

Introduction to Deep Learning

Thierry Pécot

Research Engineer

Biosit SFR UMS CNRS 3480 – Inserm 018

CZI Imaging Scientist



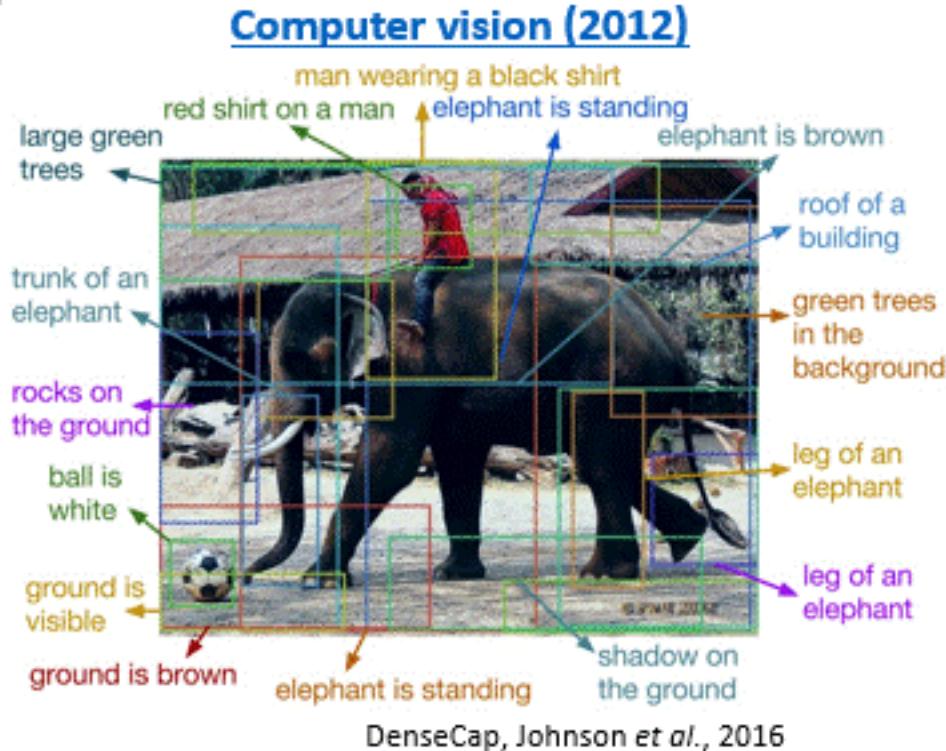
DEEP LEARNING

Speech Recognition (2011)



Tech industry

Art, music, journalism, ...

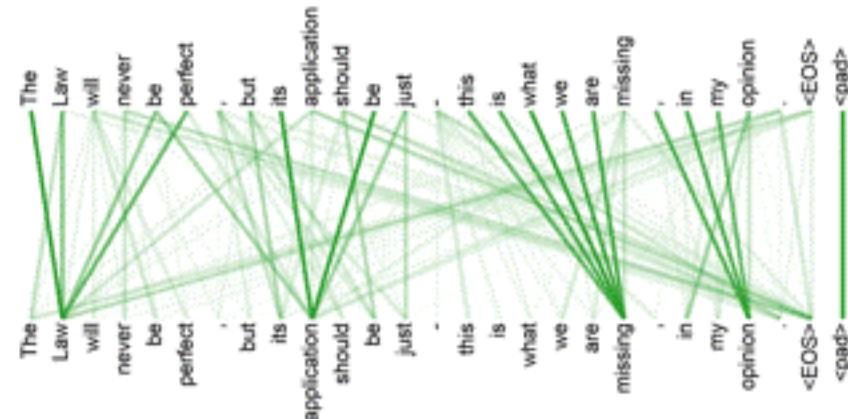


Go game (2016)

Mastering the game of Go with deep neural networks and tree search, Silver et al., Nature, 2016

Natural Language Processing for translation and chatbots

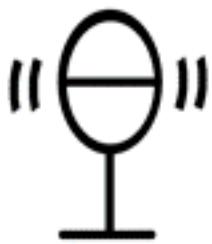
(transformers 2017)



Attention is all you need, Vaswani et al., 2017

DEEP LEARNING

Speech Recognition (2011)

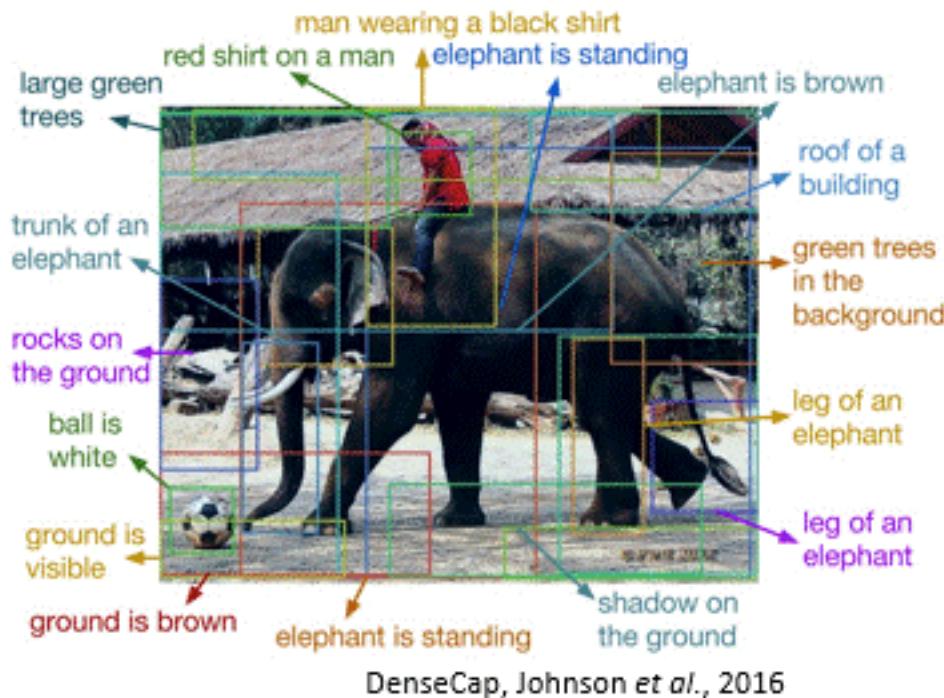


Tech industry

Art, music, journalism, ...

What is AI, more particularly deep learning?

Computer vision (2012)

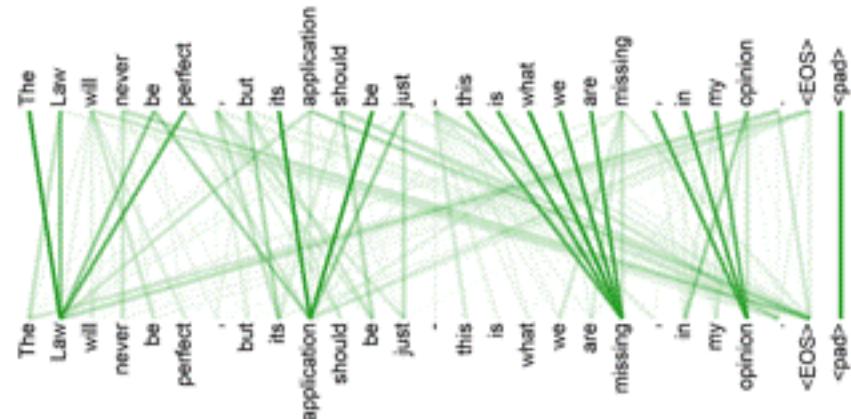


Go game (2016)

Mastering the game of Go with deep neural networks and tree search, Silver et al., Nature, 2016

Natural Language Processing for translation and chatbots

(transformers 2017)



Attention is all you need, Vaswani et al., 2017

Artificial Intelligence

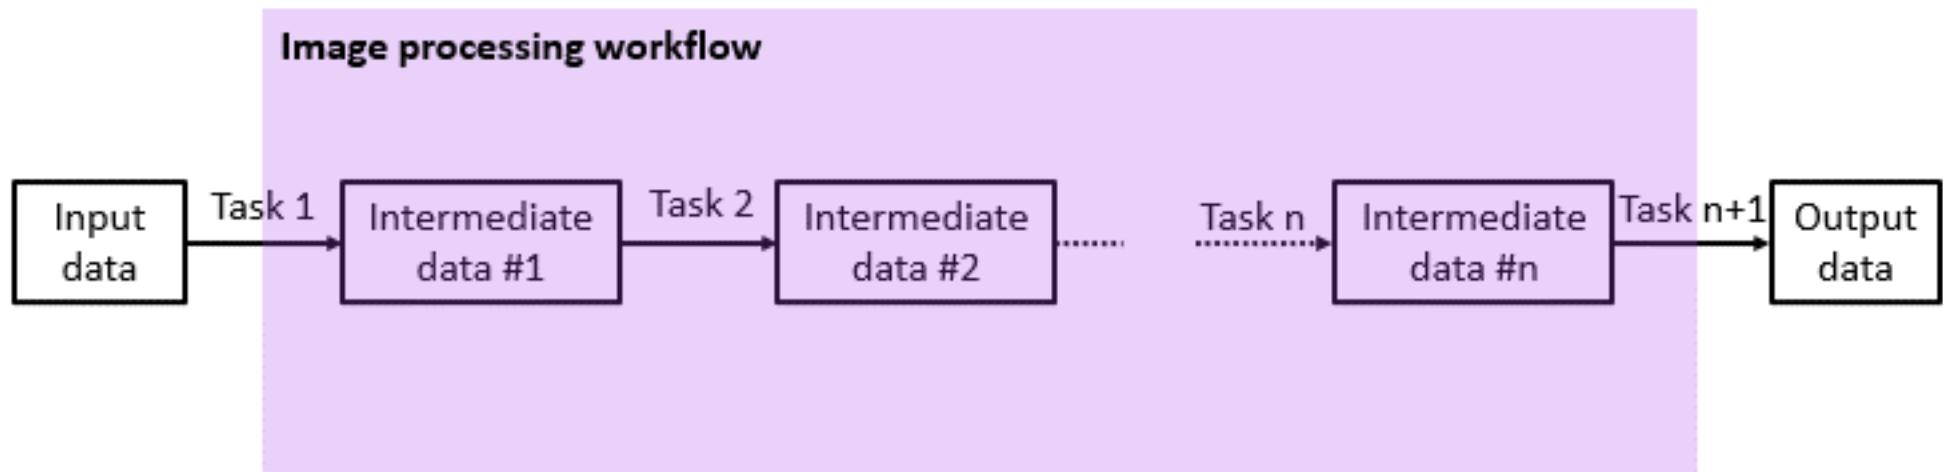
Machine Learning

- Linear/polynomial/logistic regression
- Support Vector Machines
- Decision Trees
- Random Forest
- Neural networks

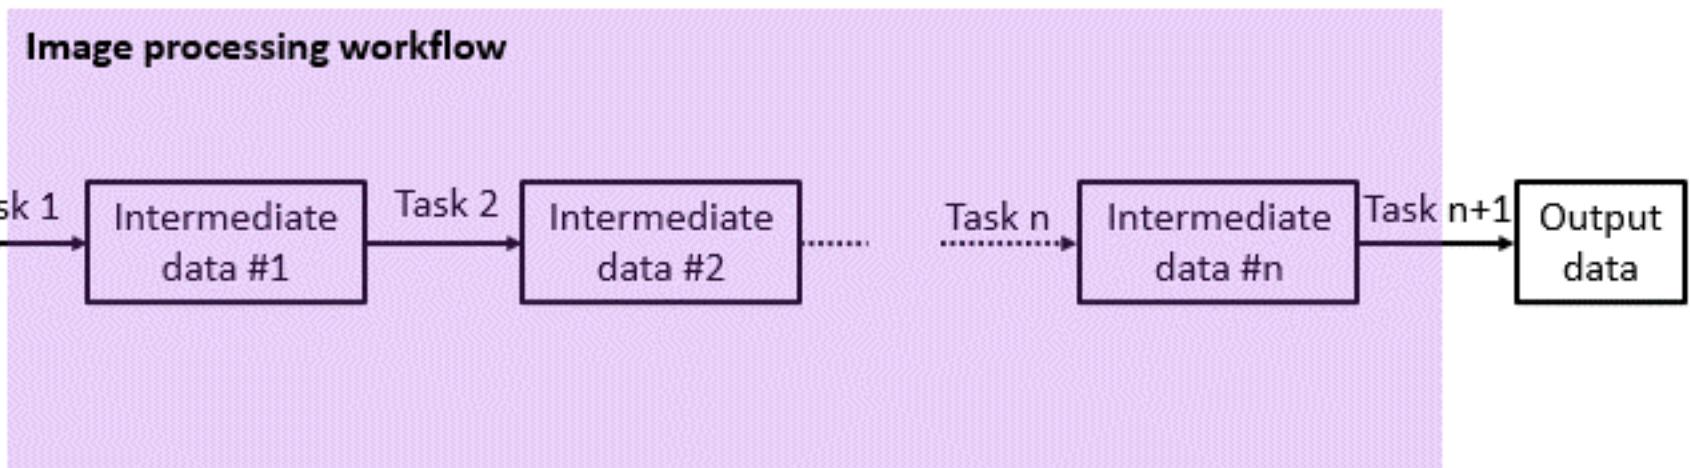
Deep Learning

- Deep Neural Networks
- Deep Convolutional Neural Networks
- Deep Recurrent Neural Networks
- Deep Reinforcement Learning
- Transformers
- Generative Adversarial Networks

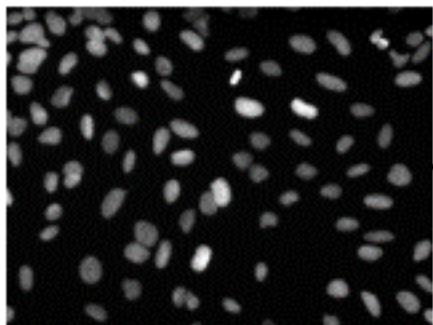
EXPLICIT PROGRAMMING



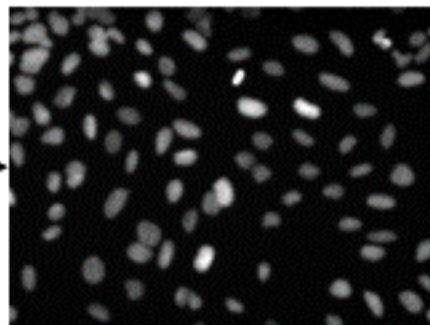
EXPLICIT PROGRAMMING



Input image

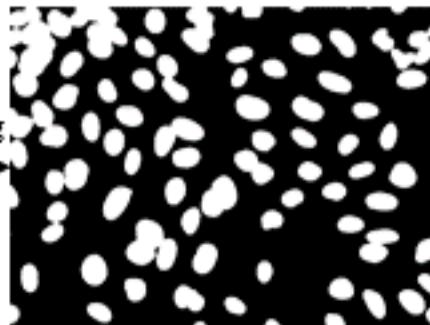


Preprocessed image



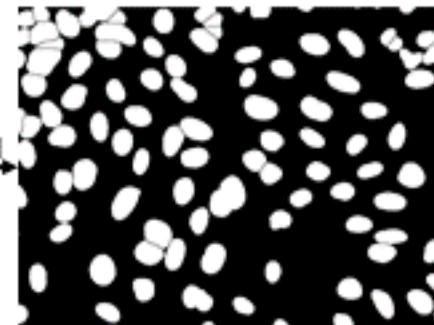
Gaussian
blurring

Segmented image



Thresholding

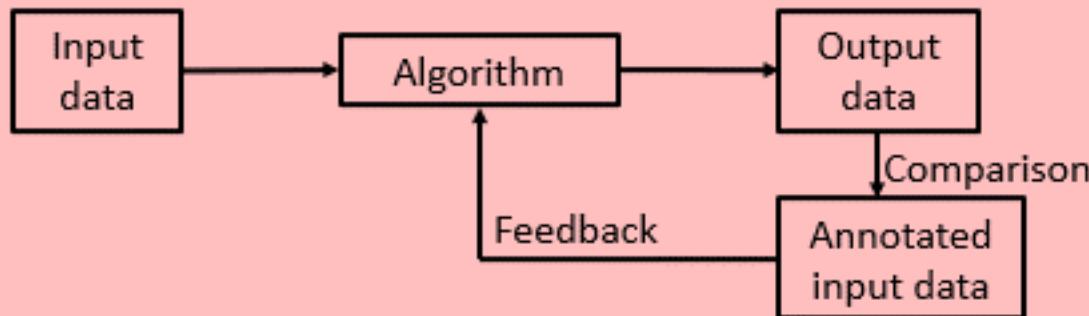
Individualized objects



Binary
watershed

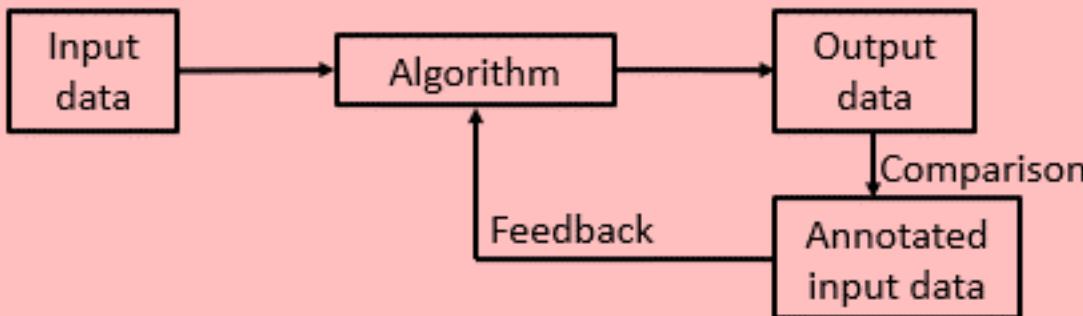
MACHINE LEARNING

Supervised Learning



MACHINE LEARNING

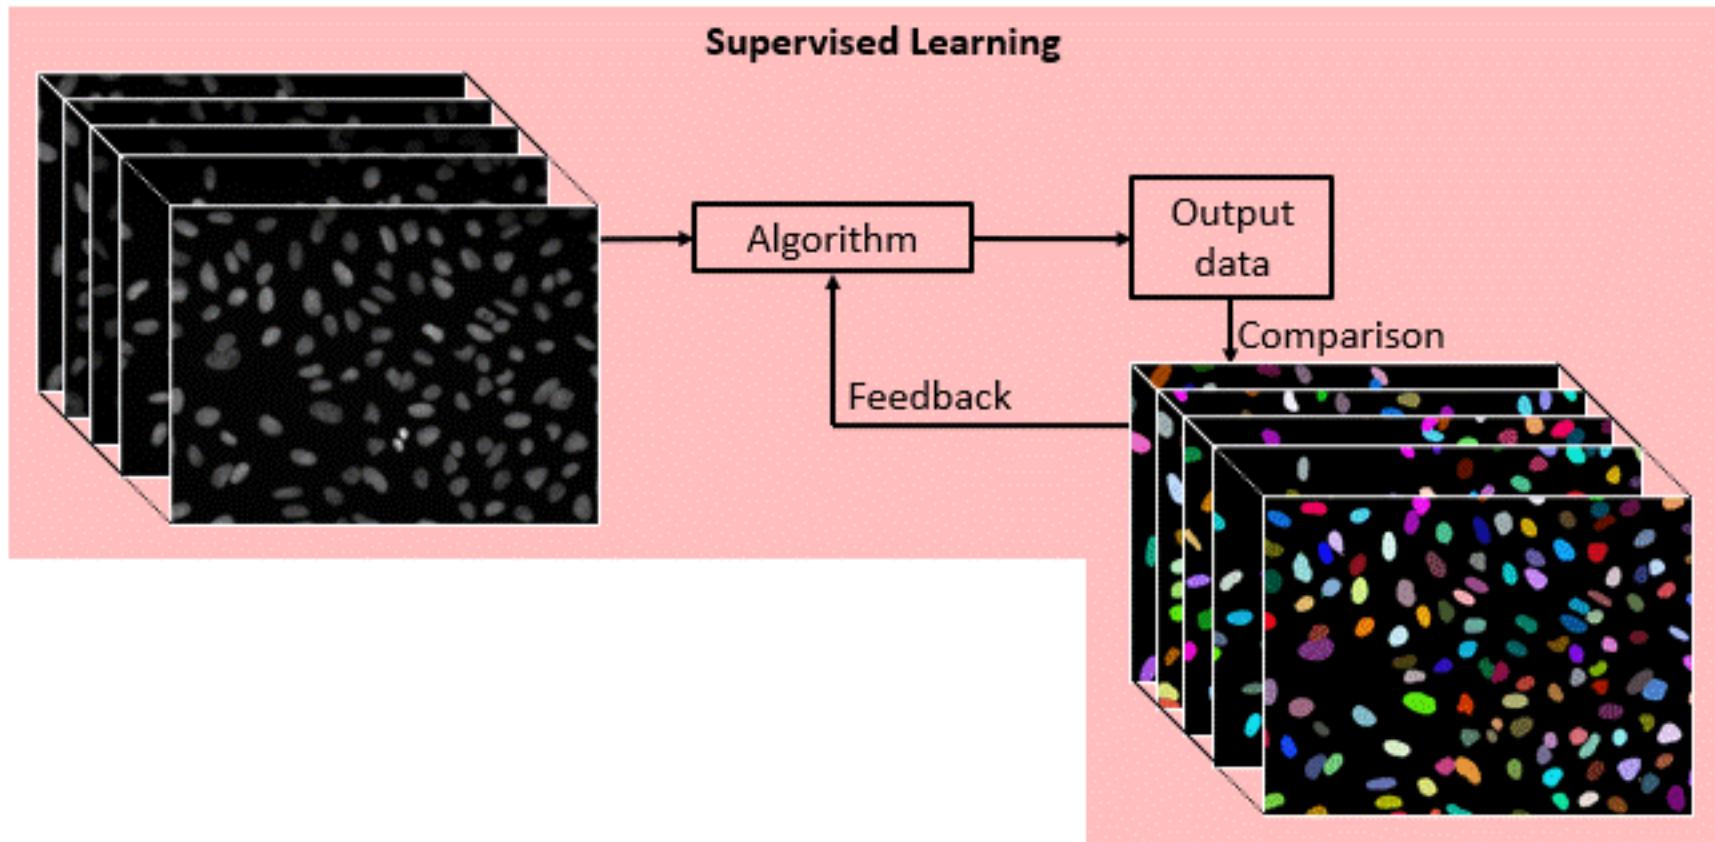
Supervised Learning



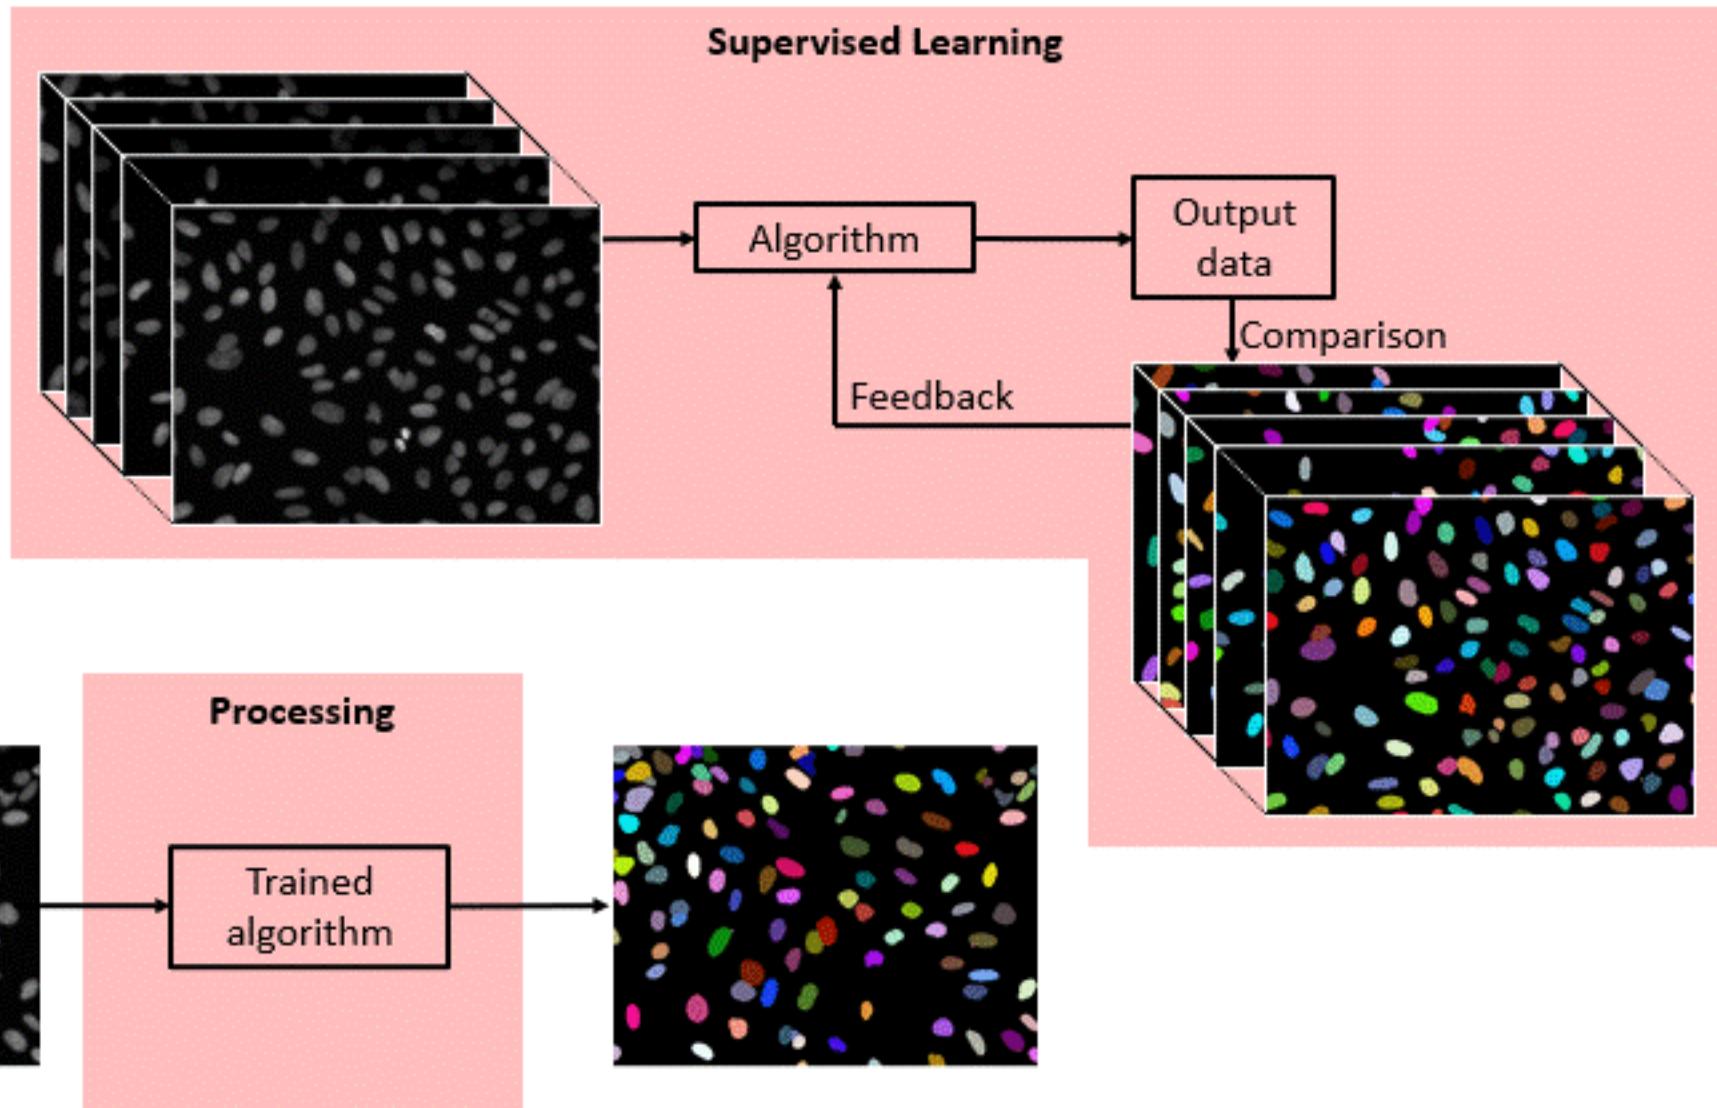
Processing



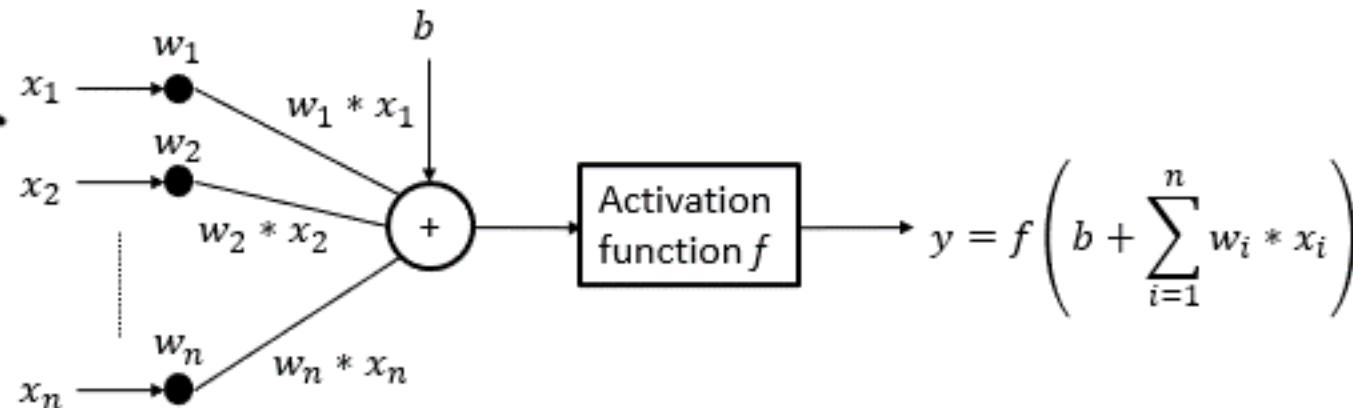
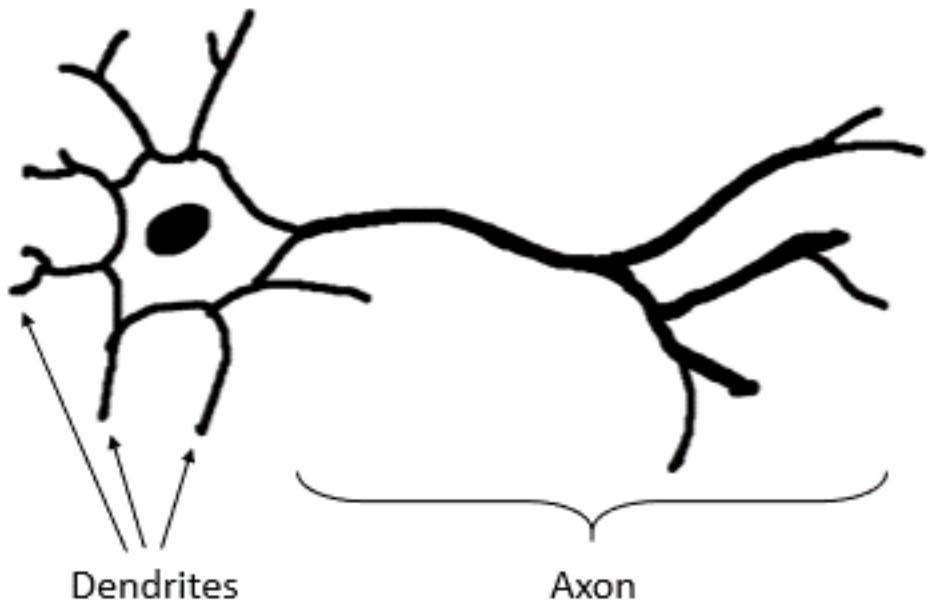
MACHINE LEARNING



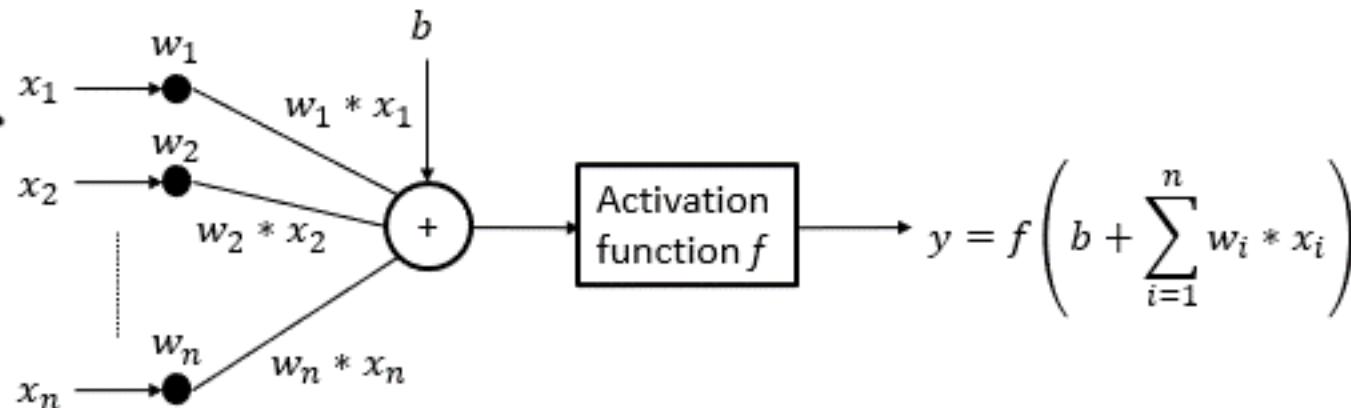
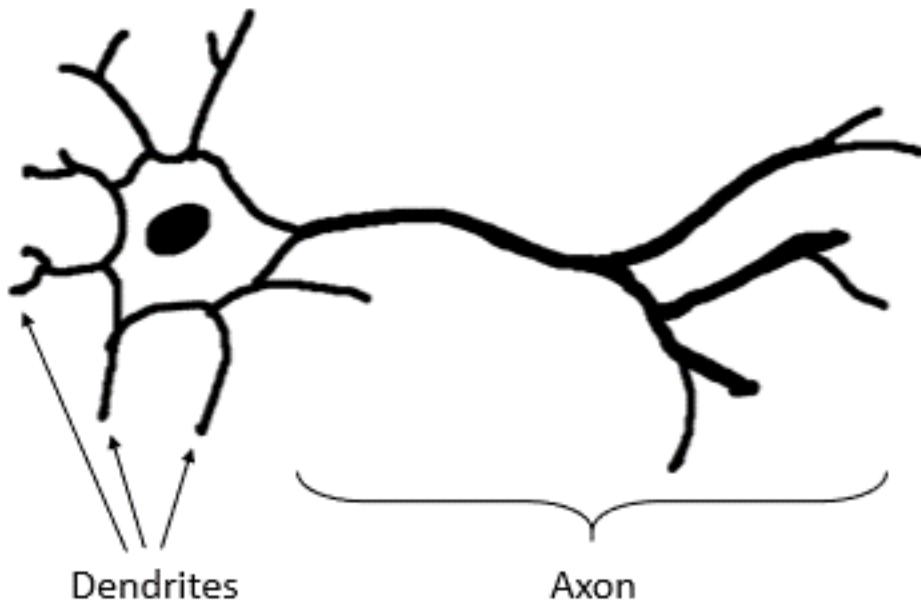
MACHINE LEARNING



ARTIFICIAL NEURON (McCULLOCH AND PITTS, 1943)



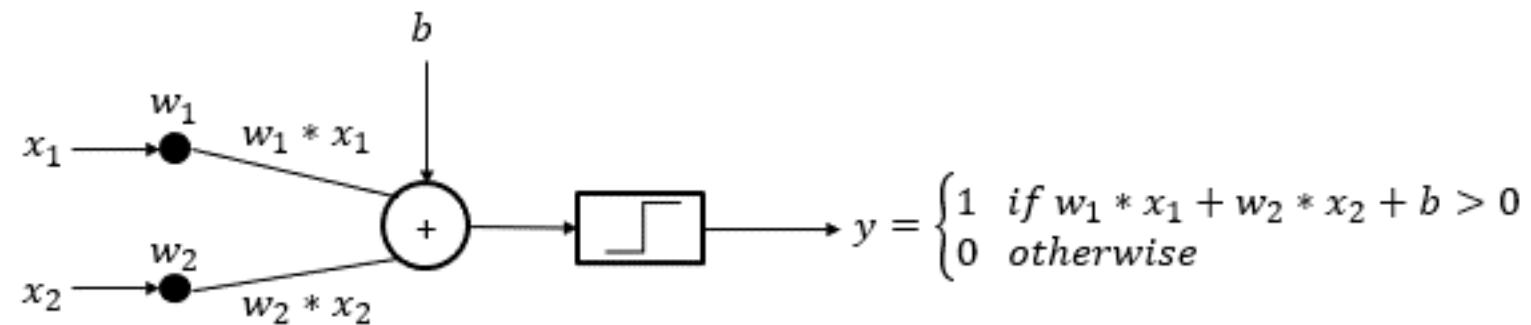
ARTIFICIAL NEURON (McCULLOCH AND PITTS, 1943)



Analogy with biological neuron:

- x_i : axons from other neurons (input axons)
- w_i : dendrites
- y : output axon

PERCEPTRON (ROSENBLATT, 1958)



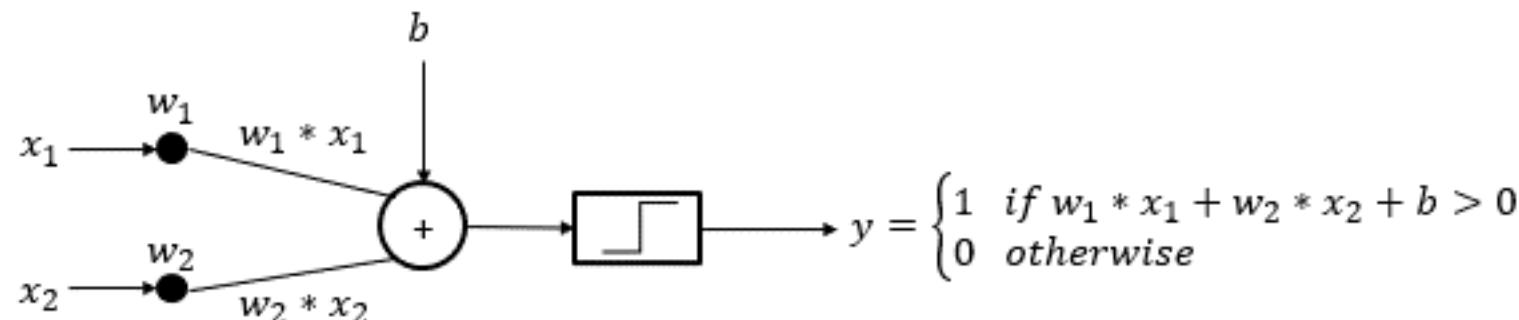
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. **Compute** y
 2. **Update** weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**



$$y = \begin{cases} 1 & \text{if } w_1 * x_1 + w_2 * x_2 + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

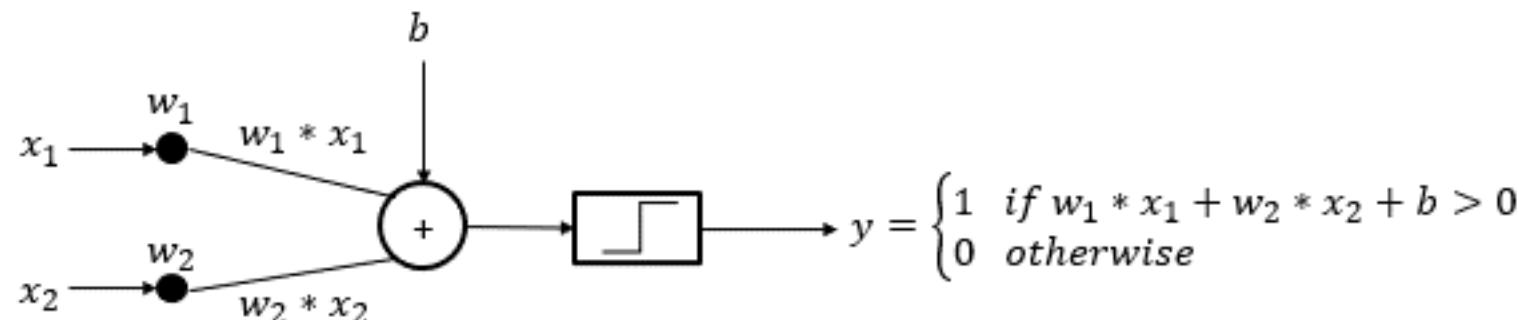
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**



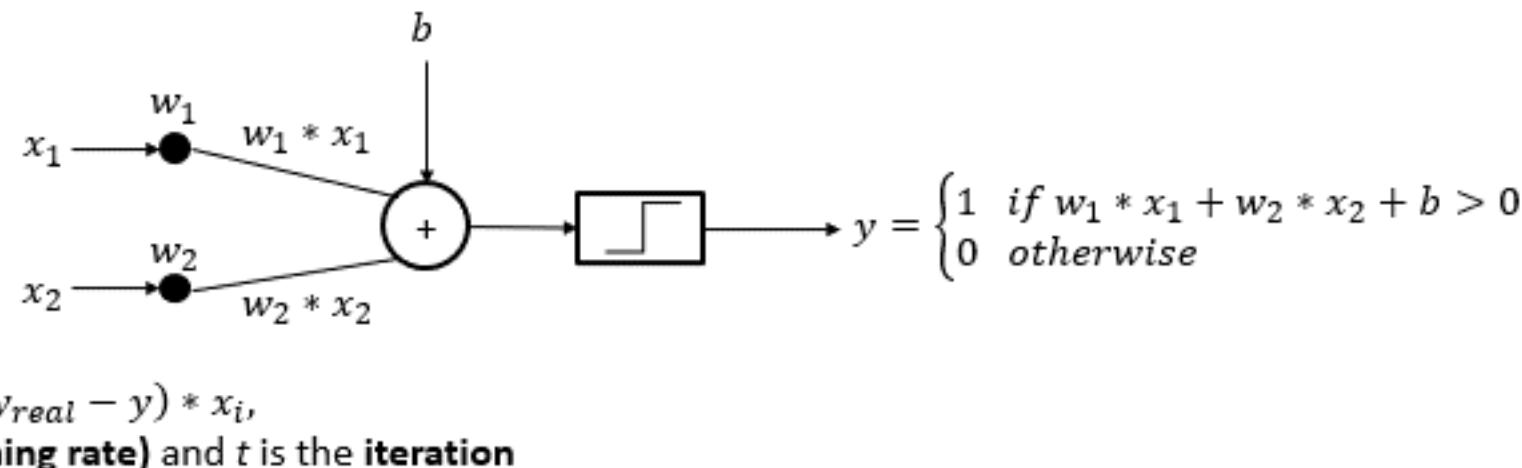
OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

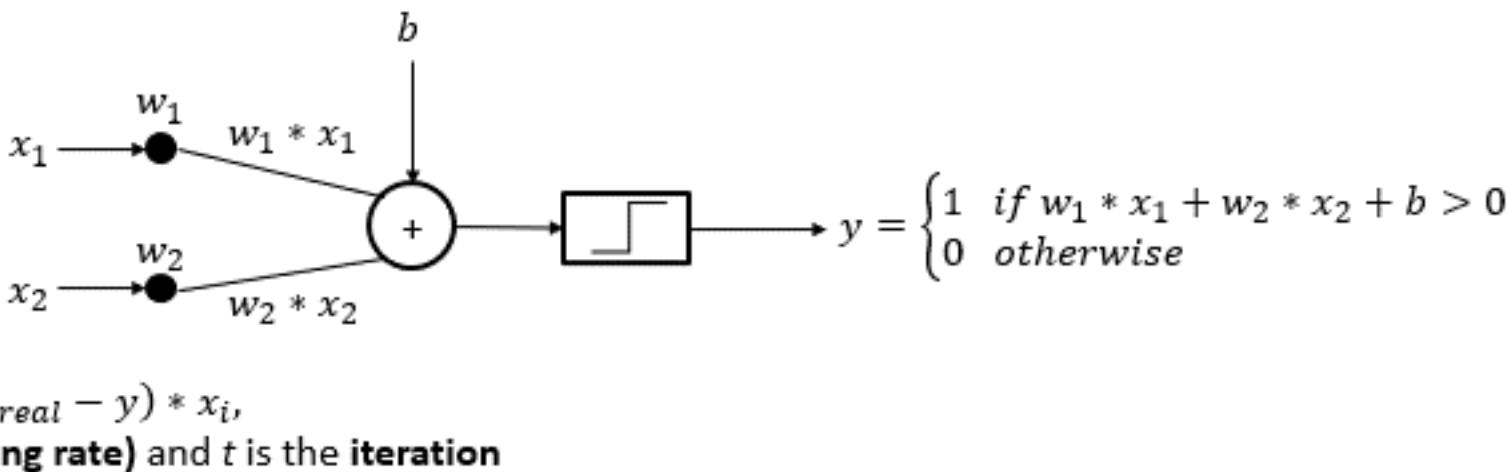
Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. **Compute** y
 2. **Update** weights such that:



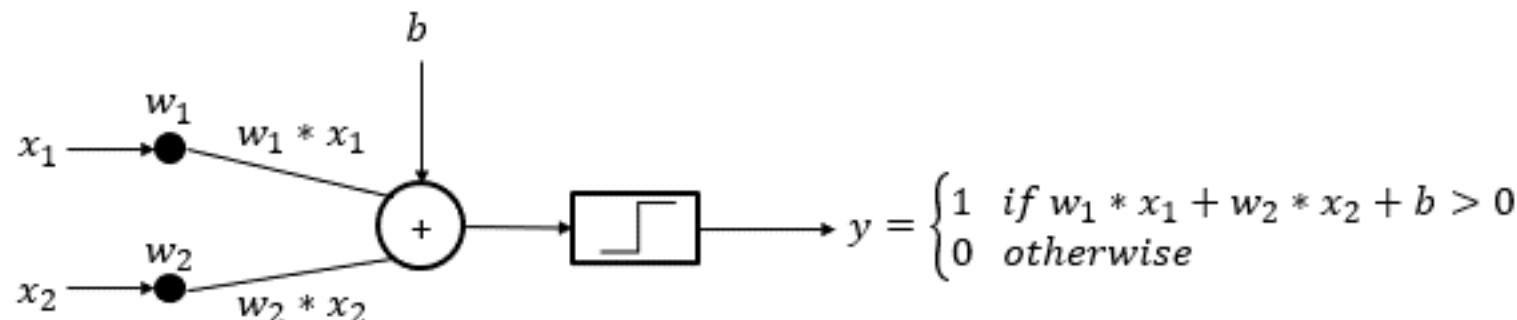
OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

PERCEPTRON (ROSENBLATT, 1958)



$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**

OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

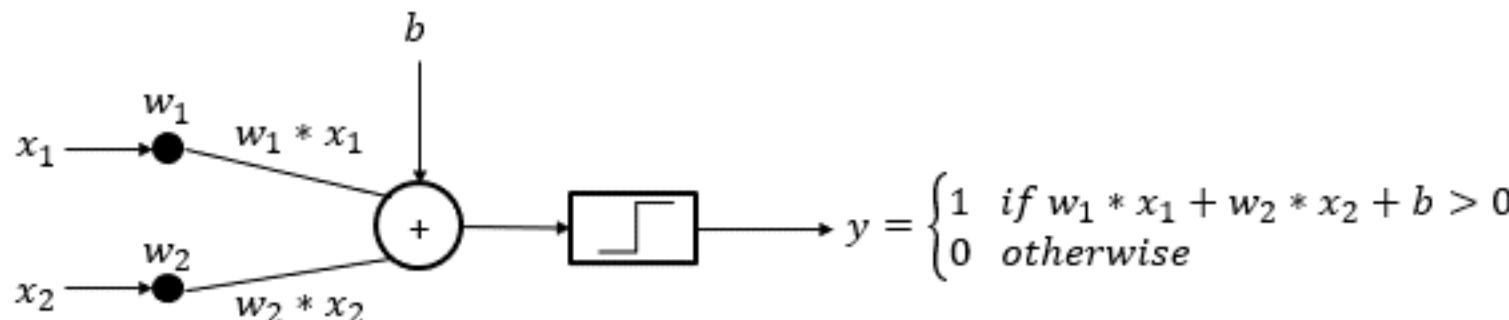
$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

PERCEPTRON (ROSENBLATT, 1958)



$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**

OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = 0 + 0.5 * 0 * 0 = 0$$

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

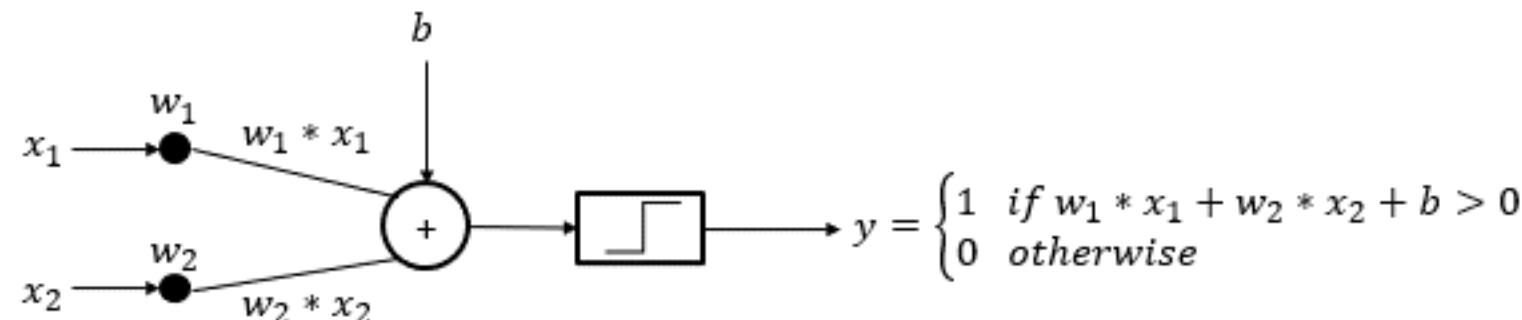
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(1) = w_2(0) + lr * (y_{real} - y) * x_2$$

$$w_2(1) = 0 + 0.5 * 0 * 0 = 0$$

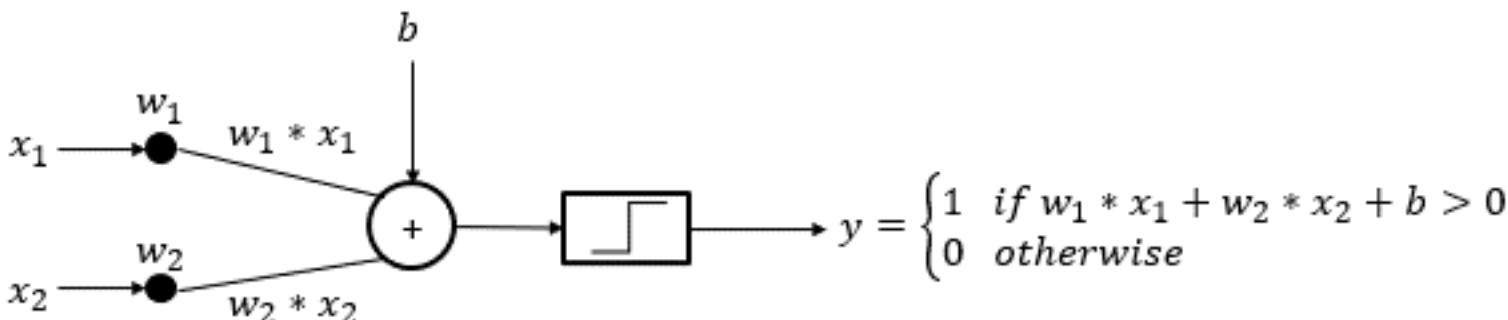
Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(\textcolor{red}{0} * \textcolor{green}{0} + \textcolor{red}{0} * \textcolor{green}{0} + \textcolor{red}{0}) = f(0) = \textcolor{purple}{0}$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = \textcolor{red}{0} + 0.5 * \textcolor{purple}{0} * \textcolor{green}{0} = 0$$

$$w_2(1) = w_2(0) + lr * (y_{real} - y) * x_2$$

$$w_2(1) = \textcolor{red}{0} + 0.5 * \textcolor{purple}{0} * \textcolor{green}{0} = 0$$

$$b(1) = b(0) + lr * (y_{real} - y)$$

$$b(1) = \textcolor{red}{0} + 0.5 * \textcolor{purple}{0} = 0$$

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

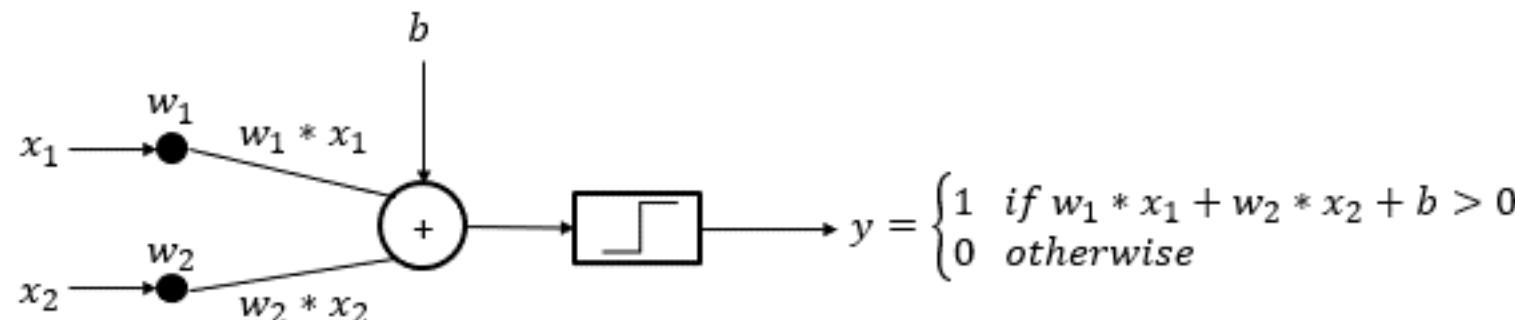
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (learning rate) and t is the iteration



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(1) = w_2(0) + lr * (y_{real} - y) * x_2$$

$$w_2(1) = 0 + 0.5 * 0 * 0 = 0$$

$$b(1) = b(0) + lr * (y_{real} - y)$$

$$b(1) = 0 + 0.5 * 0 = 0$$

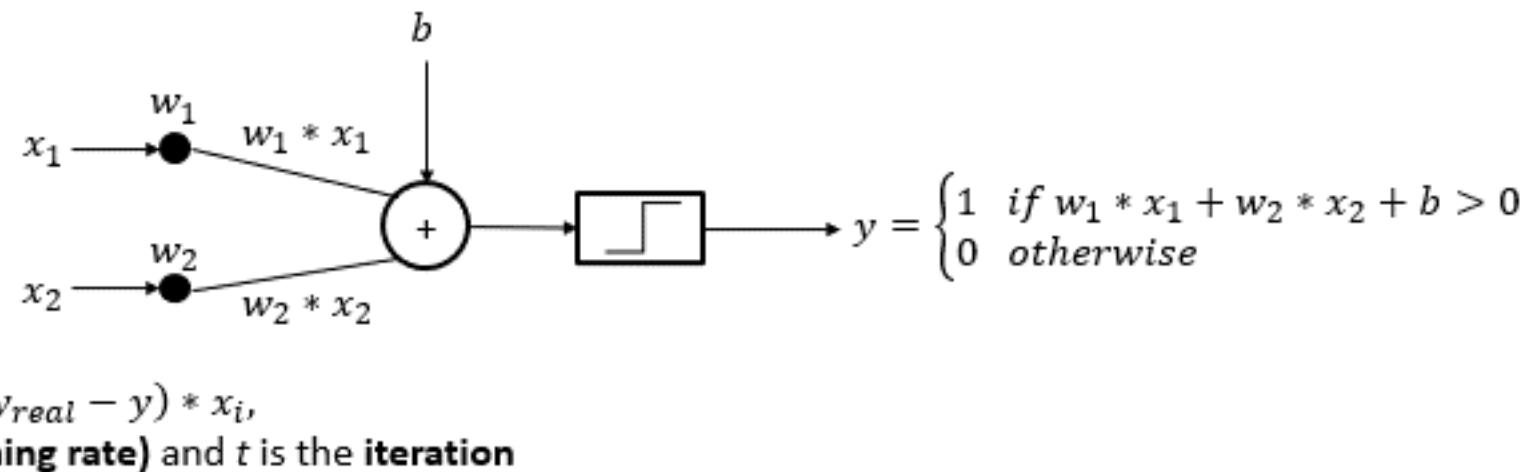
Iteration 1:

$w_1(1)$	$w_2(1)$	$b(1)$
0	0	0

PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

Iteration 1:

$w_1(1)$	$w_2(1)$	$b(1)$
0	0	0

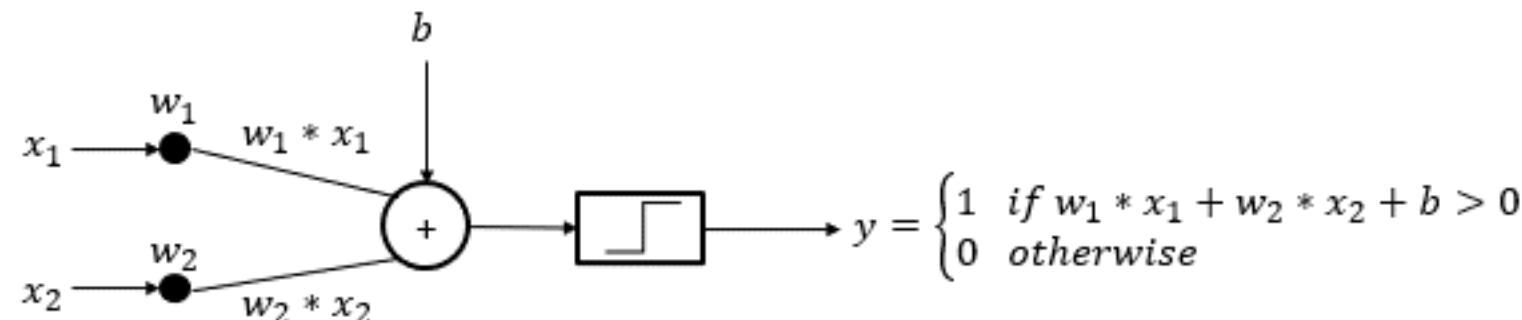
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (learning rate) and t is the iteration



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(1) * x_1 + w_2(1) * x_2 + b(1))$$

$$y = f(0 * 1 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(2) = w_1(1) + lr * (y_{real} - y) * x_1$$

$$w_1(2) = 0 + 0.5 * 1 * 1 = 0.5$$

$$w_2(2) = w_2(1) + lr * (y_{real} - y) * x_2$$

$$w_2(2) = 0 + 0.5 * 0 * 0 = 0$$

$$b(2) = b(1) + lr * (y_{real} - y)$$

$$b(2) = 0 + 0.5 * 1 = 0.5$$

Iteration 1:

$w_1(1)$	$w_2(1)$	$b(1)$
0	0	0

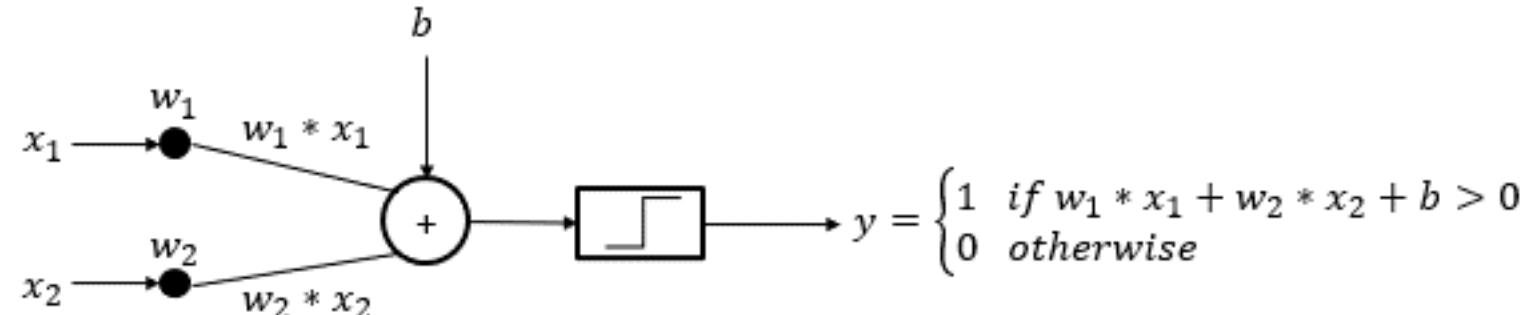
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (learning rate) and t is the iteration



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(1) * x_1 + w_2(1) * x_2 + b(1))$$

$$y = f(0 * 1 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(2) = w_1(1) + lr * (y_{real} - y) * x_1$$

$$w_1(2) = 0 + 0.5 * 1 * 1 = 0.5$$

$$w_2(2) = w_2(1) + lr * (y_{real} - y) * x_2$$

$$w_2(2) = 0 + 0.5 * 0 * 0 = 0$$

$$b(2) = b(1) + lr * (y_{real} - y)$$

$$b(2) = 0 + 0.5 * 1 = 0.5$$

Iteration 2:

$w_1(2)$	$w_2(2)$	$b(2)$
0.5	0	0.5

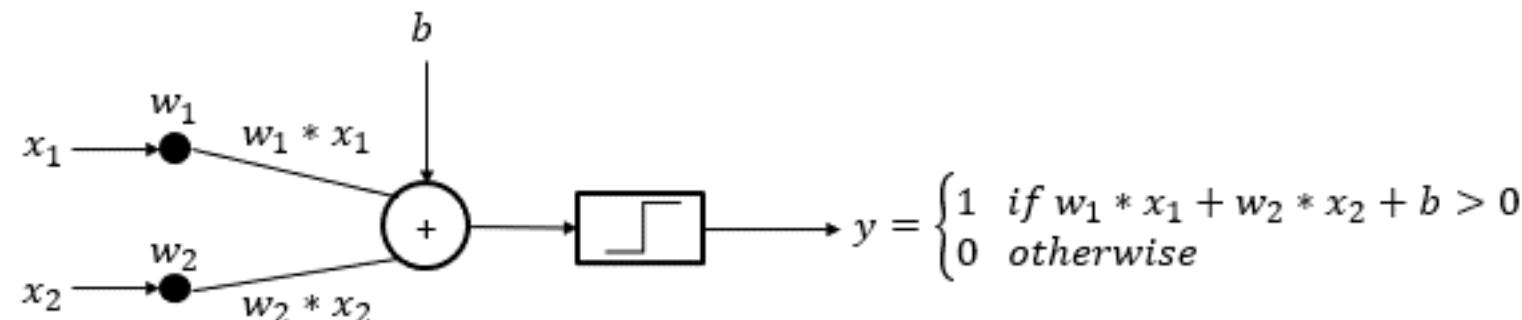
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (learning rate) and t is the iteration



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(2) * x_1 + w_2(2) * x_2 + b(2))$$

$$y = f(0.5 * 0 + 0 * 1 + 0.5) = f(0.5) = 1$$

$$w_1(3) = w_1(2) + lr * (y_{real} - y) * x_1$$

$$w_1(3) = 0.5 + 0.5 * 0 * 0 = 0.5$$

$$w_2(3) = w_2(2) + lr * (y_{real} - y) * x_2$$

$$w_2(3) = 0 + 0.5 * 0 * 1 = 0$$

$$b(3) = b(2) + lr * (y_{real} - y)$$

$$b(3) = 0.5 + 0.5 * 0 = 0.5$$

Iteration 2:

$w_1(2)$	$w_2(2)$	$b(2)$
0.5	0	0.5

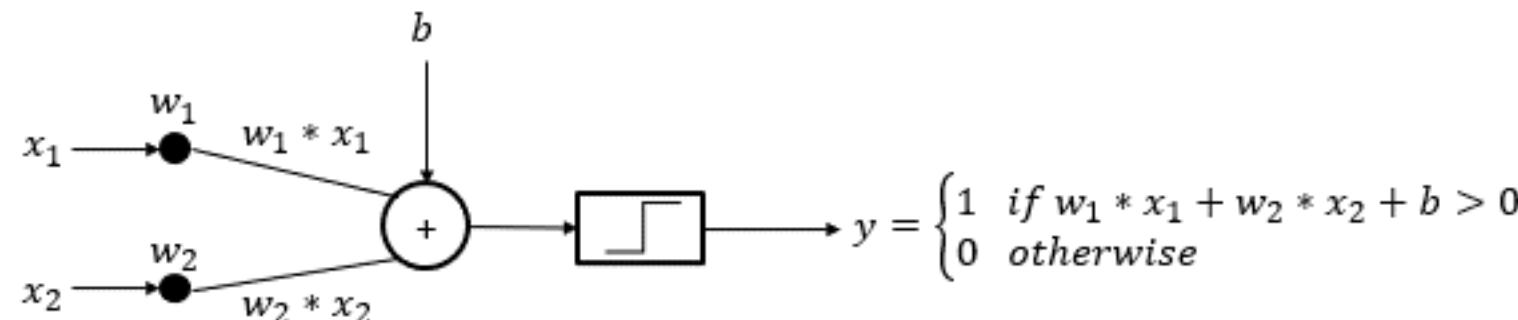
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (learning rate) and t is the iteration



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(2) * x_1 + w_2(2) * x_2 + b(2))$$

$$y = f(0.5 * 0 + 0 * 1 + 0.5) = f(0.5) = 1$$

$$w_1(3) = w_1(2) + lr * (y_{real} - y) * x_1$$

$$w_1(3) = 0.5 + 0.5 * 0 * 0 = 0.5$$

$$w_2(3) = w_2(2) + lr * (y_{real} - y) * x_2$$

$$w_2(3) = 0 + 0.5 * 0 * 1 = 0$$

$$b(3) = b(2) + lr * (y_{real} - y)$$

$$b(3) = 0.5 + 0.5 * 0 = 0.5$$

Iteration 3:

$w_1(3)$	$w_2(3)$	$b(3)$
0.5	0	0.5

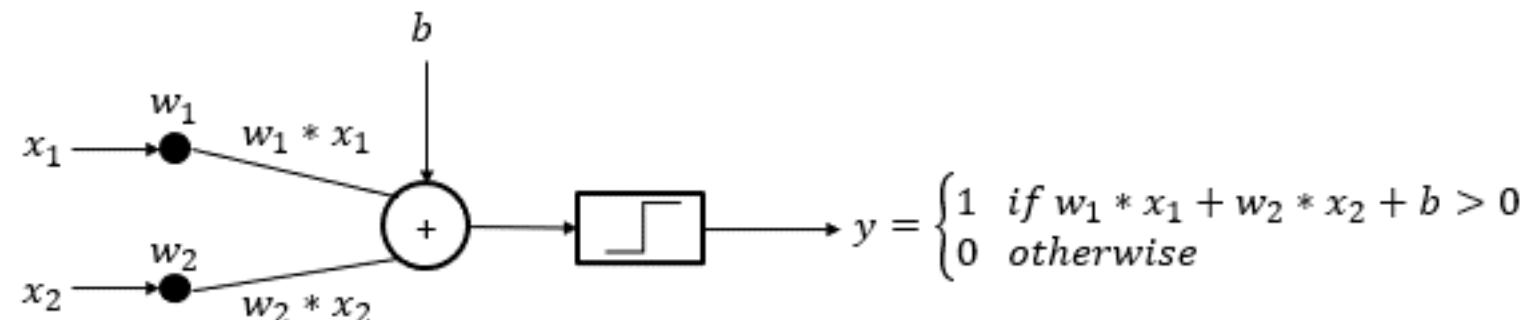
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is a constant (**learning rate**) and t is the **iteration**



OR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	1

$$y = f(w_1(3) * x_1 + w_2(3) * x_2 + b(3))$$

$$y = f(0.5 * 1 + 0 * 1 + 0.5) = f(1) = 1$$

$$w_1(4) = w_1(3) + lr * (y_{real} - y) * x_1$$

$$w_1(4) = 0.5 + 0.5 * 0 * 1 = 0.5$$

$$w_2(4) = w_2(3) + lr * (y_{real} - y) * x_2$$

$$w_2(4) = 0 + 0.5 * 0 * 1 = 0$$

$$b(4) = b(3) + lr * (y_{real} - y)$$

$$b(4) = 0.5 + 0.5 * 0 = 0.5$$

Iteration 3:

$w_1(3)$	$w_2(3)$	$b(3)$
0.5	0	0.5

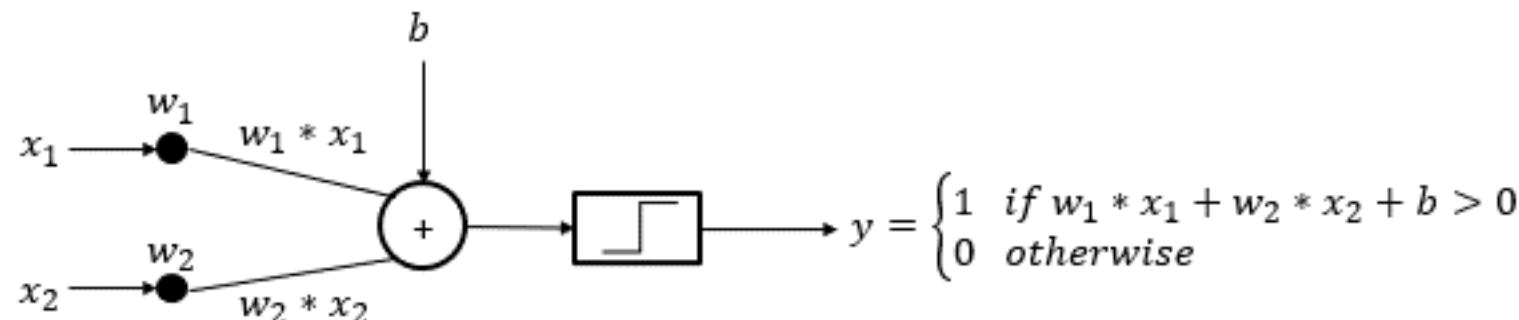
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

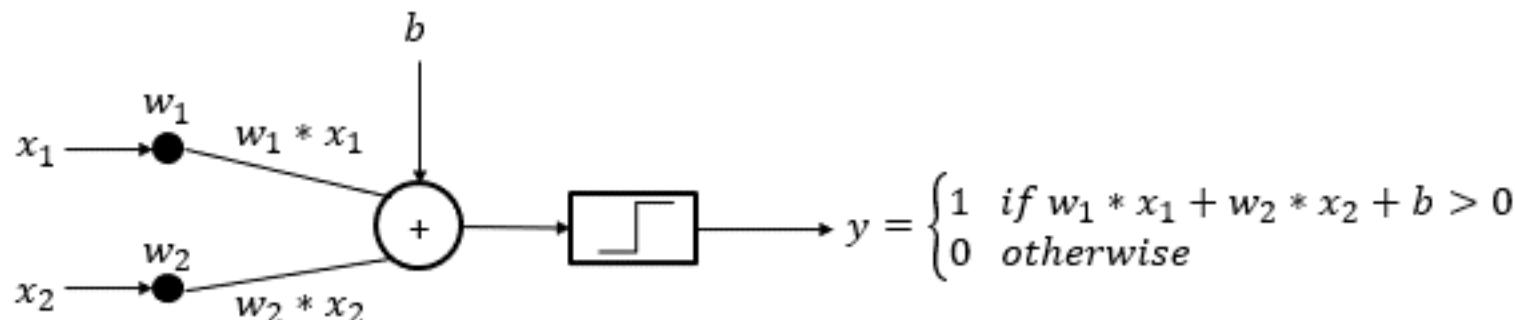
where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

PERCEPTRON (ROSENBLATT, 1958)



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

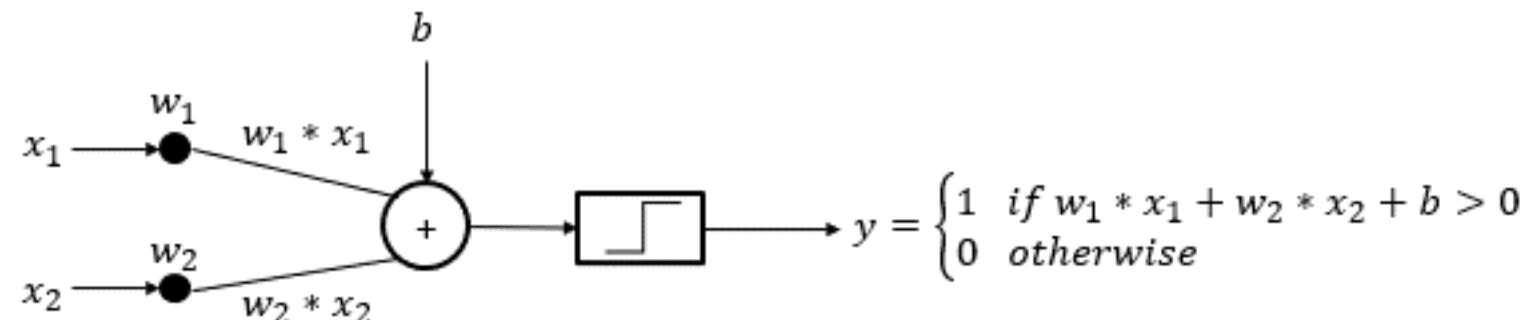
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(1) = w_2(0) + lr * (y_{real} - y) * x_2$$

$$w_2(1) = 0 + 0.5 * 0 * 0 = 0$$

$$b(1) = b(0) + lr * (y_{real} - y)$$

$$b(1) = 0 + 0.5 * 0 = 0$$

Initialization:

$w_1(0)$	$w_2(0)$	$b(0)$
0	0	0

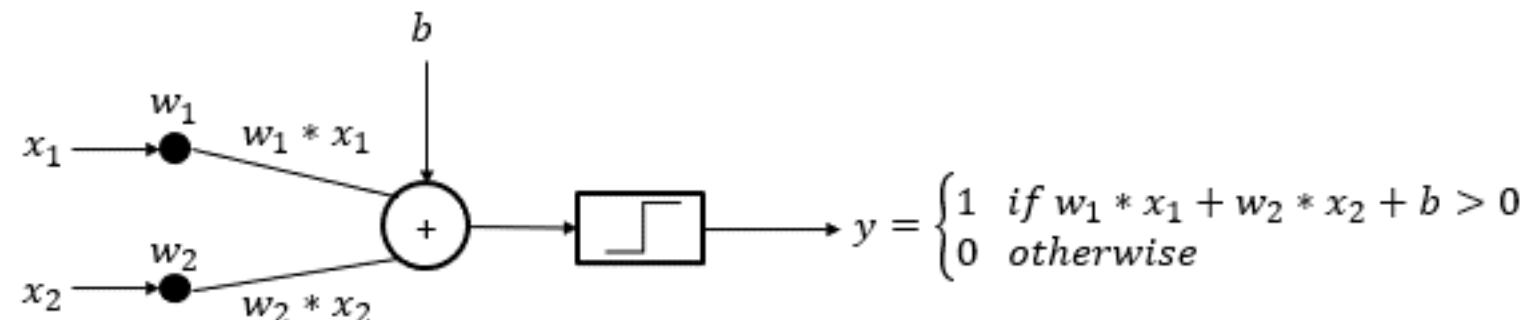
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(0) * x_1 + w_2(0) * x_2 + b(0))$$

$$y = f(0 * 0 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(1) = w_1(0) + lr * (y_{real} - y) * x_1$$

$$w_1(1) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(1) = w_2(0) + lr * (y_{real} - y) * x_2$$

$$w_2(1) = 0 + 0.5 * 0 * 0 = 0$$

$$b(1) = b(0) + lr * (y_{real} - y)$$

$$b(1) = 0 + 0.5 * 0 = 0$$

Iteration 1:

$w_1(1)$	$w_2(1)$	$b(1)$
0	0	0

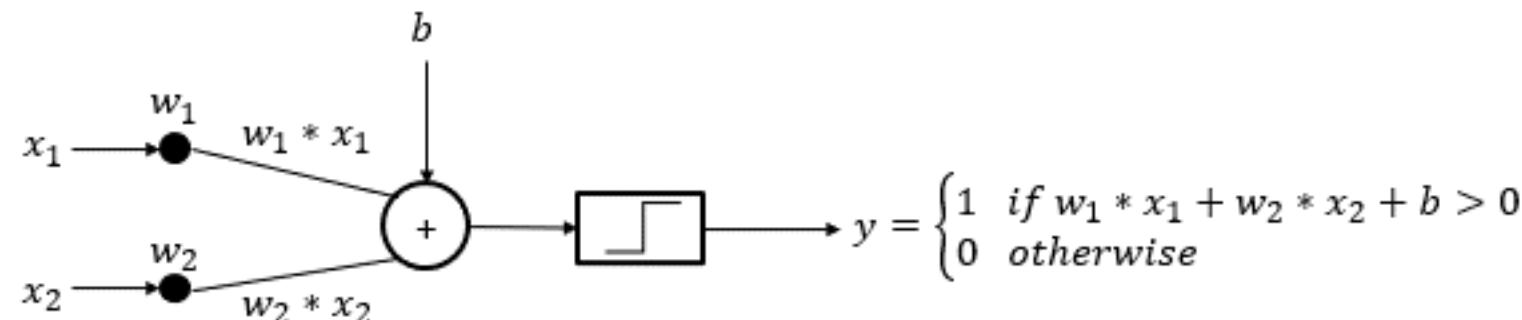
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(1) * x_1 + w_2(1) * x_2 + b(1))$$

$$y = f(0 * 1 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(2) = w_1(1) + lr * (y_{real} - y) * x_1$$

$$w_1(2) = 0 + 0.5 * 0 * 1 = 0$$

$$w_2(2) = w_2(1) + lr * (y_{real} - y) * x_2$$

$$w_2(2) = 0 + 0.5 * 0 * 0 = 0$$

$$b(2) = b(1) + lr * (y_{real} - y)$$

$$b(2) = 0 + 0.5 * 0 = 0$$

Iteration 1:

$w_1(1)$	$w_2(1)$	$b(1)$
0	0	0

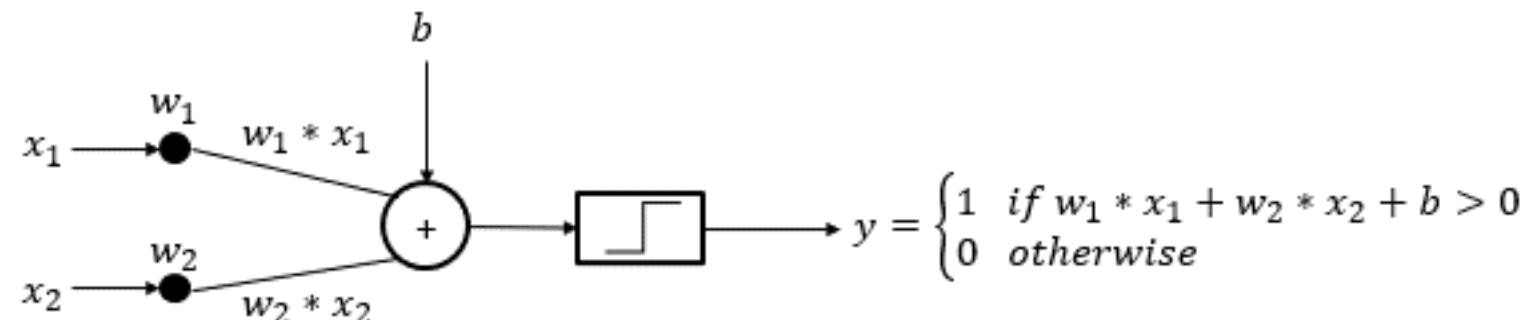
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(1) * x_1 + w_2(1) * x_2 + b(1))$$

$$y = f(0 * 1 + 0 * 0 + 0) = f(0) = 0$$

$$w_1(2) = w_1(1) + lr * (y_{real} - y) * x_1$$

$$w_1(2) = 0 + 0.5 * 0 * 1 = 0$$

$$w_2(2) = w_2(1) + lr * (y_{real} - y) * x_2$$

$$w_2(2) = 0 + 0.5 * 0 * 0 = 0$$

$$b(2) = b(1) + lr * (y_{real} - y)$$

$$b(2) = 0 + 0.5 * 0 = 0$$

Iteration 2:

$w_1(2)$	$w_2(2)$	$b(2)$
0	0	0

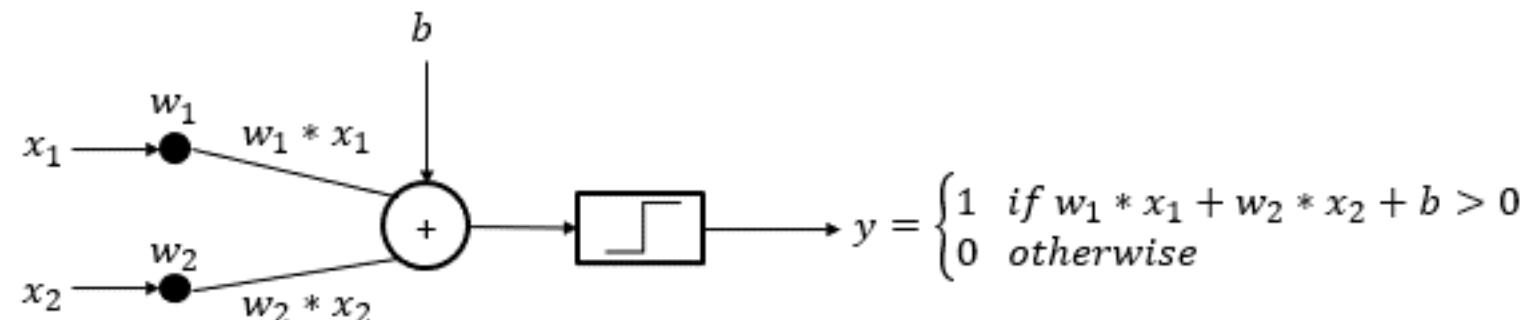
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(2) * x_1 + w_2(2) * x_2 + b(2))$$

$$y = f(0 * 0 + 0 * 1 + 0) = f(0) = 0$$

$$w_1(3) = w_1(2) + lr * (y_{real} - y) * x_1$$

$$w_1(3) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(3) = w_2(2) + lr * (y_{real} - y) * x_2$$

$$w_2(3) = 0 + 0.5 * 0 * 1 = 0$$

$$b(3) = b(2) + lr * (y_{real} - y)$$

$$b(3) = 0 + 0.5 * 0 = 0$$

Iteration 2:

$w_1(2)$	$w_2(2)$	$b(2)$
0	0	0

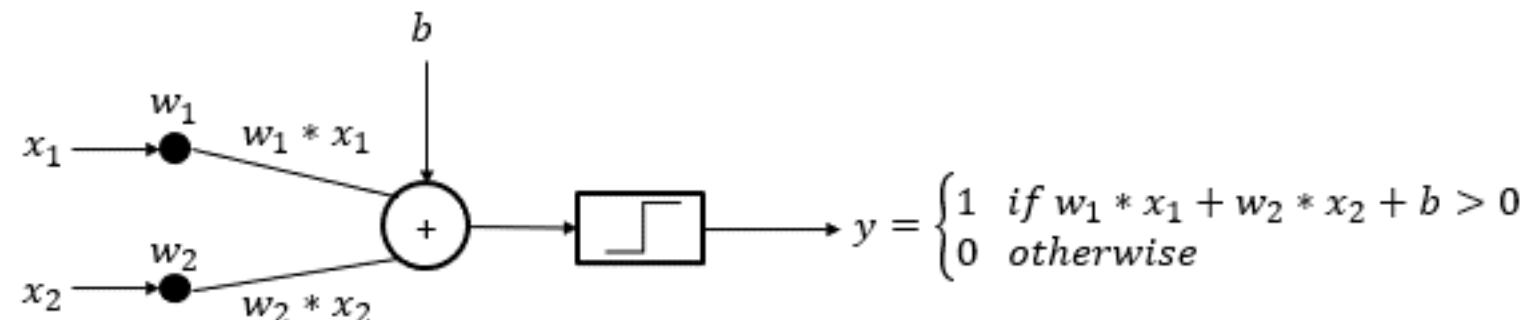
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(2) * x_1 + w_2(2) * x_2 + b(2))$$

$$y = f(0 * 0 + 0 * 1 + 0) = f(0) = 0$$

$$w_1(3) = w_1(2) + lr * (y_{real} - y) * x_1$$

$$w_1(3) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(3) = w_2(2) + lr * (y_{real} - y) * x_2$$

$$w_2(3) = 0 + 0.5 * 0 * 1 = 0$$

$$b(3) = b(2) + lr * (y_{real} - y)$$

$$b(3) = 0 + 0.5 * 0 = 0$$

Iteration 3:

$w_1(3)$	$w_2(3)$	$b(3)$
0	0	0

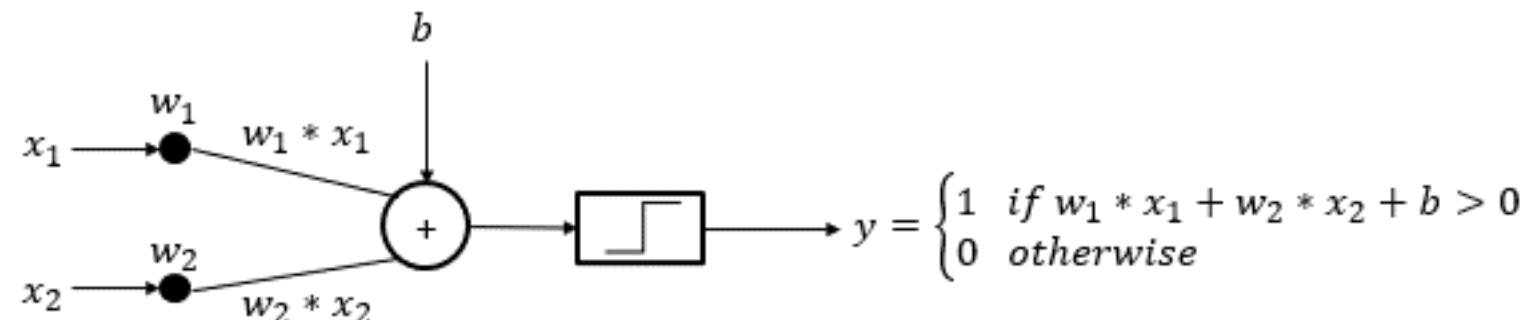
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(3) * x_1 + w_2(3) * x_2 + b(3))$$

$$y = f(0 * 1 + 0 * 1 + 0) = f(0) = 0$$

$$w_1(4) = w_1(3) + lr * (y_{real} - y) * x_1$$

$$w_1(4) = 0 + 0.5 * 1 * 1 = 0.5$$

$$w_2(4) = w_2(3) + lr * (y_{real} - y) * x_2$$

$$w_2(4) = 0 + 0.5 * 1 * 1 = 0.5$$

$$b(4) = b(3) + lr * (y_{real} - y)$$

$$b(4) = 0 + 0.5 * 1 = 0.5$$

Iteration 3:

$w_1(3)$	$w_2(3)$	$b(3)$
0	0	0

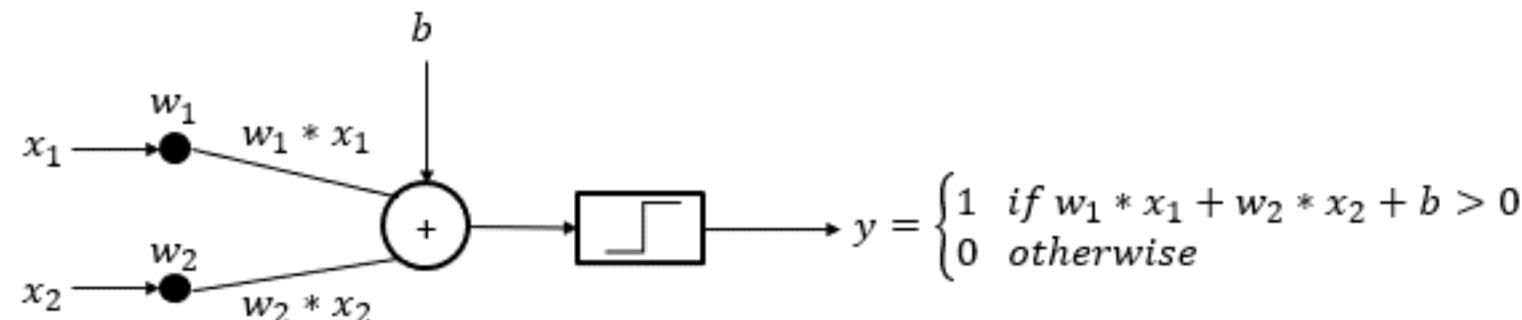
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(3) * x_1 + w_2(3) * x_2 + b(3))$$

$$y = f(0 * 1 + 0 * 1 + 0) = f(0) = 0$$

$$w_1(4) = w_1(3) + lr * (y_{real} - y) * x_1$$

$$w_1(4) = 0 + 0.5 * 1 * 1 = 0.5$$

$$w_2(4) = w_2(3) + lr * (y_{real} - y) * x_2$$

$$w_2(4) = 0 + 0.5 * 1 * 1 = 0.5$$

$$b(4) = b(3) + lr * (y_{real} - y)$$

$$b(4) = 0 + 0.5 * 1 = 0.5$$

Iteration 4:

$w_1(4)$	$w_2(4)$	$b(4)$
0.5	0.5	0.5

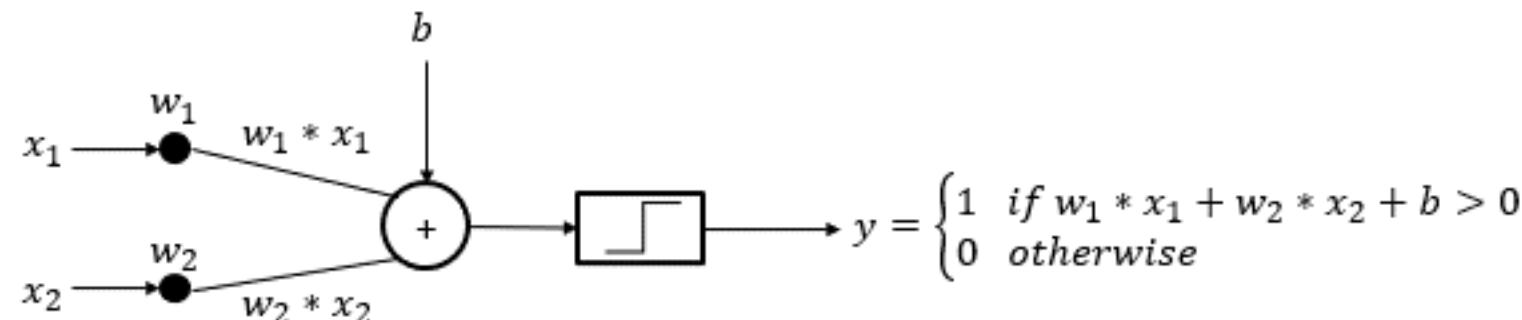
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(4) * x_1 + w_2(4) * x_2 + b(4))$$

$$y = f(0.5 * 0 + 0.5 * 0 + 0.5) = f(0.5) = 1$$

$$w_1(5) = w_1(4) + lr * (y_{real} - y) * x_1$$

$$w_1(5) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

$$w_2(5) = w_2(4) + lr * (y_{real} - y) * x_2$$

$$w_2(5) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

$$b(5) = b(4) + lr * (y_{real} - y)$$

$$b(5) = 0.5 + 0.5 * (-1) = 0$$

Iteration 4:

$w_1(4)$	$w_2(4)$	$b(4)$
0.5	0.5	0.5

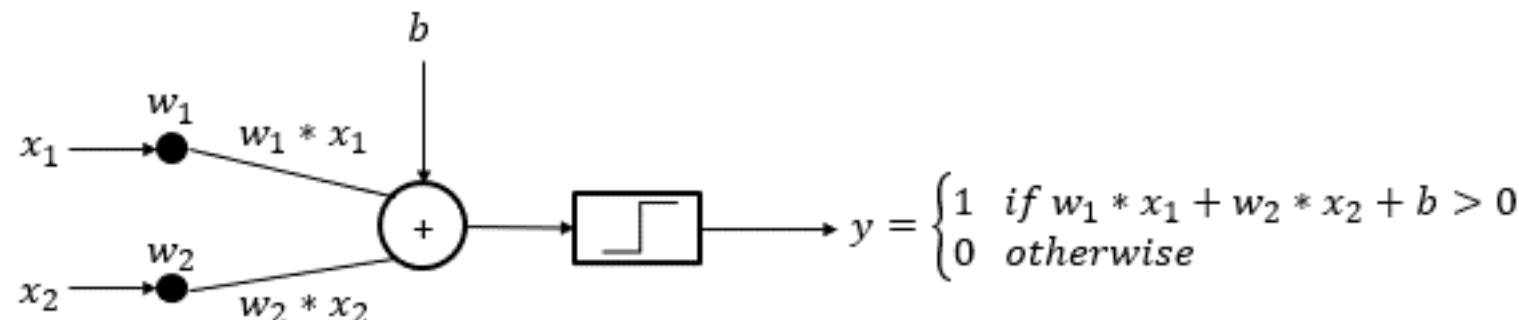
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(4) * x_1 + w_2(4) * x_2 + b(4))$$

$$y = f(0.5 * 0 + 0.5 * 0 + 0.5) = f(0.5) = 1$$

$$w_1(5) = w_1(4) + lr * (y_{real} - y) * x_1$$

$$w_1(5) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

$$w_2(5) = w_2(4) + lr * (y_{real} - y) * x_2$$

$$w_2(5) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

Iteration 5:

$w_1(5)$	$w_2(5)$	$b(5)$
0.5	0.5	0

$$b(5) = b(4) + lr * (y_{real} - y)$$

$$b(5) = 0.5 + 0.5 * (-1) = 0$$

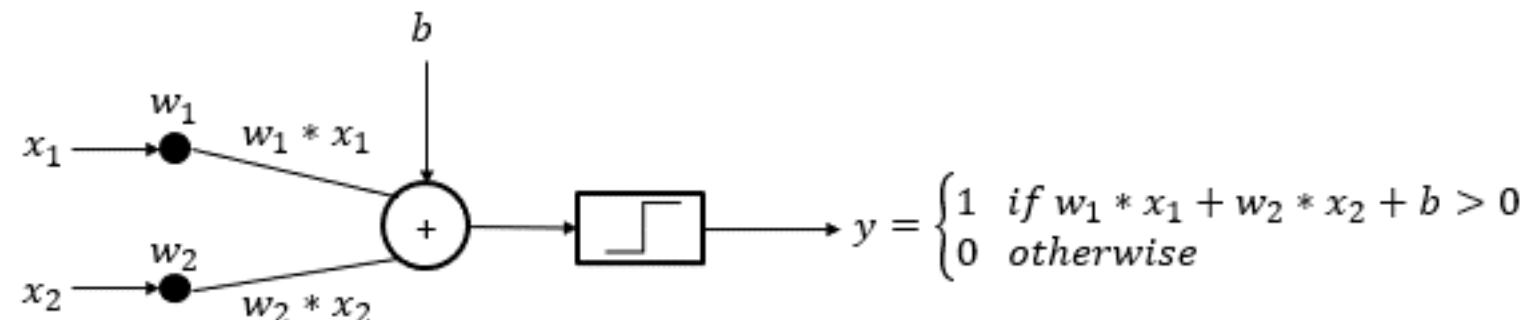
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(5) * x_1 + w_2(5) * x_2 + b(5))$$

$$y = f(0.5 * 1 + 0.5 * 0 + 0) = f(0.5) = 1$$

$$w_1(6) = w_1(5) + lr * (y_{real} - y) * x_1$$

$$w_1(6) = 0.5 + 0.5 * (-1) * 1 = 0$$

$$w_2(6) = w_2(5) + lr * (y_{real} - y) * x_2$$

$$w_2(6) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

$$b(6) = b(5) + lr * (y_{real} - y)$$

$$b(6) = 0 + 0.5 * (-1) = -0.5$$

Iteration 5:

$w_1(5)$	$w_2(5)$	$b(5)$
0.5	0.5	0

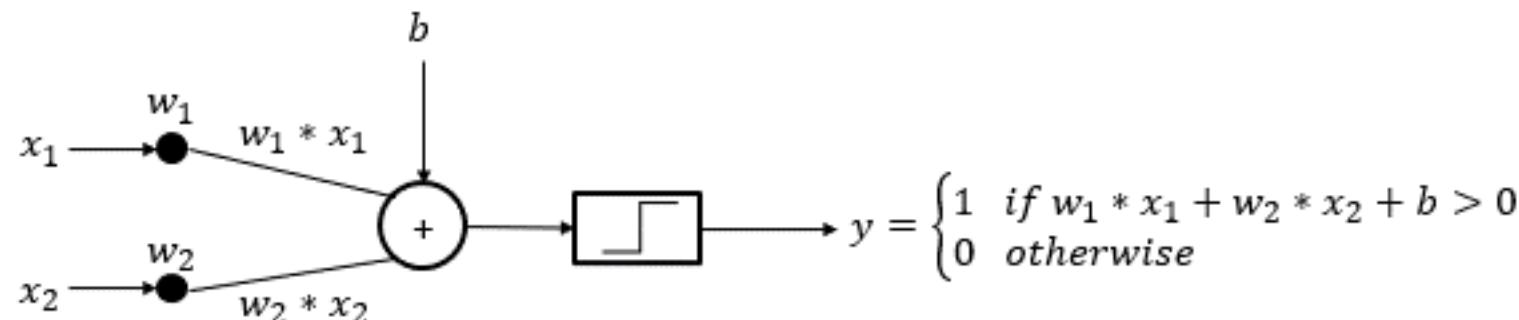
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(5) * x_1 + w_2(5) * x_2 + b(5))$$

$$y = f(0.5 * 1 + 0.5 * 0 + 0) = f(0.5) = 1$$

$$w_1(6) = w_1(5) + lr * (y_{real} - y) * x_1$$

$$w_1(6) = 0.5 + 0.5 * (-1) * 1 = 0$$

$$w_2(6) = w_2(5) + lr * (y_{real} - y) * x_2$$

$$w_2(6) = 0.5 + 0.5 * (-1) * 0 = 0.5$$

$$b(6) = b(5) + lr * (y_{real} - y)$$

$$b(6) = 0 + 0.5 * (-1) = -0.5$$

Iteration 6:

$w_1(6)$	$w_2(6)$	$b(6)$
0	0.5	-0.5

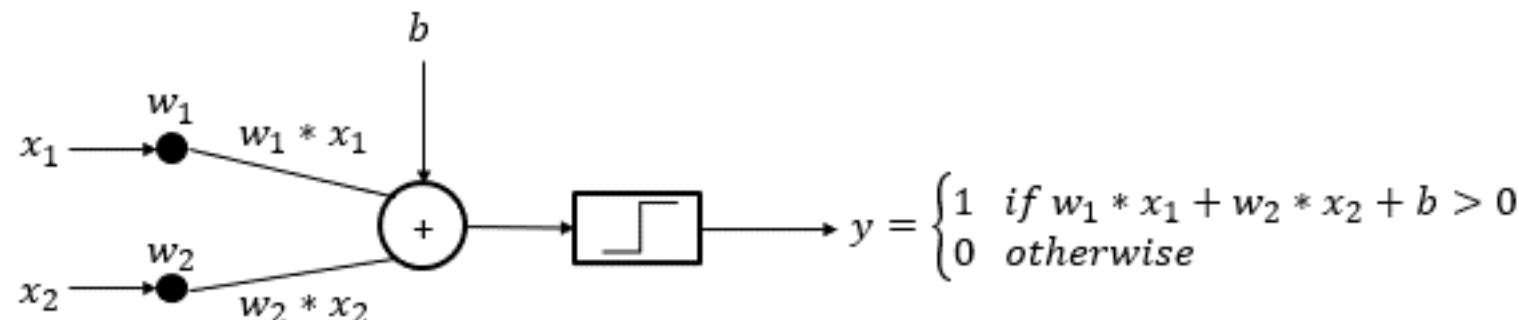
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(6) * x_1 + w_2(6) * x_2 + b(6))$$

$$y = f(0 * 0 + 0.5 * 1 - 0.5) = f(0) = 0$$

$$w_1(7) = w_1(6) + lr * (y_{real} - y) * x_1$$

$$w_1(7) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(7) = w_2(6) + lr * (y_{real} - y) * x_2$$

$$w_2(7) = 0.5 + 0.5 * 0 * 1 = 0.5$$

$$b(7) = b(6) + lr * (y_{real} - y)$$

$$b(7) = -0.5 + 0.5 * 0 = -0.5$$

Iteration 6:

$w_1(6)$	$w_2(6)$	$b(6)$
0	0.5	-0.5

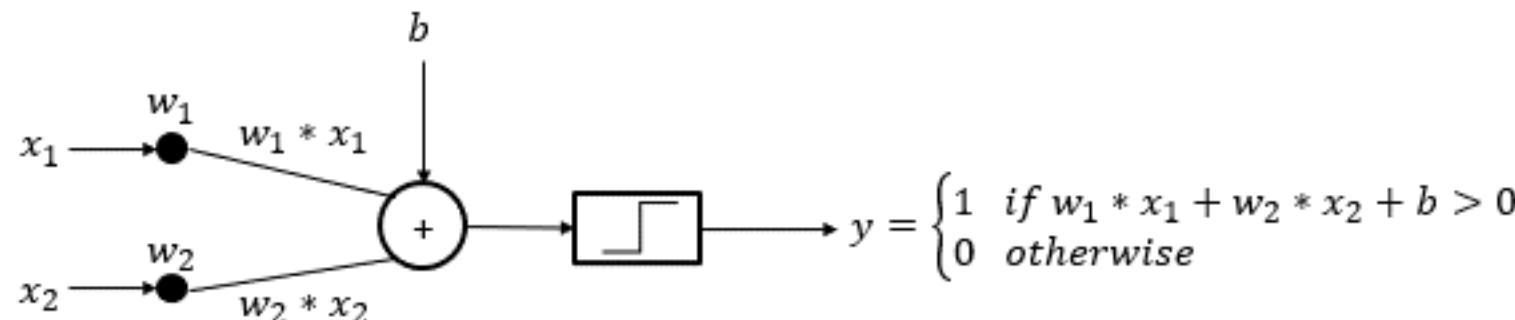
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight initialization
- Until convergence:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is learning rate and t is the iteration



AND gate:

x_1	x_2	y_{real}
0	0	0
1	0	0
0	1	0
1	1	1

$$y = f(w_1(6) * x_1 + w_2(6) * x_2 + b(6))$$

$$y = f(0 * 0 + 0.5 * 1 - 0.5) = f(0) = 0$$

$$w_1(7) = w_1(6) + lr * (y_{real} - y) * x_1$$

$$w_1(7) = 0 + 0.5 * 0 * 0 = 0$$

$$w_2(7) = w_2(6) + lr * (y_{real} - y) * x_2$$

$$w_2(7) = 0.5 + 0.5 * 0 * 1 = 0.5$$

$$b(7) = b(6) + lr * (y_{real} - y)$$

$$b(7) = -0.5 + 0.5 * 0 = -0.5$$

Iteration 7:

$w_1(7)$	$w_2(7)$	$b(7)$
0	0.5	-0.5

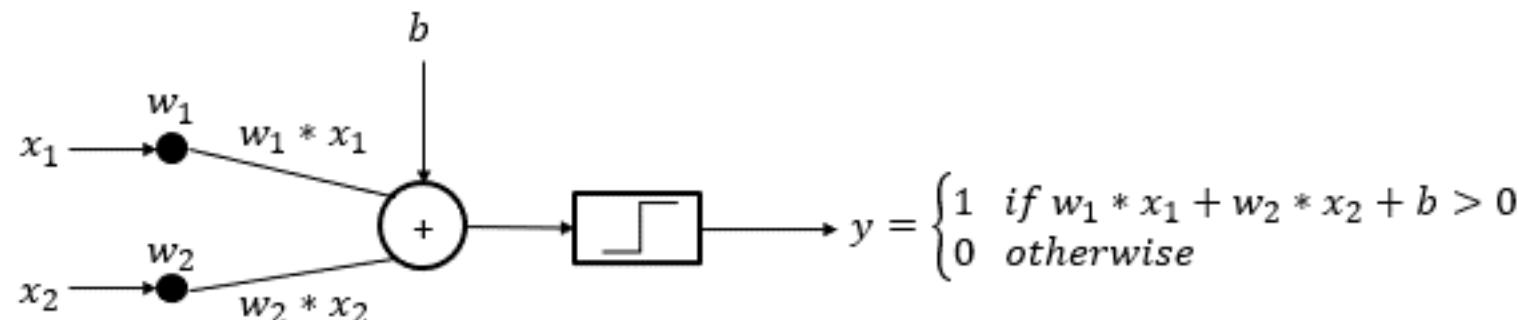
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is **learning rate** and t is the **iteration**



XOR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	0

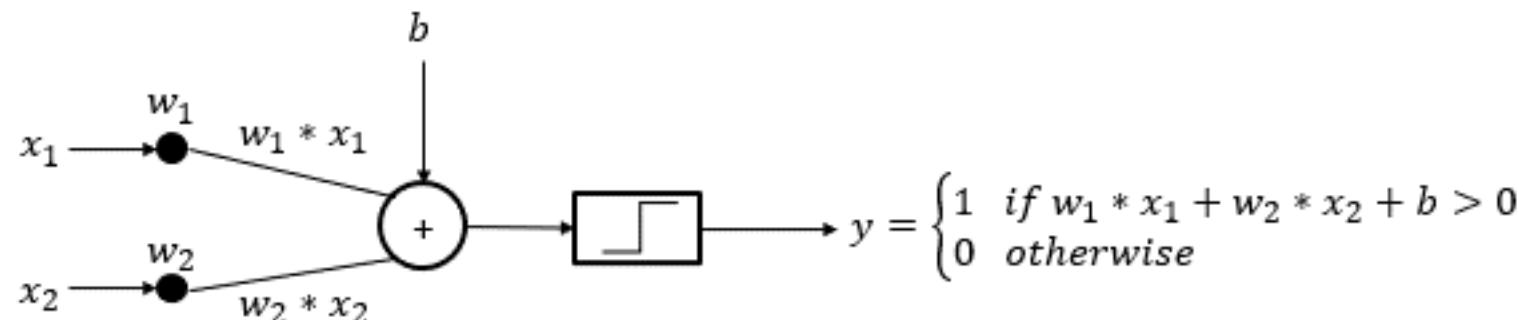
PERCEPTRON (ROSENBLATT, 1958)

Training algorithm:

- Weight **initialization**
- Until **convergence**:
 1. Compute y
 2. Update weights such that:

$$w_i(t+1) = w_i(t) + lr * (y_{real} - y) * x_i,$$

where lr is **learning rate** and t is the **iteration**

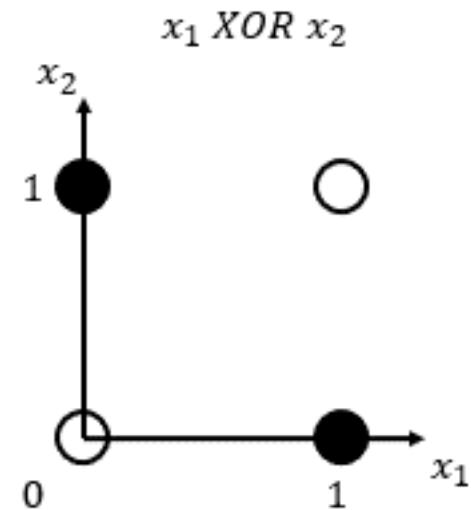
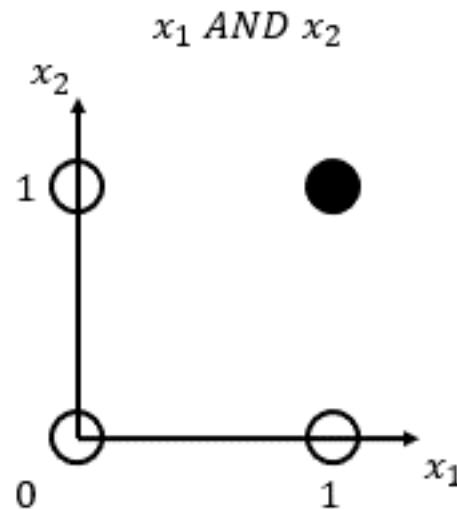
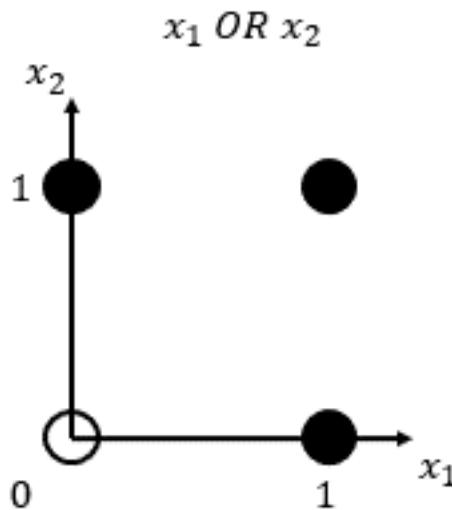


XOR gate:

x_1	x_2	y_{real}
0	0	0
1	0	1
0	1	1
1	1	0

→ Training does not converge

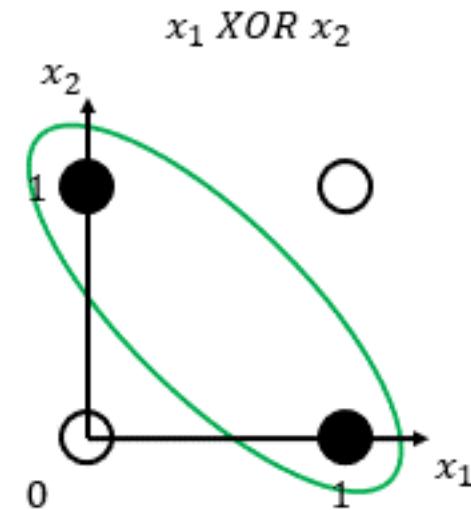
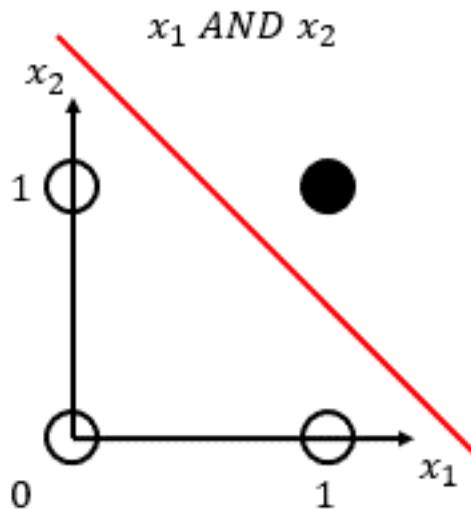
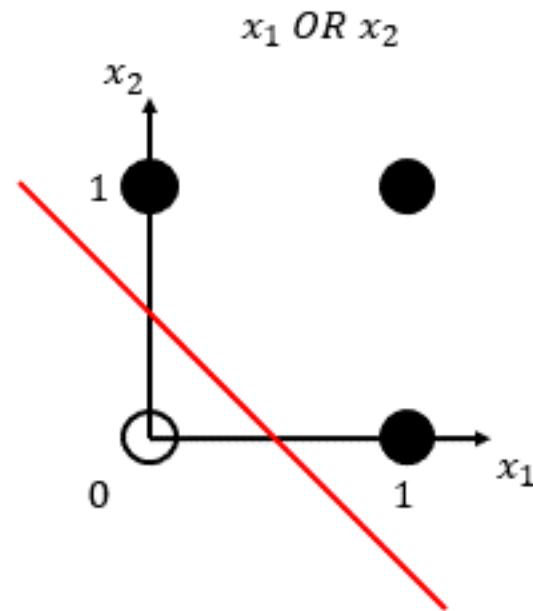
PERCEPTRON (ROSENBLATT, 1958)



○ $y = 0$
● $y = 1$

PERCEPTRON (ROSENBLATT, 1958)

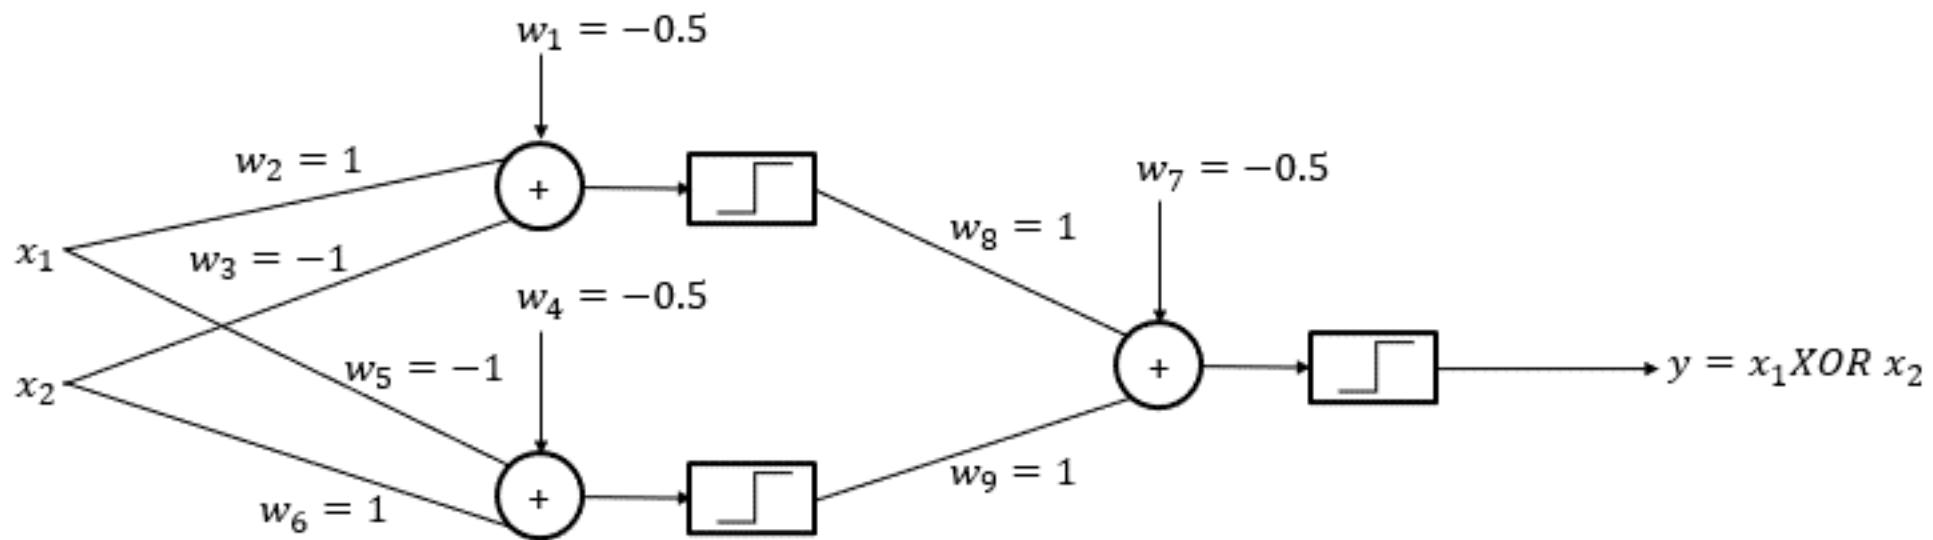
Linearly separable



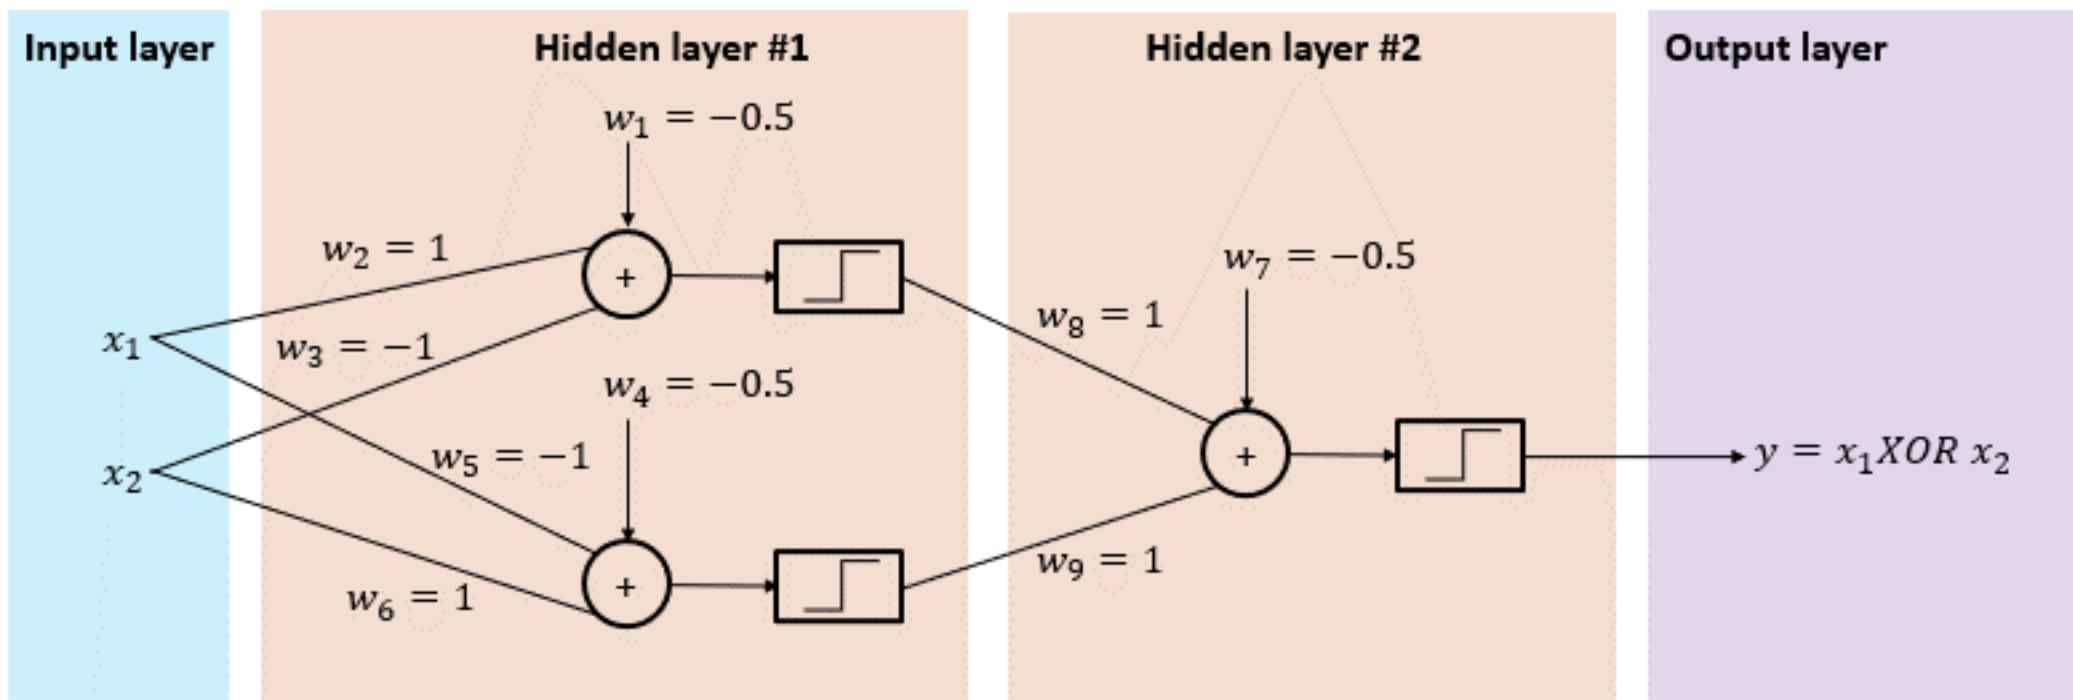
- $y = 0$
- $y = 1$

Non linearly separable

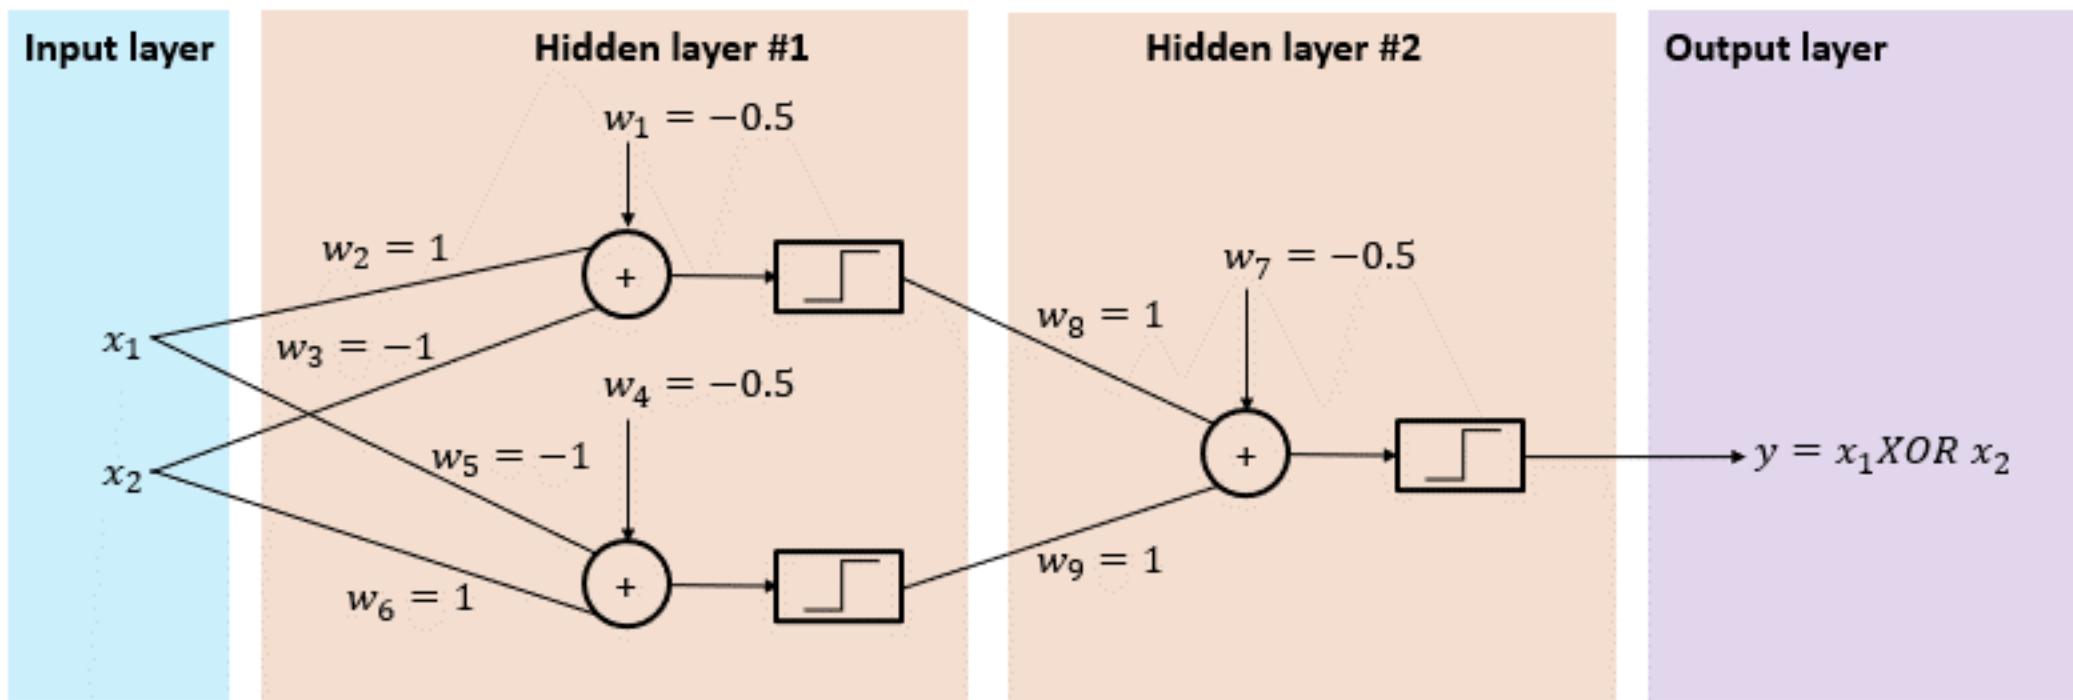
PERCEPTRON (ROSENBLATT, 1958)



PERCEPTRON (ROSENBLATT, 1958)

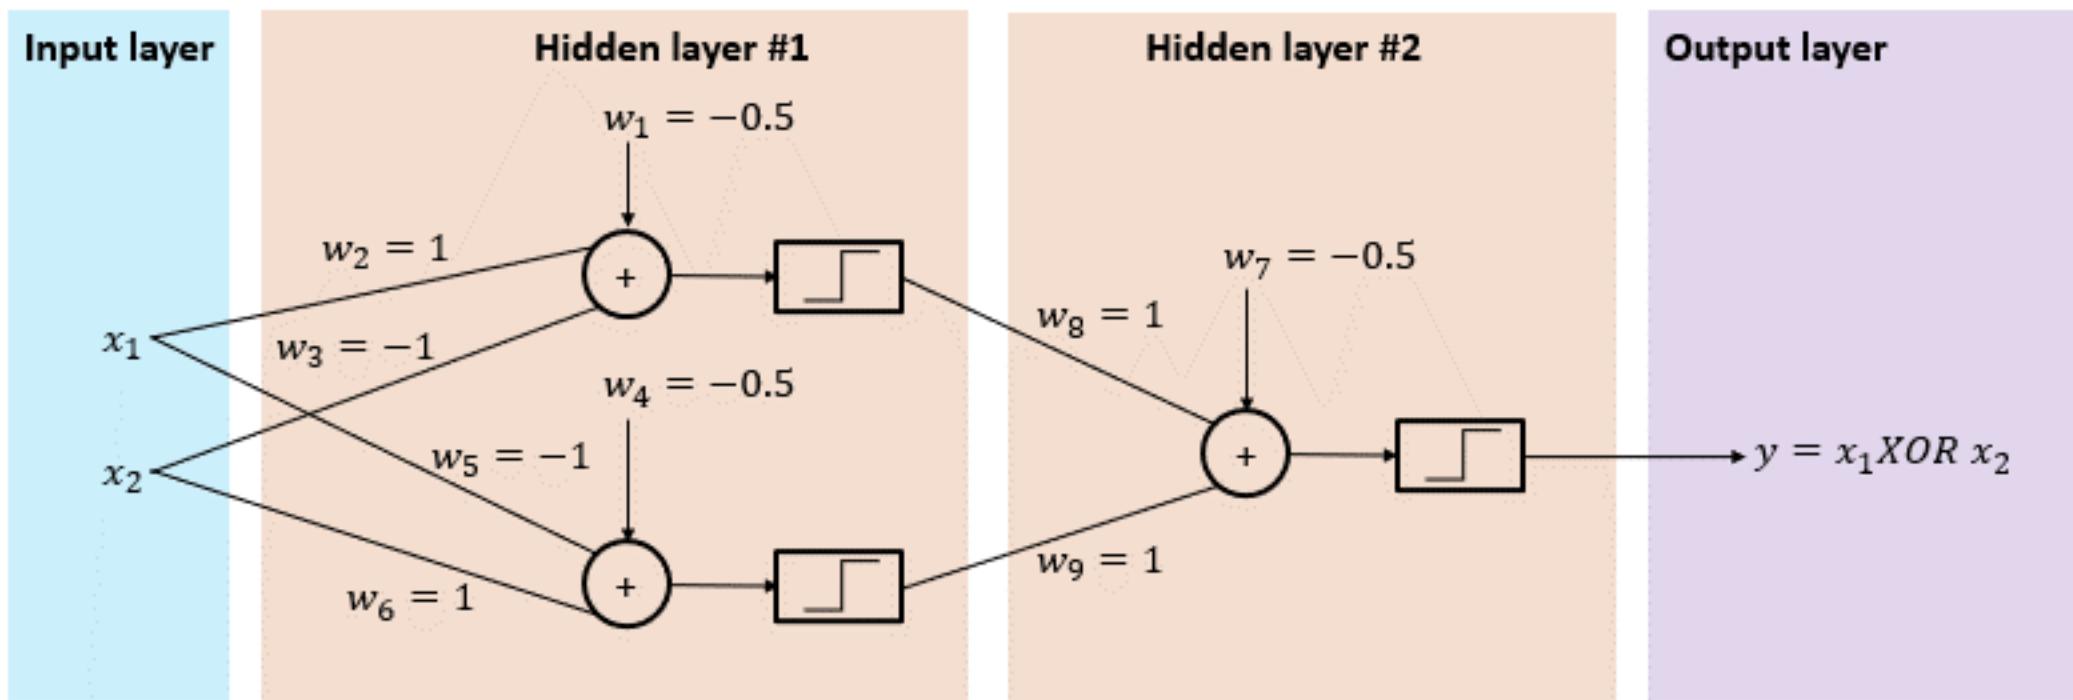


PERCEPTRON (ROSENBLATT, 1958)



2-hidden layer perceptron can **solve non linear** problems

PERCEPTRON (ROSENBLATT, 1958)

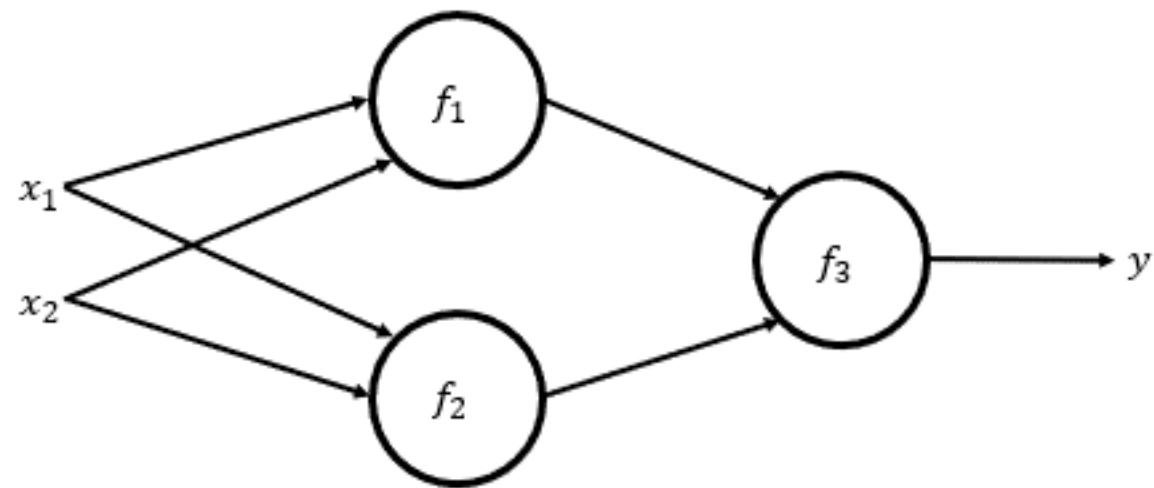


2-hidden layer perceptron can **solve non linear** problems

How to train such network?

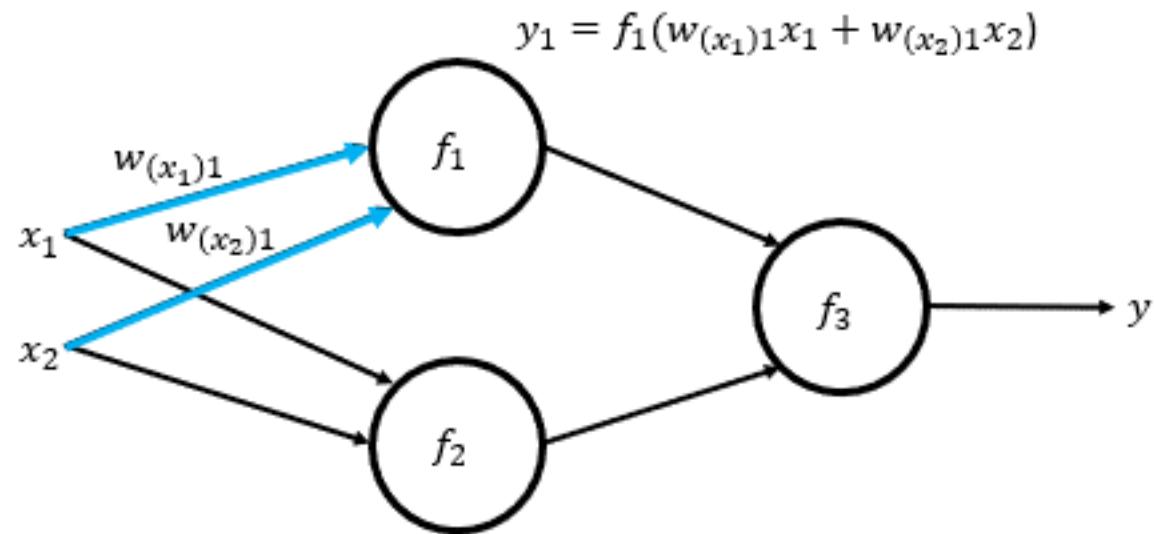
BACKPROPAGATION

Forward propagation:



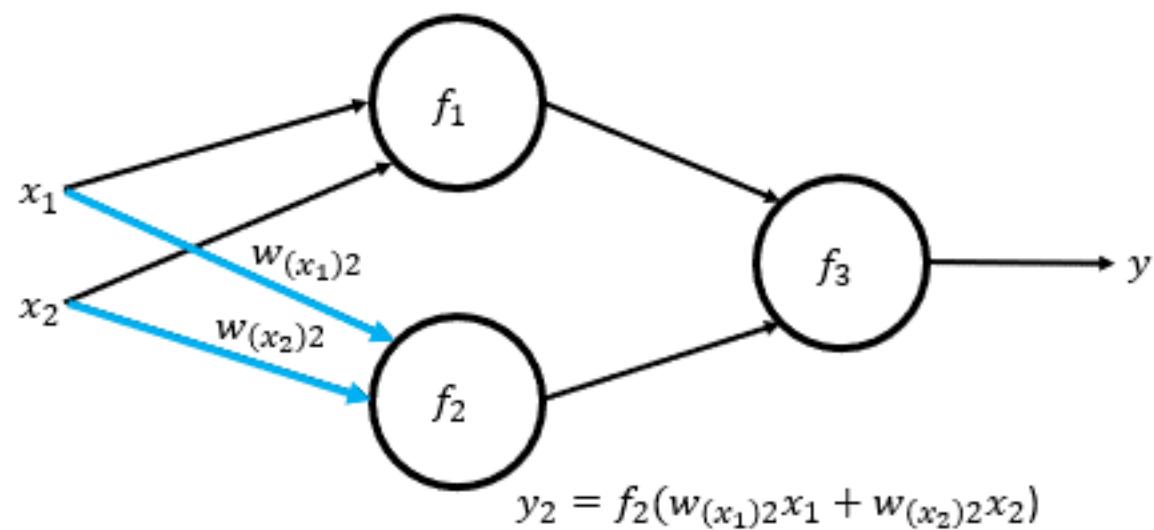
BACKPROPAGATION

Forward propagation:



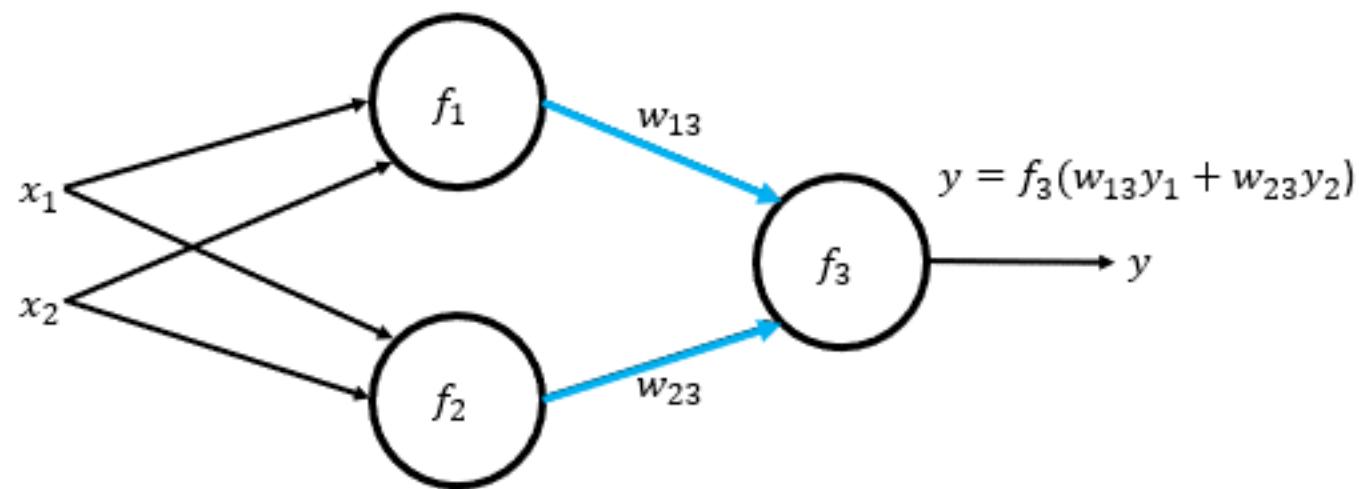
BACKPROPAGATION

Forward propagation:



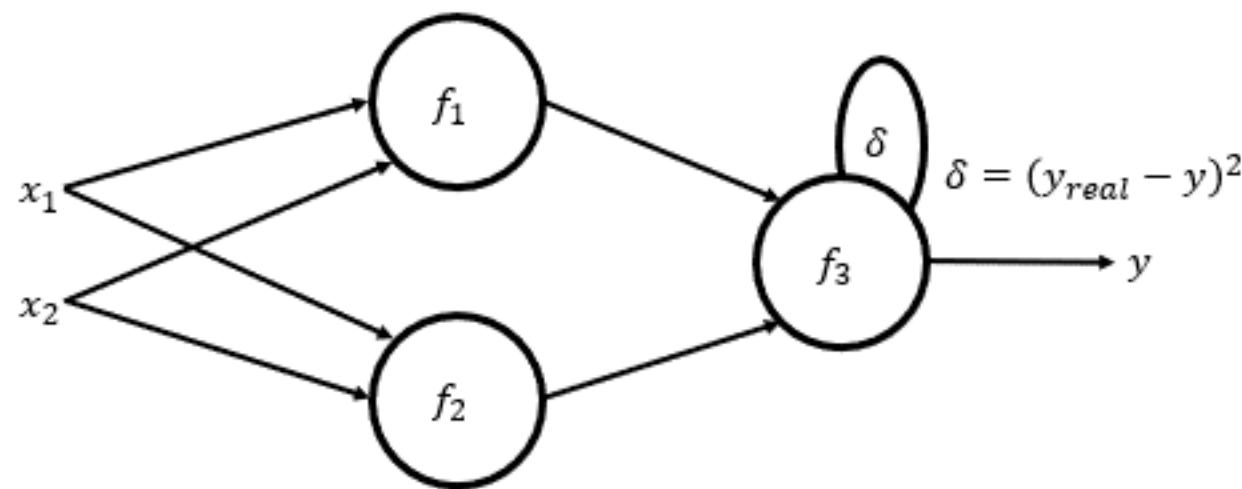
BACKPROPAGATION

Forward propagation:



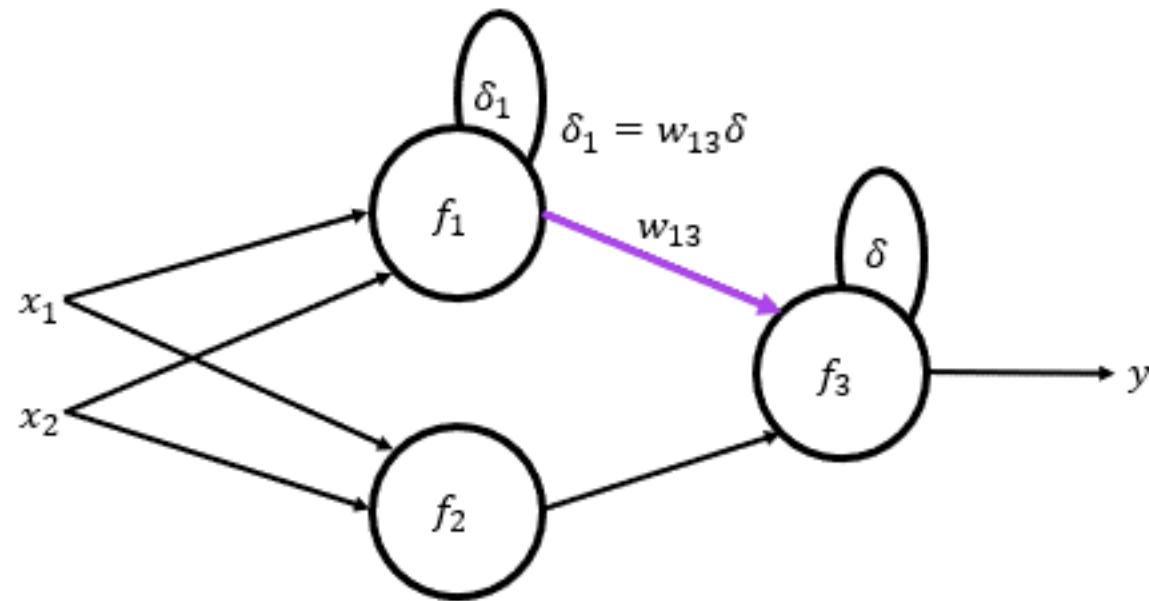
BACKPROPAGATION

Forward propagation:



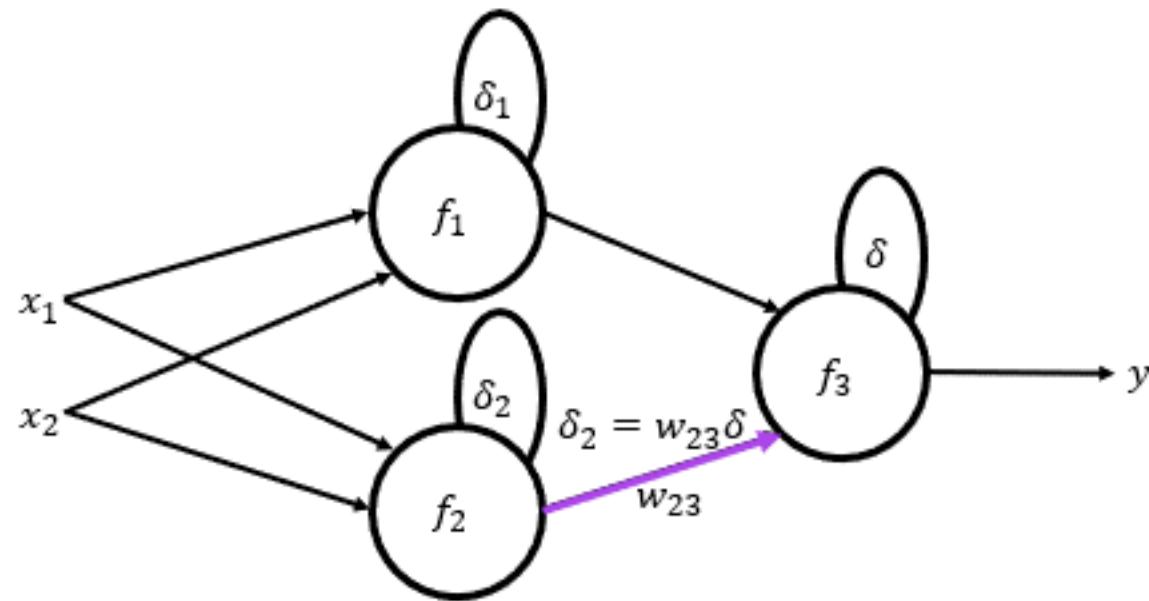
BACKPROPAGATION

Backpropagation:



BACKPROPAGATION

Backpropagation:

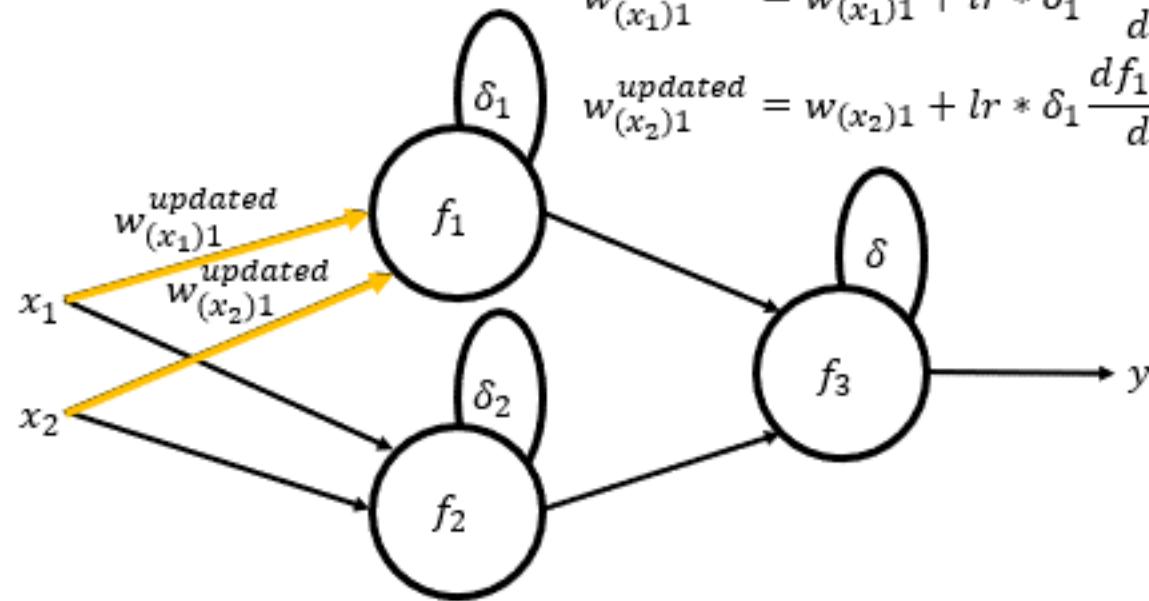


BACKPROPAGATION

Backpropagation:

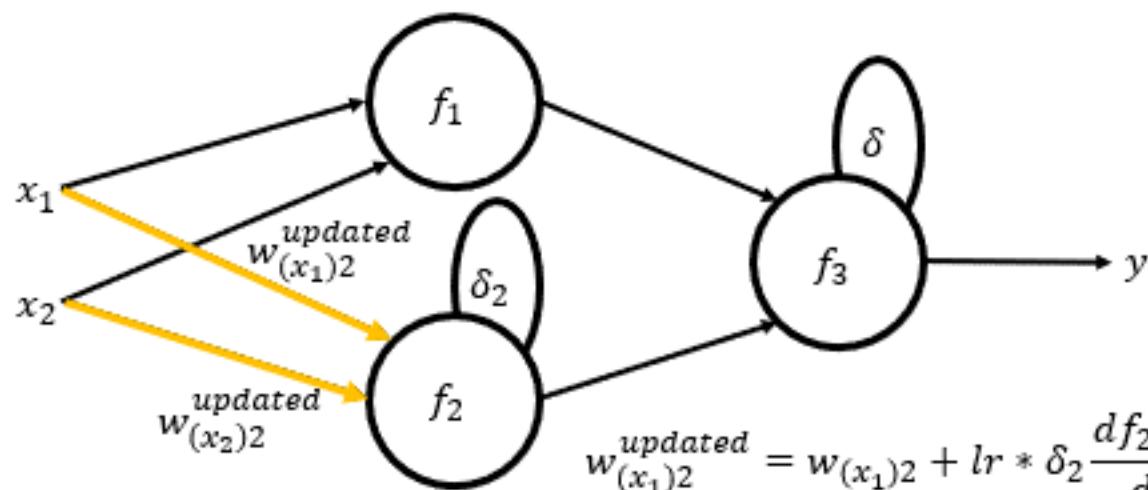
$$w_{(x_1)1}^{updated} = w_{(x_1)1} + lr * \delta_1 \frac{df_1(u)}{du} * x_1$$

$$w_{(x_2)1}^{updated} = w_{(x_2)1} + lr * \delta_1 \frac{df_1(u)}{du} * x_2$$



BACKPROPAGATION

Backpropagation:

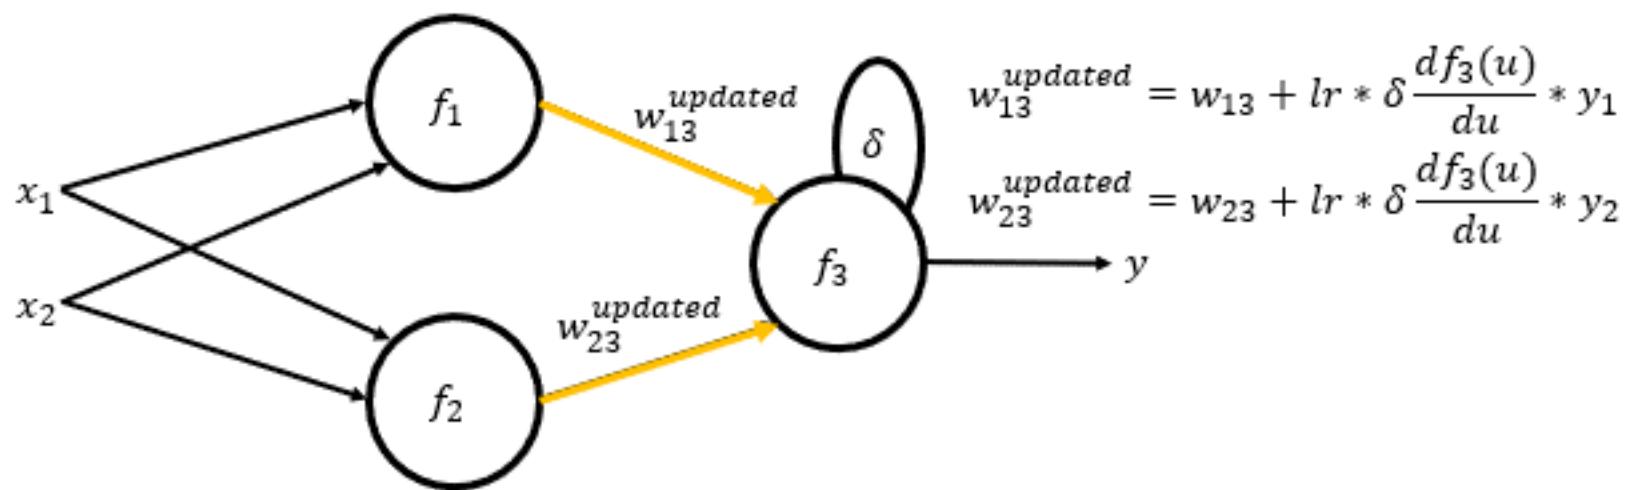


$$w_{(x_1)2}^{updated} = w_{(x_1)2} + lr * \delta_2 \frac{df_2(u)}{du} * x_1$$

$$w_{(x_2)2}^{updated} = w_{(x_2)2} + lr * \delta_2 \frac{df_2(u)}{du} * x_2$$

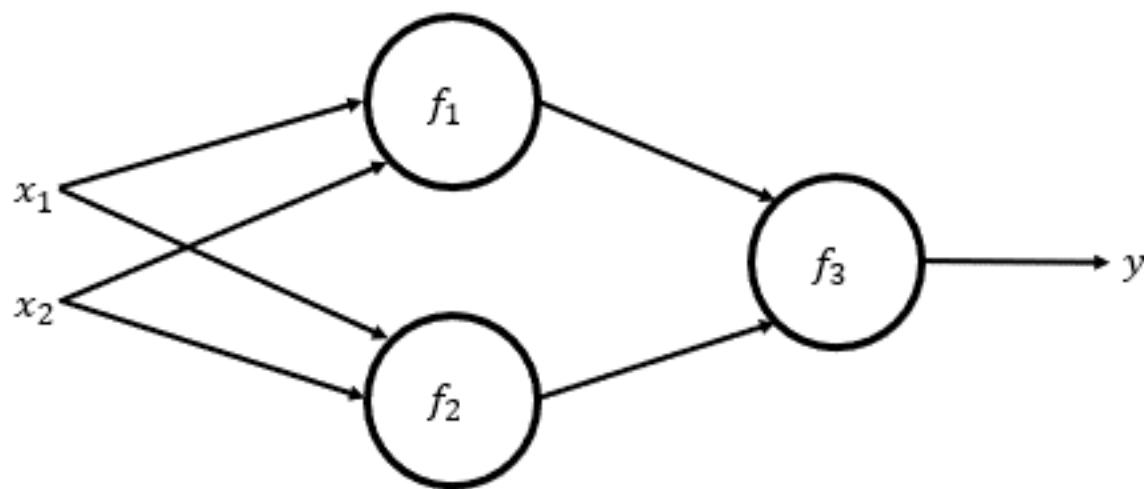
BACKPROPAGATION

Backpropagation:



BACKPROPAGATION

Backpropagation:



Critical point: $\frac{df(u)}{du}$ is known while the Heaviside function is **not differentiable**

Examples of **activation functions**: sigmoid, hyperbolic tangent, rectified linear unit (ReLU), softmax

OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{real}, y)$ (Mean Squared, error, Cross-Entropy, ...)

OPTIMIZATION

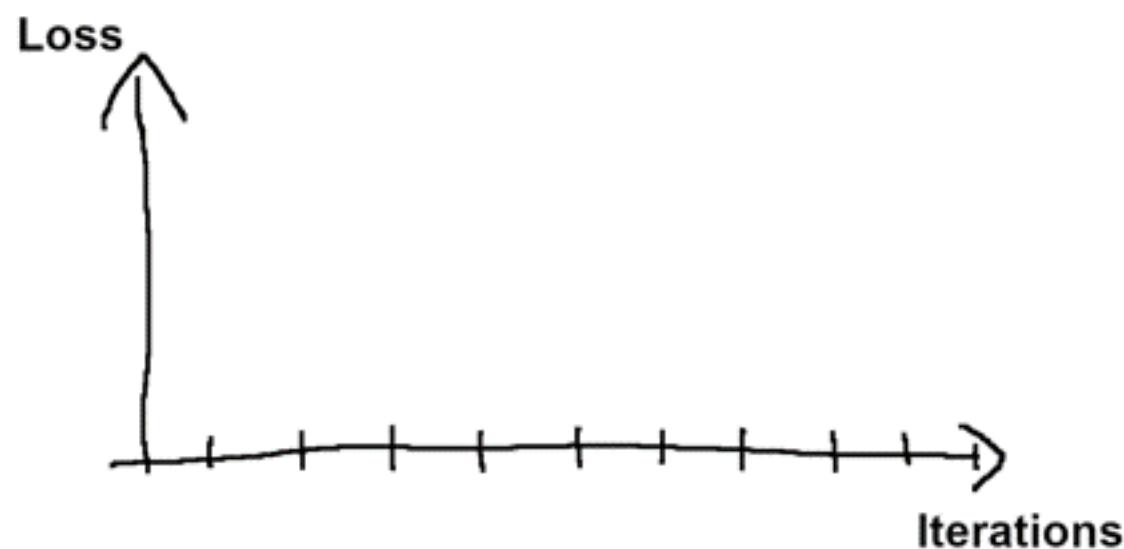
- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{real}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)

OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)

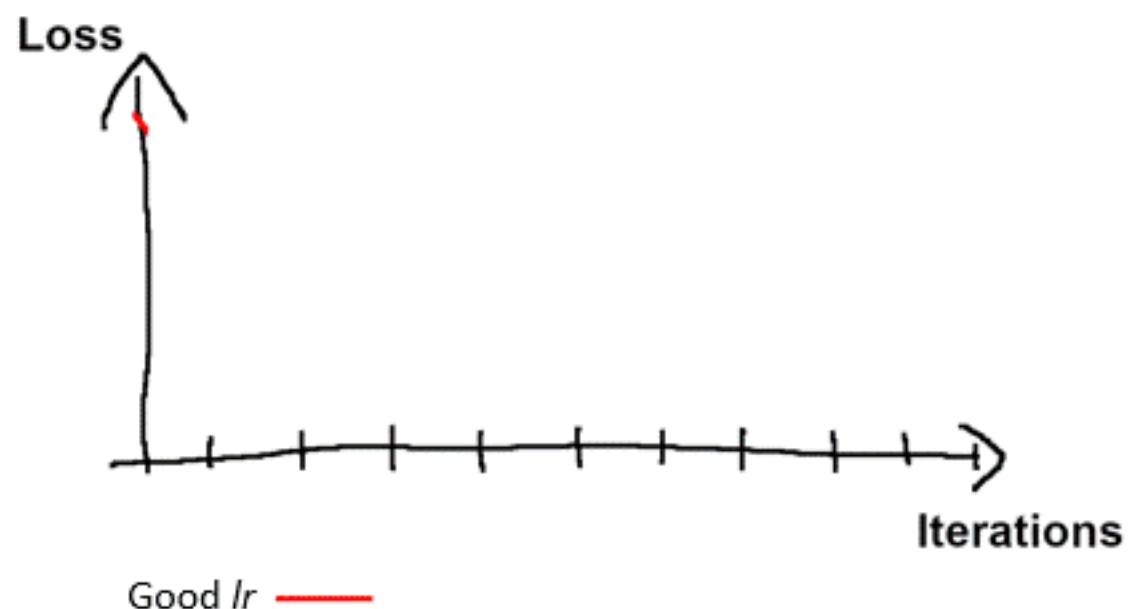
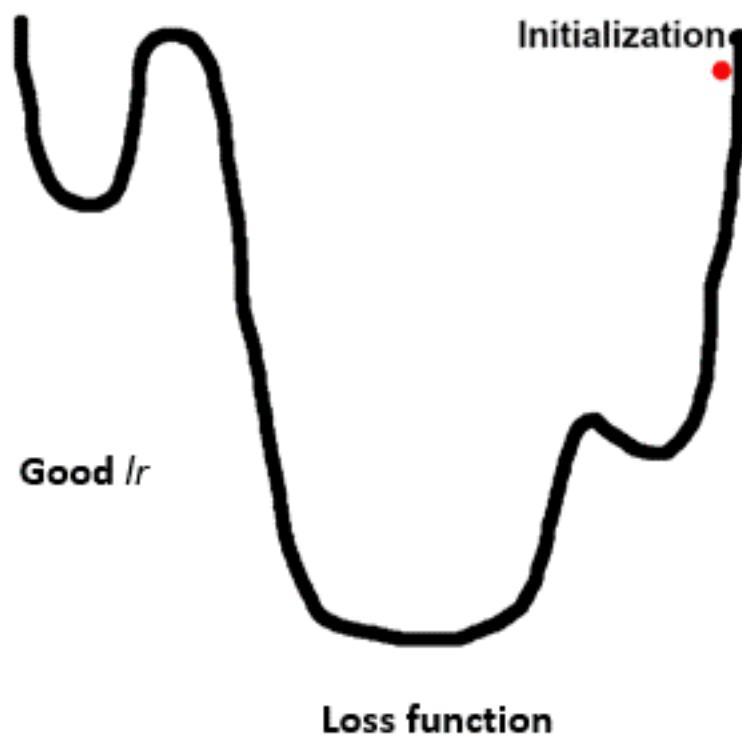


Loss function



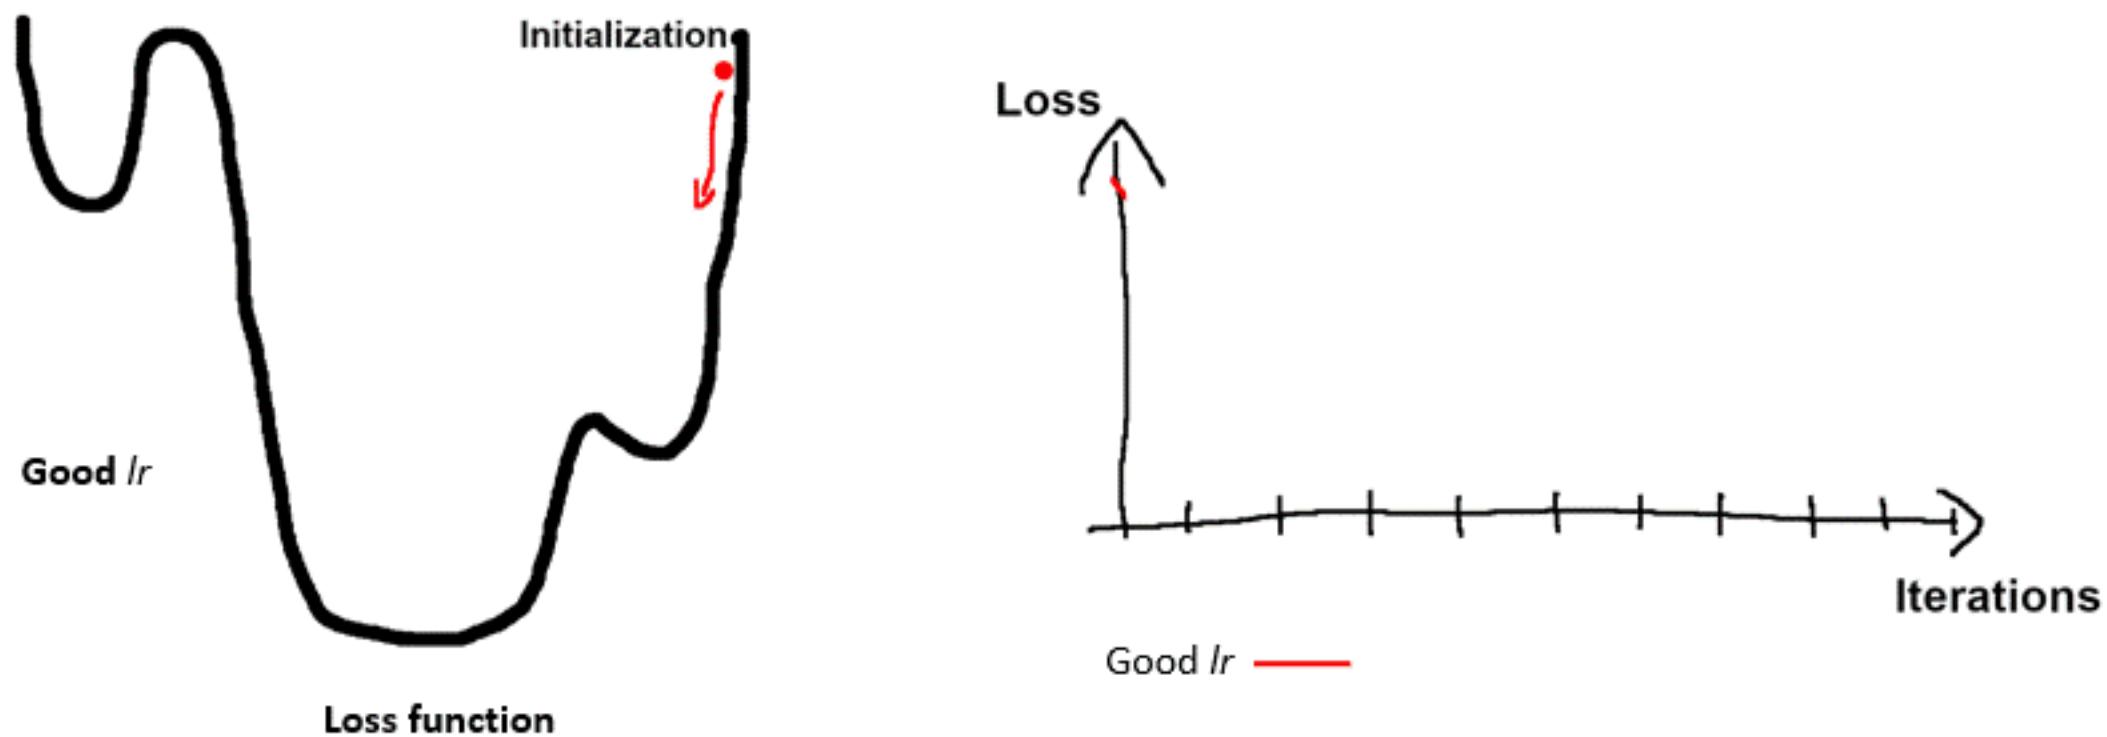
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



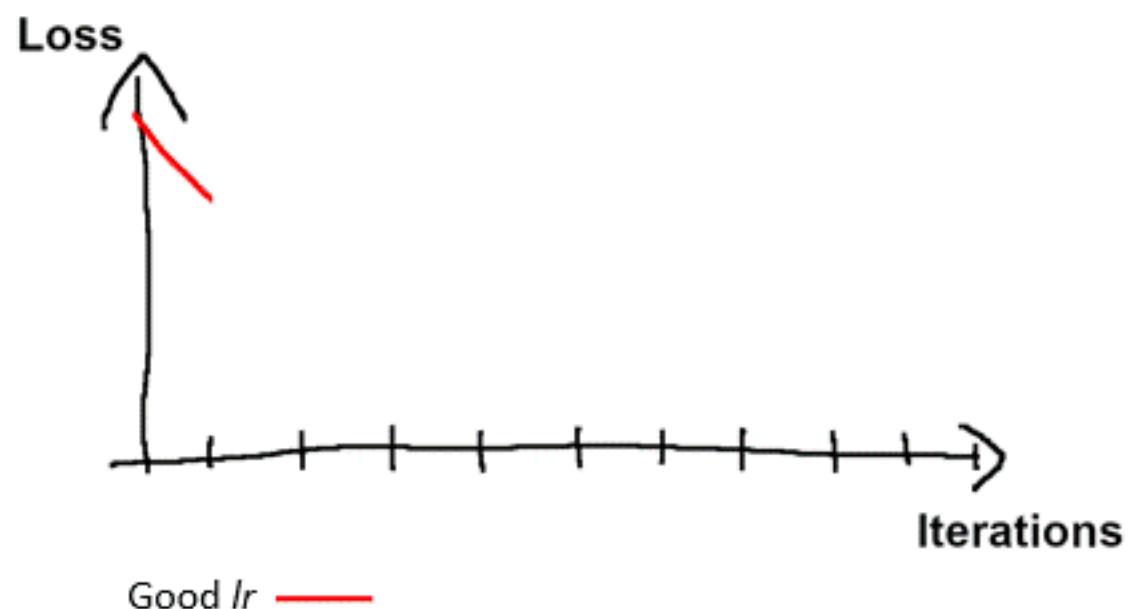
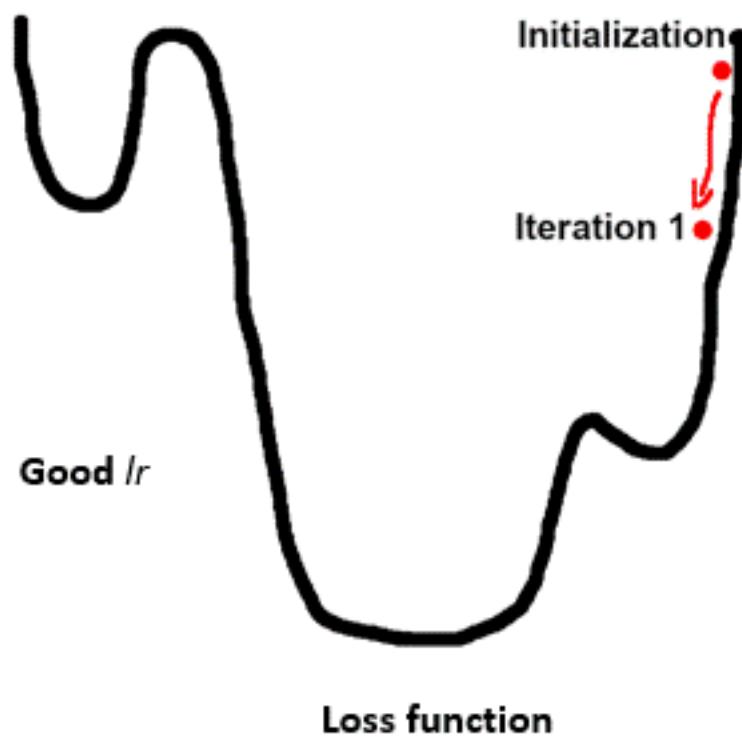
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



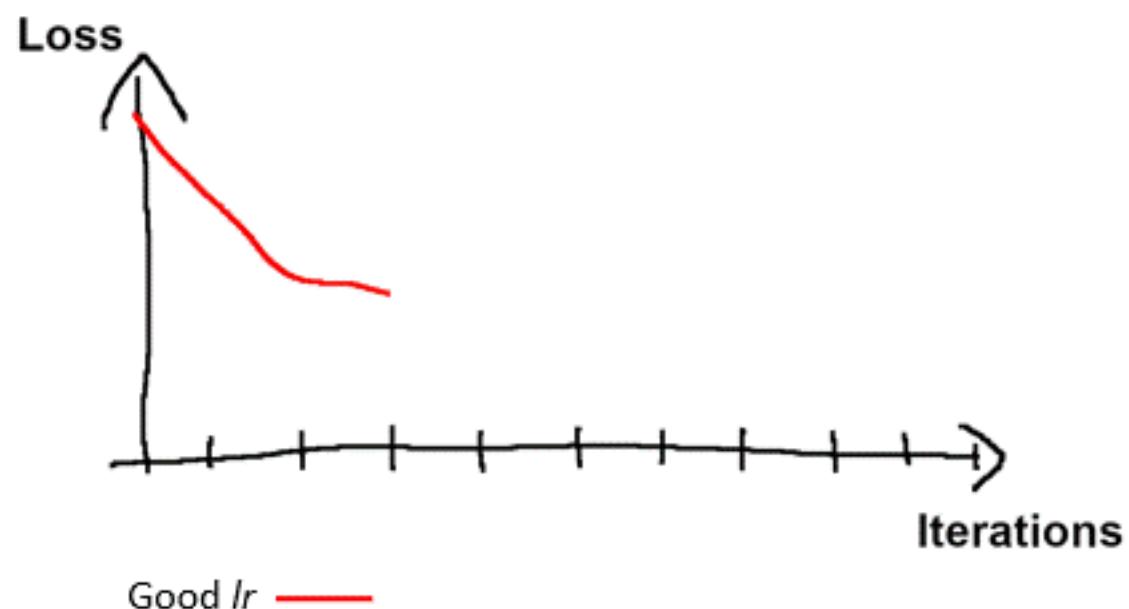
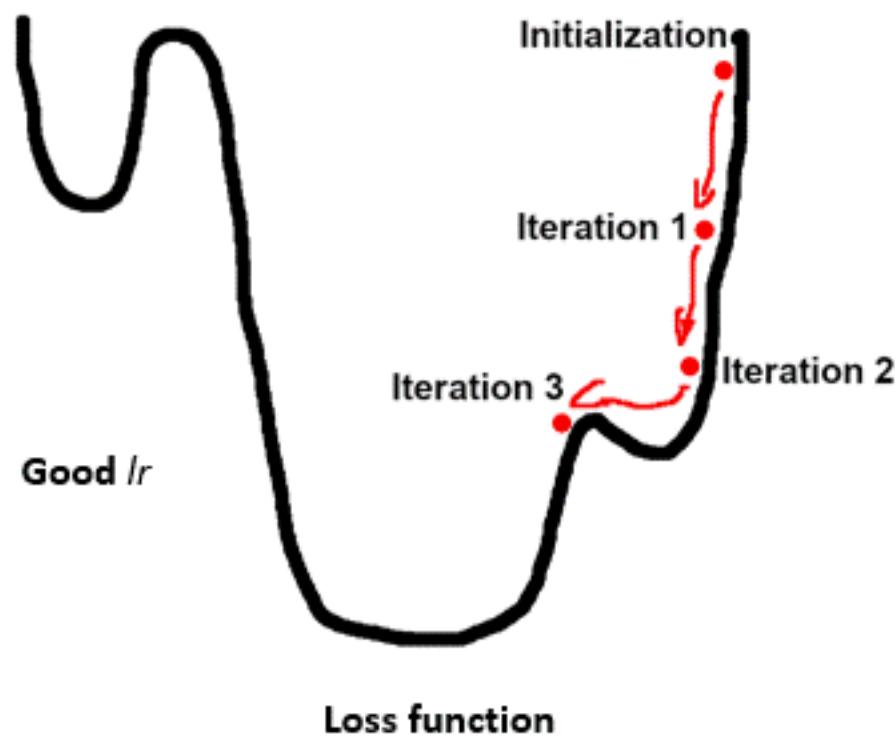
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



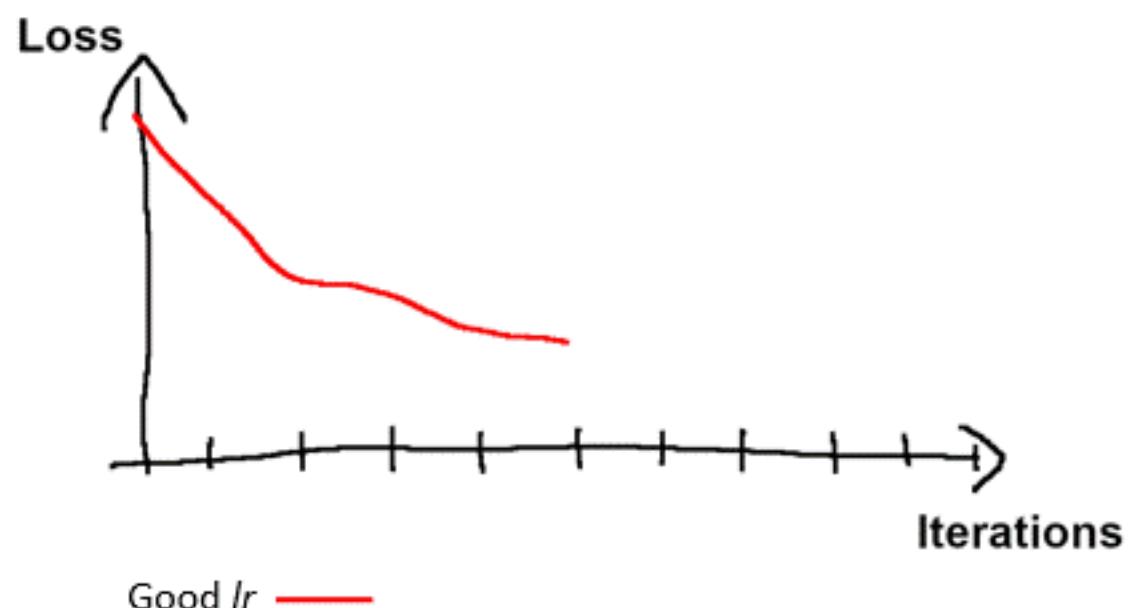
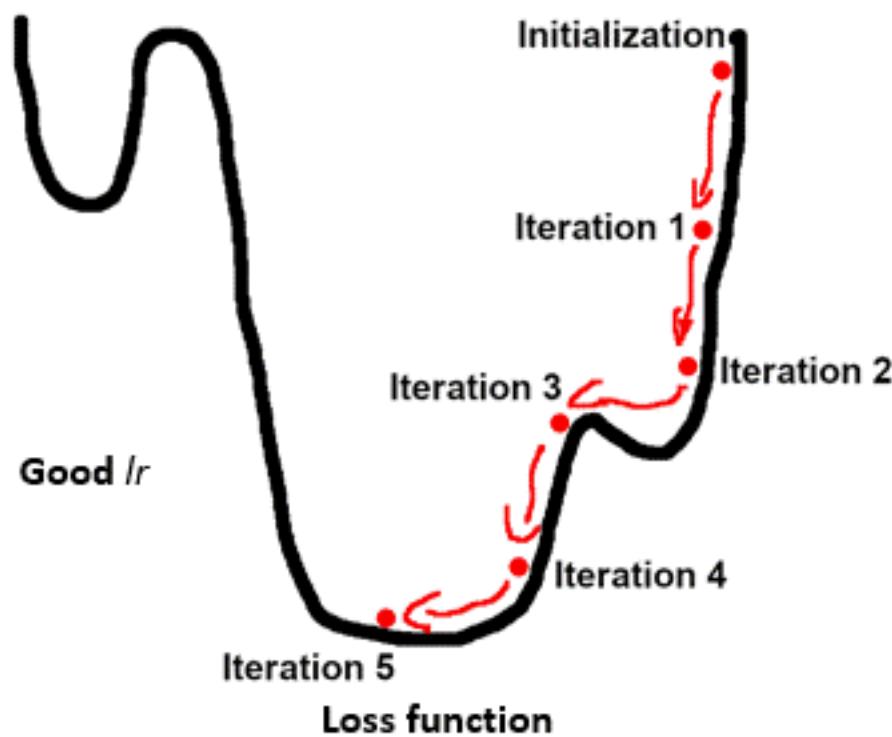
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



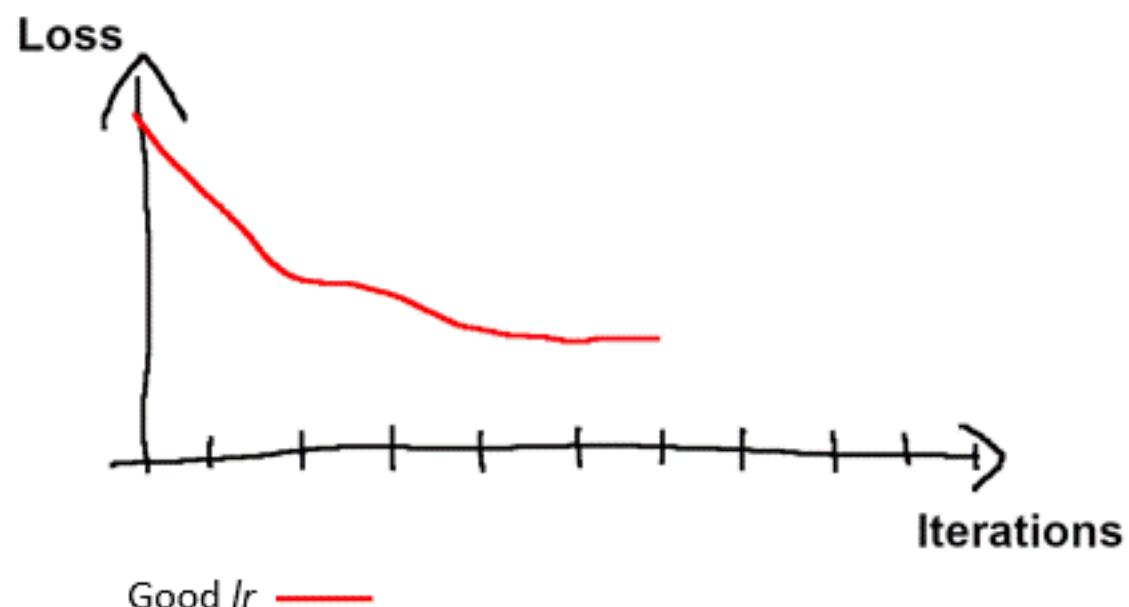
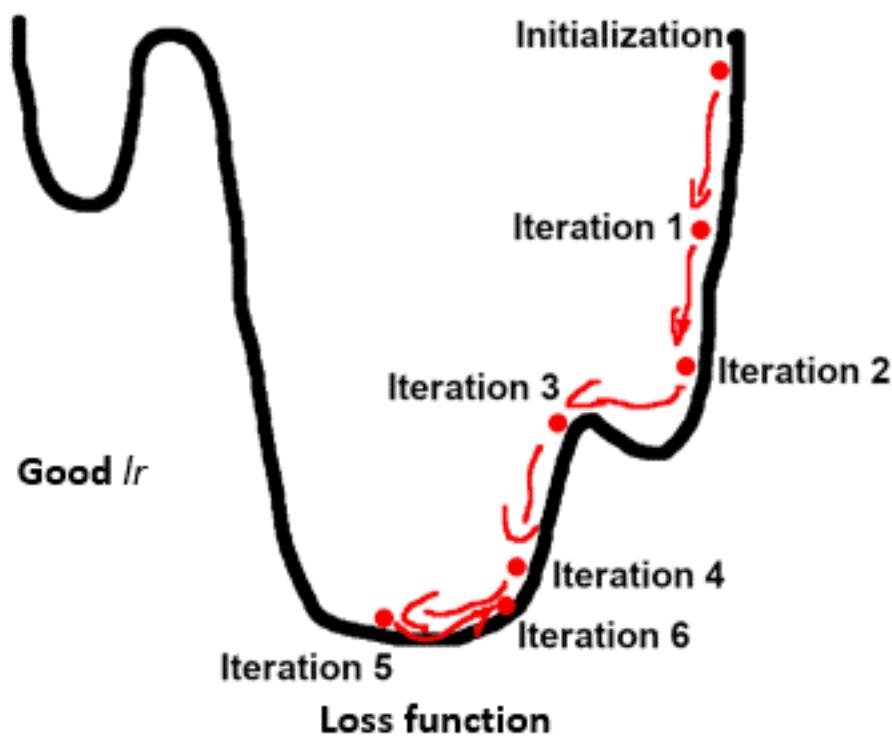
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



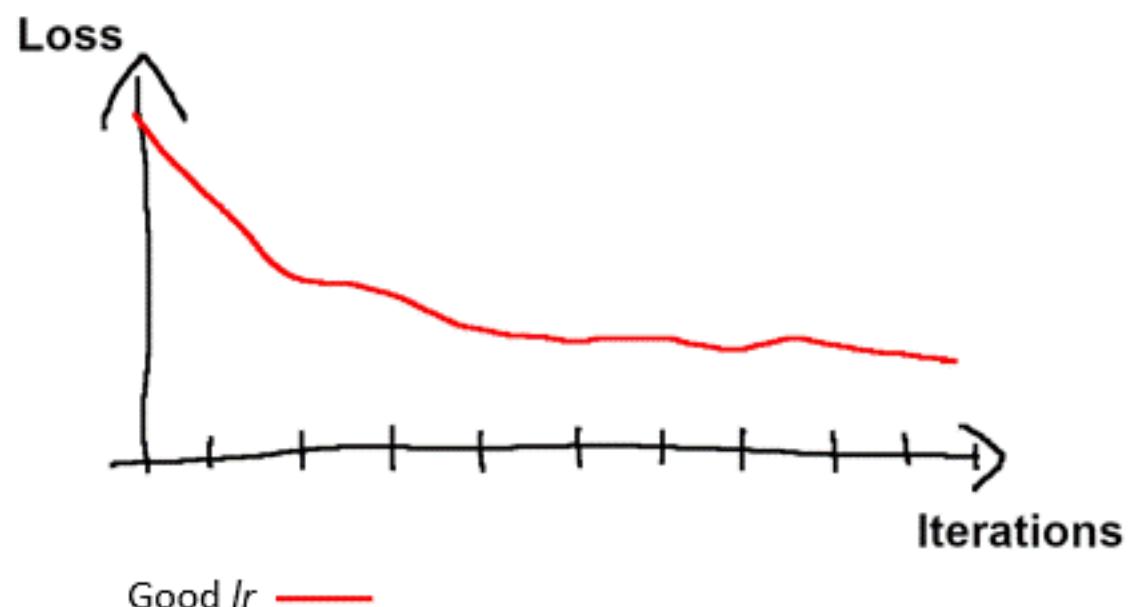
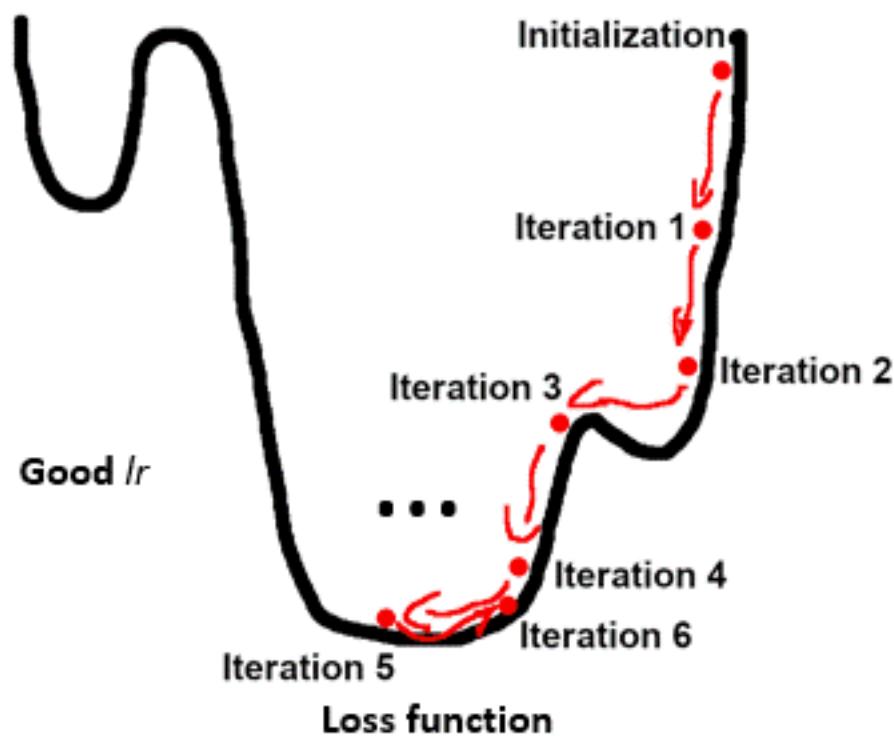
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



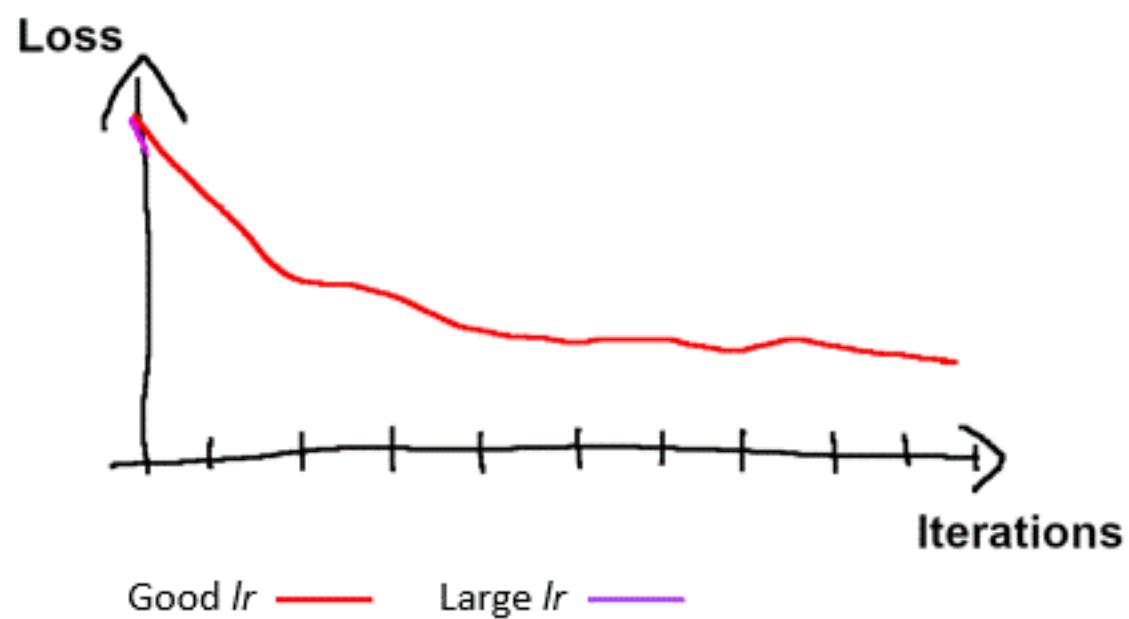
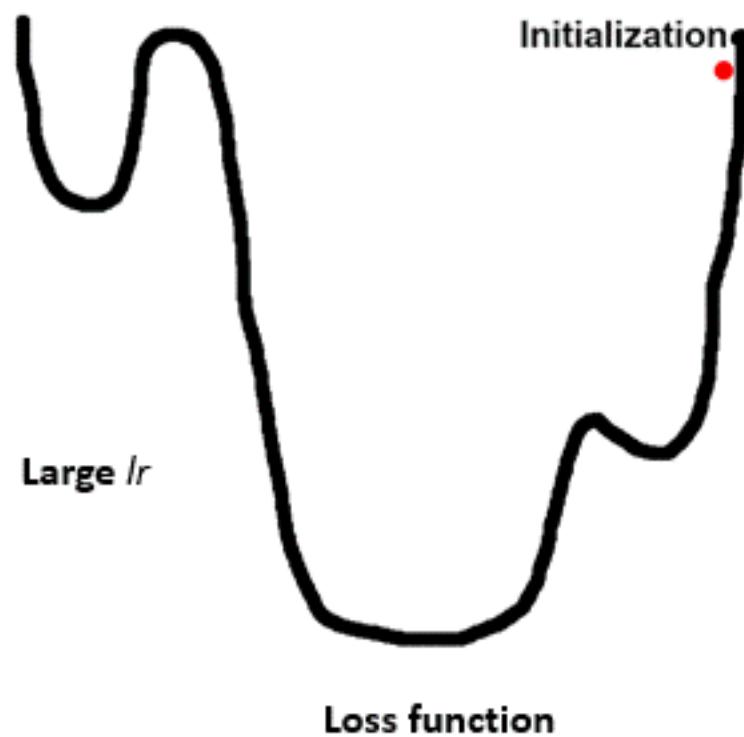
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



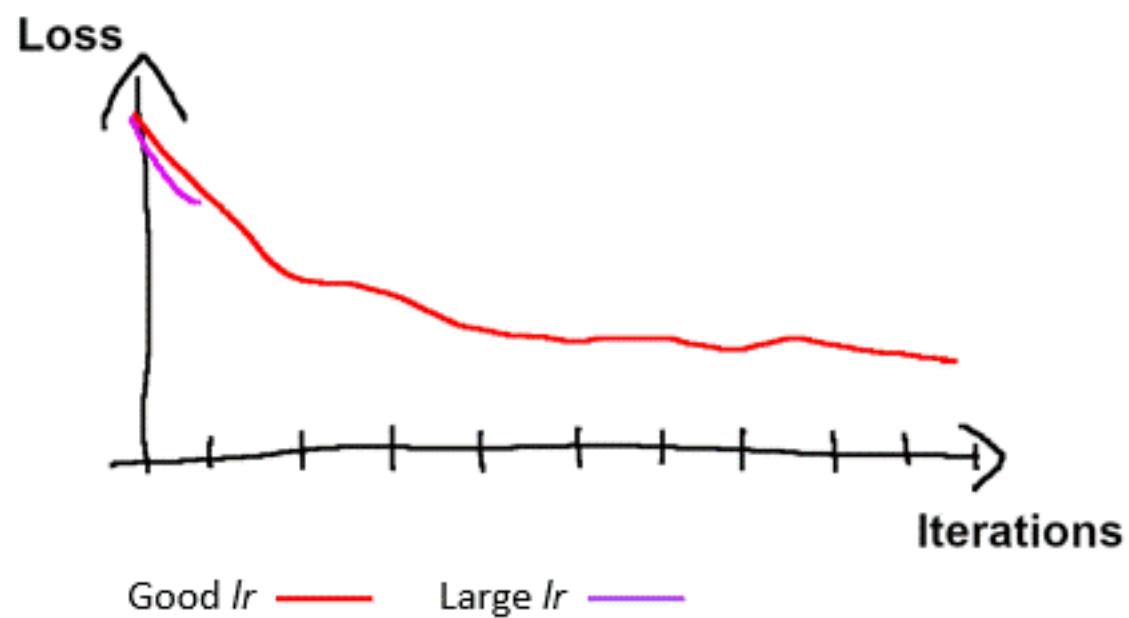
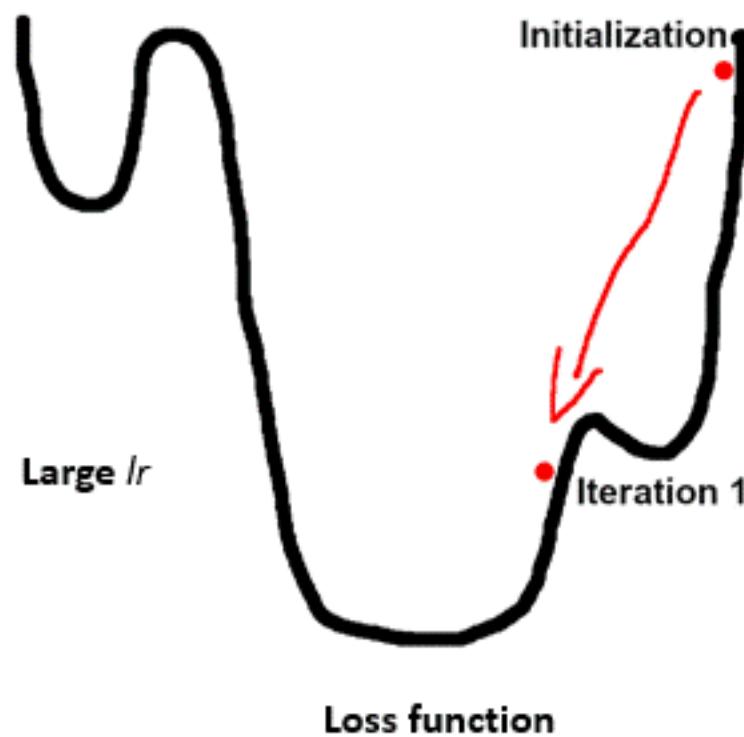
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



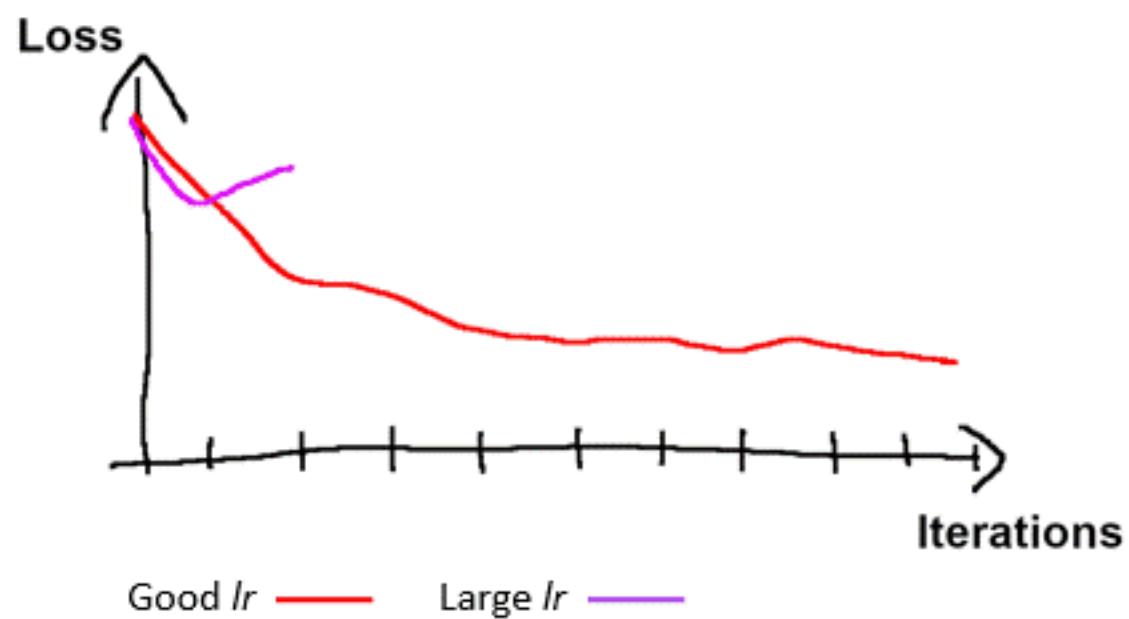
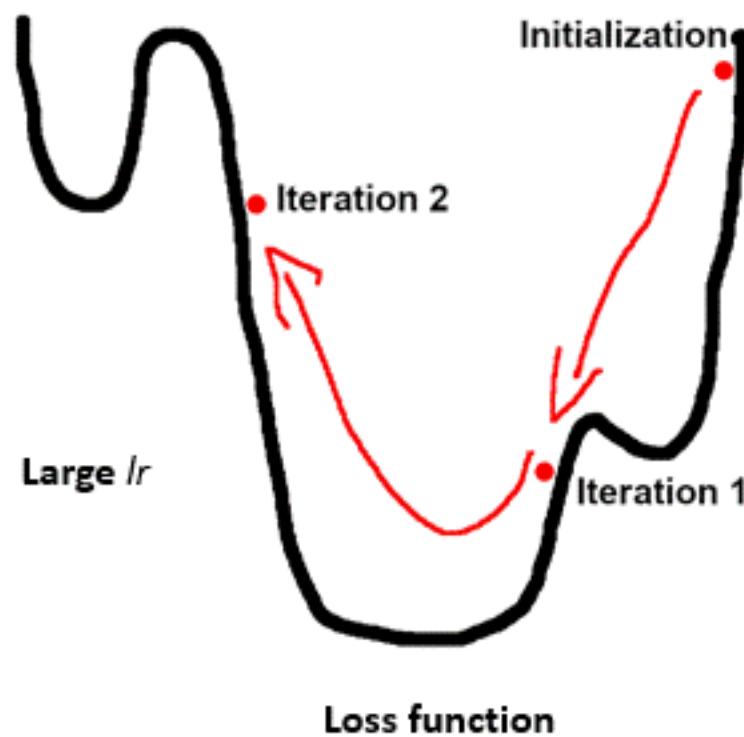
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



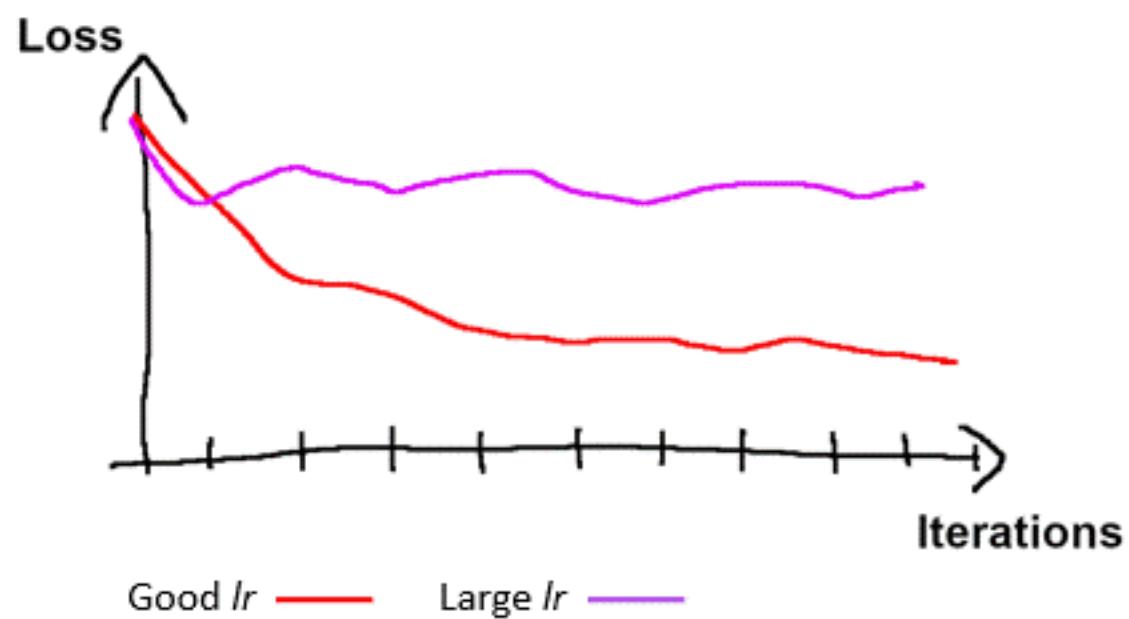
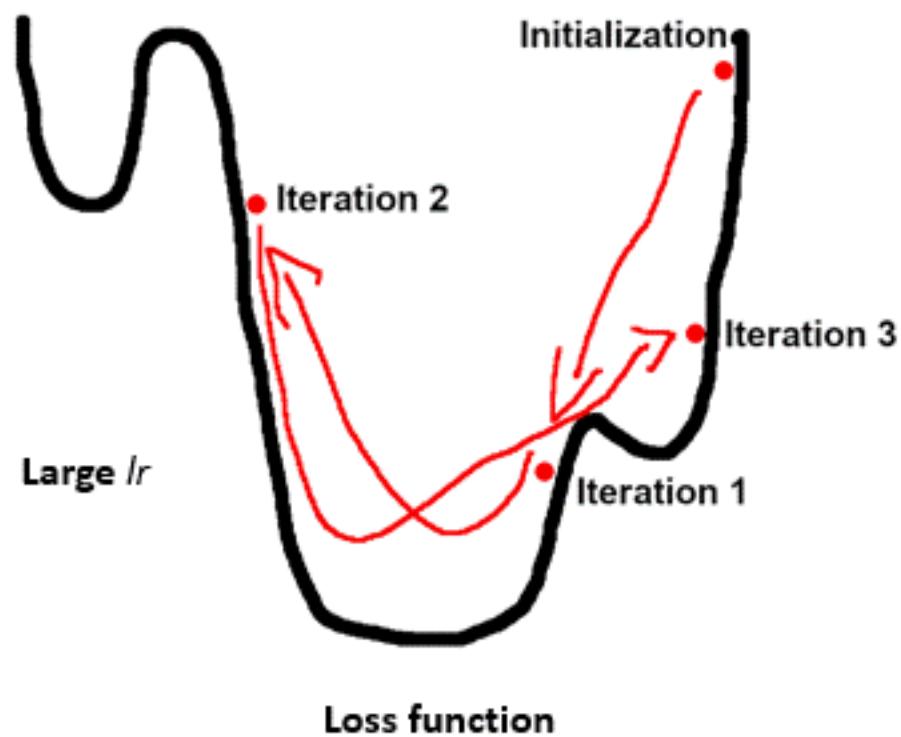
OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)



OPTIMIZATION

- The goal is to **minimize the loss (cost) function** $\text{Loss}(y_{\text{real}}, y)$ (Mean Squared error, Cross-Entropy, ...)
- This is achieved via a **gradient descent** (stochastic gradient descent, RMSprop, ADAM, ...)

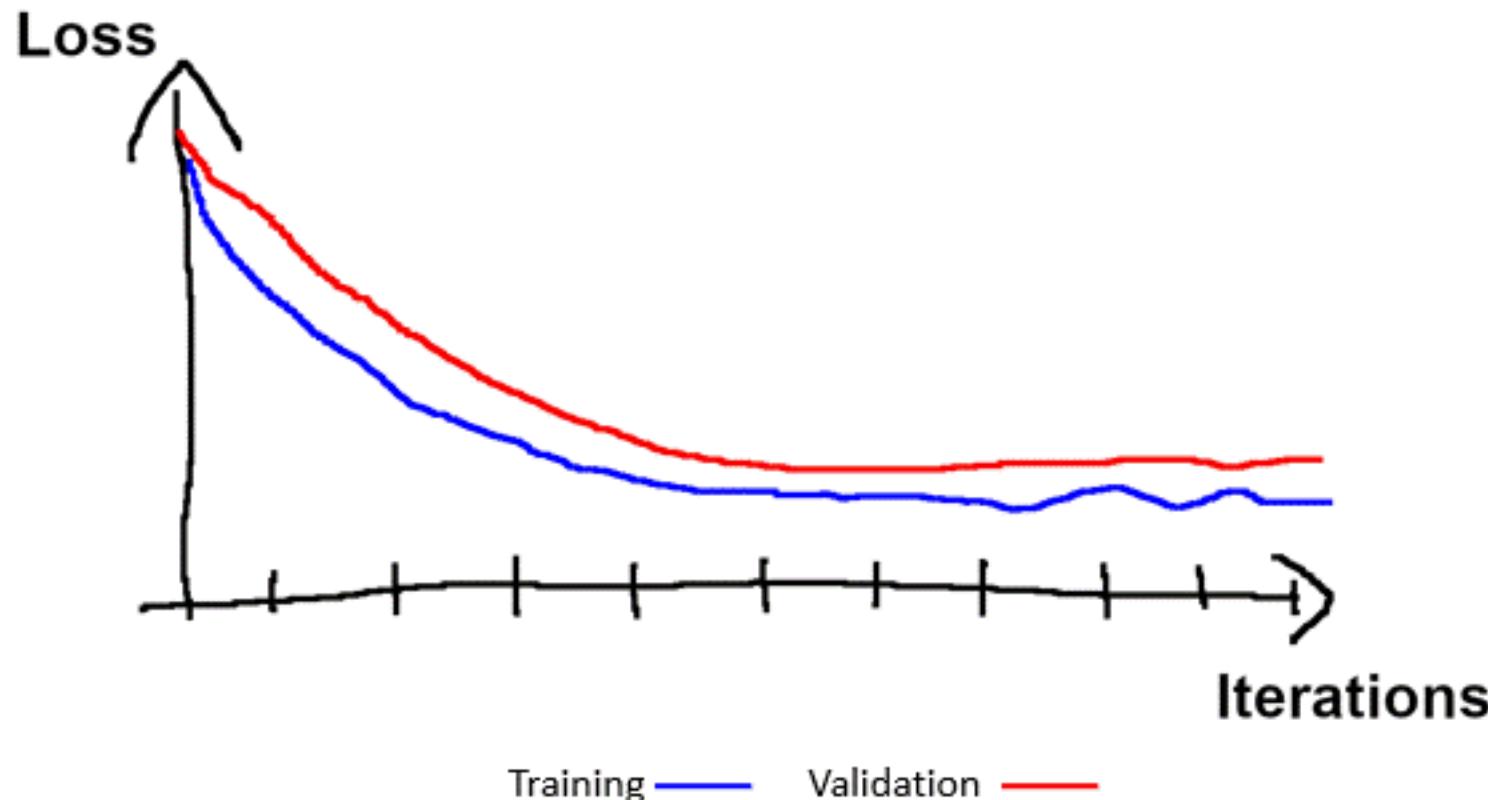


TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**

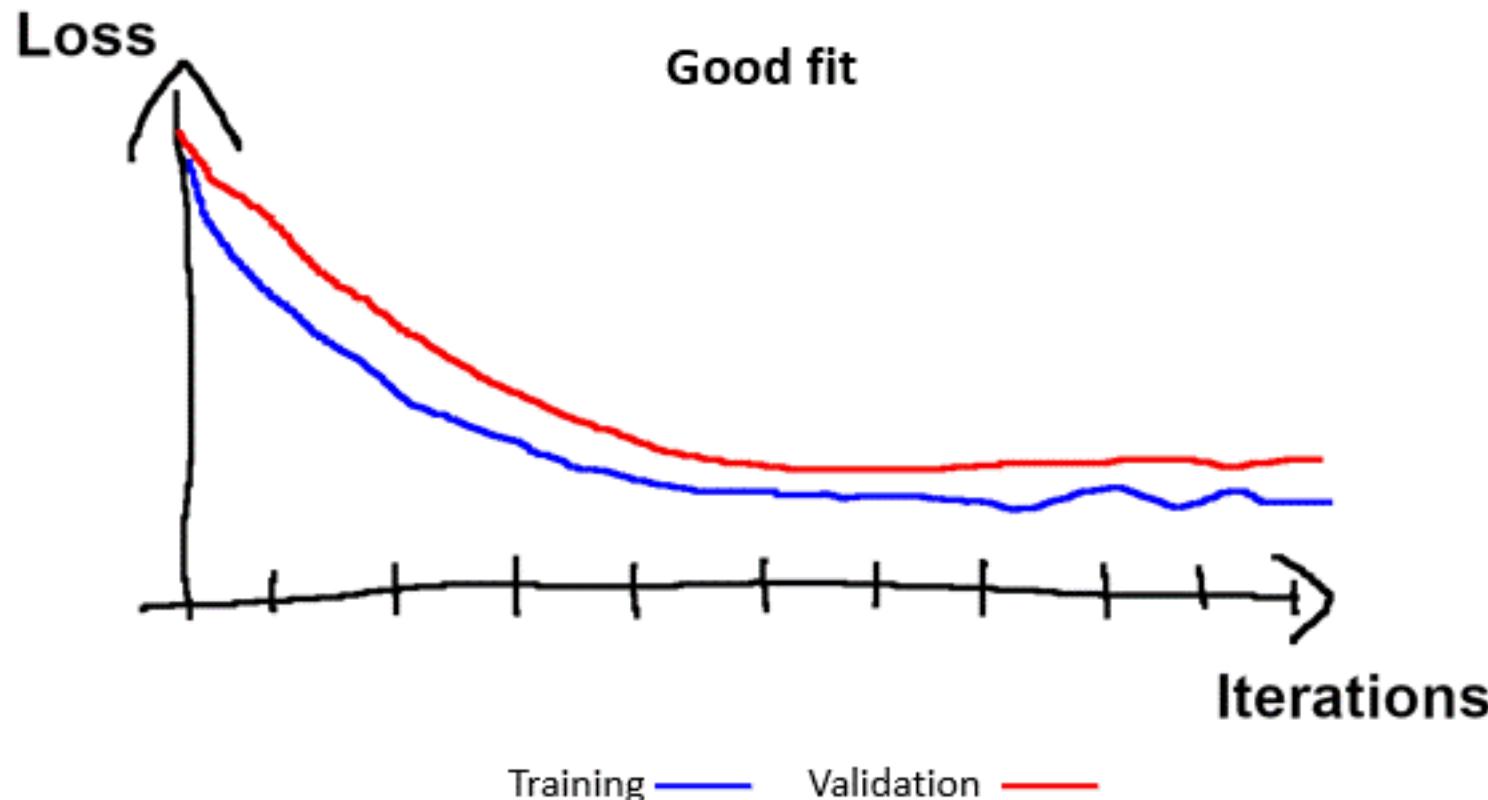
TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



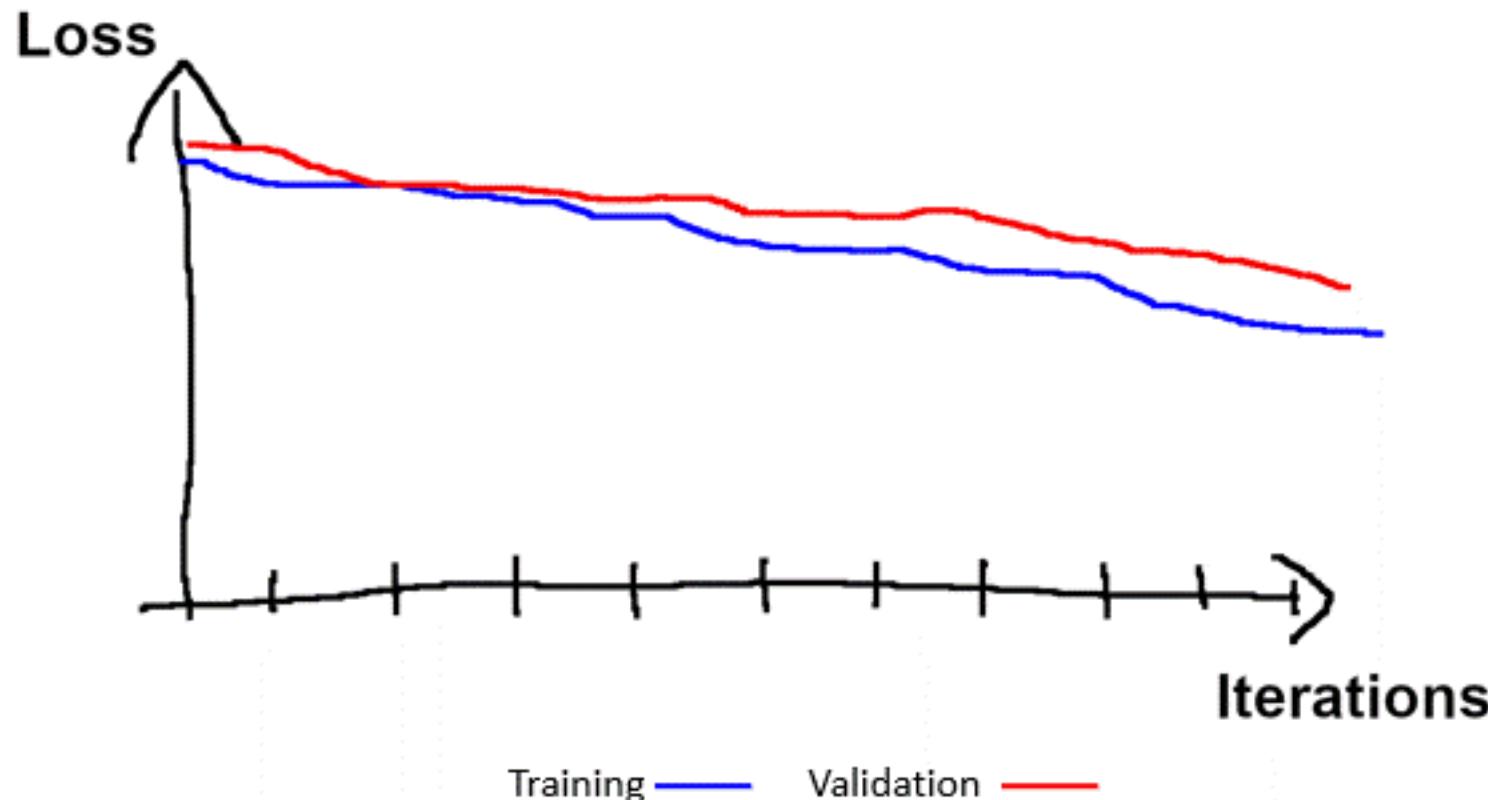
TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



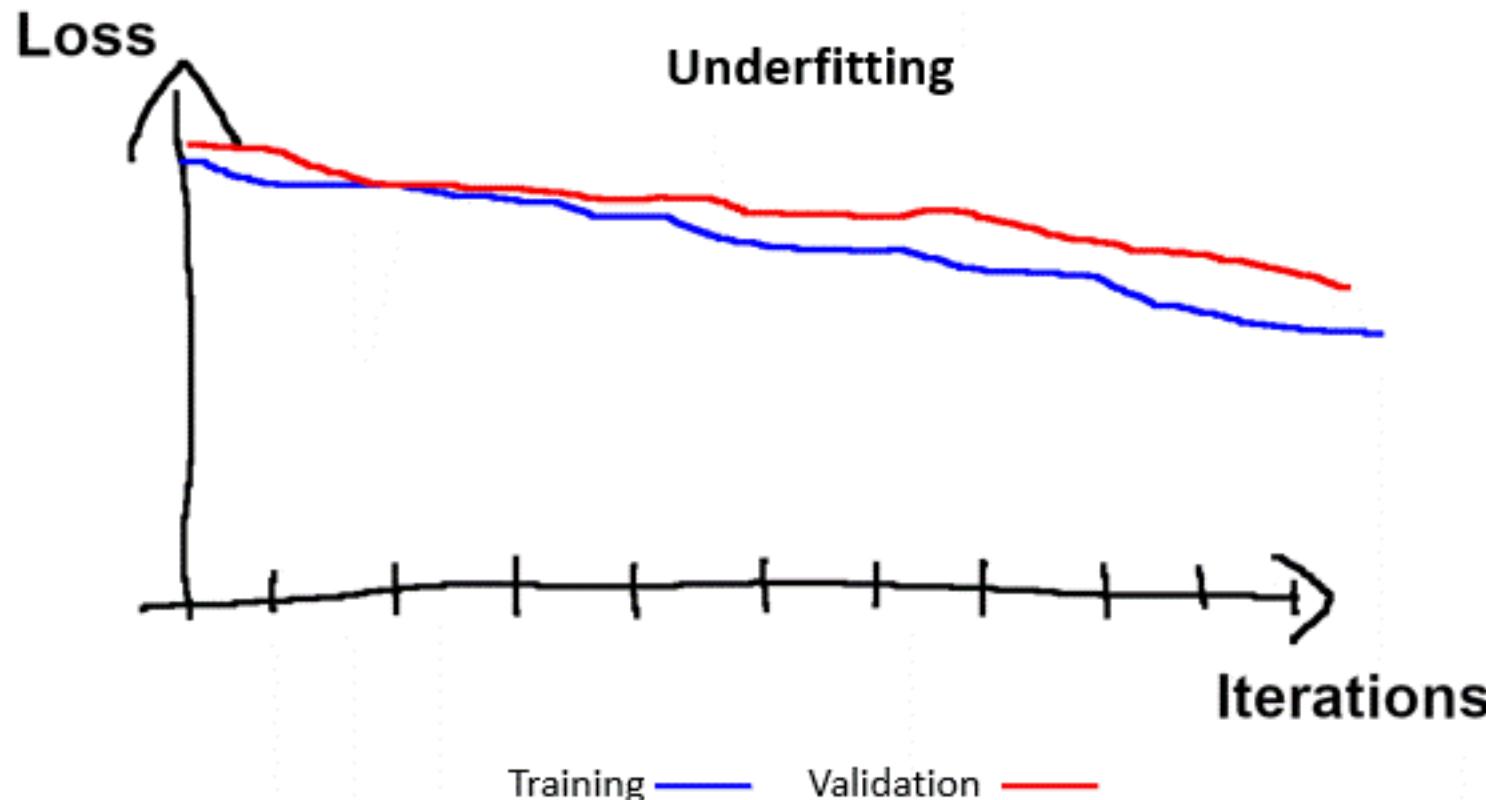
TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



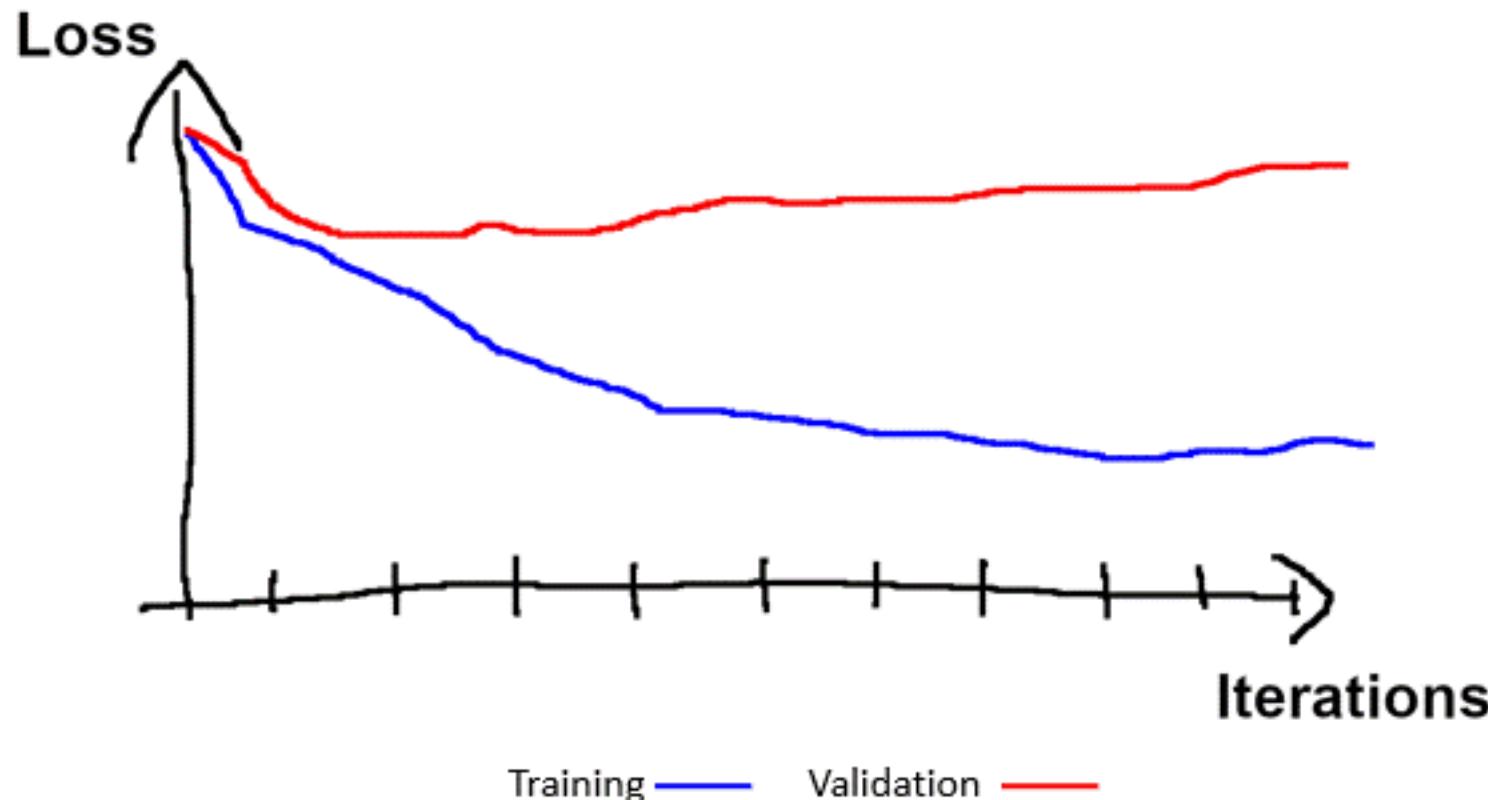
TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



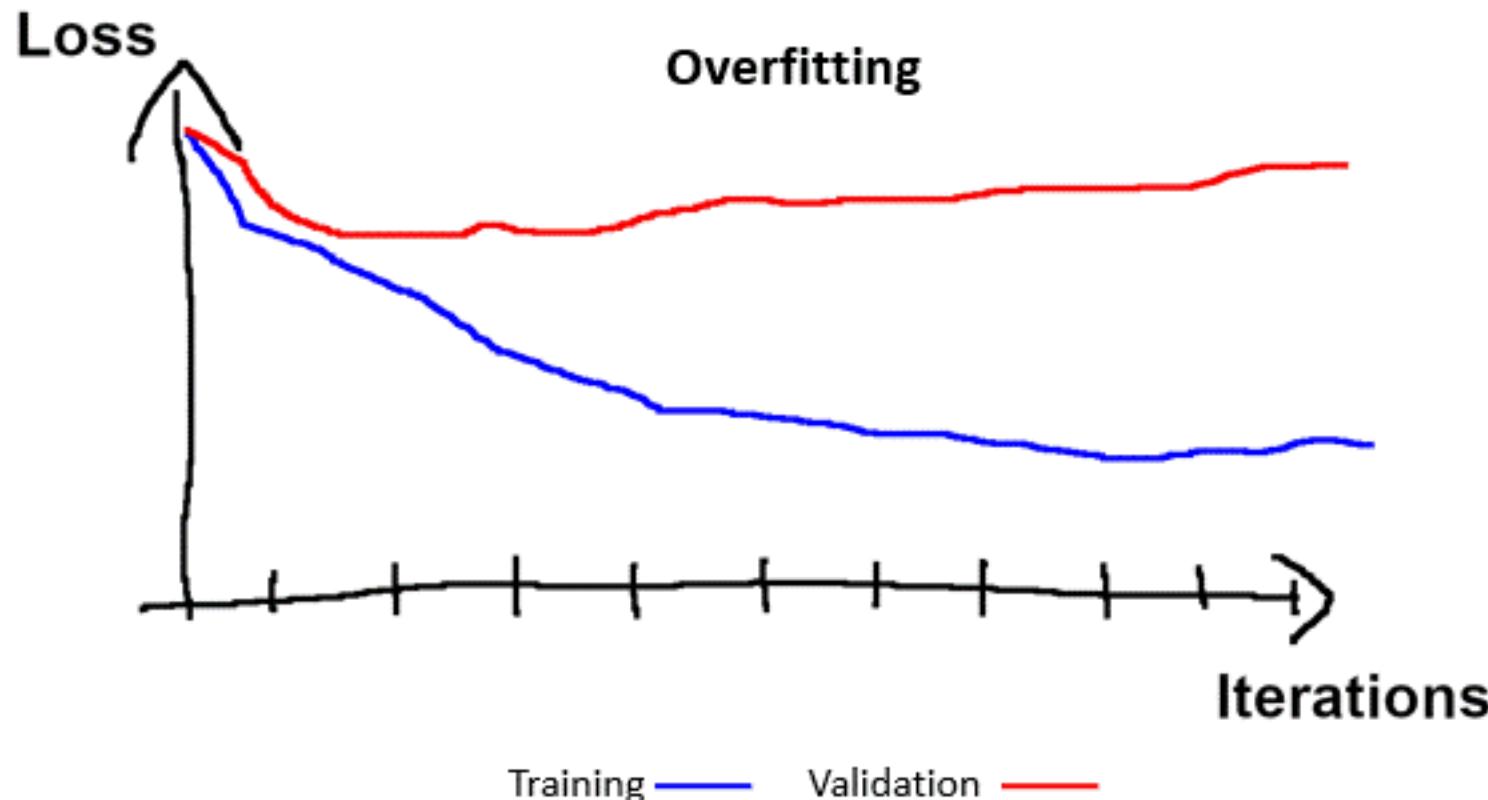
TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



TRAINING AND VALIDATION

Training dataset is divided into two parts: **training** and **validation**



FULLY CONNECTED NEURONS

Input layer:
5x5 image patch

10	15	22	18	16
29	77	76	74	17
24	79	98	72	19
59	74	81	66	18
13	14	18	17	16

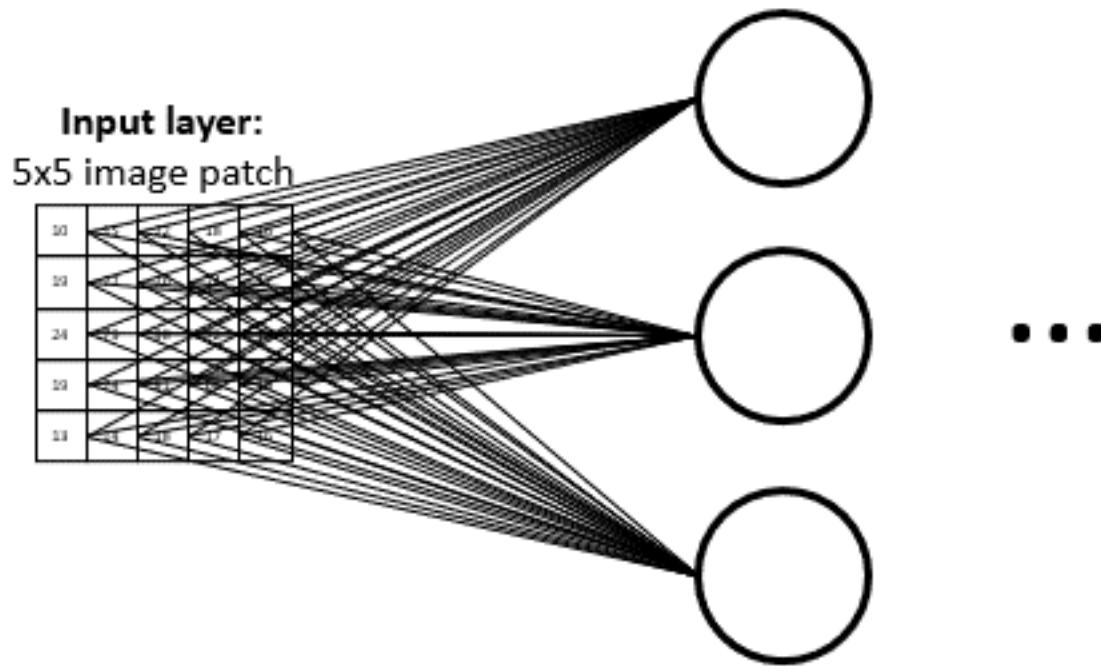
FULLY CONNECTED NEURONS

Input layer:
5x5 image patch

10	15	22	18	16
29	77	76	74	17
24	79	96	72	19
19	74	81	68	18
13	14	18	17	16



FULLY CONNECTED NEURONS



CONVOLUTION

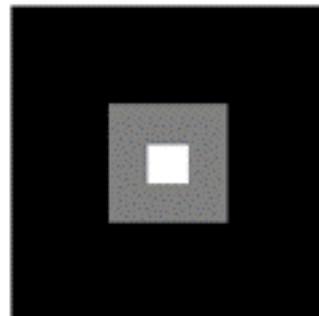
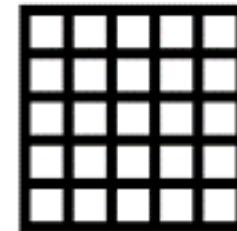
Image

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	5	5	0	0
0	0	5	10	5	0	0
0	0	5	5	5	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Filter

1	1	1
1	1	1
1	1	1

Filtered
image



CONVOLUTION

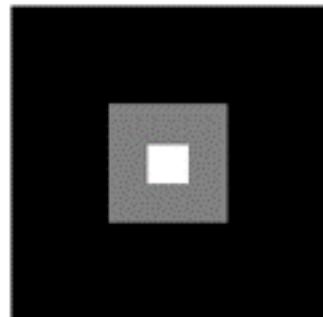
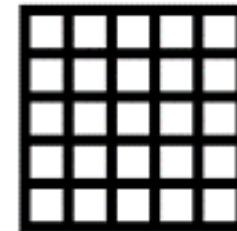
Image

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	5	5	5	0	0	0
0	0	5	10	5	0	0	0
0	0	5	5	5	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

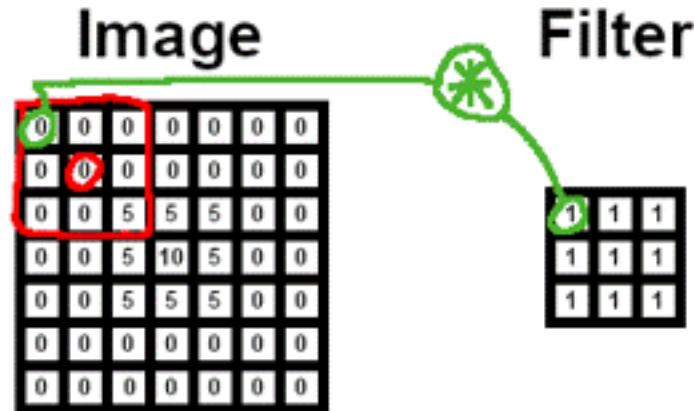
Filter

1	1	1
1	1	1
1	1	1

Filtered
image

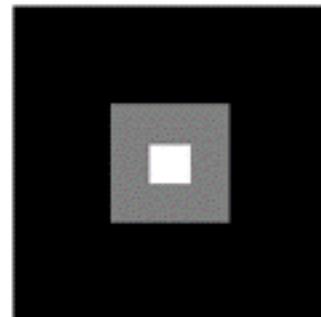
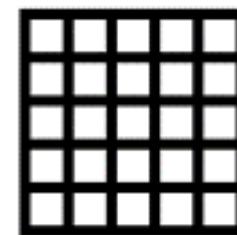


CONVOLUTION

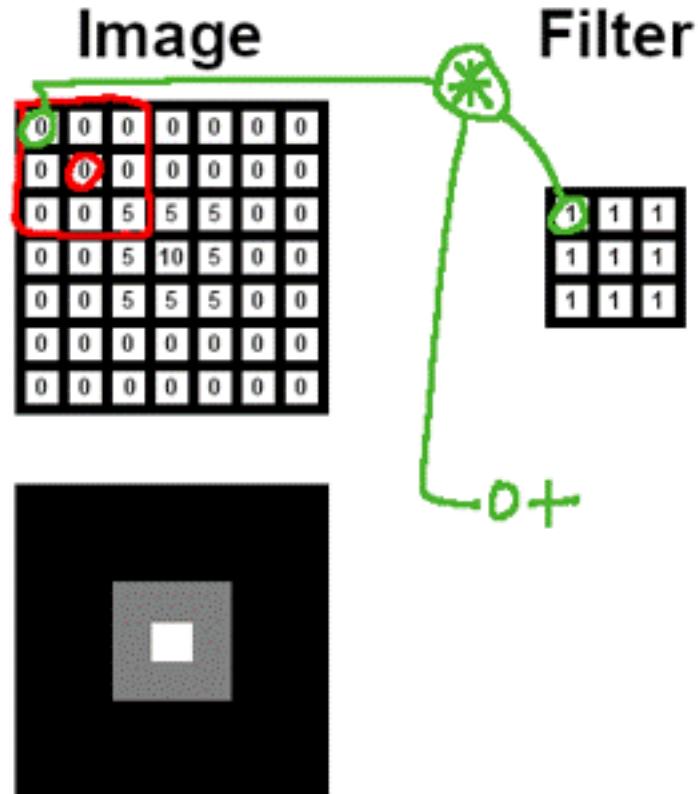


Filter

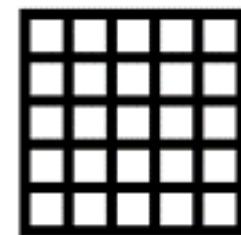
**Filtered
image**



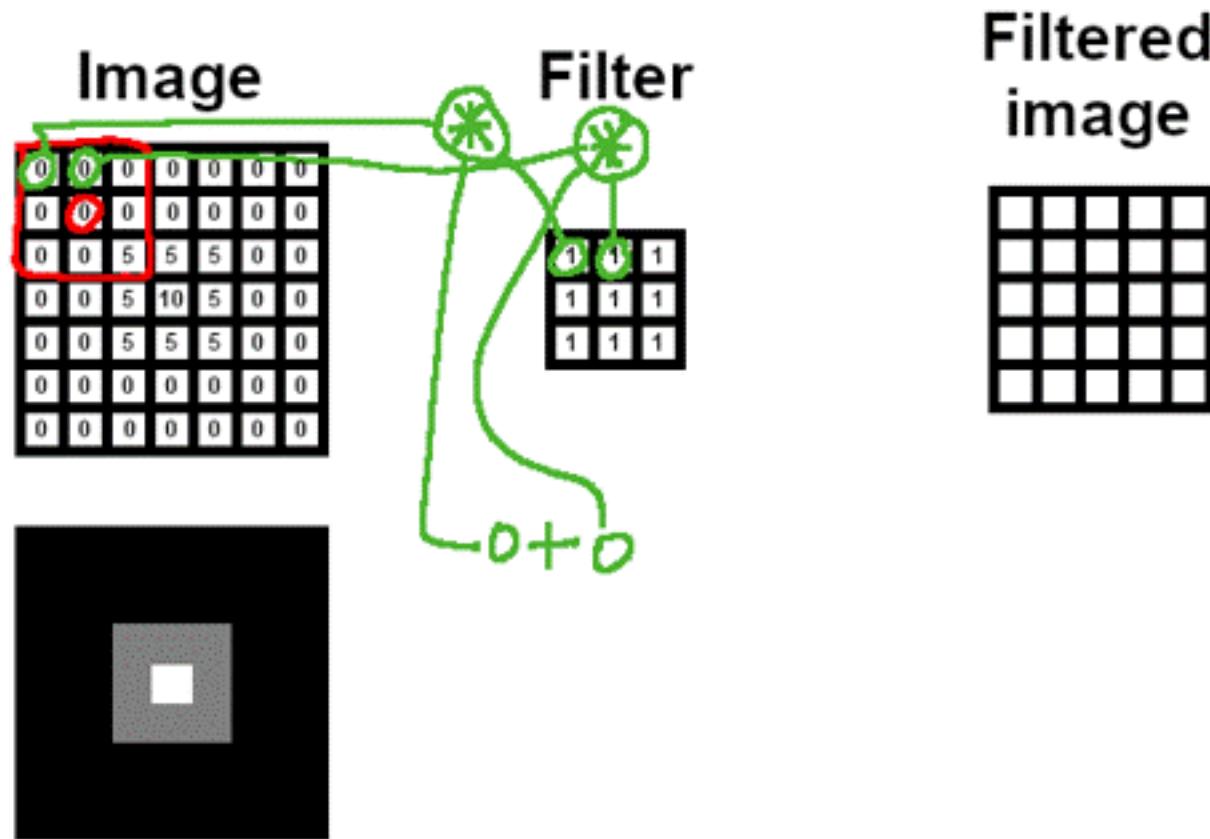
CONVOLUTION



Filtered
image

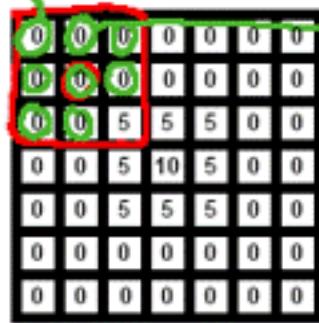


CONVOLUTION



CONVOLUTION

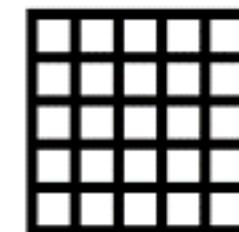
Image



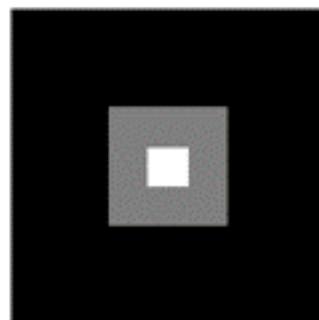
Filter



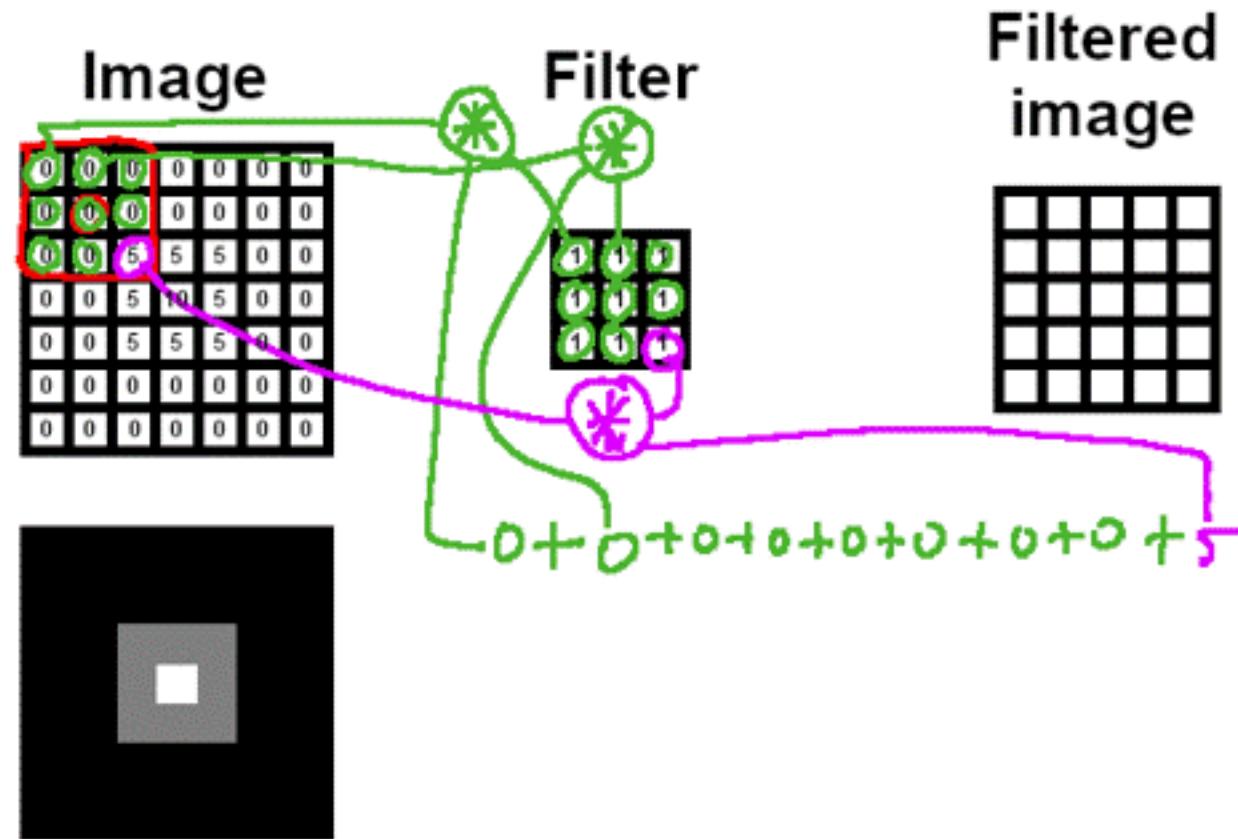
Filtered image



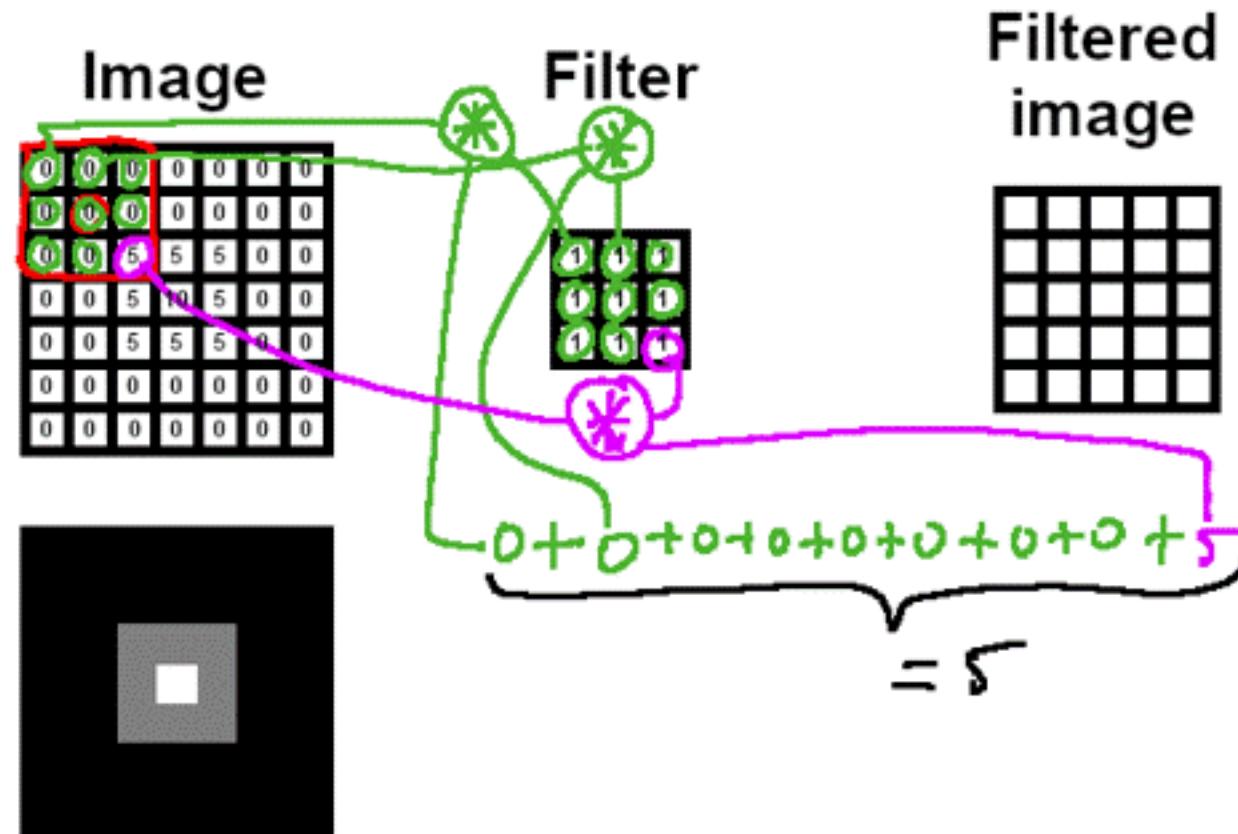
$\cdot 0 + \overline{0} + 0 + 0 + 0 + 0 + 0 + 0 +$



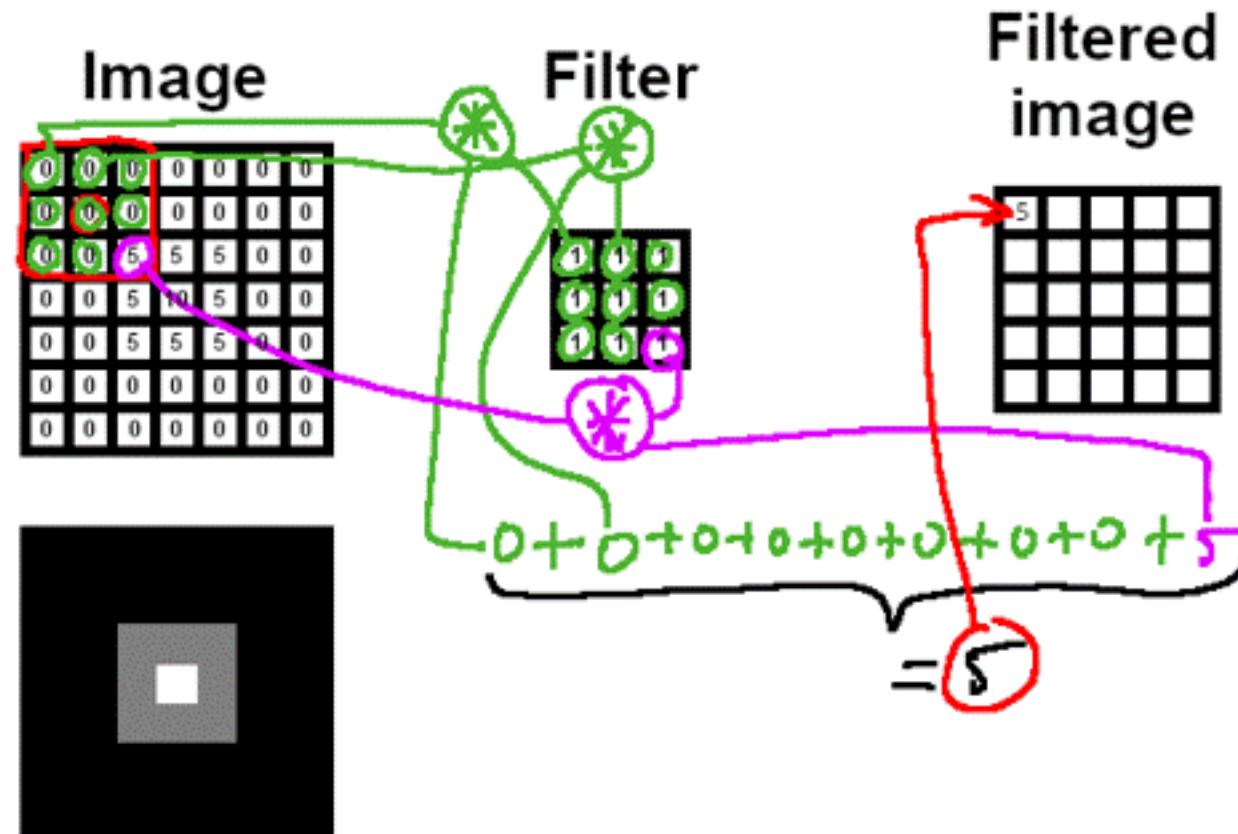
CONVOLUTION



CONVOLUTION



CONVOLUTION



CONVOLUTION

Image

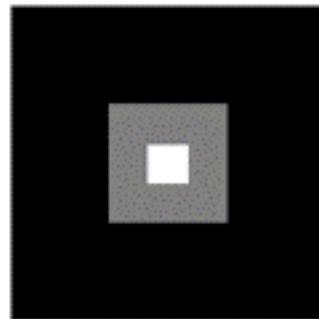
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	5	5	0	0
0	0	5	10	5	0	0
0	0	5	5	5	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Filter

1	1	1
1	1	1
1	1	1

Filtered
image

5			



CONVOLUTION

Image

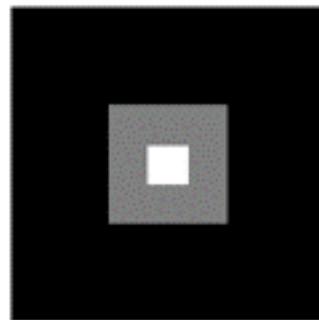
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	5	5	5	0	0	0
0	0	5	10	5	0	0	0
0	0	5	5	5	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Filter

1	1	1
1	1	1
1	1	1

Filtered
image

5			



CONVOLUTION

Image

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	5	5	5	0	0	0
0	0	5	10	5	0	0	0
0	0	5	5	5	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

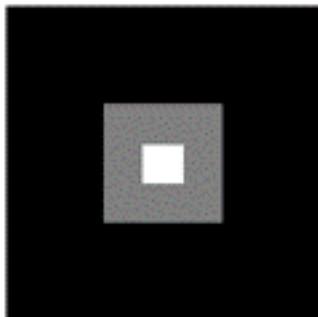
Filter

1	1	1
1	1	1
1	1	1

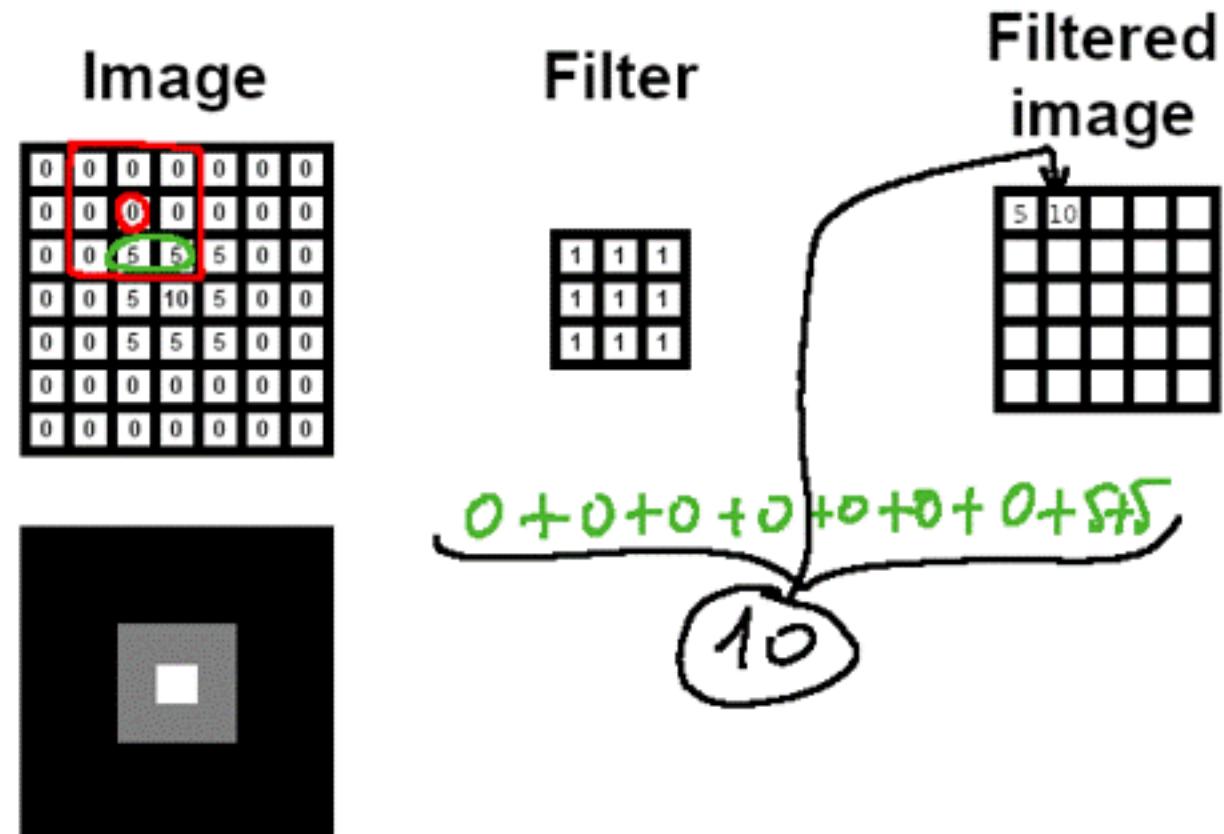
Filtered
image

5			

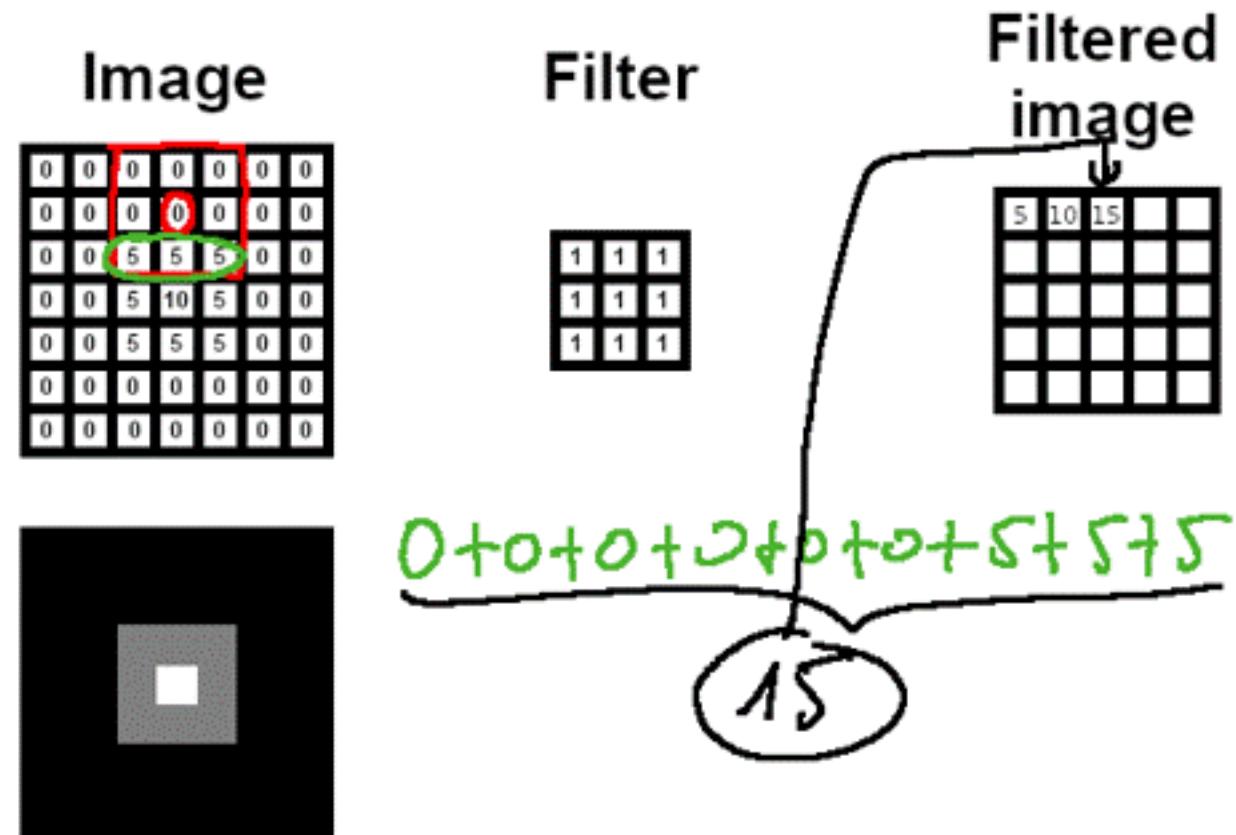
$$0+0+0+0+0+0+\cancel{0} = 5$$



CONVOLUTION



CONVOLUTION



CONVOLUTION

Image

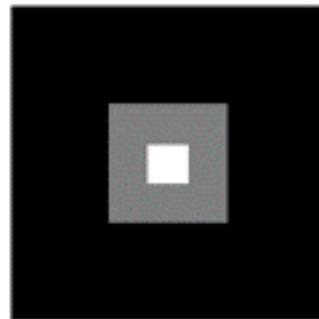
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	5	5	5	0	0
0	0	5	10	5	0	0
0	0	5	5	5	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

Filter

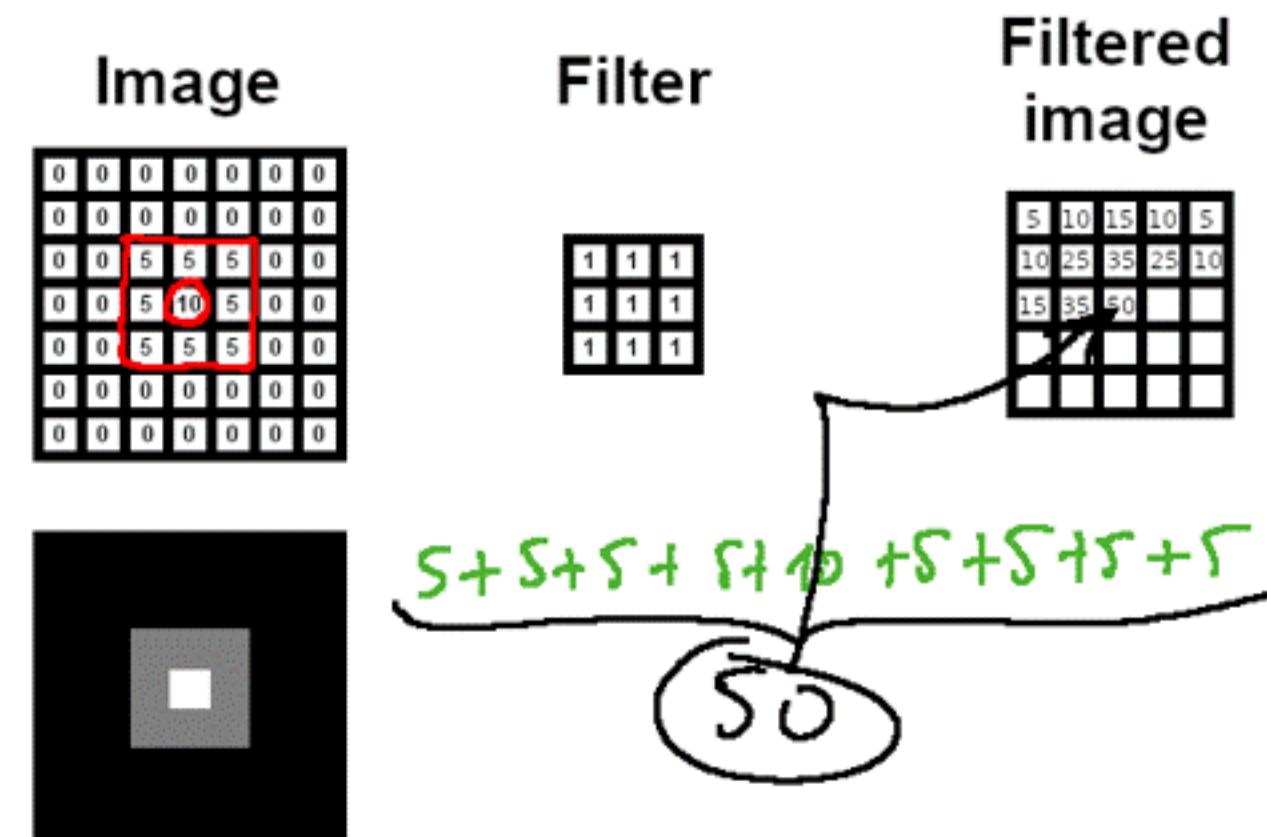
1	1	1
1	1	1
1	1	1

Filtered
image

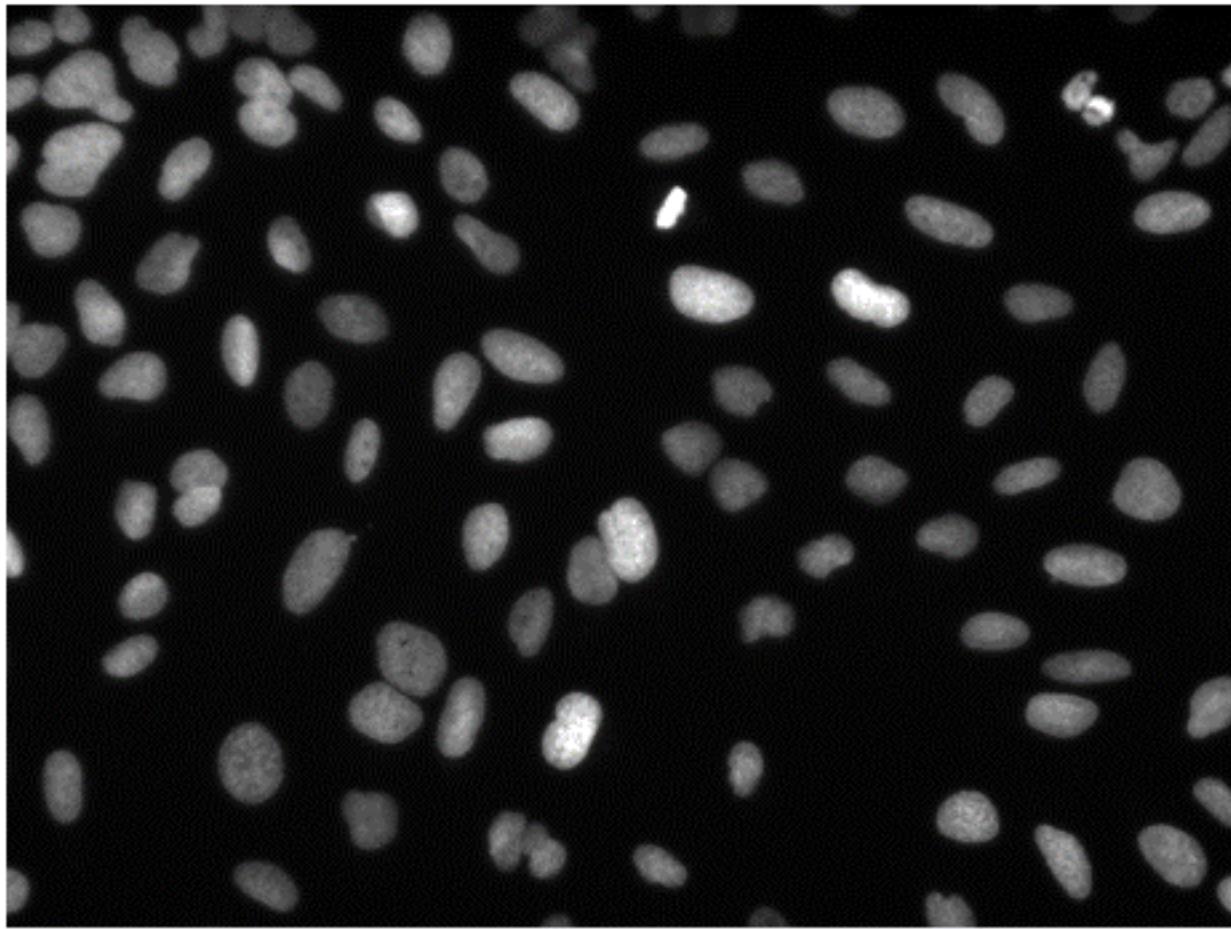
5	10	15	10	5
10	25	35	25	10
15	35			



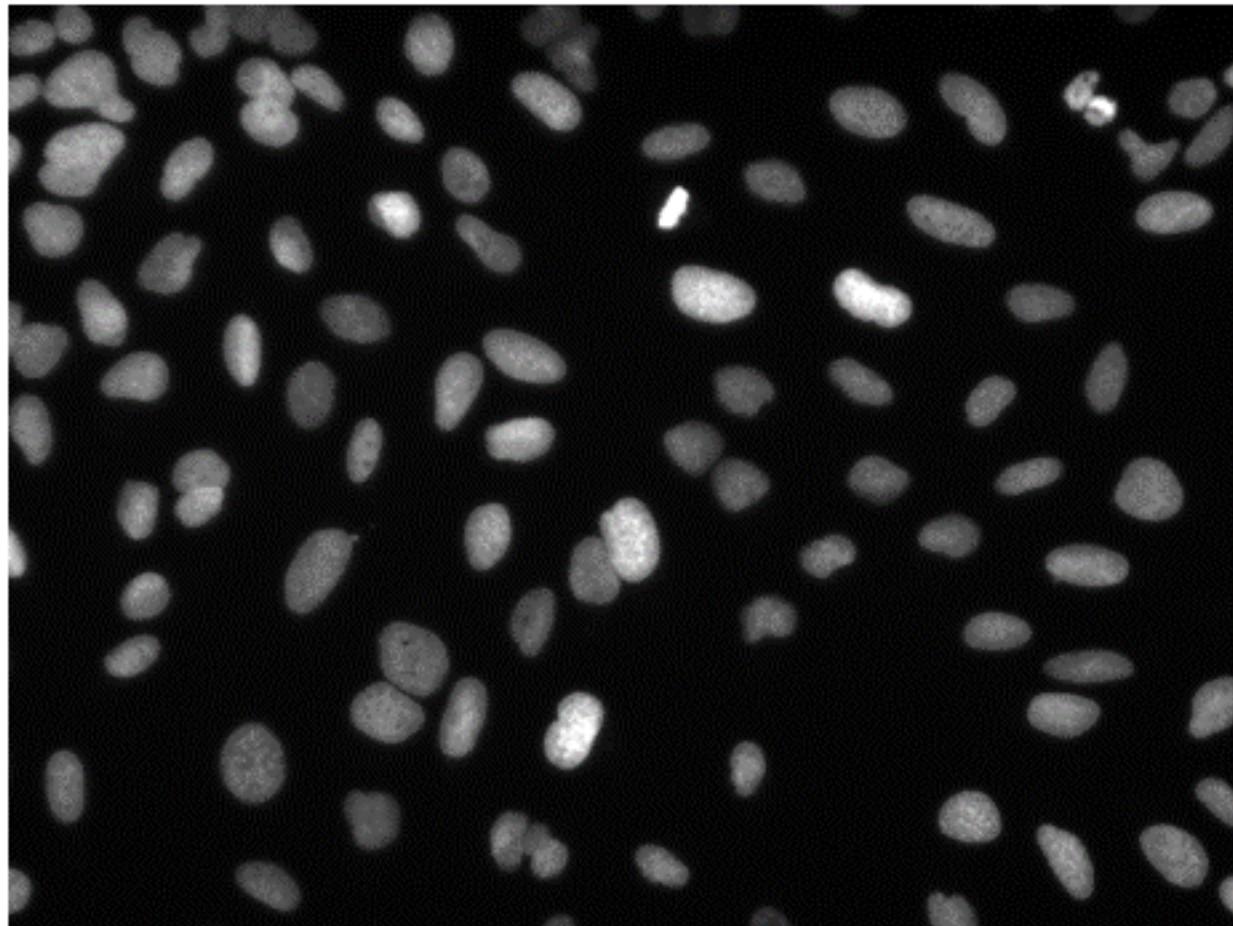
CONVOLUTION



NORMALIZATION

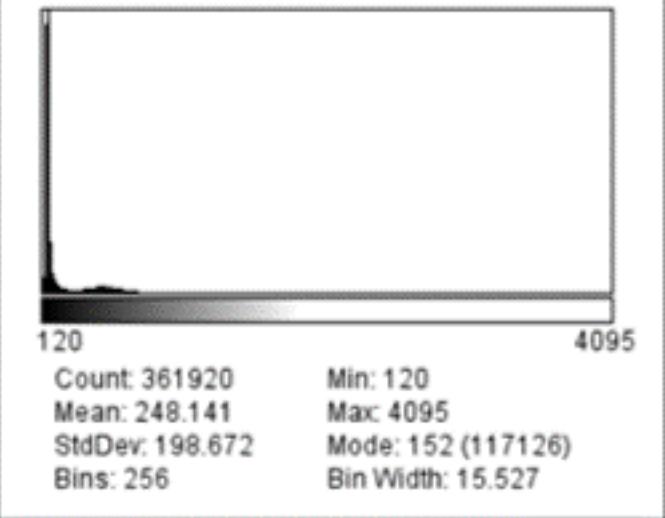


NORMALIZATION

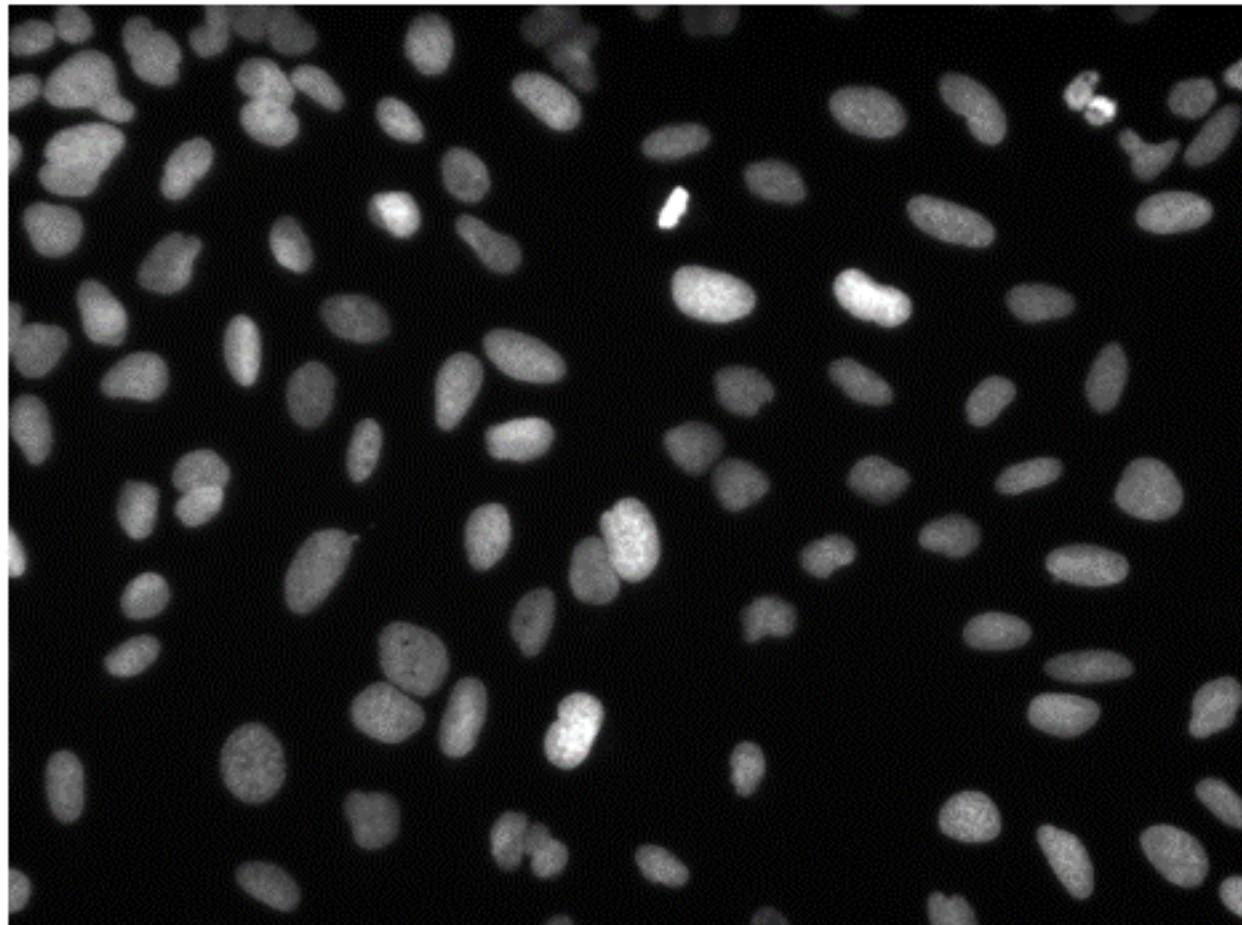


Histogram of IXMtest_A0... — □ X

300x240 pixels; RGB; 281K



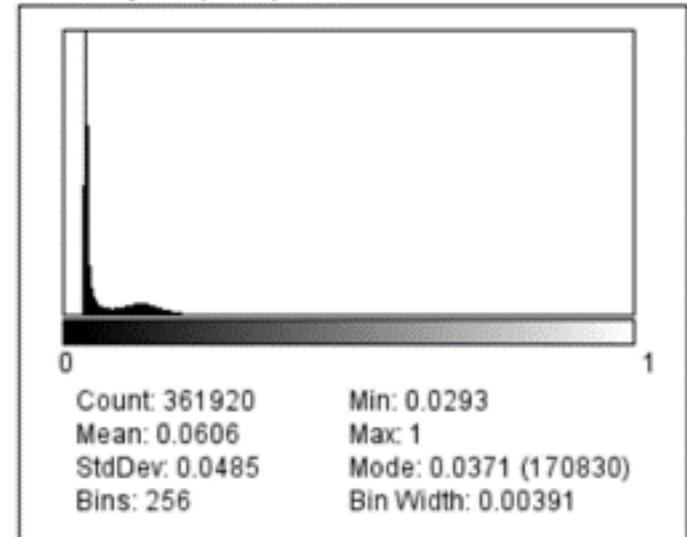
NORMALIZATION



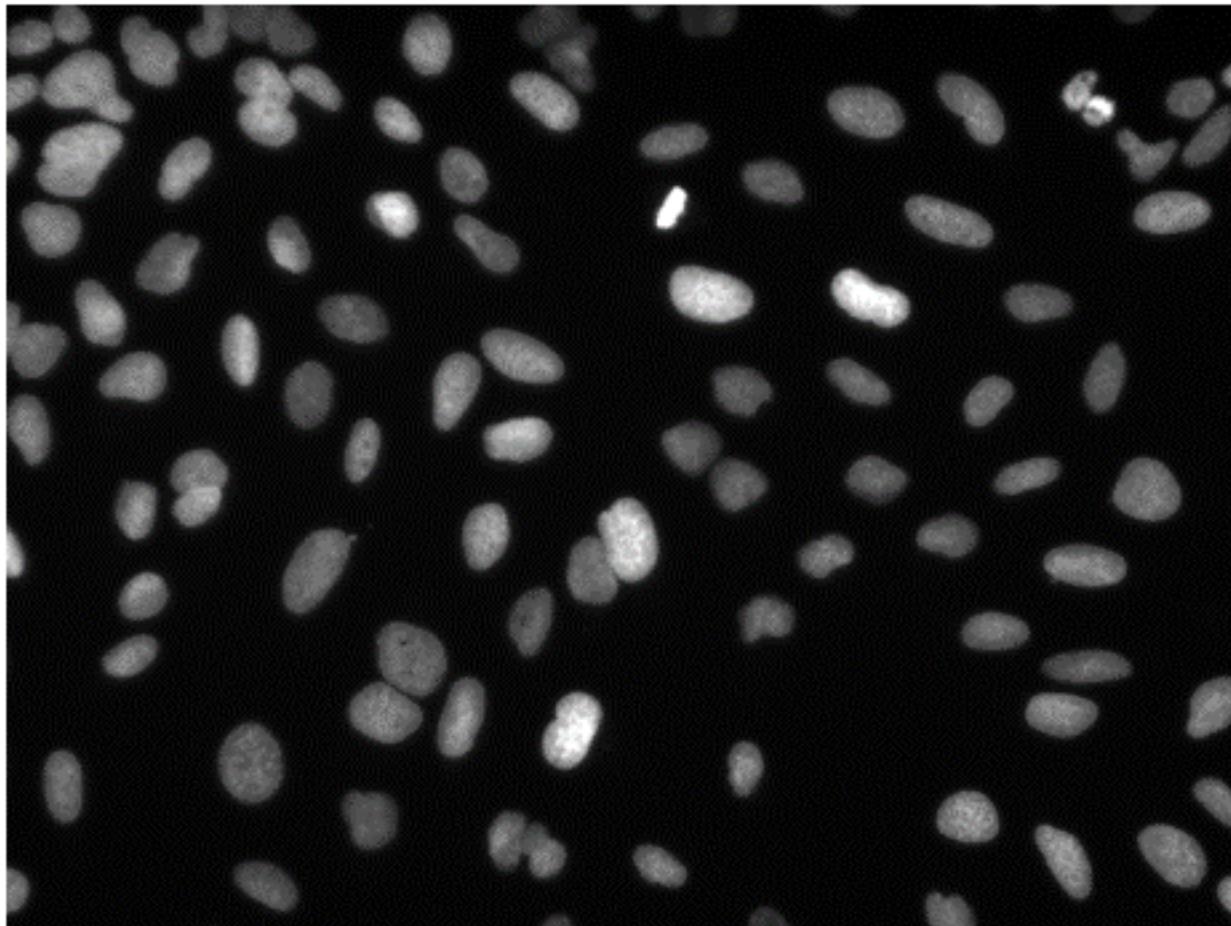
Pixel normalization

Histogram of IXMtest_A0... — □ X

300x240 pixels; RGB; 281K



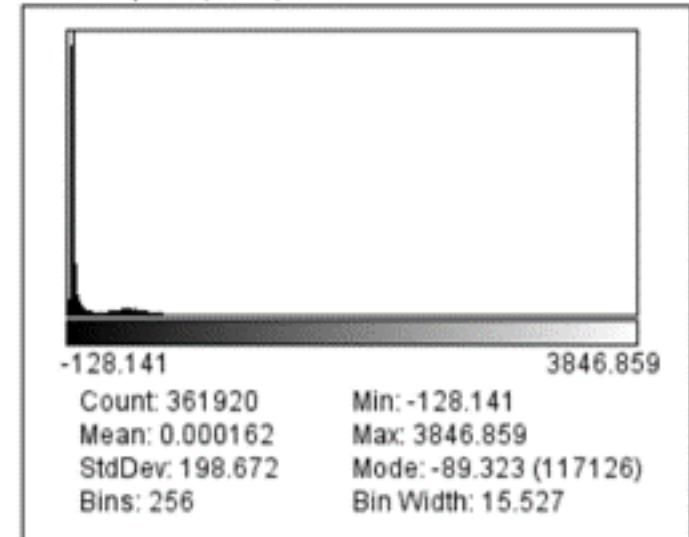
NORMALIZATION



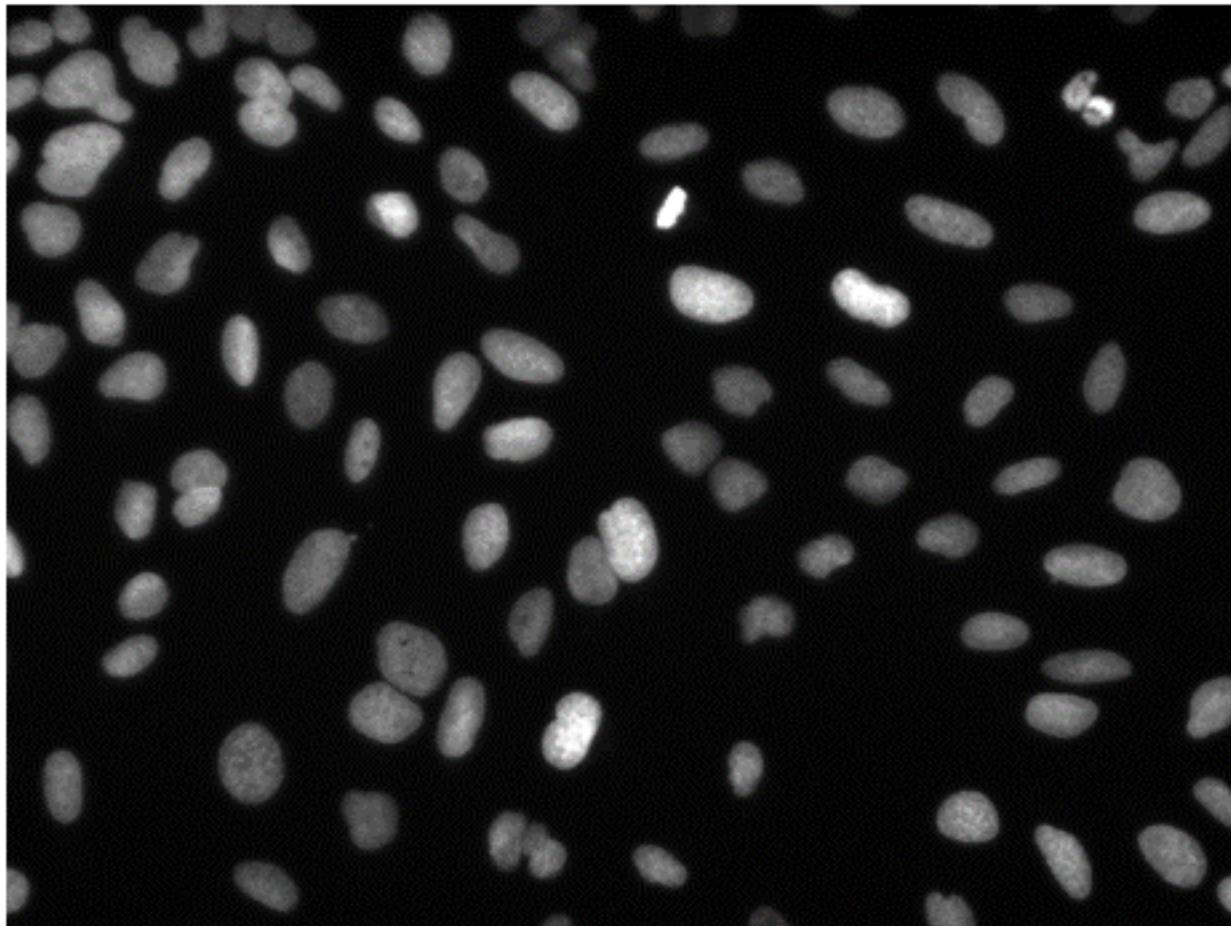
Pixel centering

Histogram of IXMtest_A0... — □ X

300x240 pixels; RGB; 281K



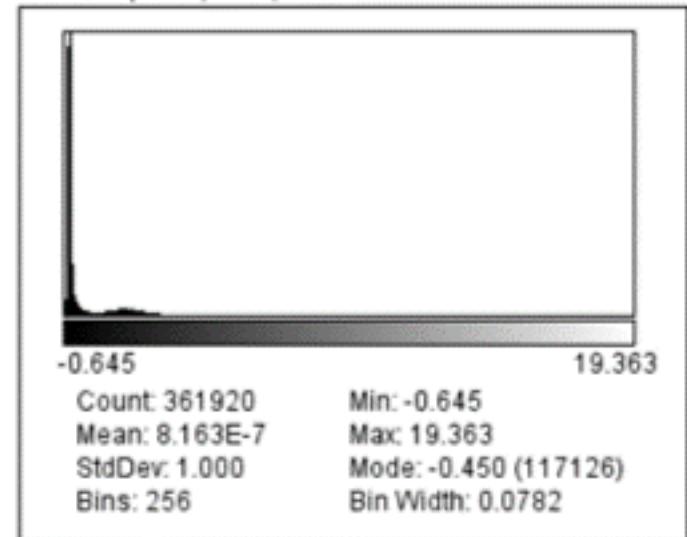
NORMALIZATION



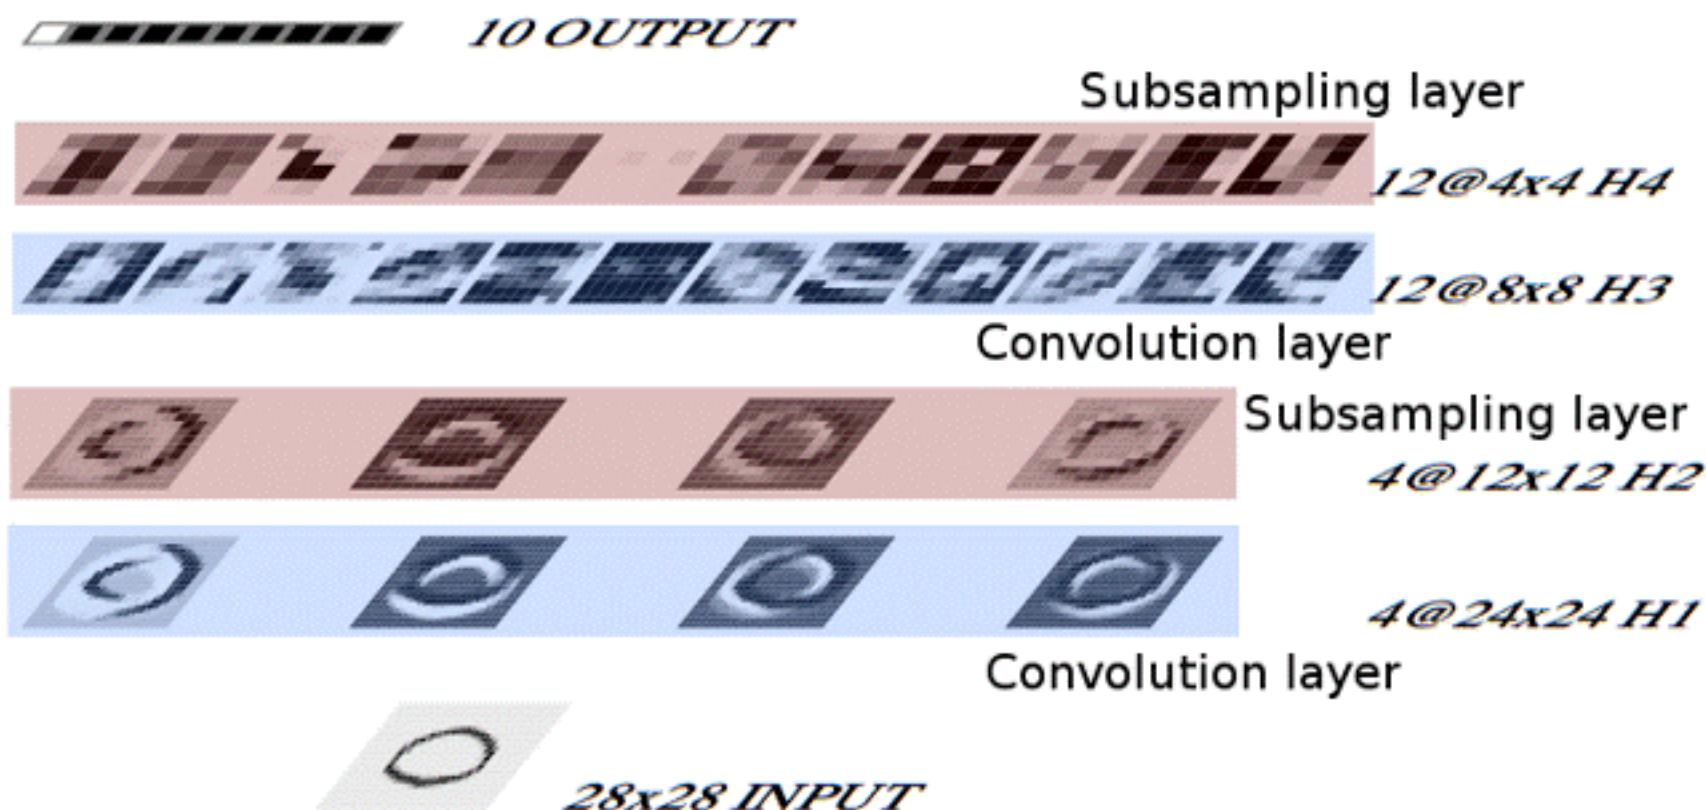
Pixel standardization

Histogram of IXMtest_A0... — □ X

300x240 pixels; RGB; 281K



CONVOLUTIONAL NEURAL NETWORK (1989)

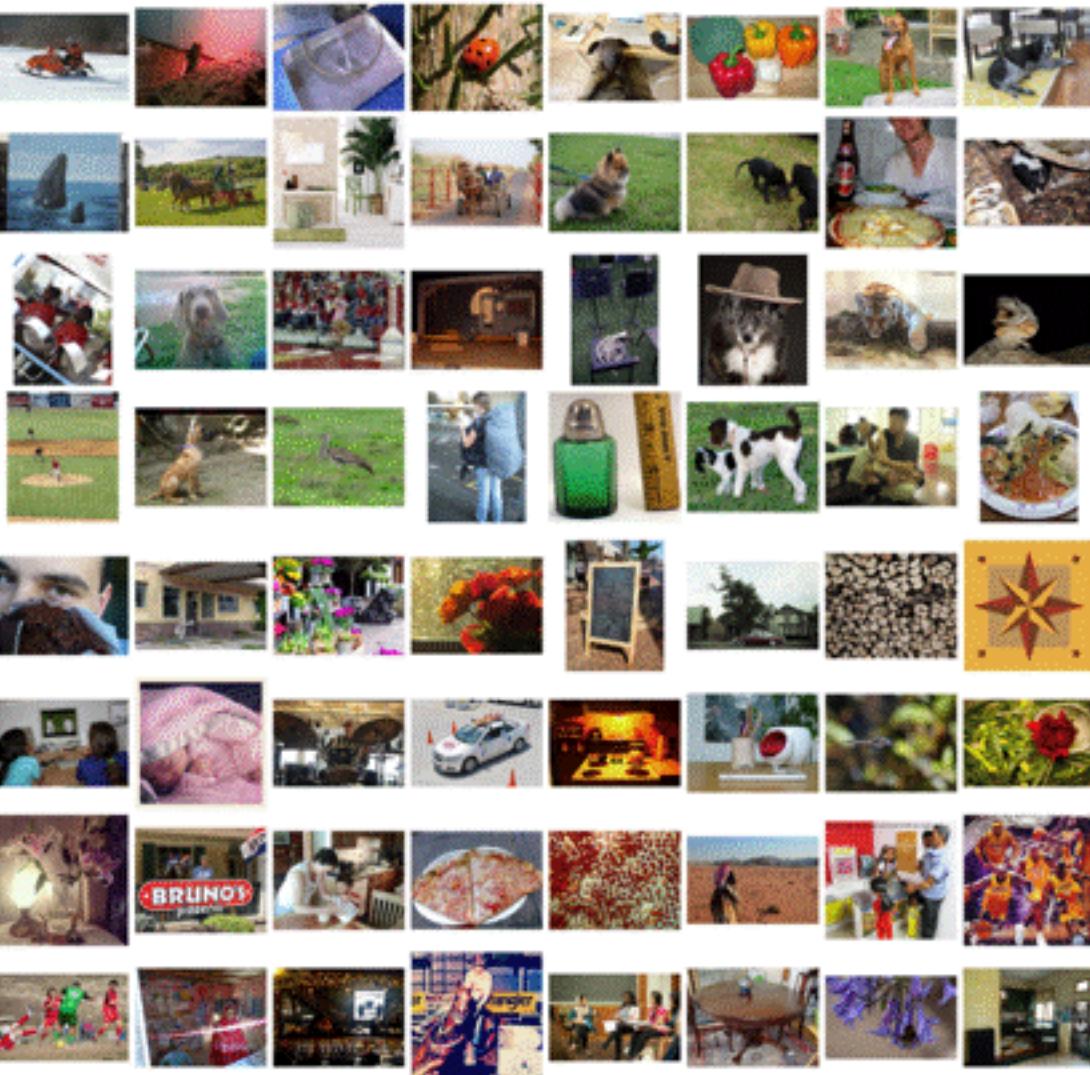


LeCun et al., Handwritten digit recognition with a back-propagation network, *Neural Information Processing Systems*, 1989

4,635 pixels and 2,578 parameters (much lower than fully connected network)

DEEP LEARNING REQUIRES

2) Large amounts of training data



ImageNet Large Scale Visual Recognition Challenge (2010 – 2017)

1000 classes – 10,000,000 images

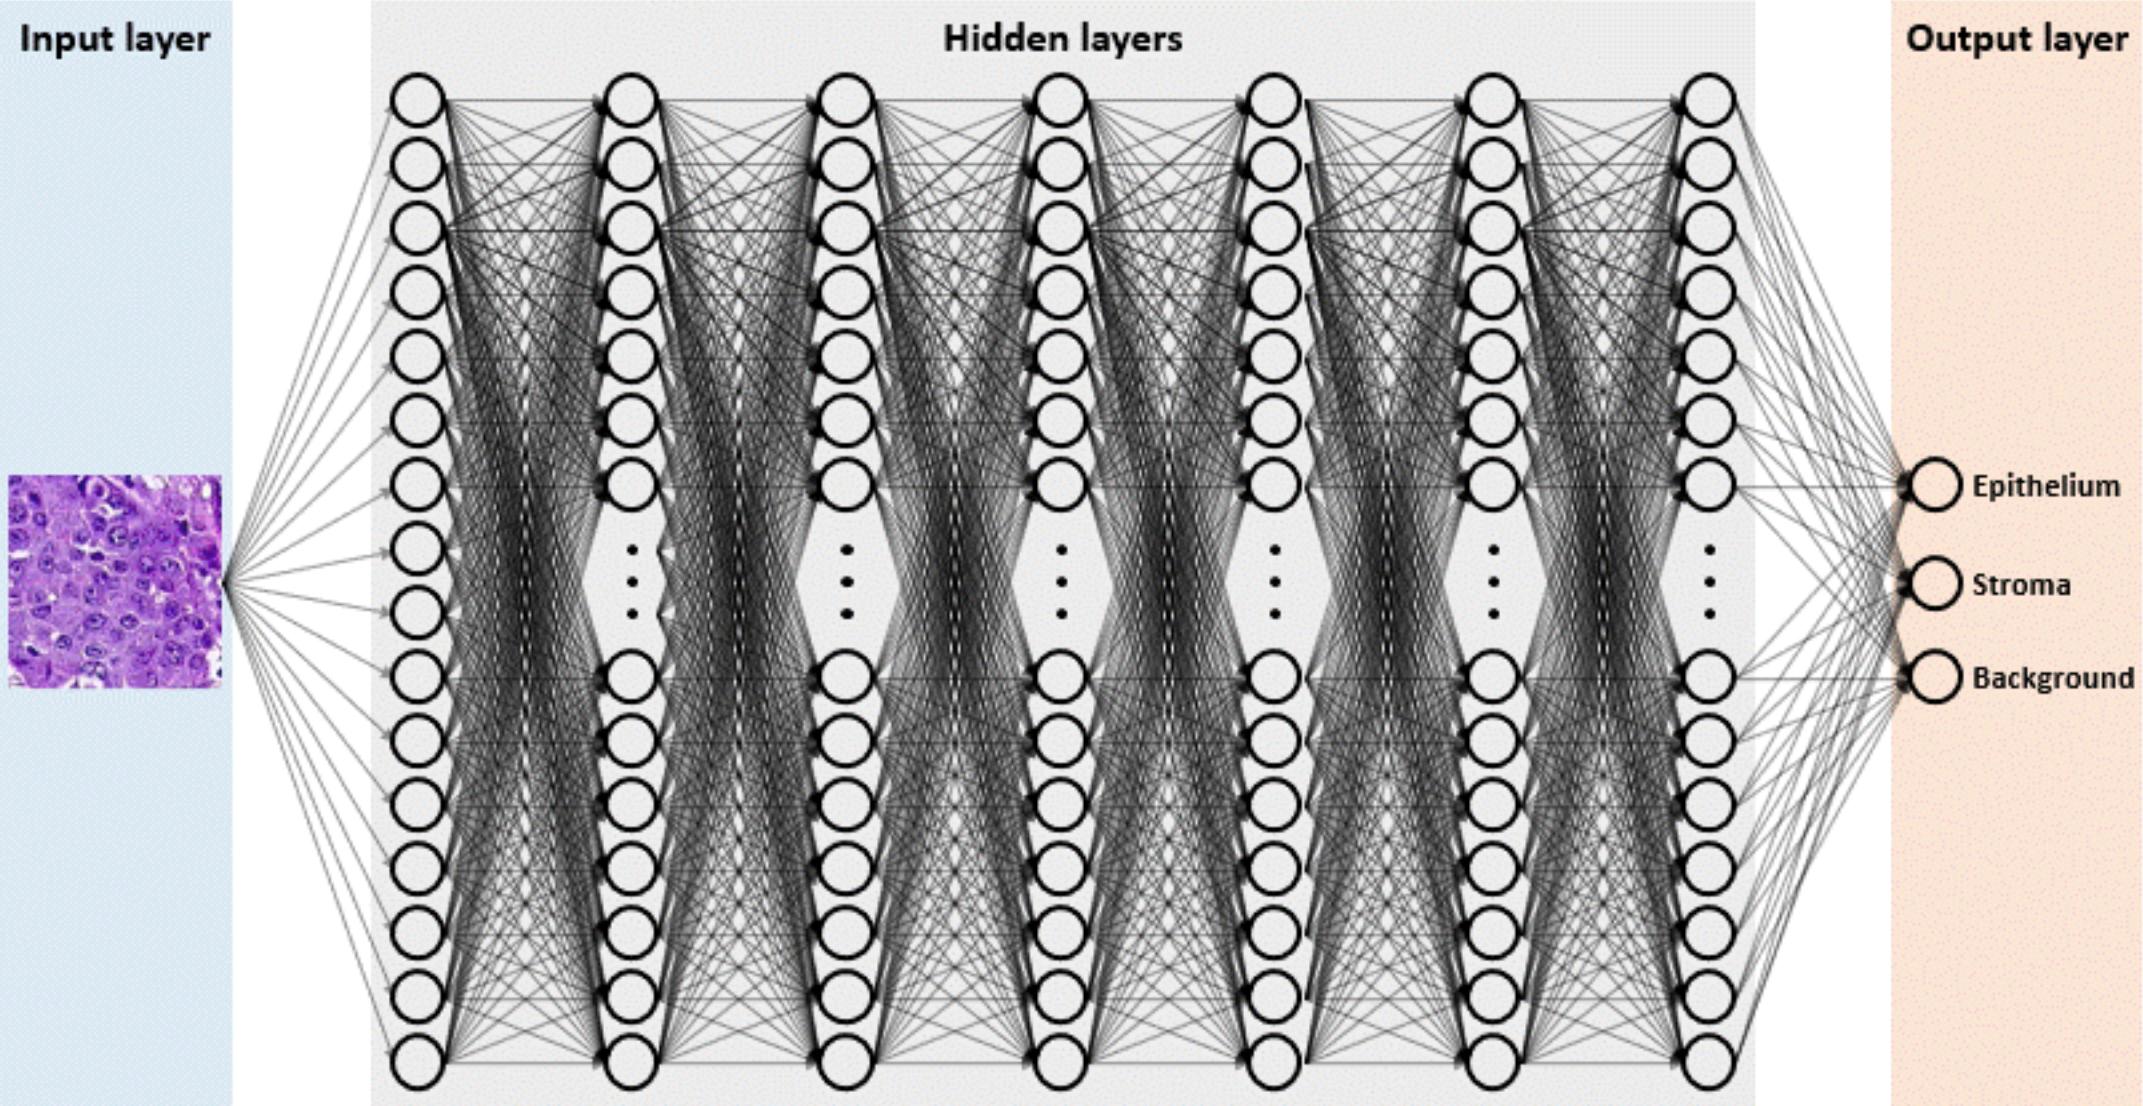
1) High computational power

4.4 Speedup factor of GPU code

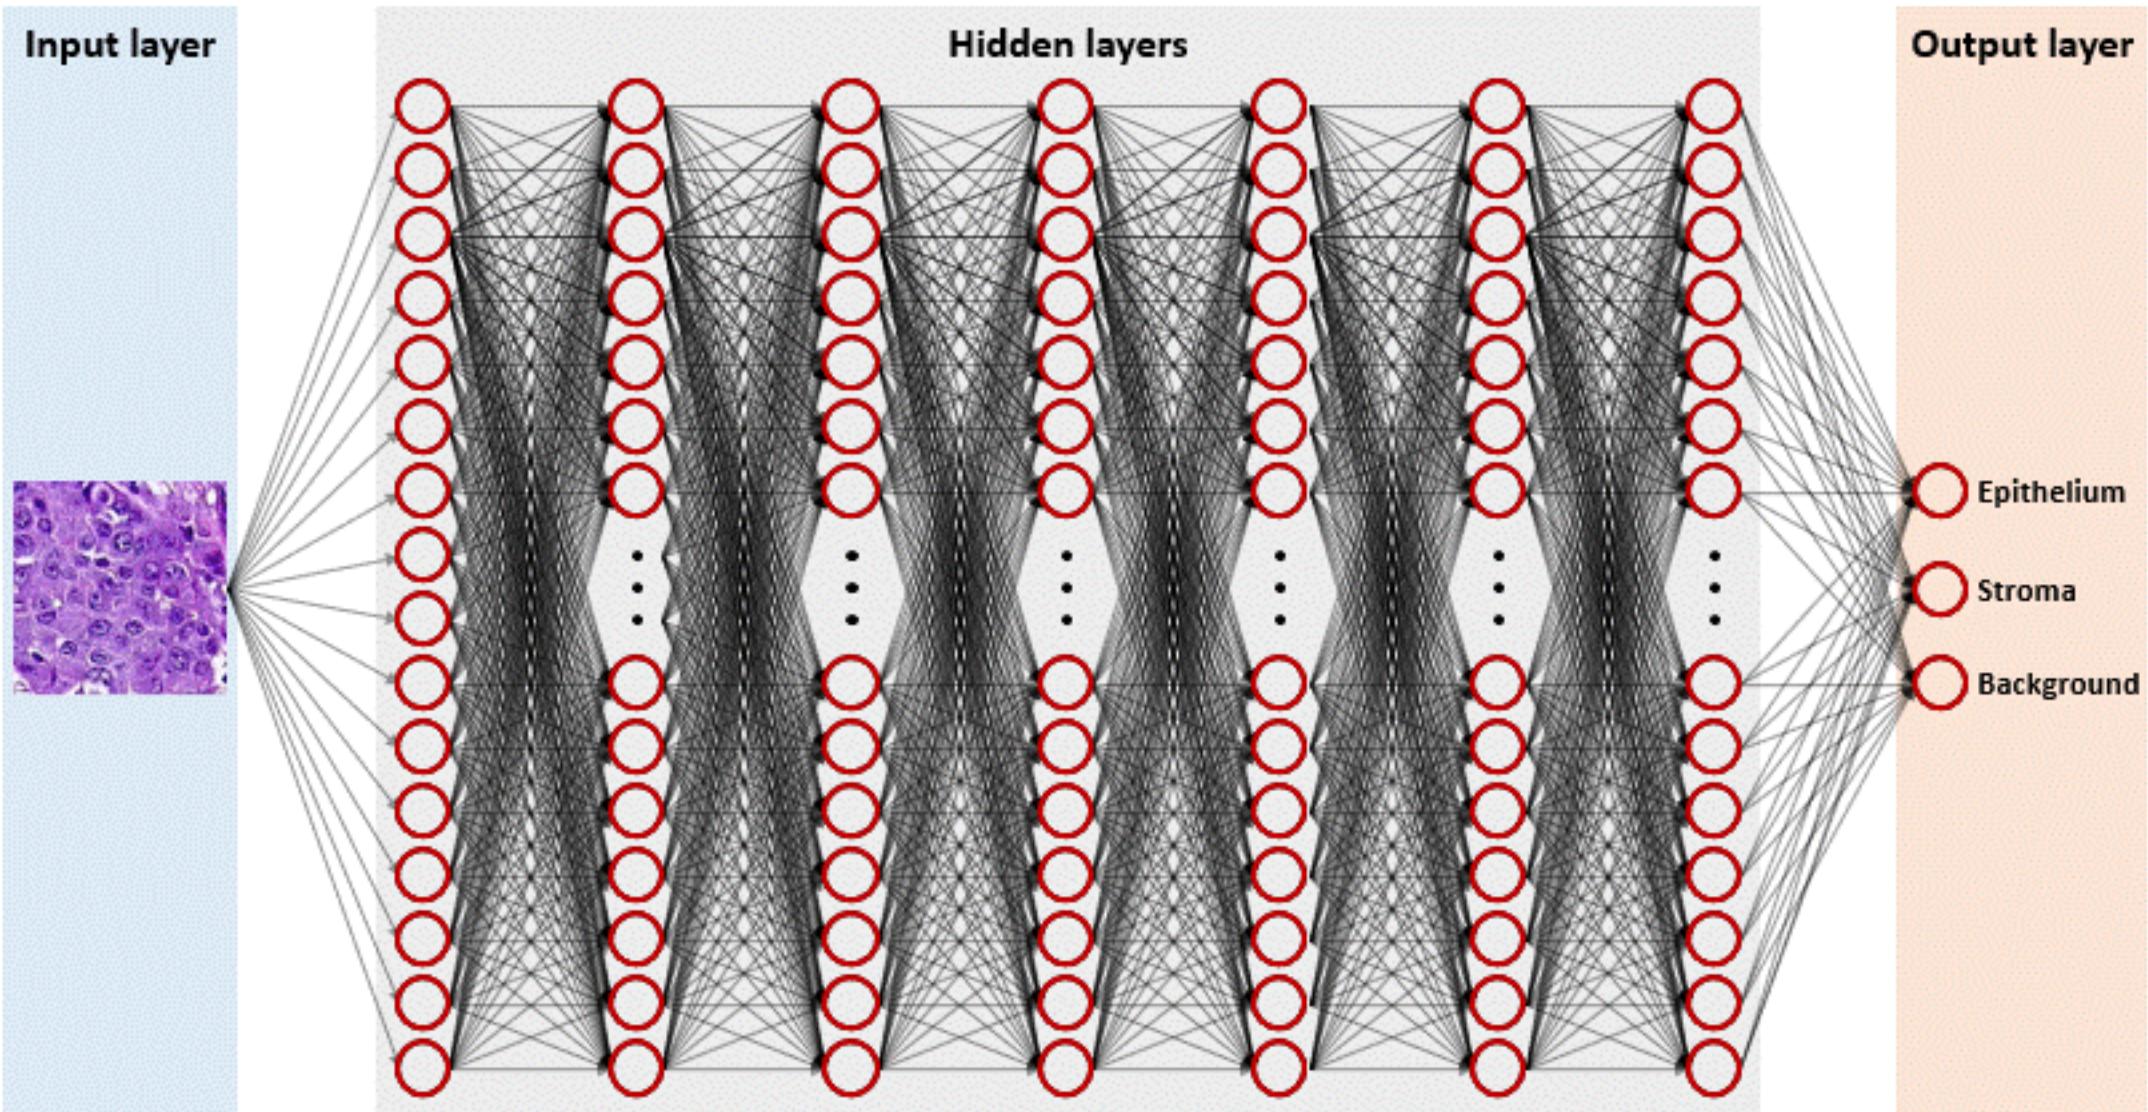
The GPU code scales well with network size. For small nets the speedup is small (but still over 10) since they fit better inside the CPU cache, and GPU resources are underutilized. For huge nets (ex: Table 2) the GPU implementation is more than 60 times faster than a compiler-optimized CPU version. Given the flexibility of our GPU version, this is a significant speedup. One epoch takes 35 GPU minutes but more than 35 CPU hours.

Cireşan *et al.*, High-Performance Neural Networks for Visual Object Classification. Tech Report, 2011.

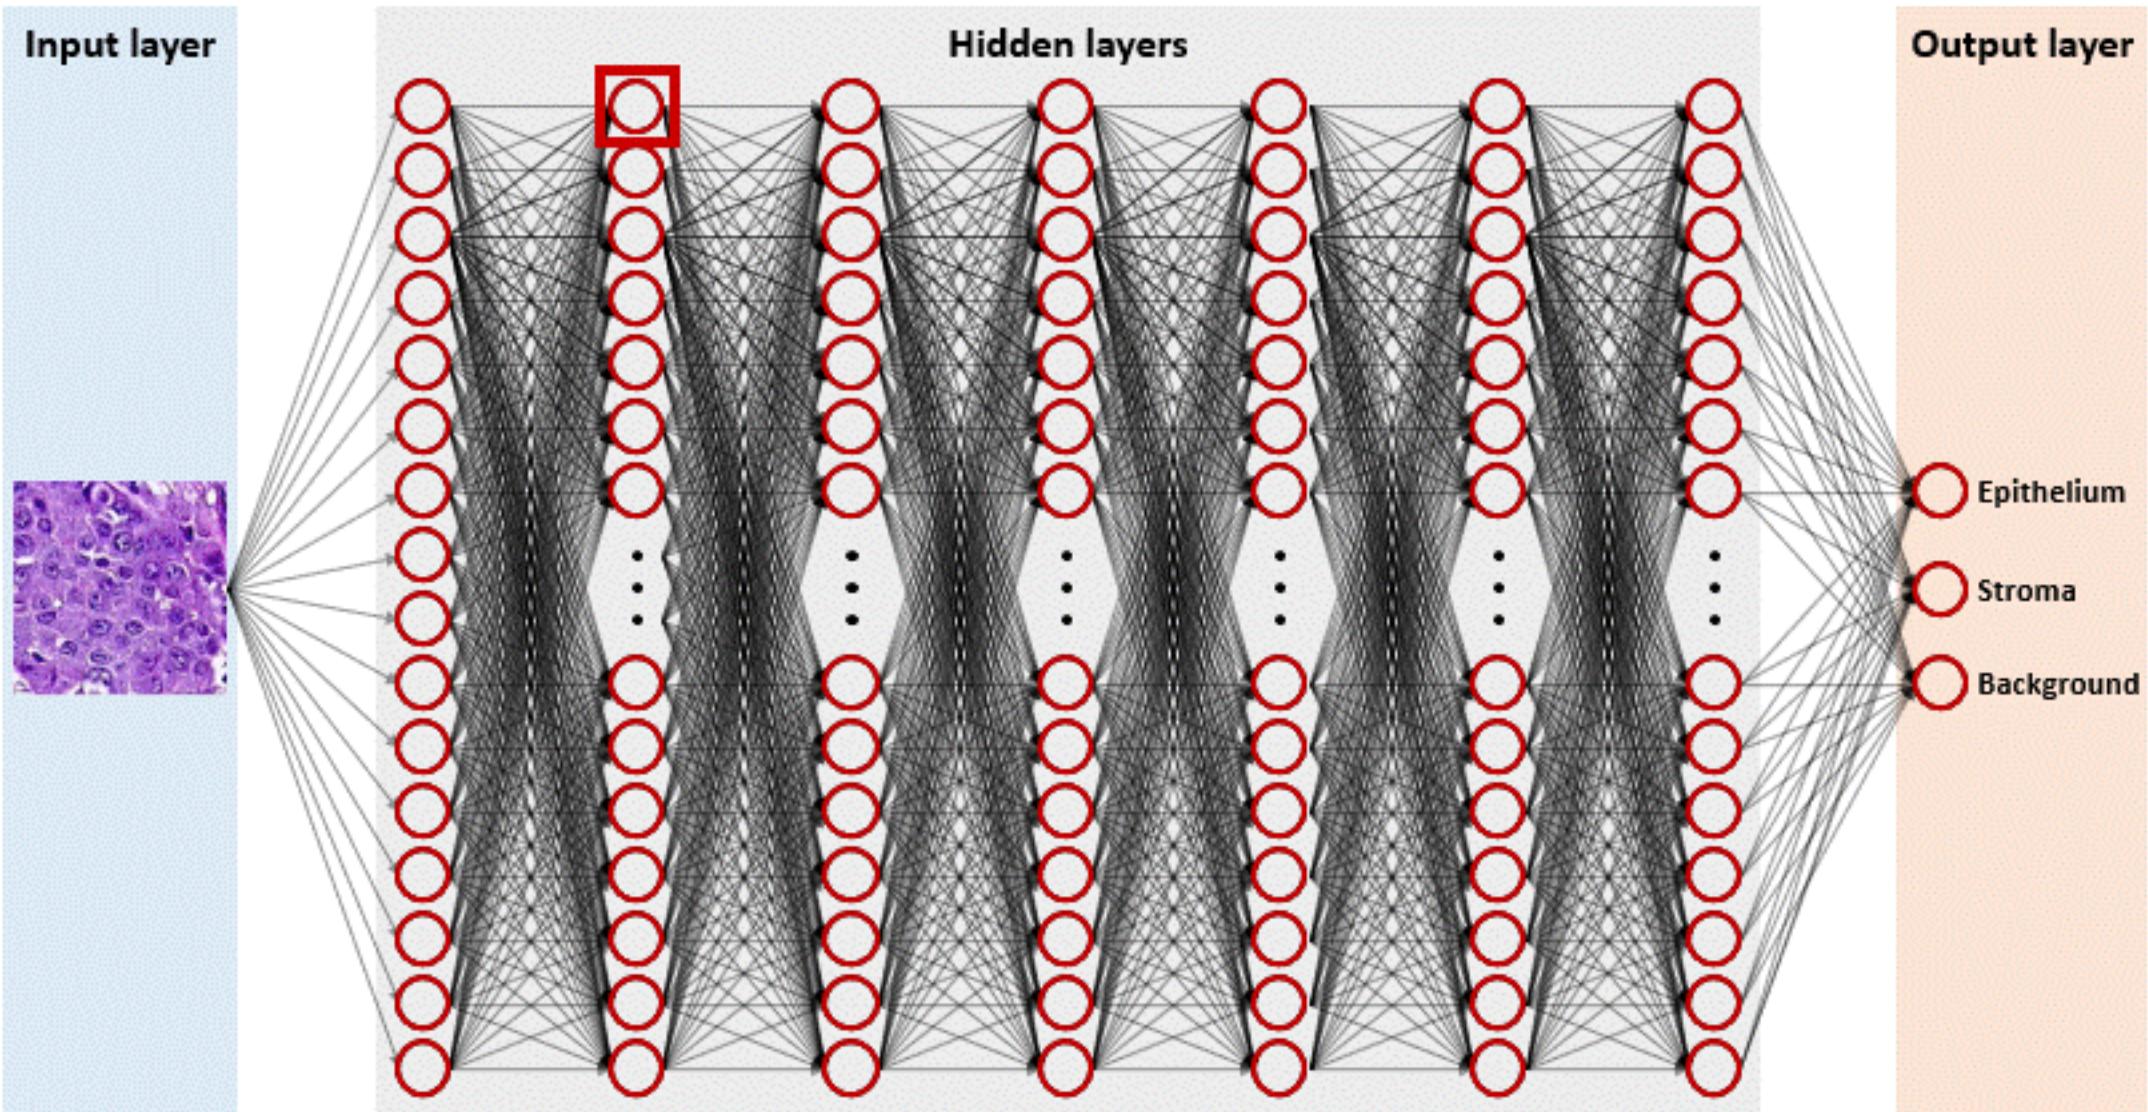
DEEP CONVOLUTIONAL NEURAL NETWORKS



DEEP CONVOLUTIONAL NEURAL NETWORKS

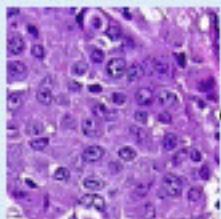


DEEP CONVOLUTIONAL NEURAL NETWORKS

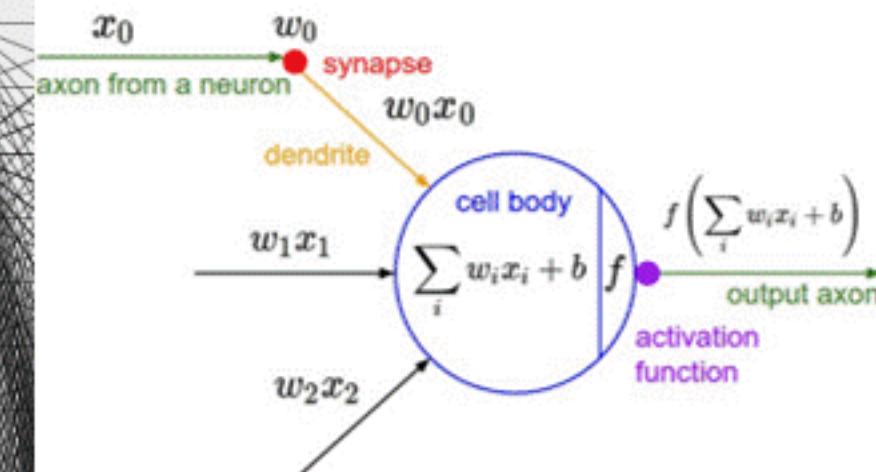


DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



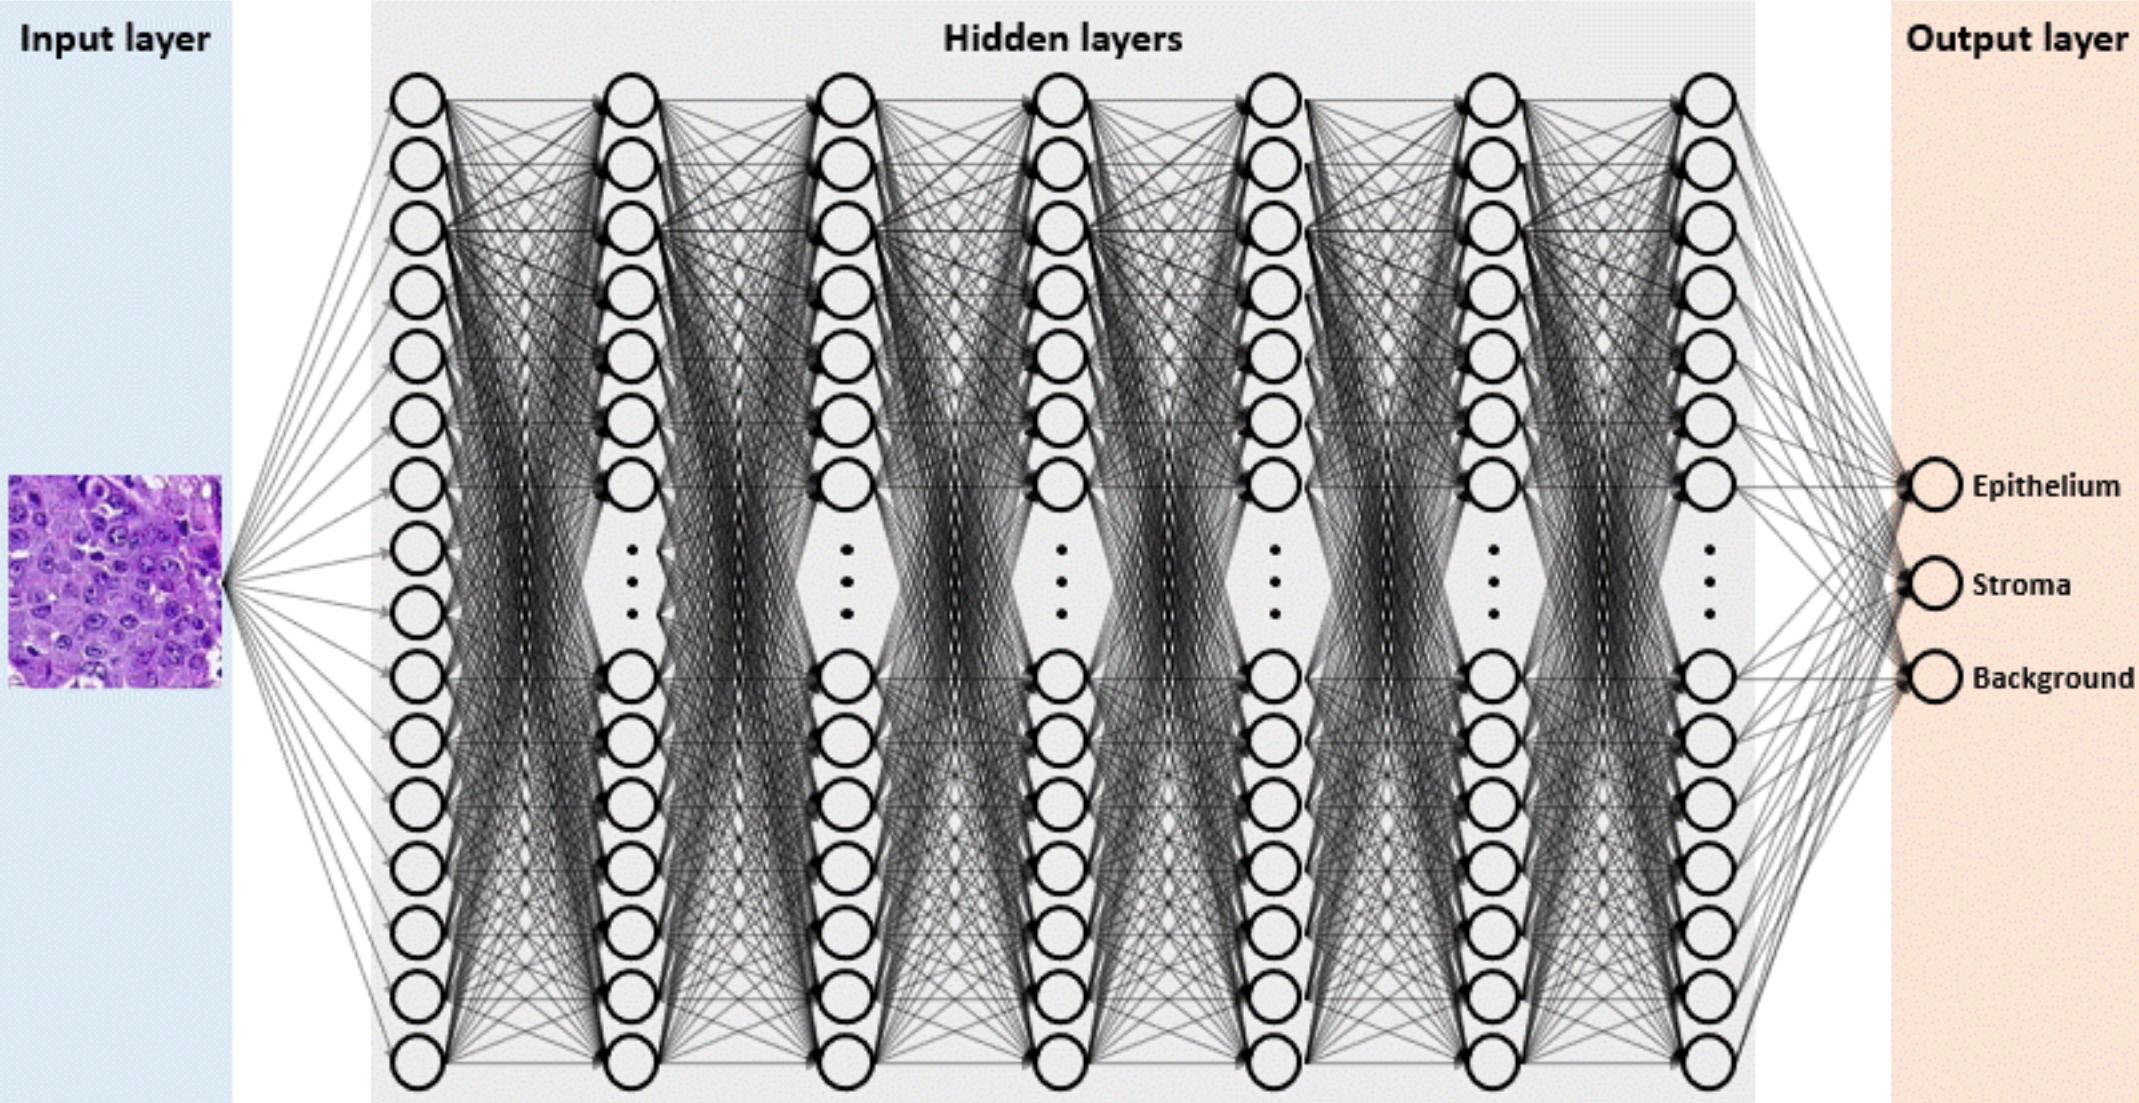
Hidden layers



Output layer

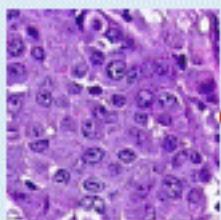
- Epithelium
- Stroma
- Background

DEEP CONVOLUTIONAL NEURAL NETWORKS

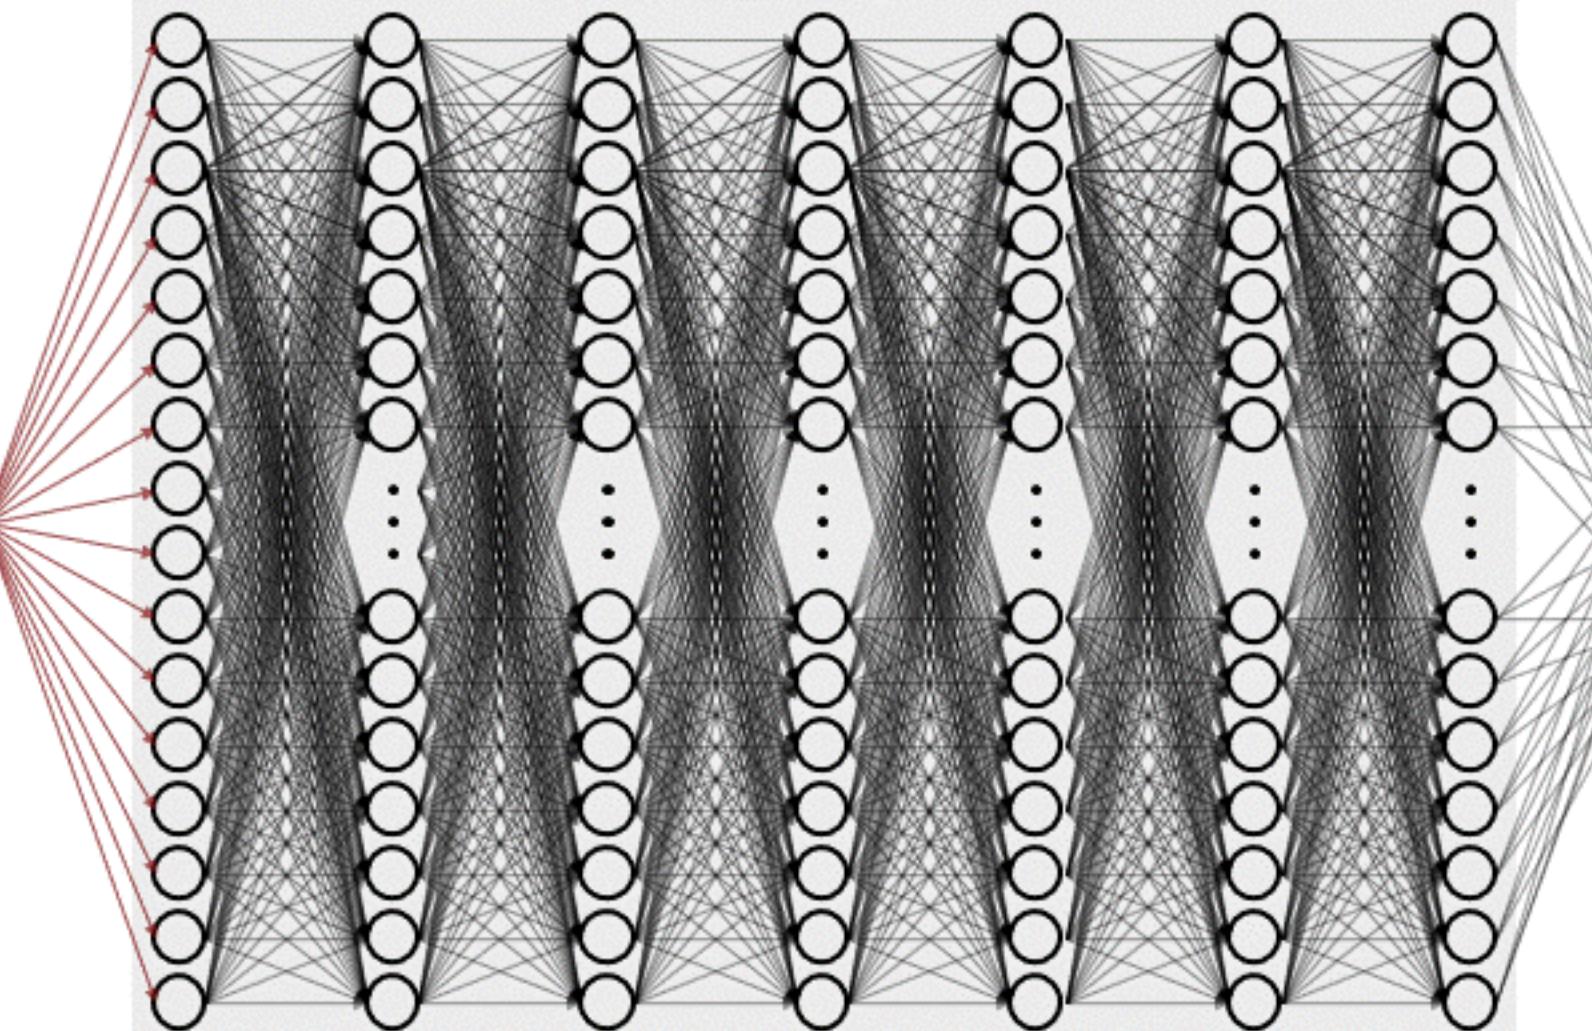


DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



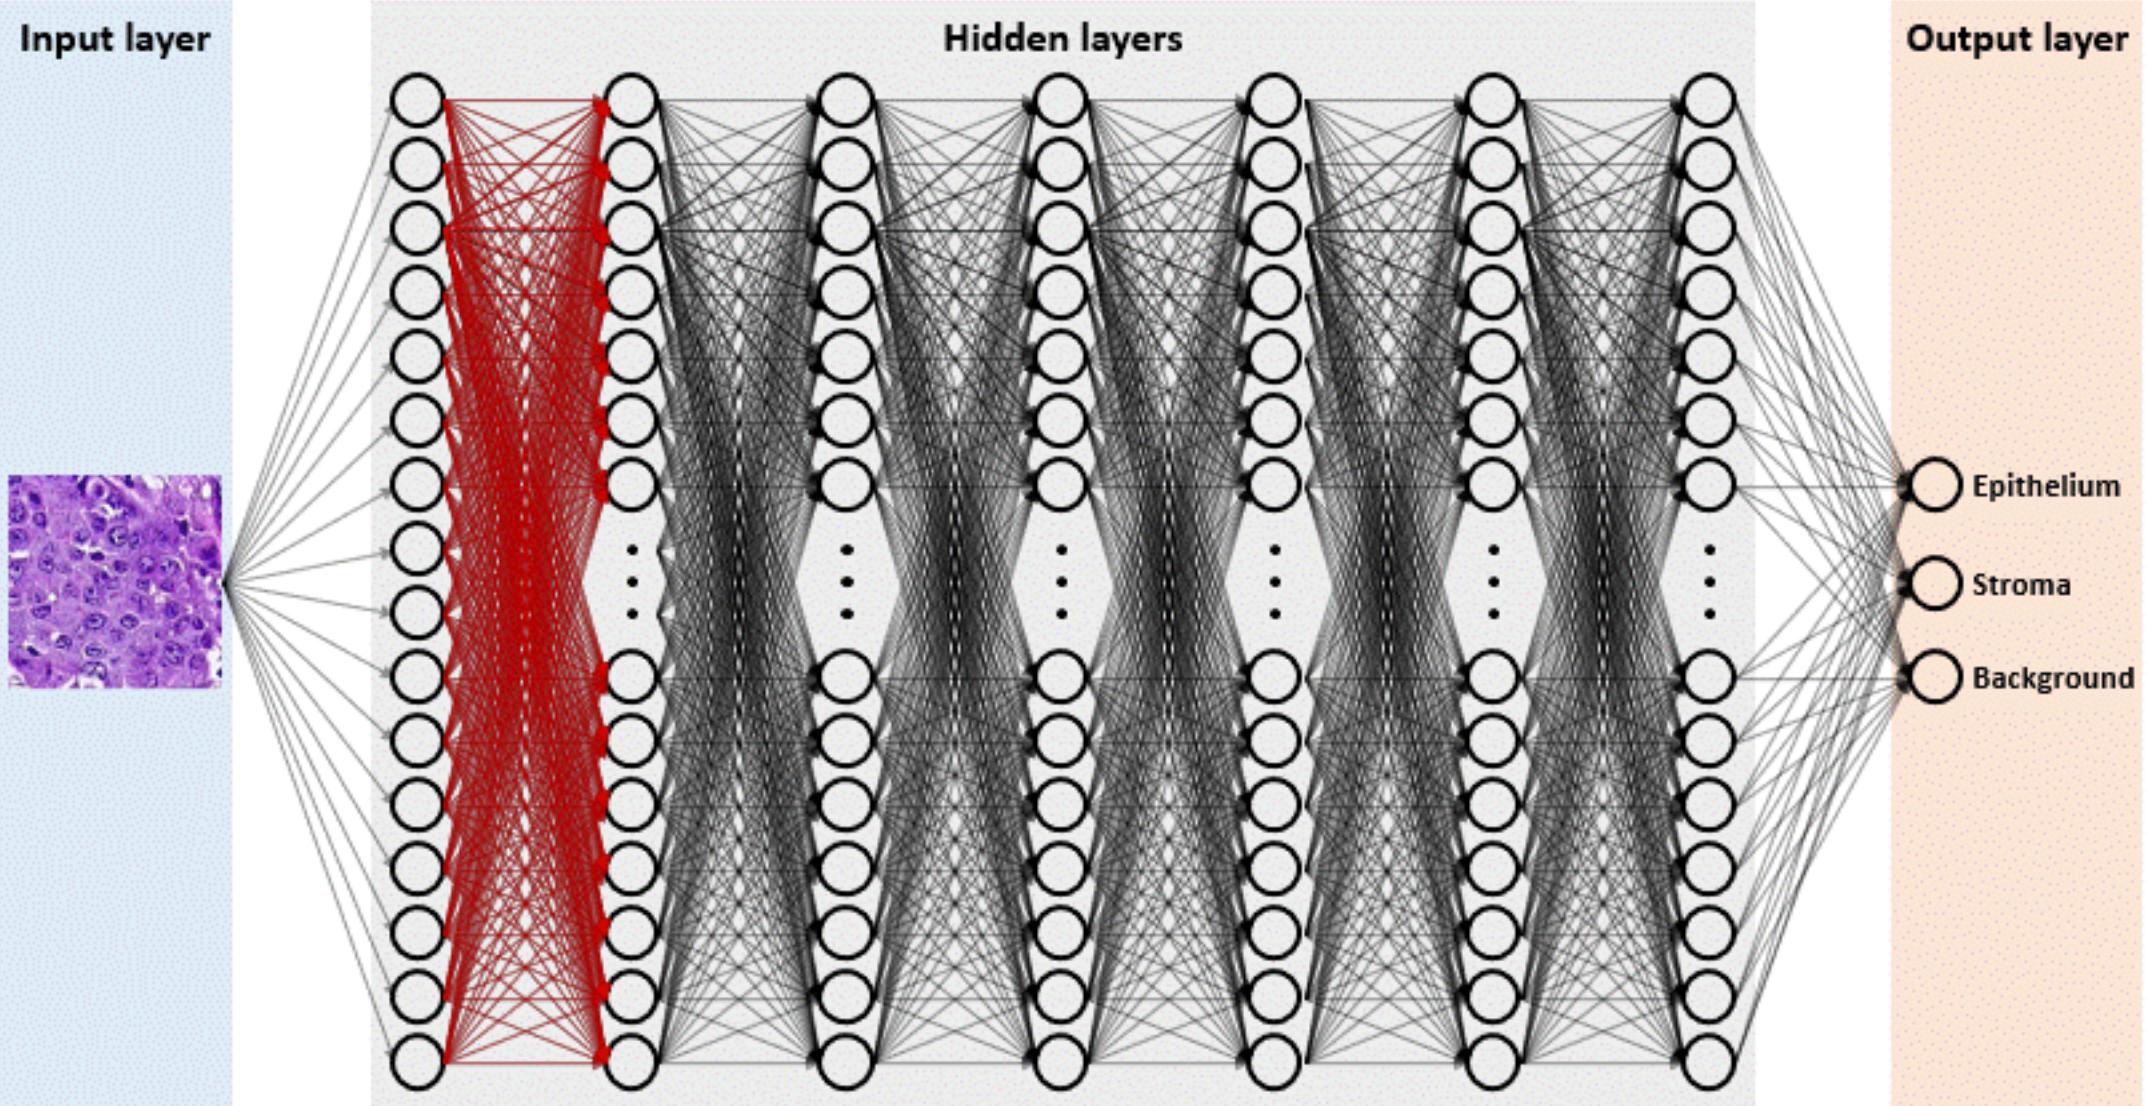
Hidden layers



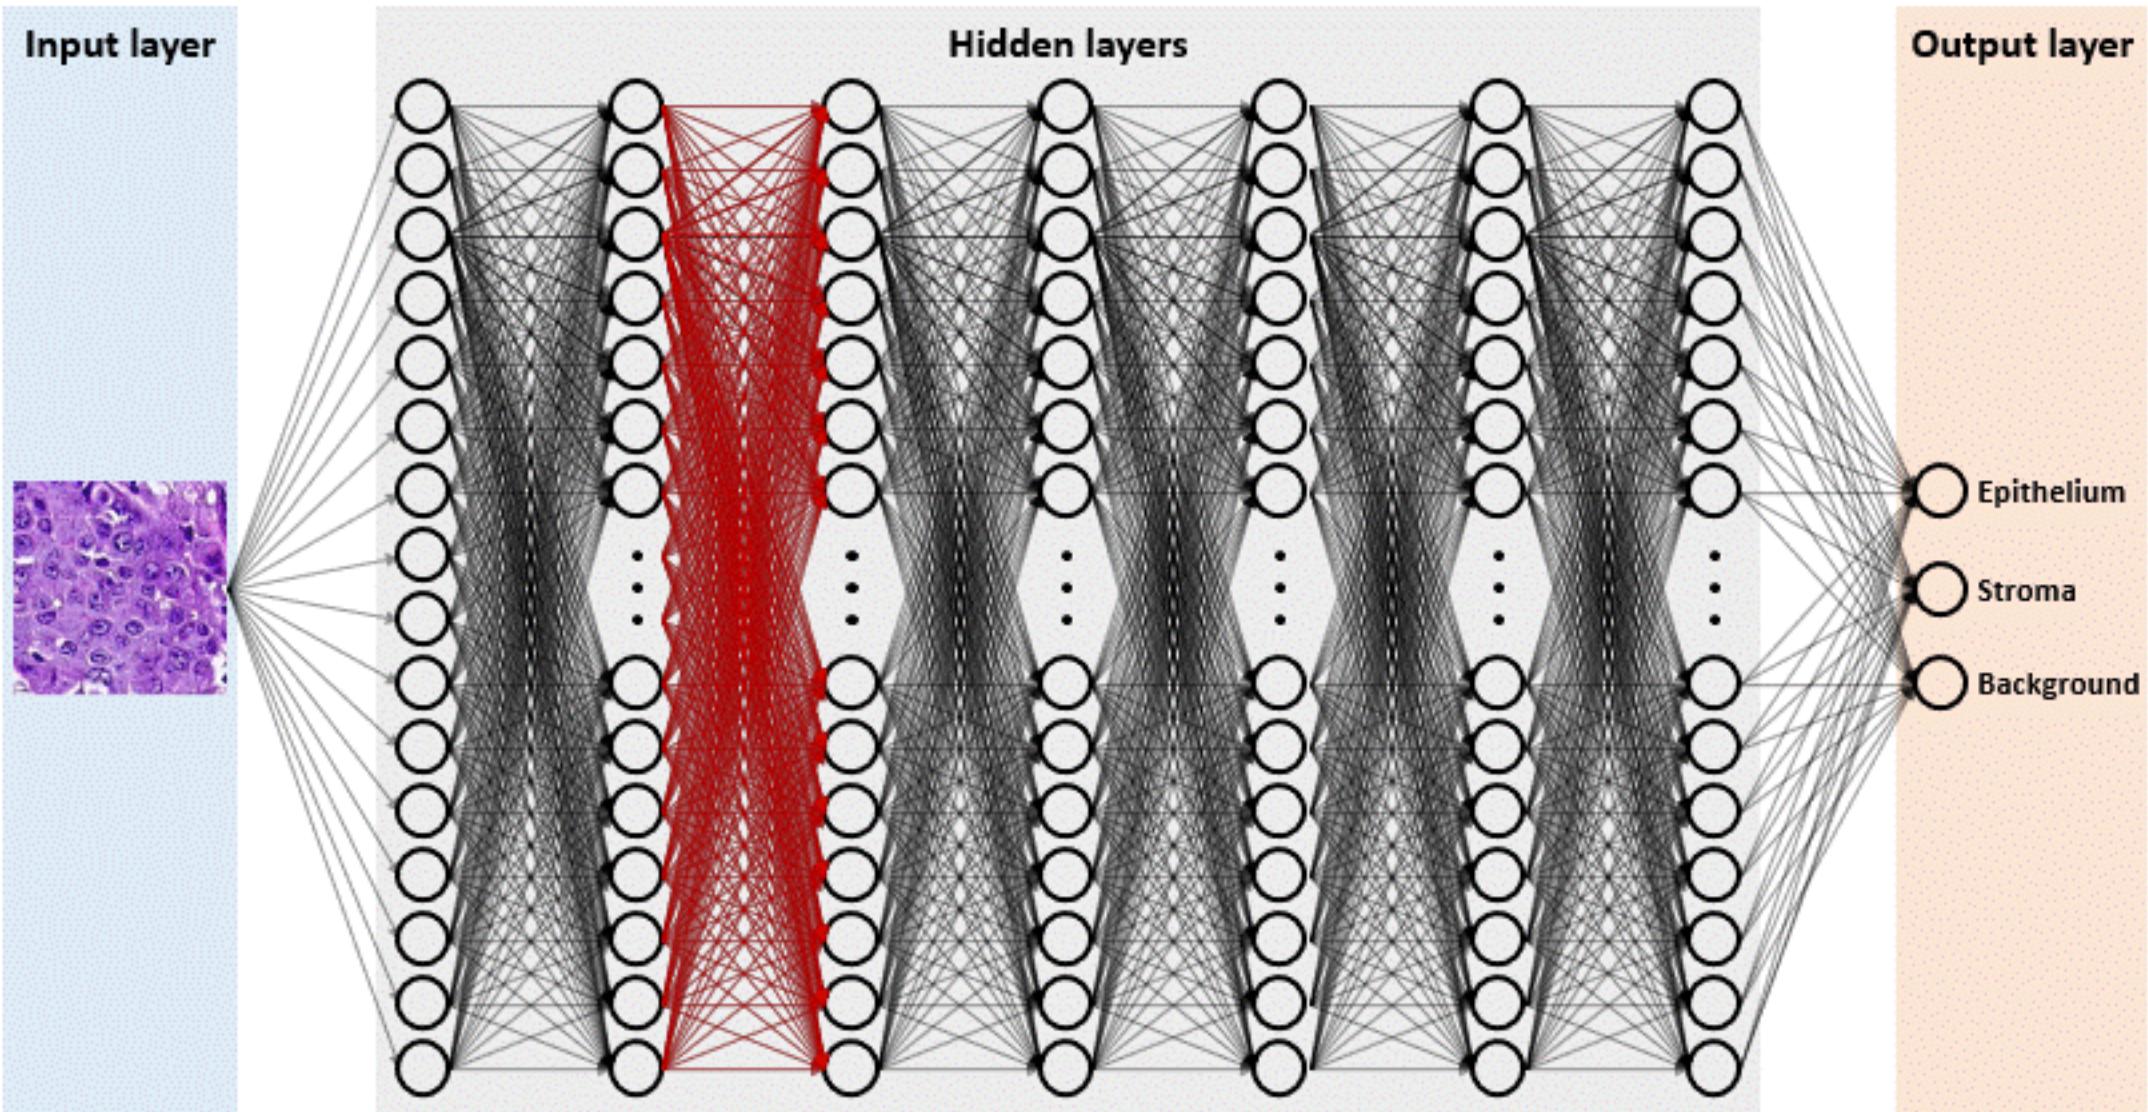
Output layer

- Epithelium
- Stroma
- Background

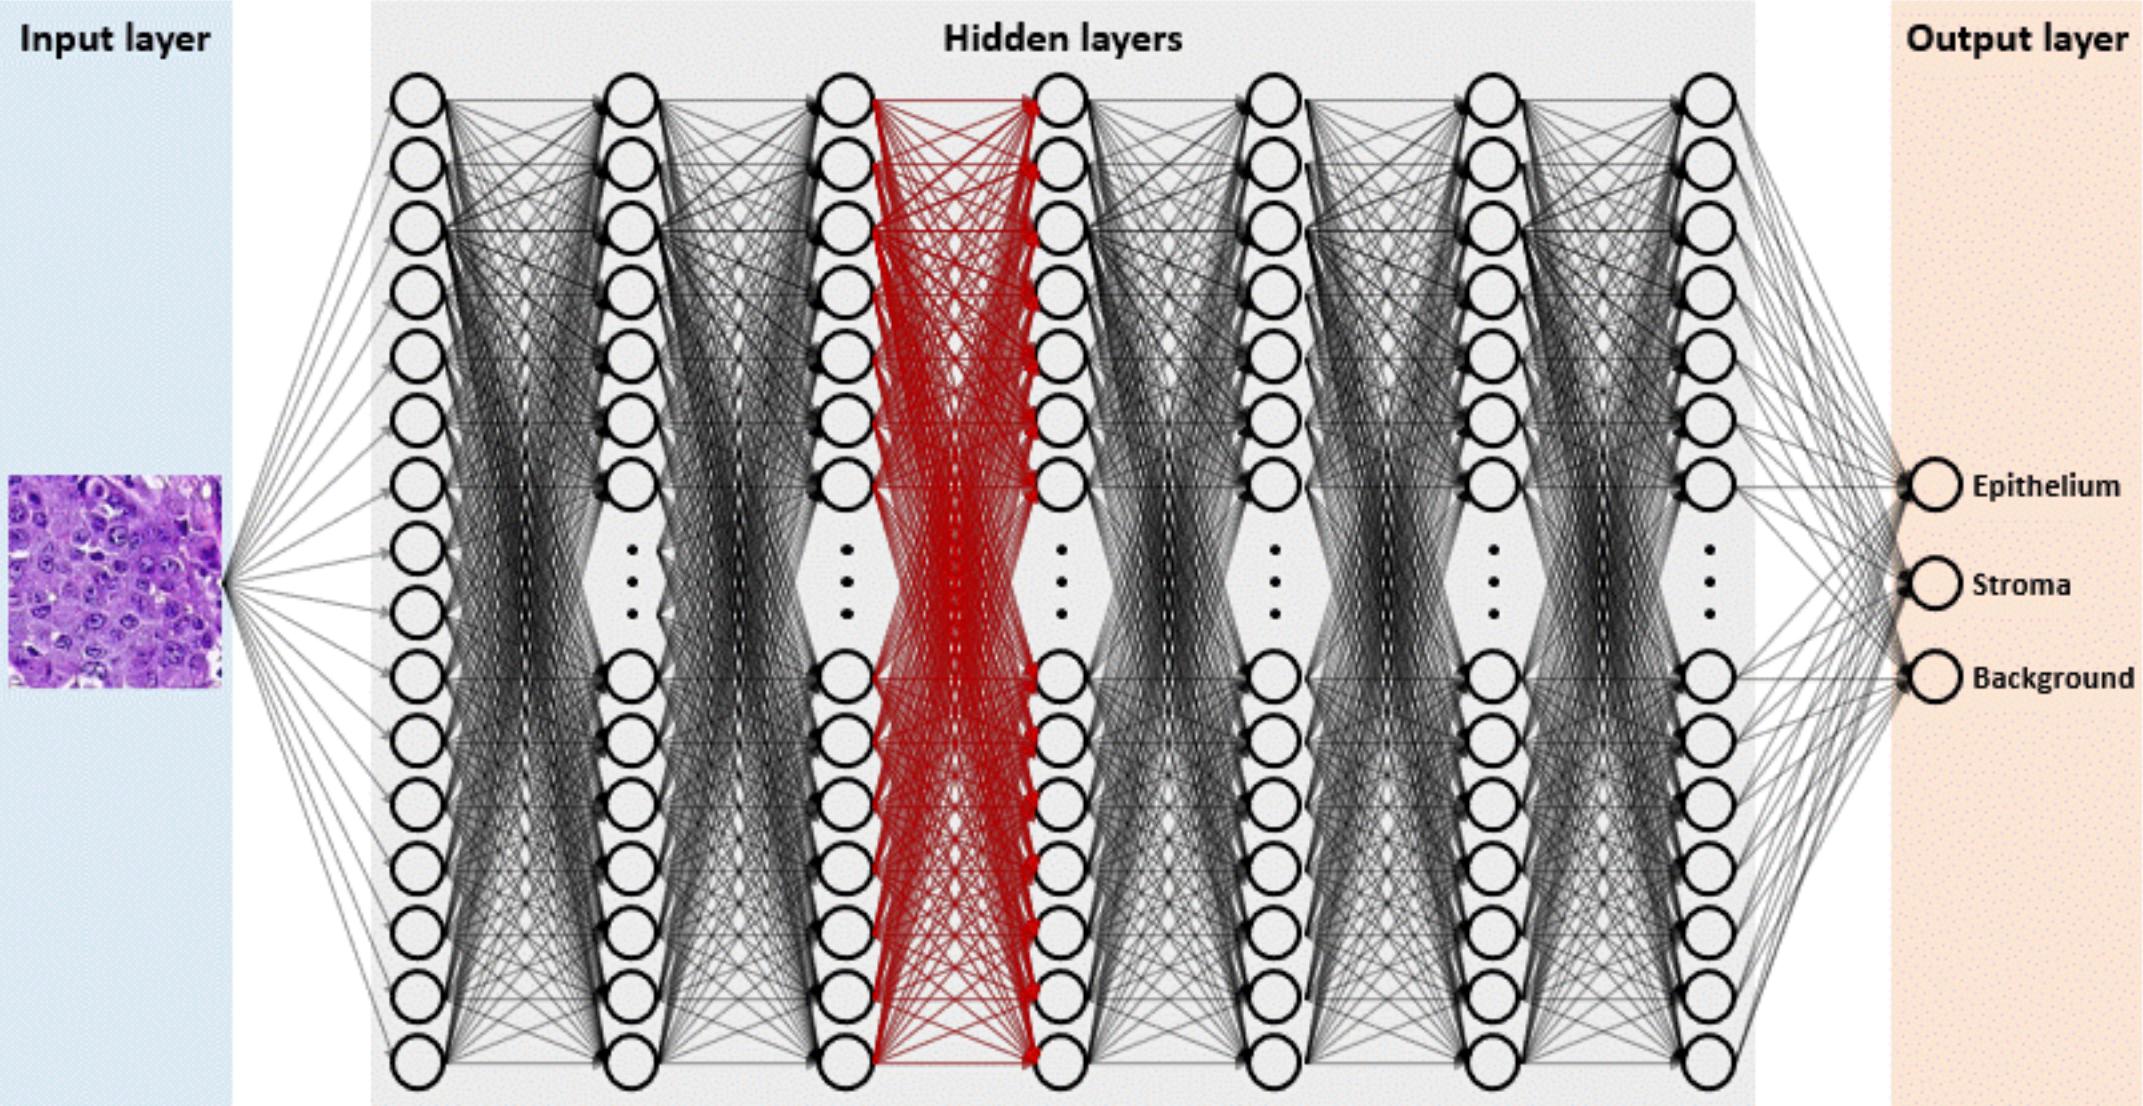
DEEP CONVOLUTIONAL NEURAL NETWORKS



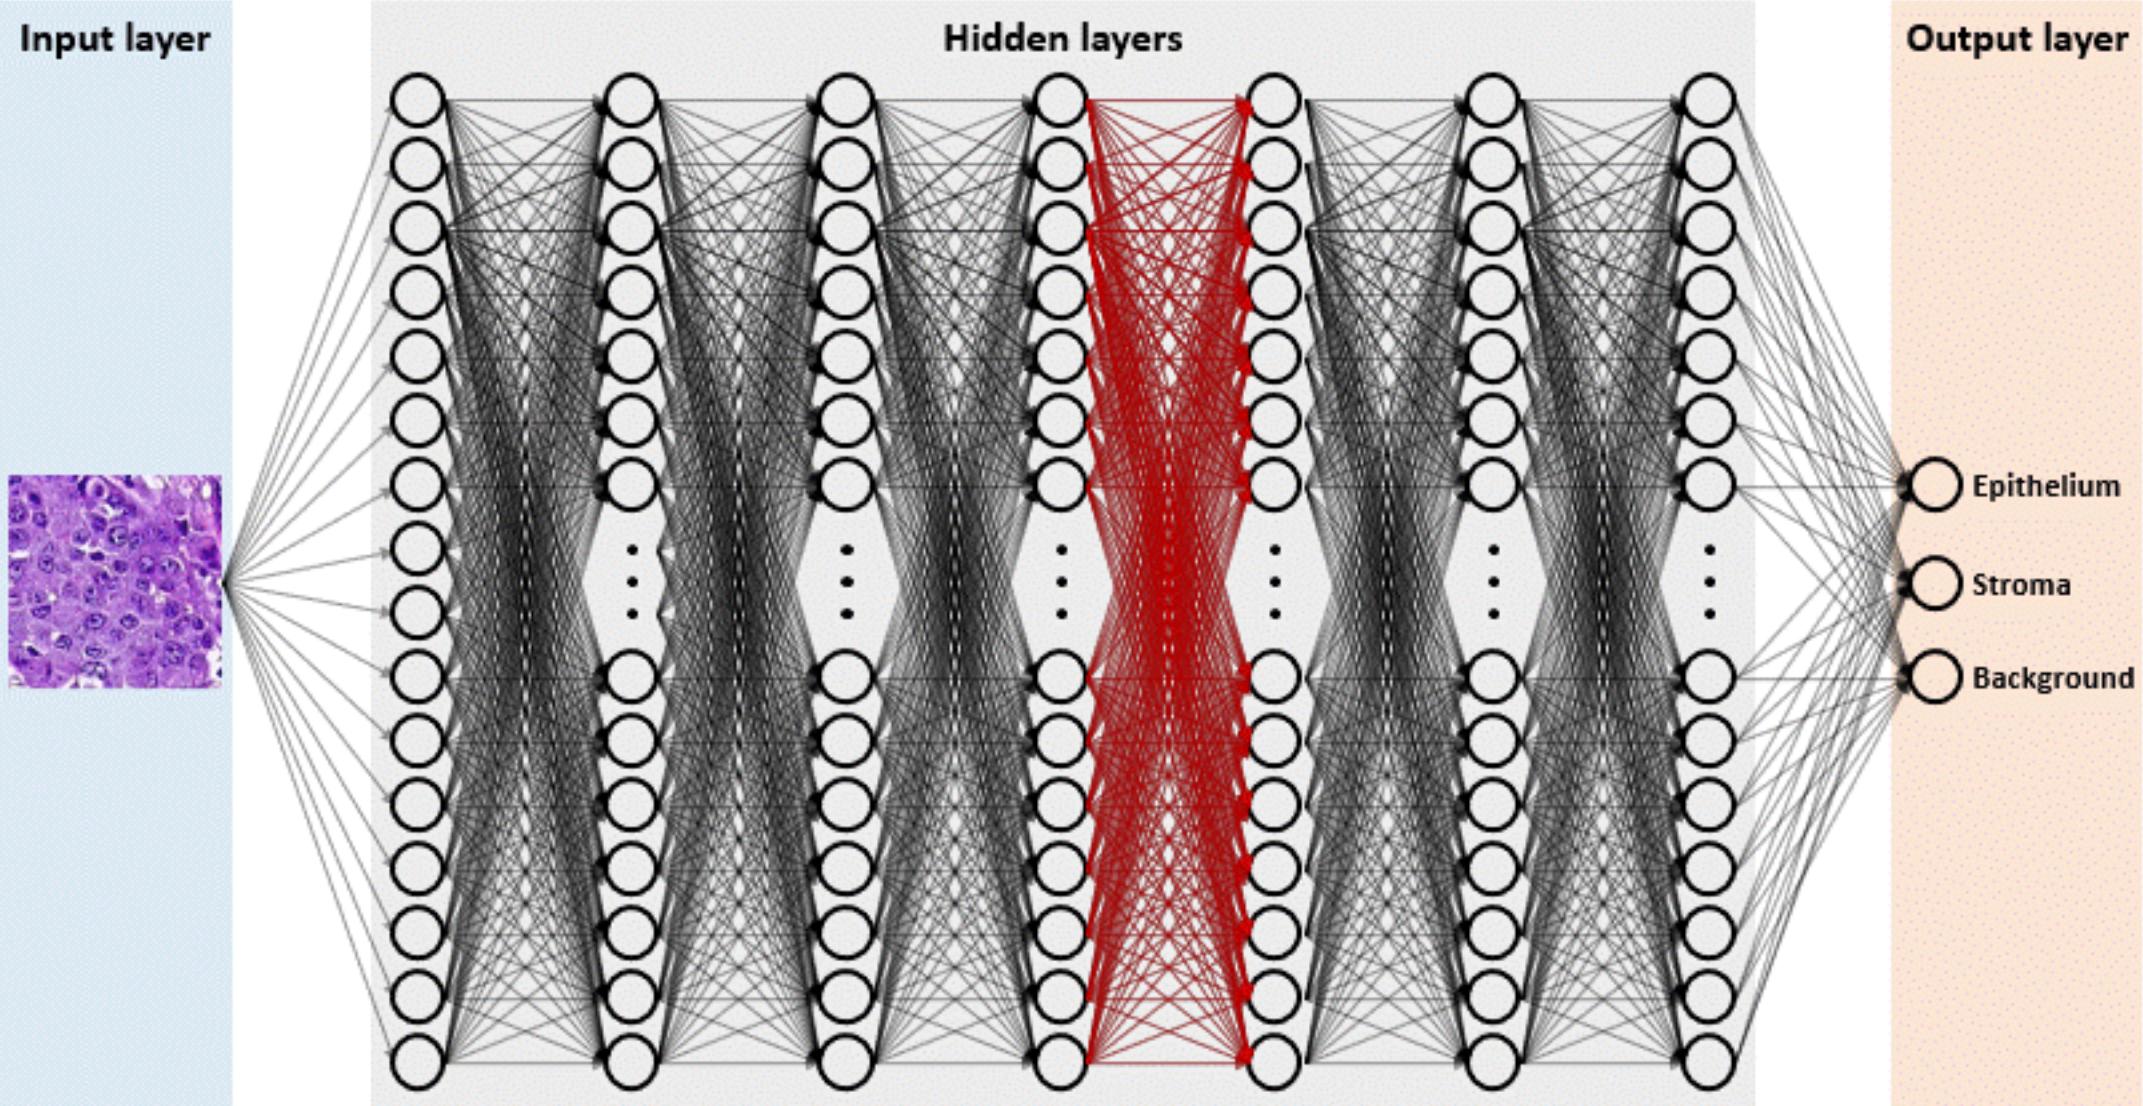
DEEP CONVOLUTIONAL NEURAL NETWORKS



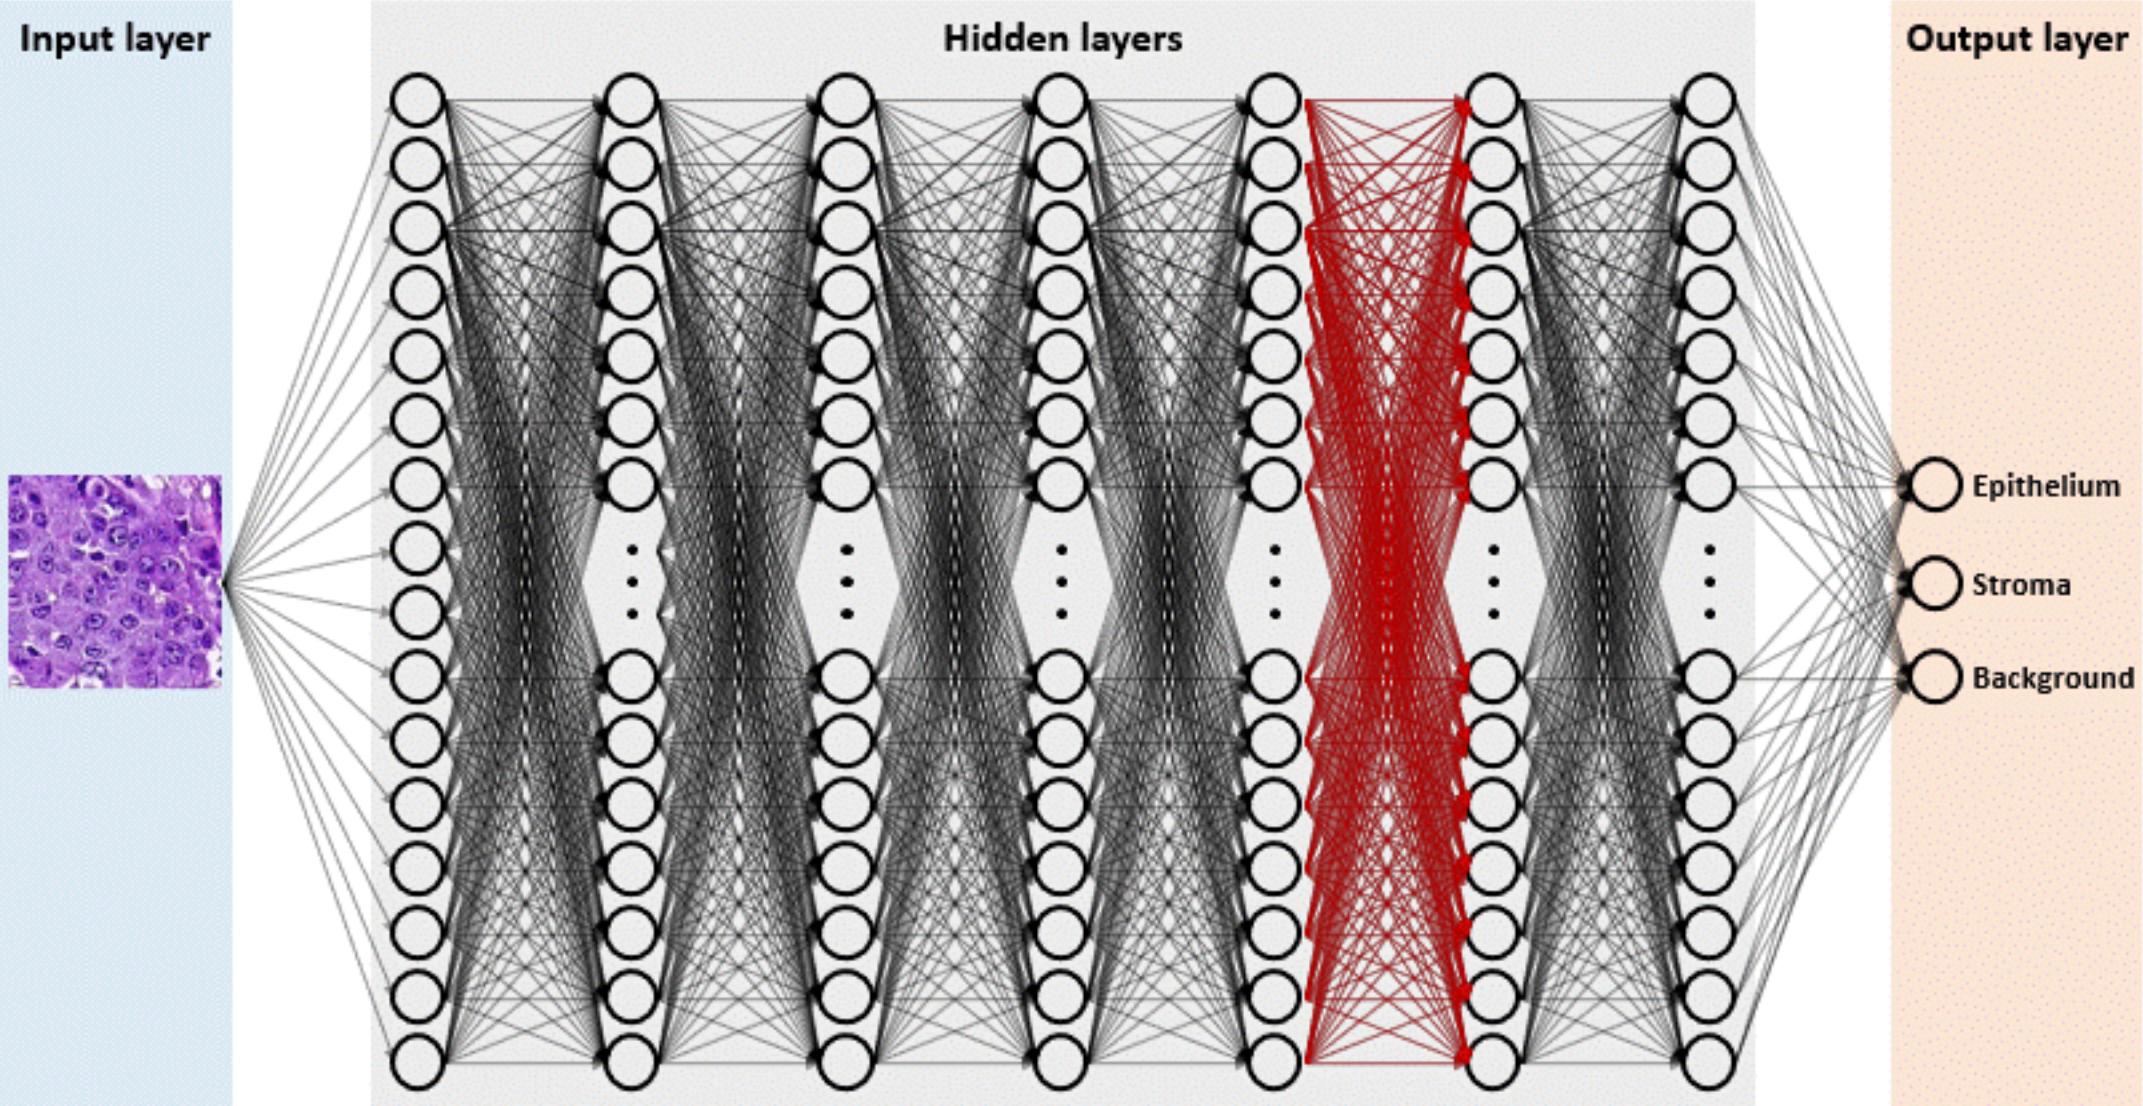
DEEP CONVOLUTIONAL NEURAL NETWORKS



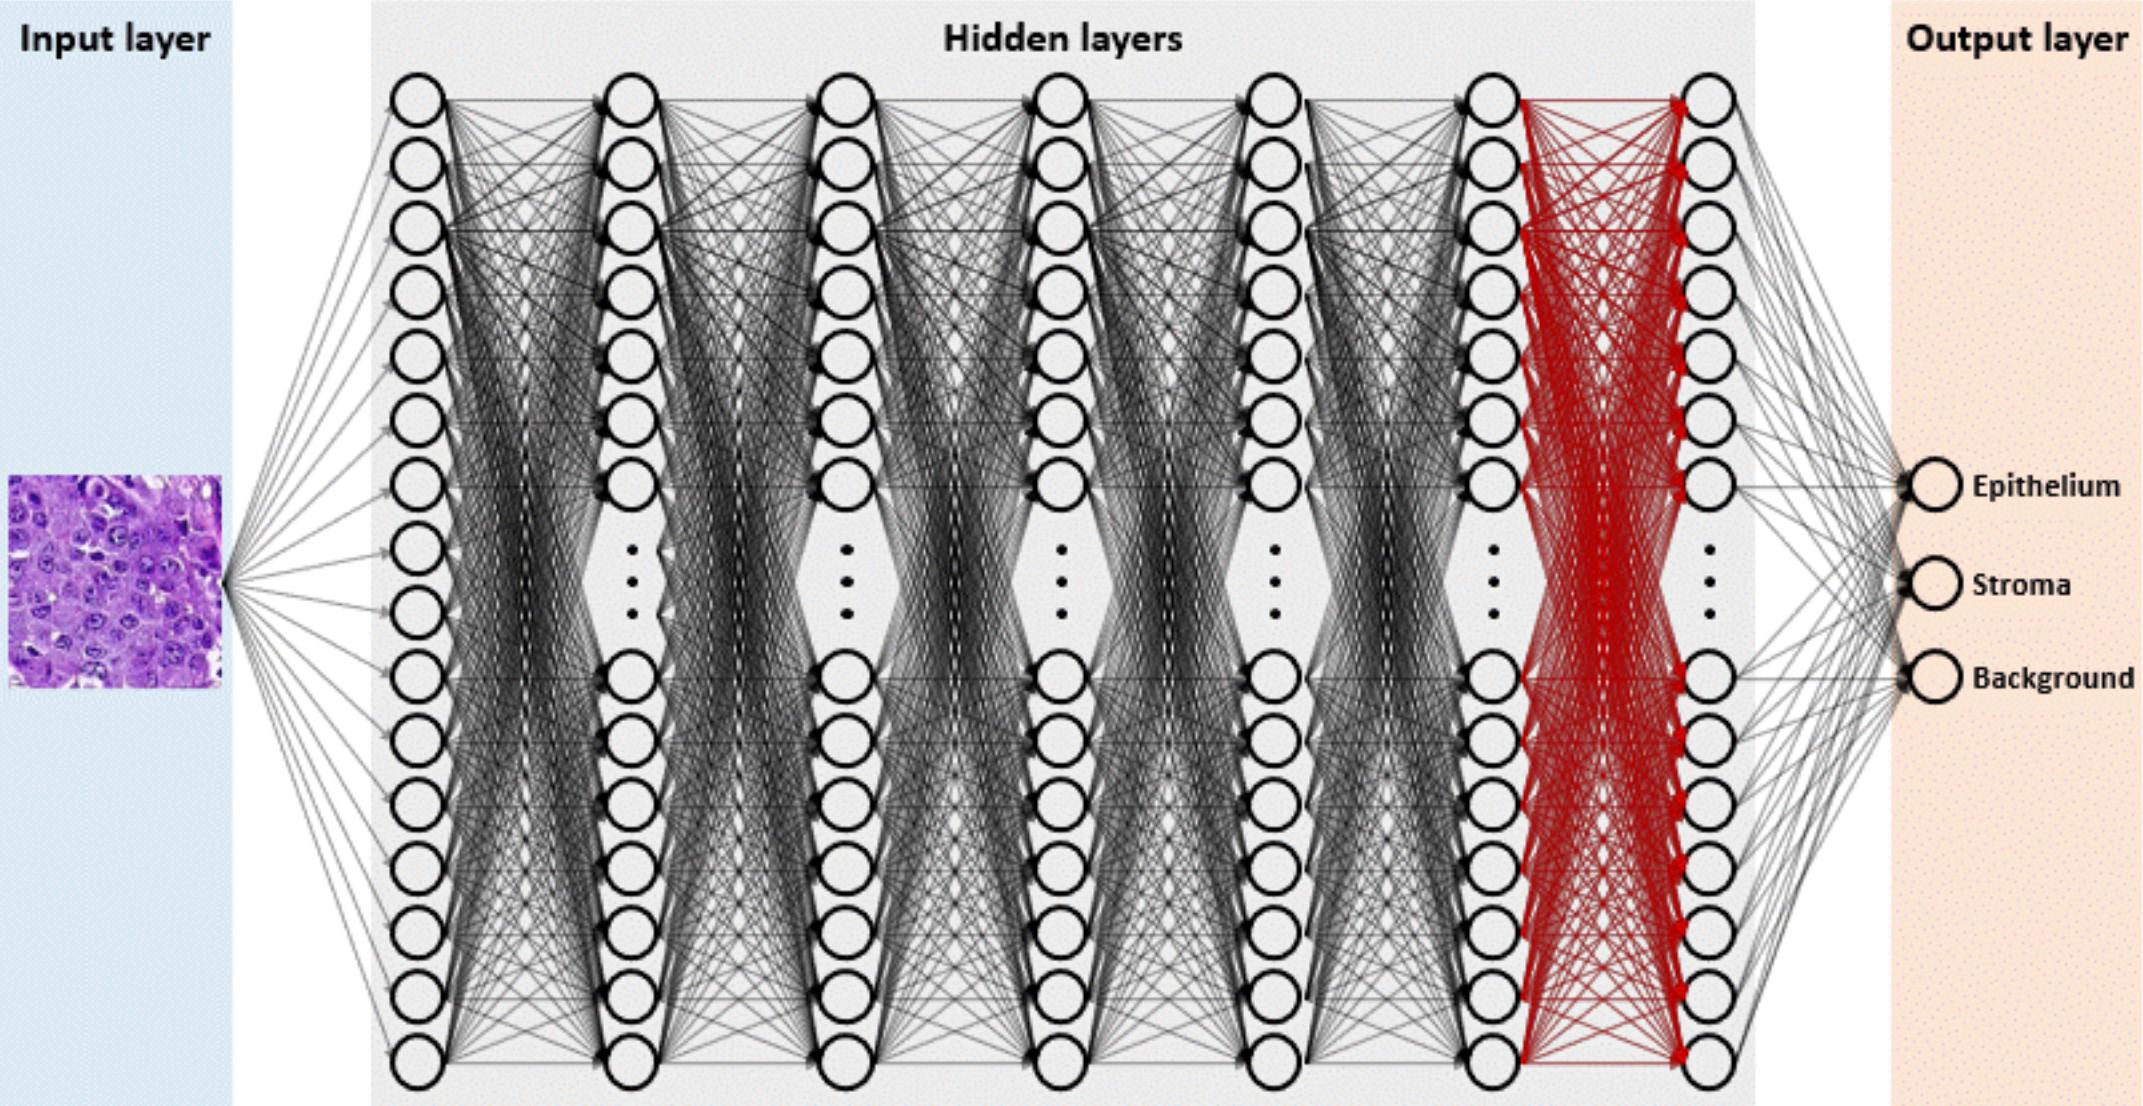
DEEP CONVOLUTIONAL NEURAL NETWORKS



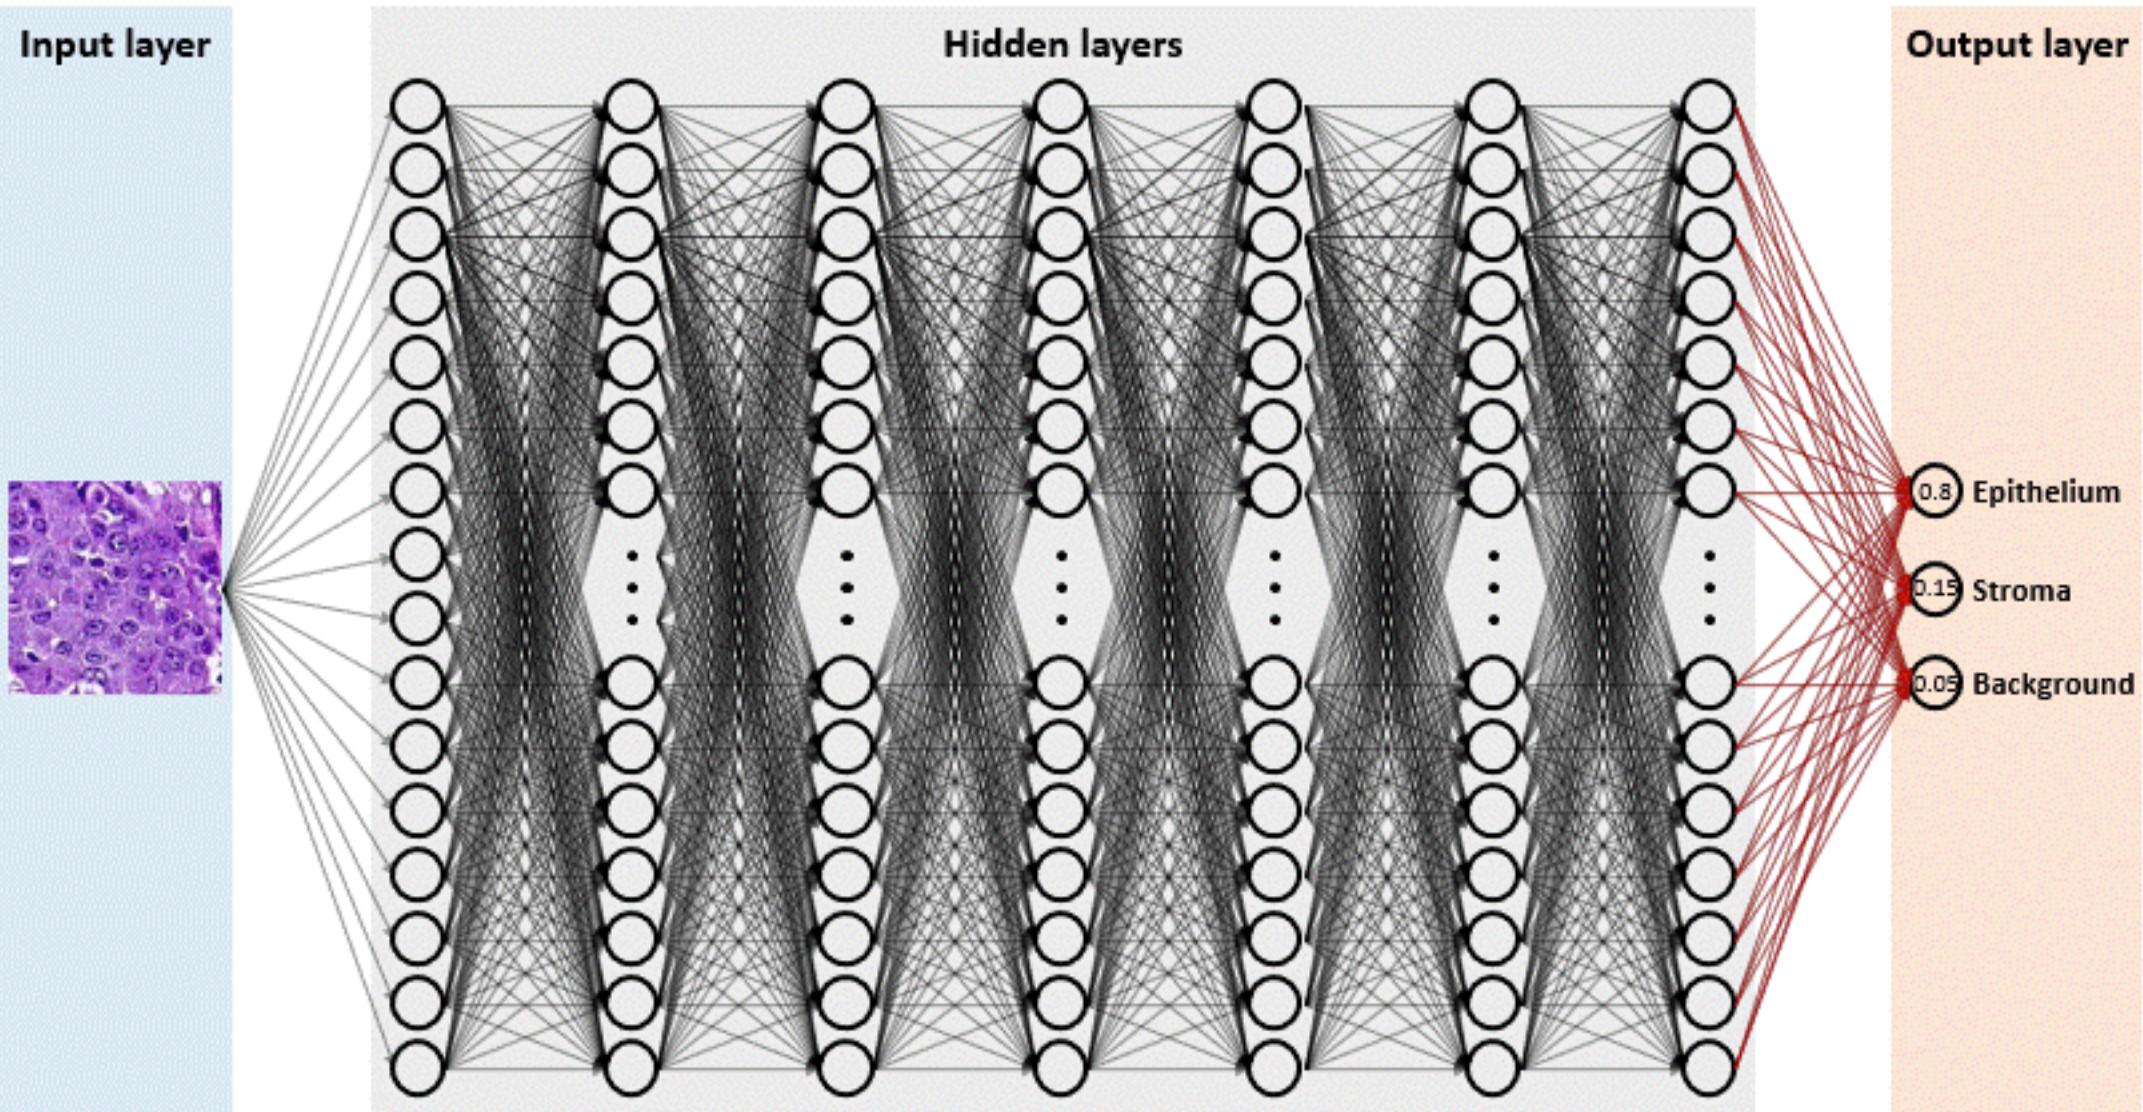
DEEP CONVOLUTIONAL NEURAL NETWORKS



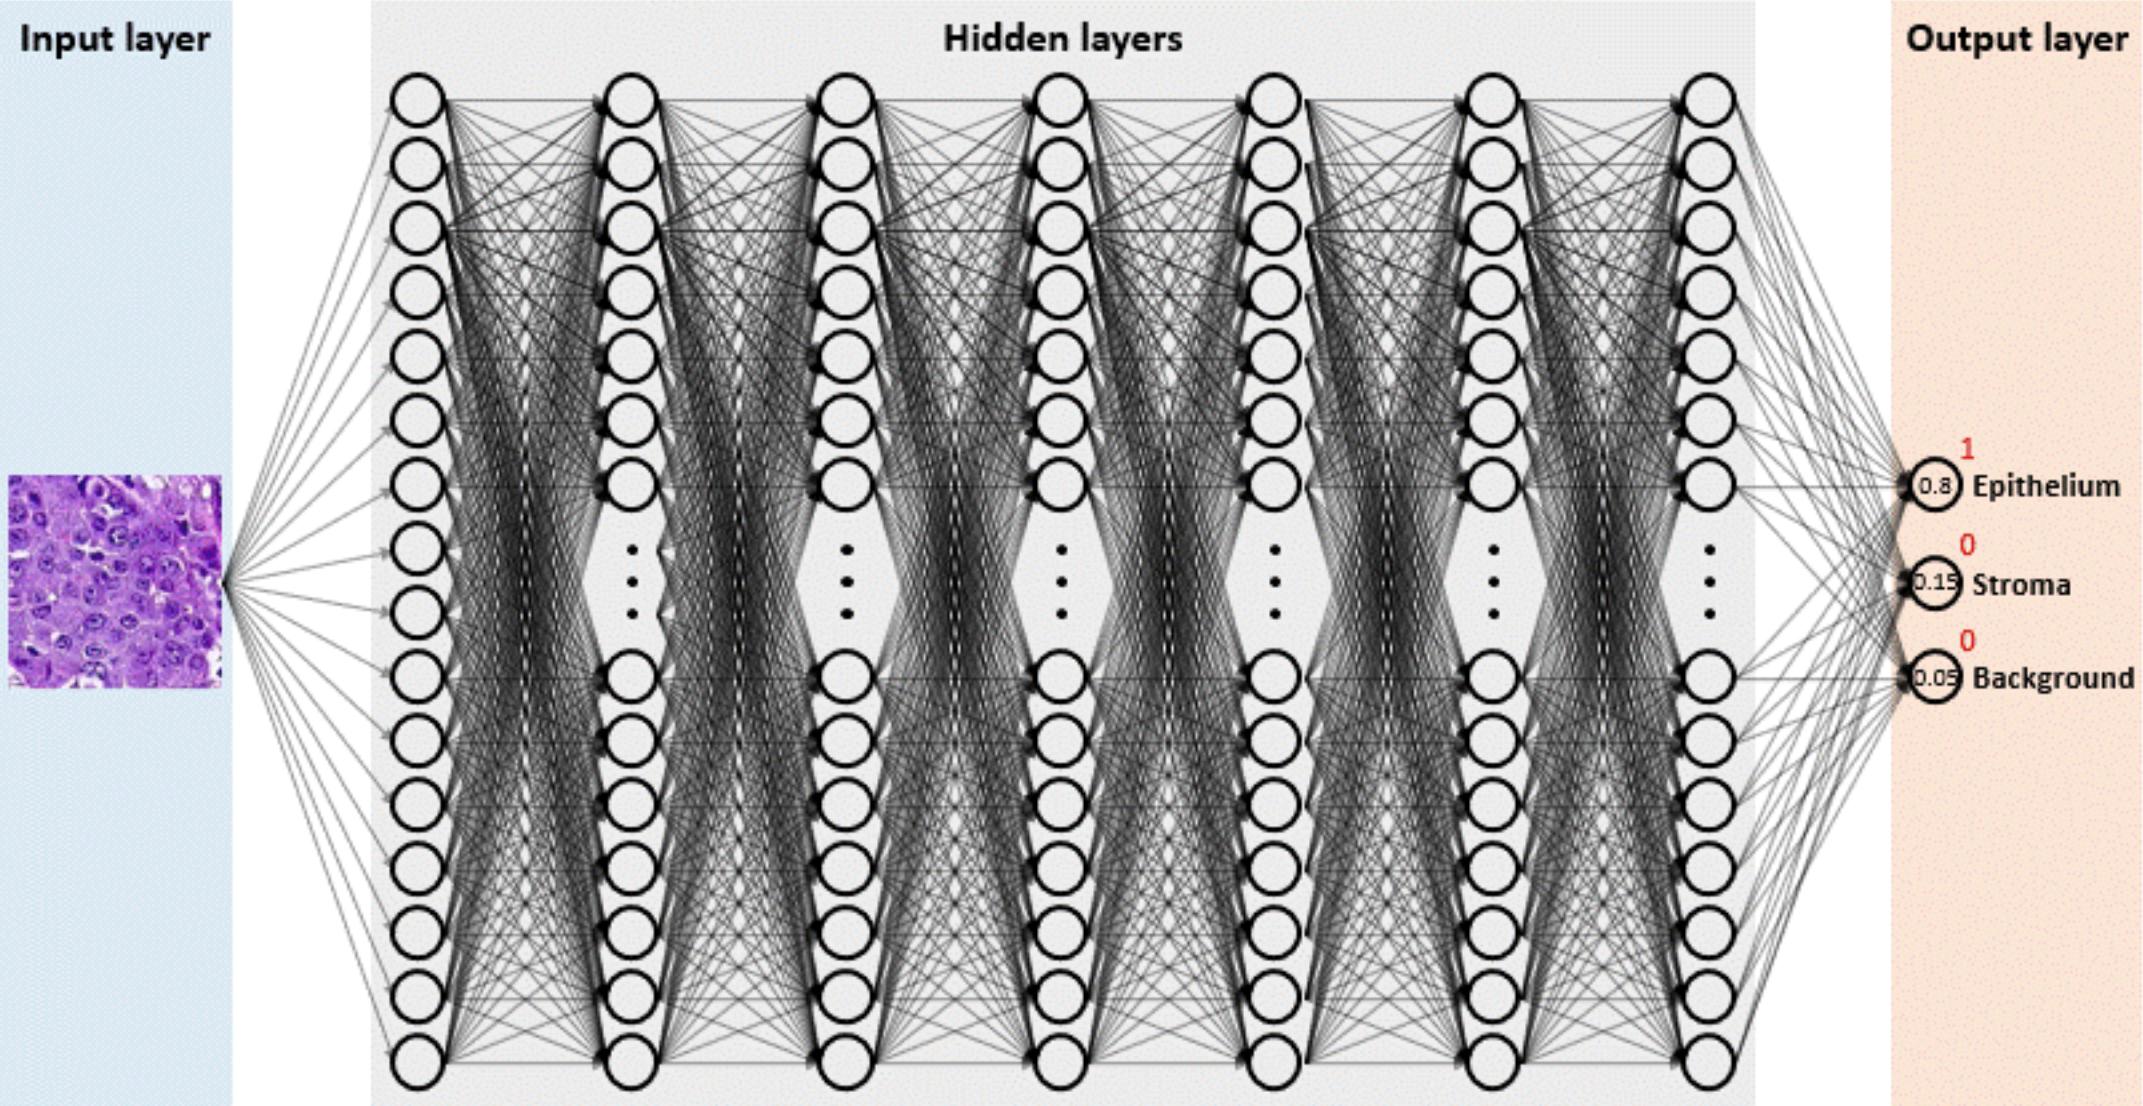
DEEP CONVOLUTIONAL NEURAL NETWORKS



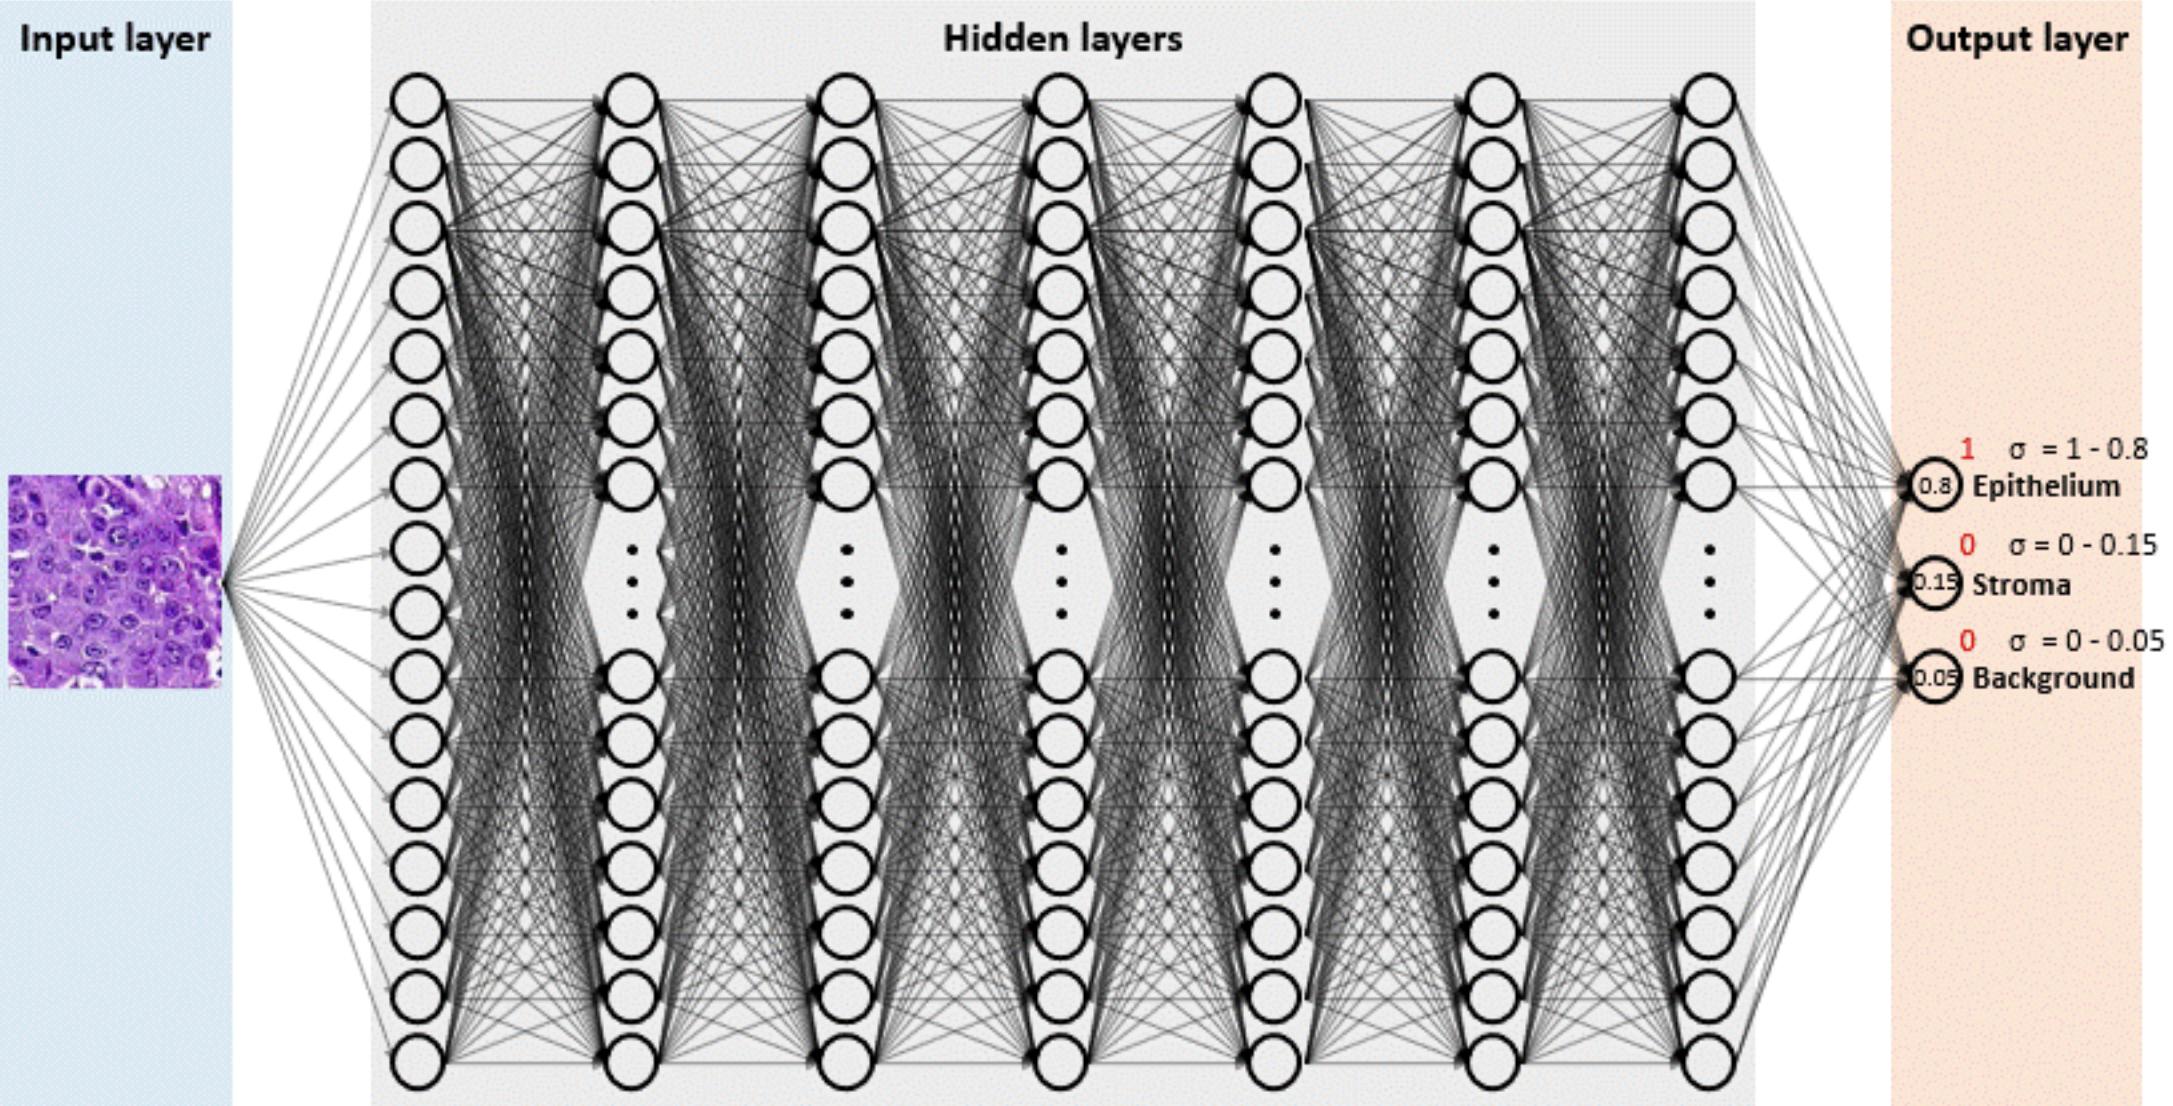
DEEP CONVOLUTIONAL NEURAL NETWORKS



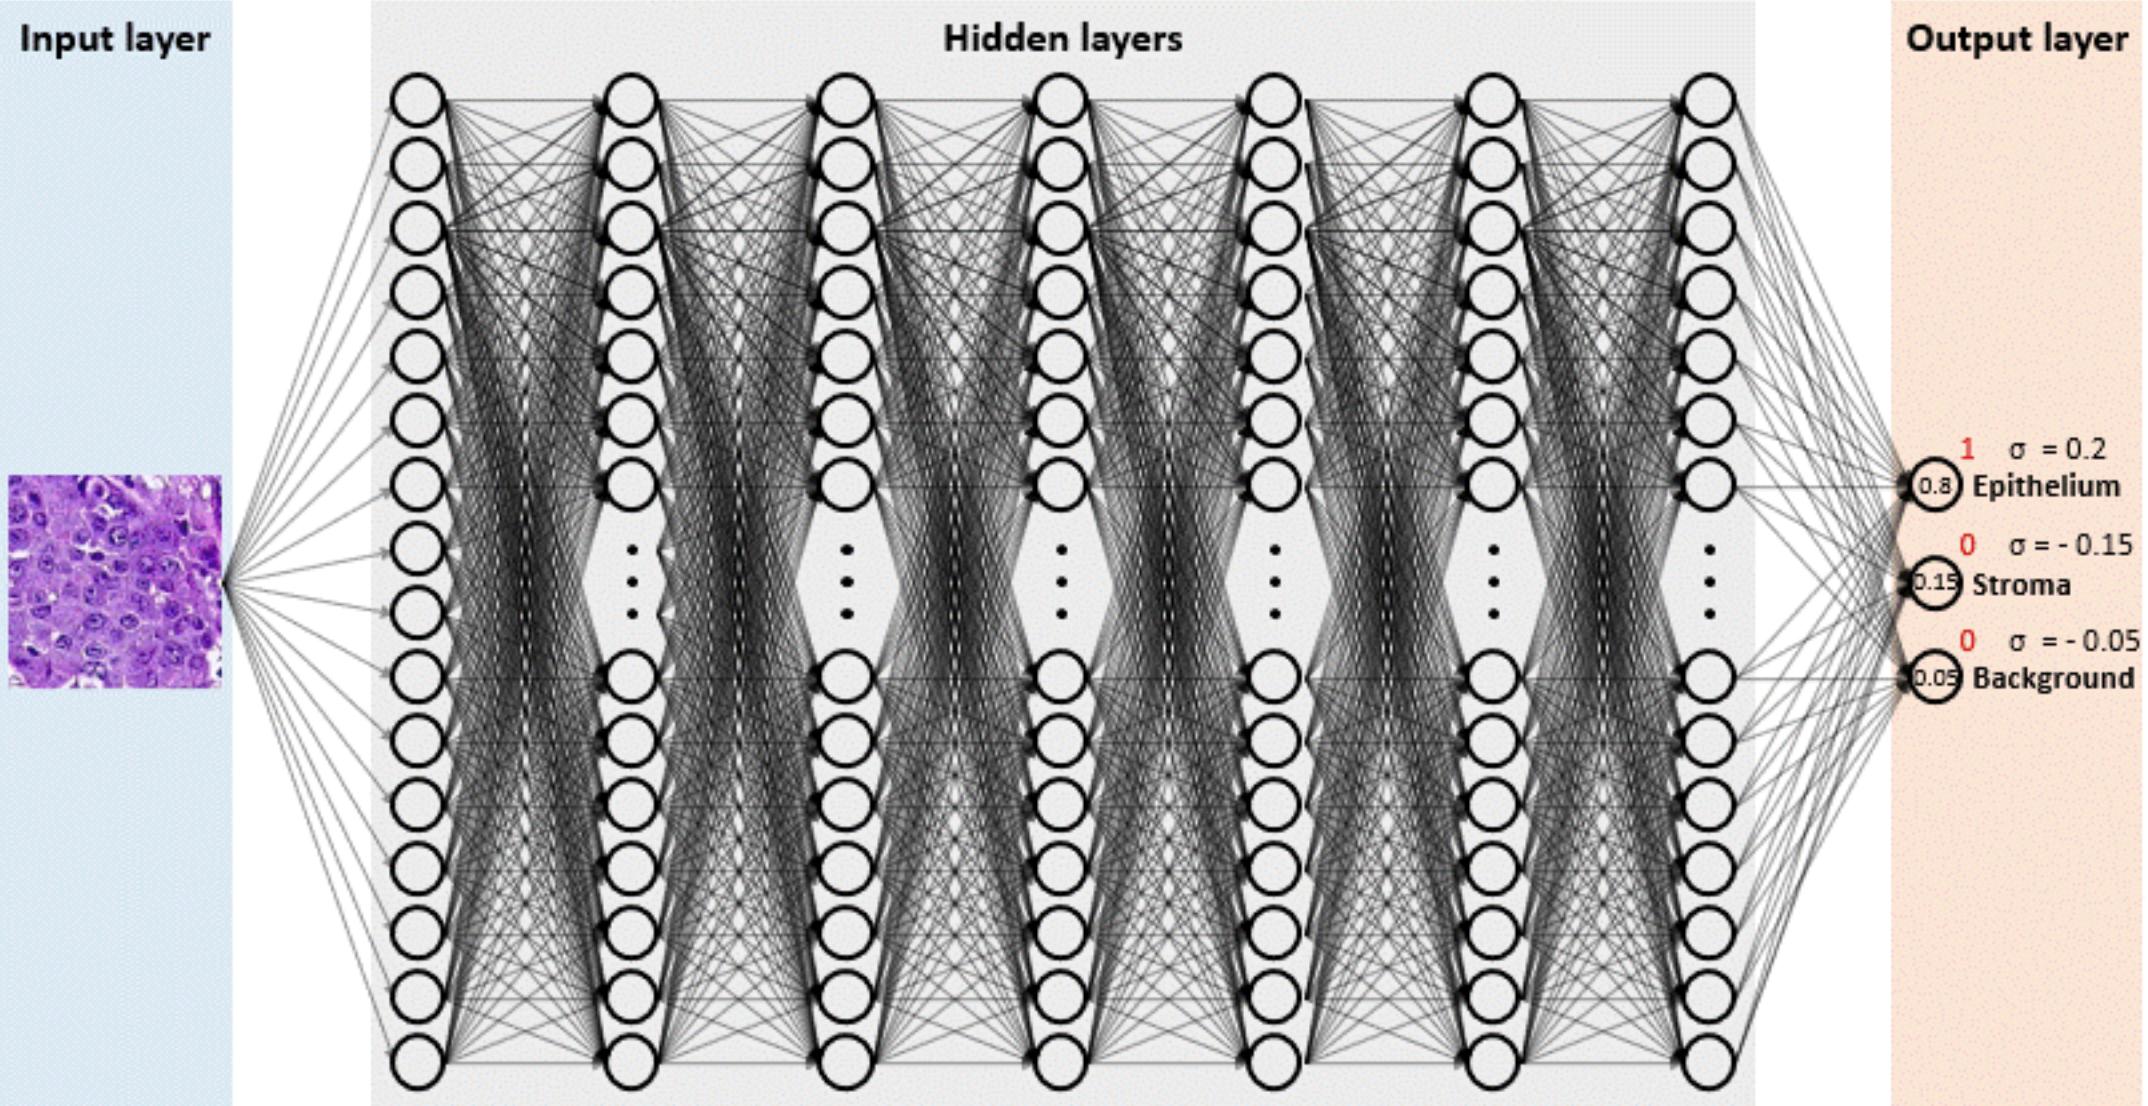
DEEP CONVOLUTIONAL NEURAL NETWORKS



DEEP CONVOLUTIONAL NEURAL NETWORKS

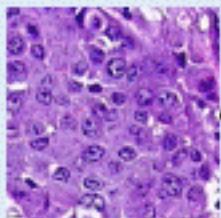


DEEP CONVOLUTIONAL NEURAL NETWORKS

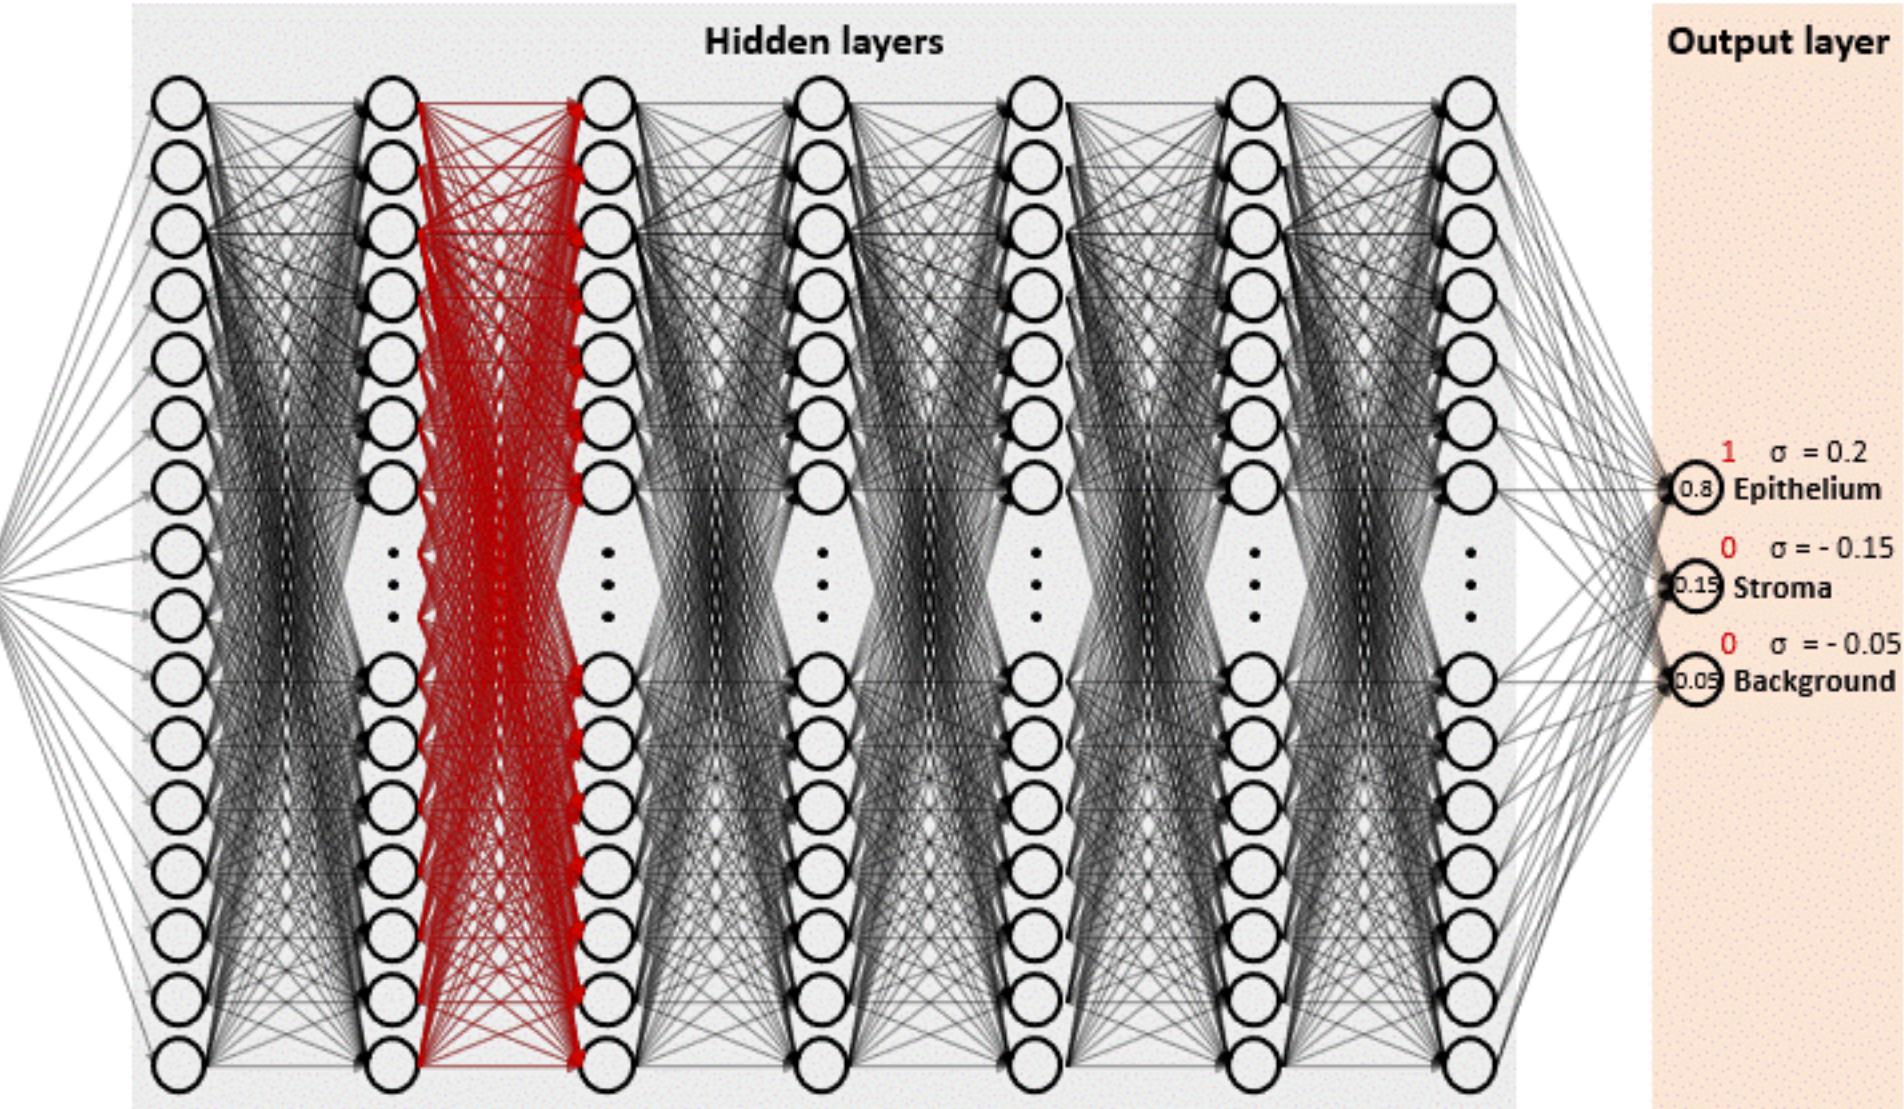


DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



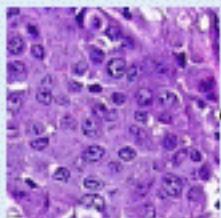
Hidden layers



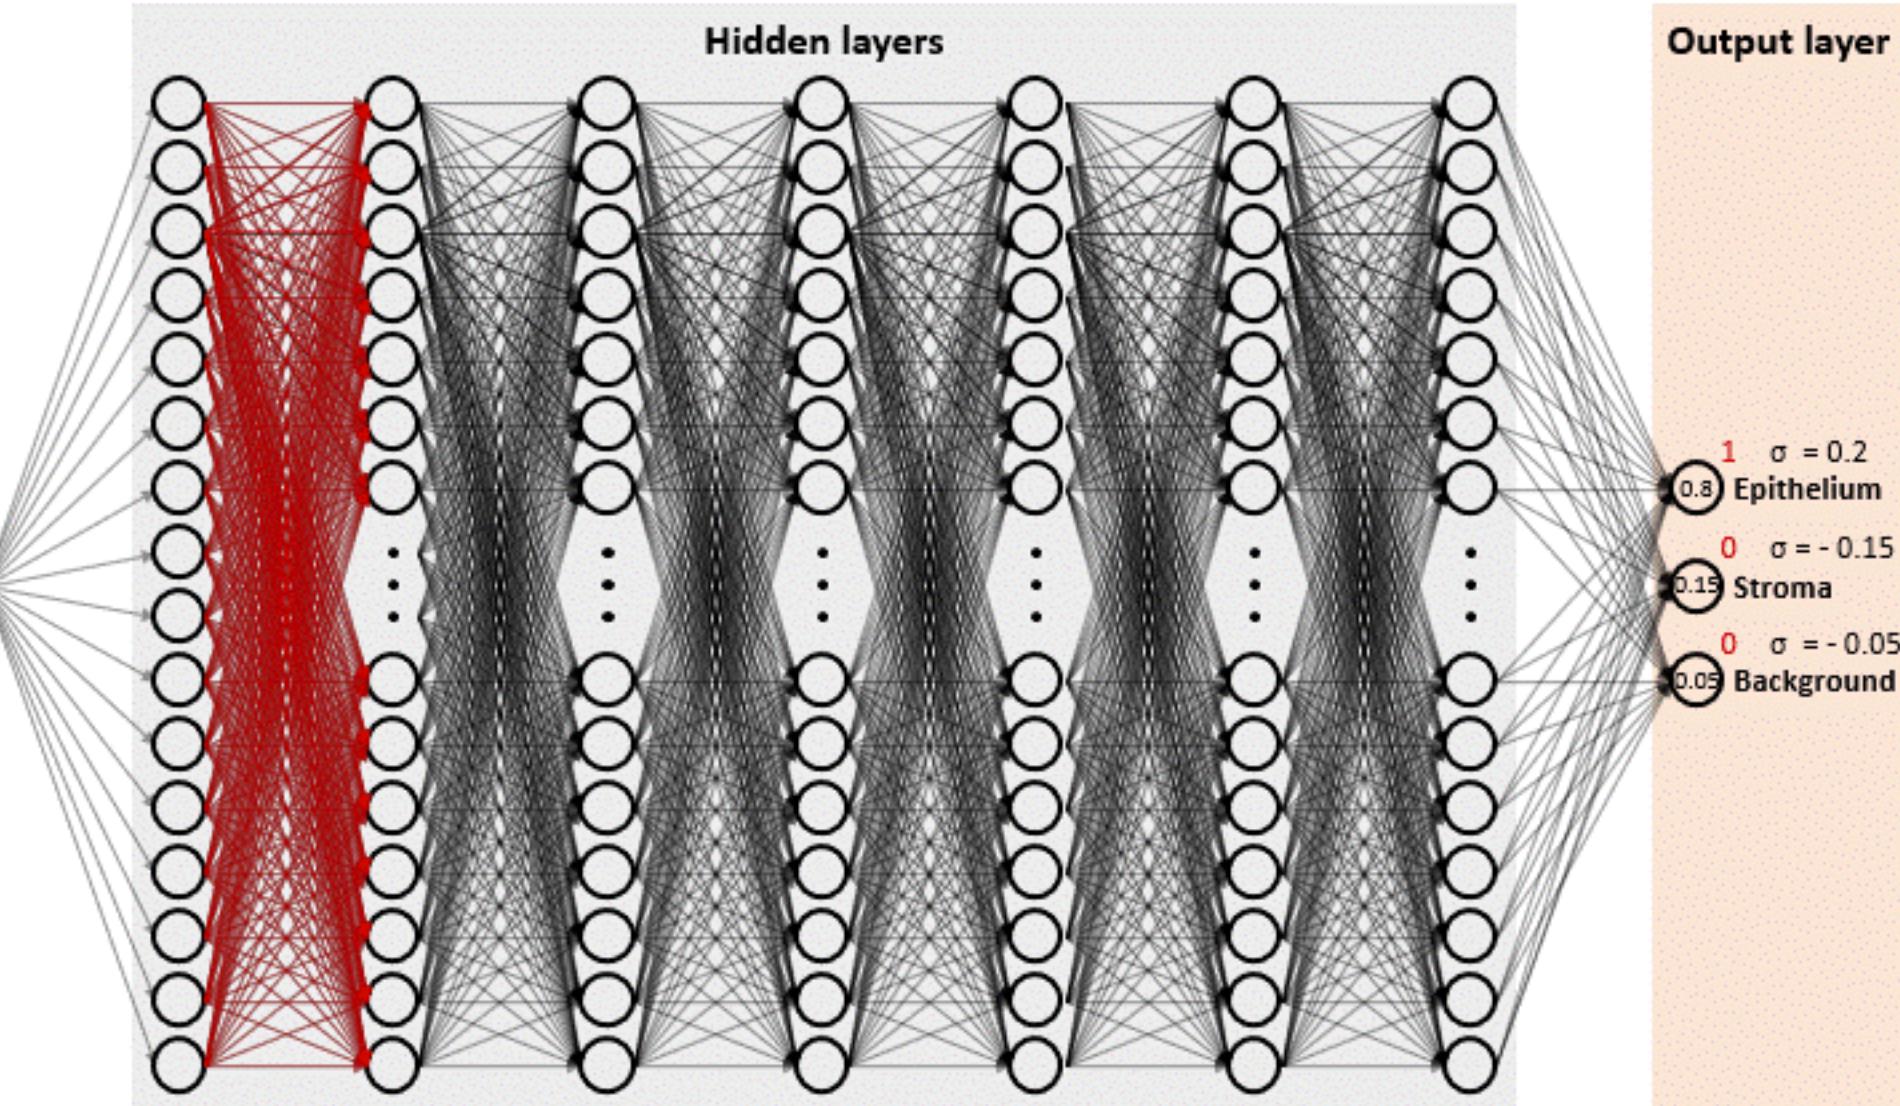
Output layer

DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



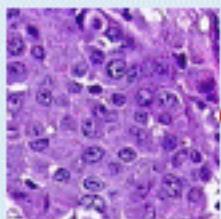
Hidden layers



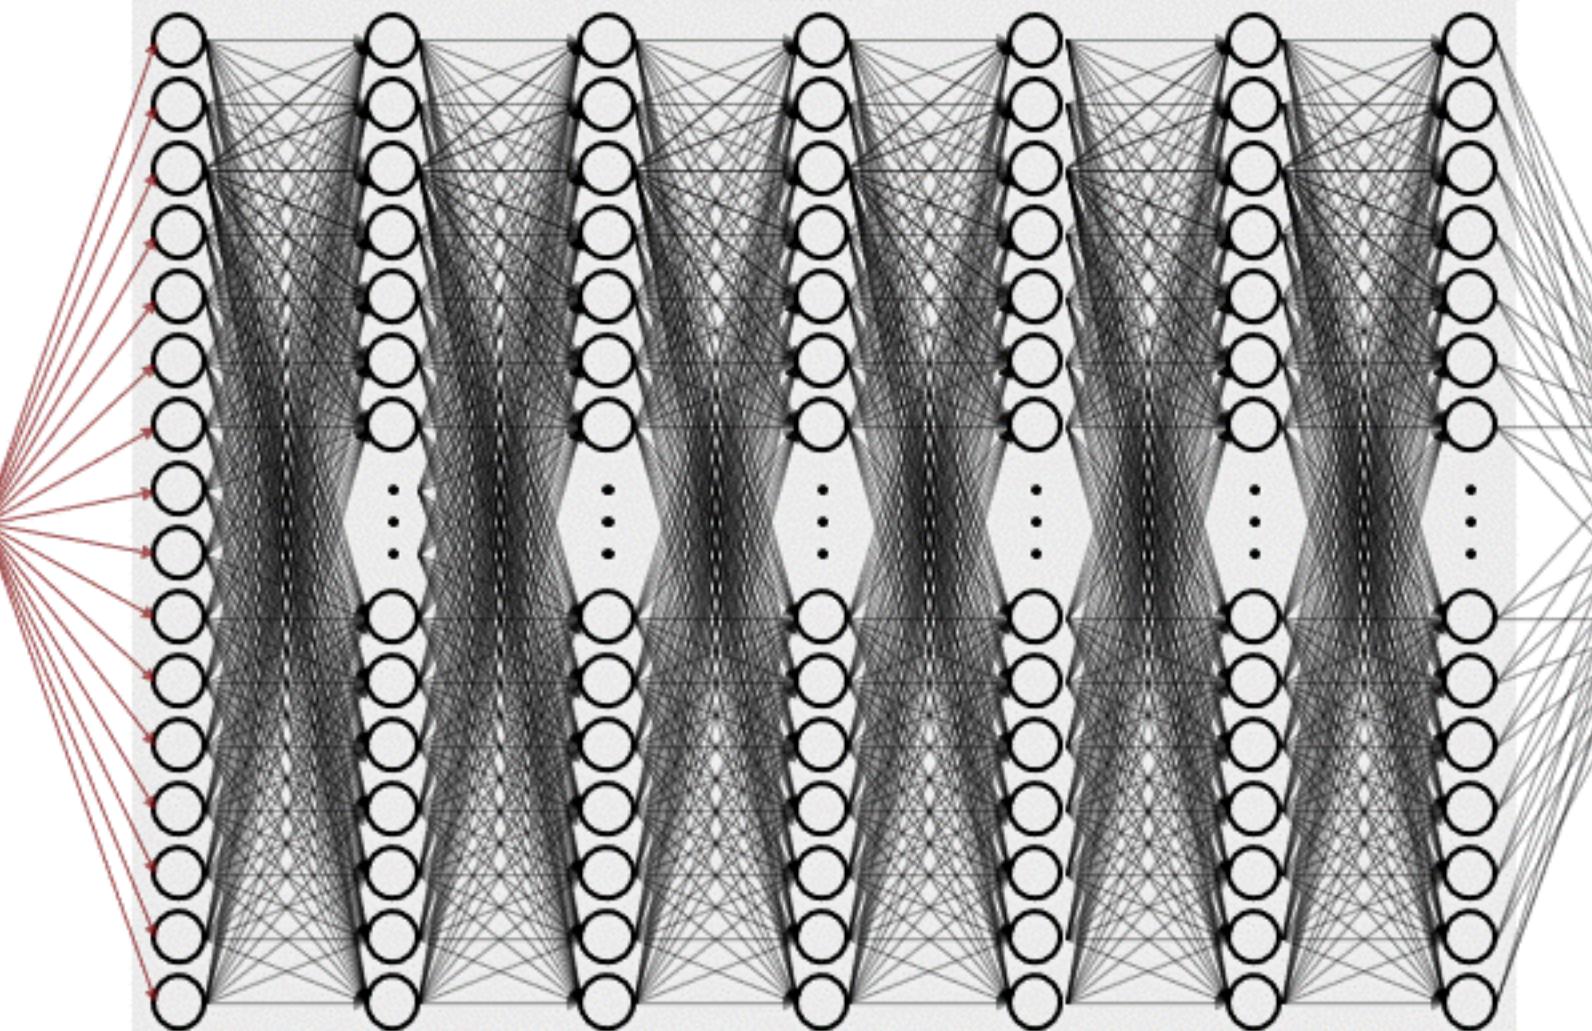
Output layer

DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



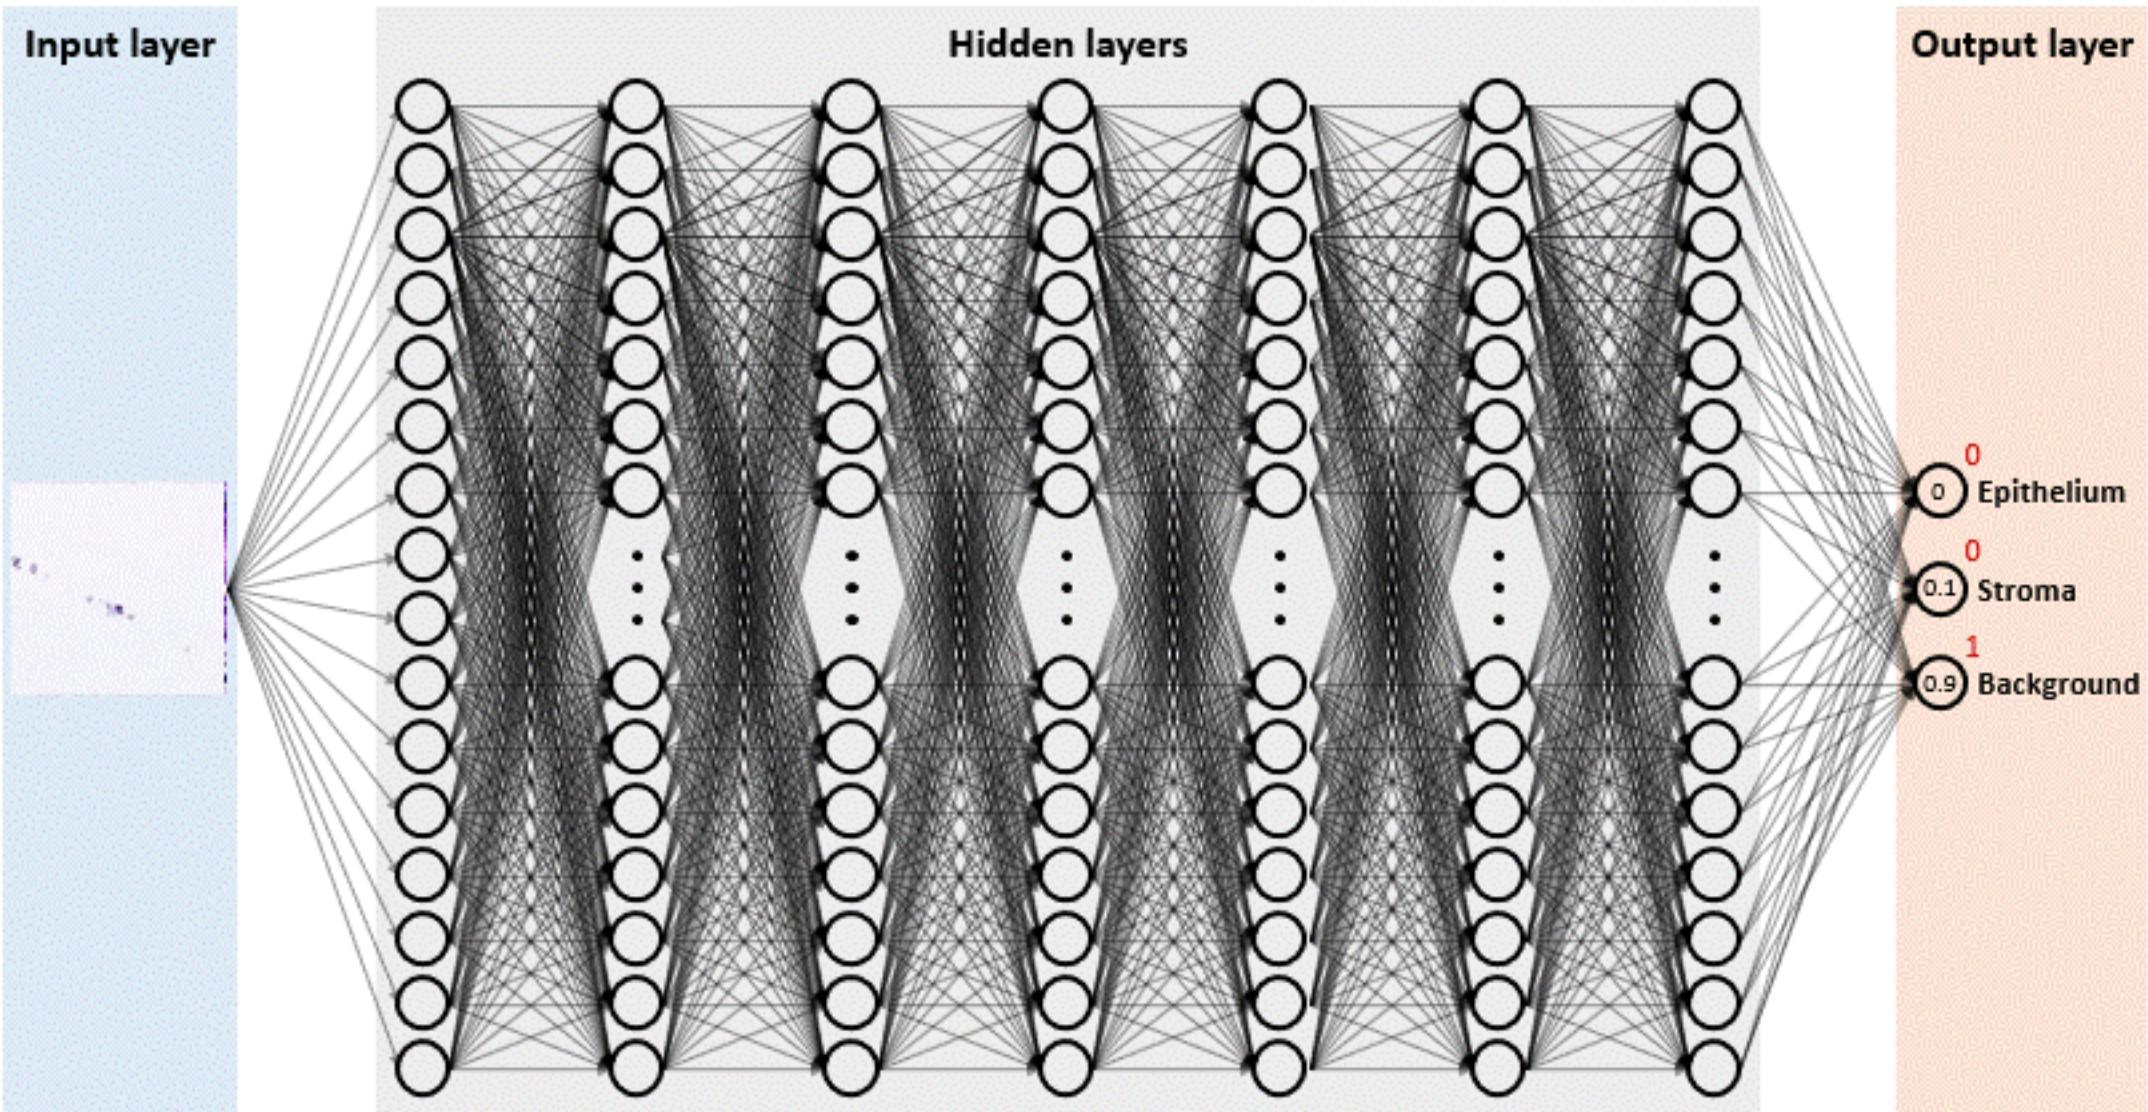
Hidden layers



Output layer

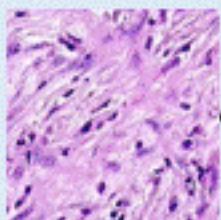
- 1 $\sigma = 0.2$ Epithelium
- 0 $\sigma = -0.15$ Stroma
- 0 $\sigma = -0.05$ Background

DEEP CONVOLUTIONAL NEURAL NETWORKS

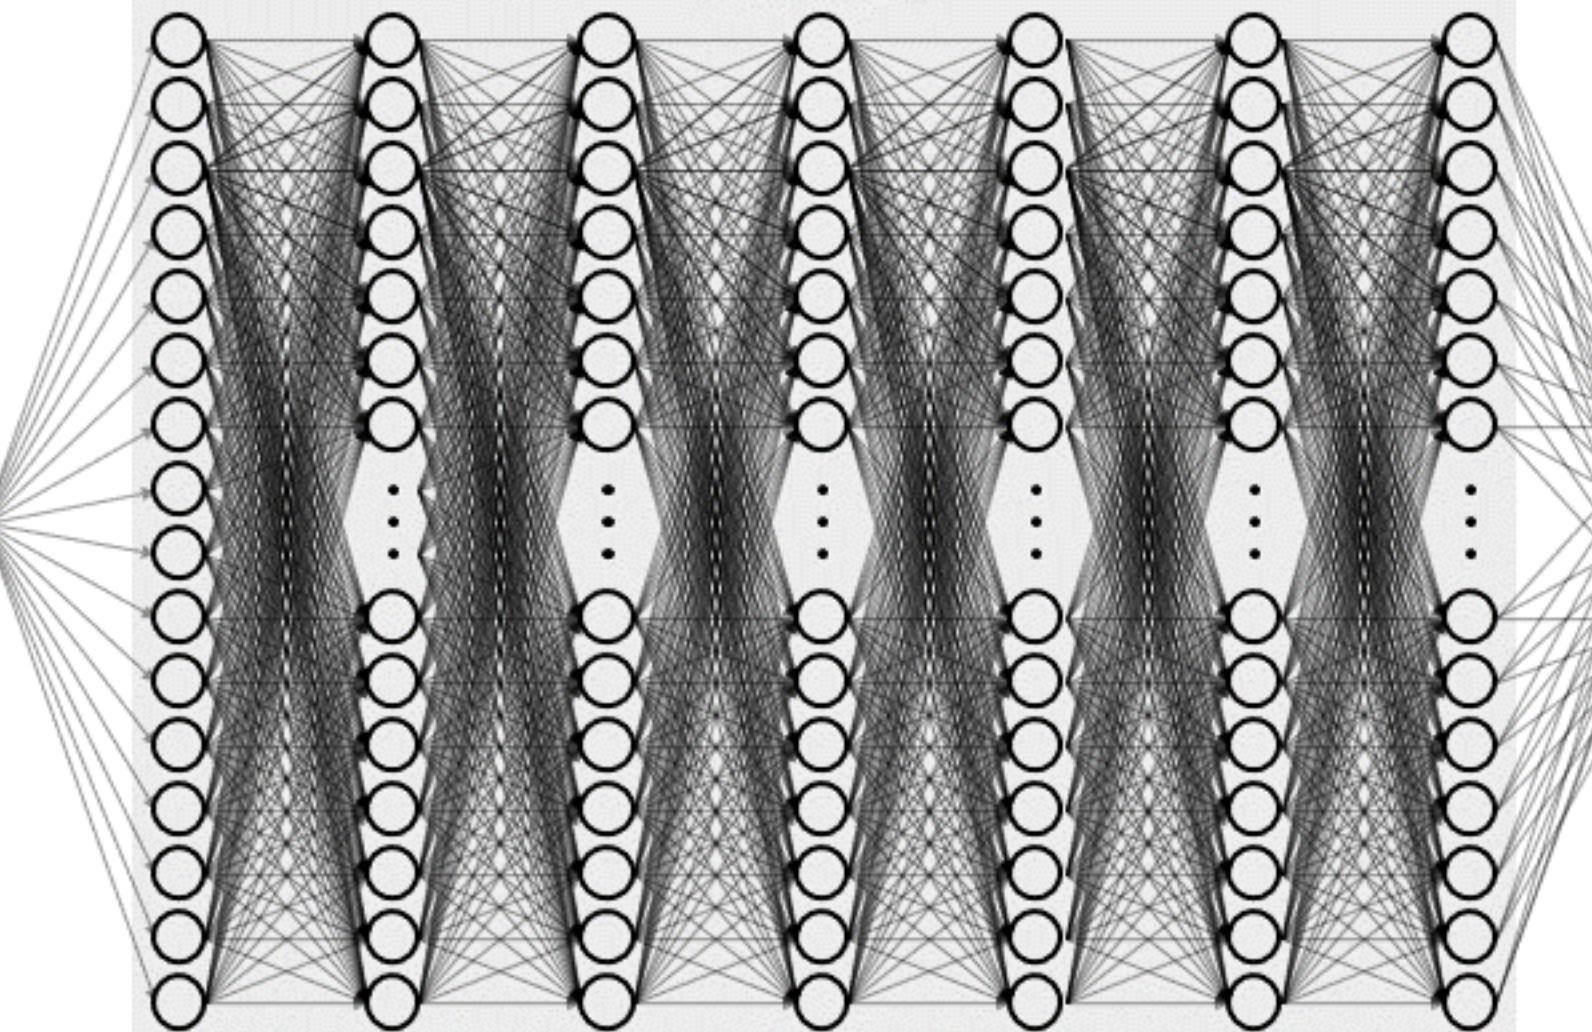


DEEP CONVOLUTIONAL NEURAL NETWORKS

Input layer



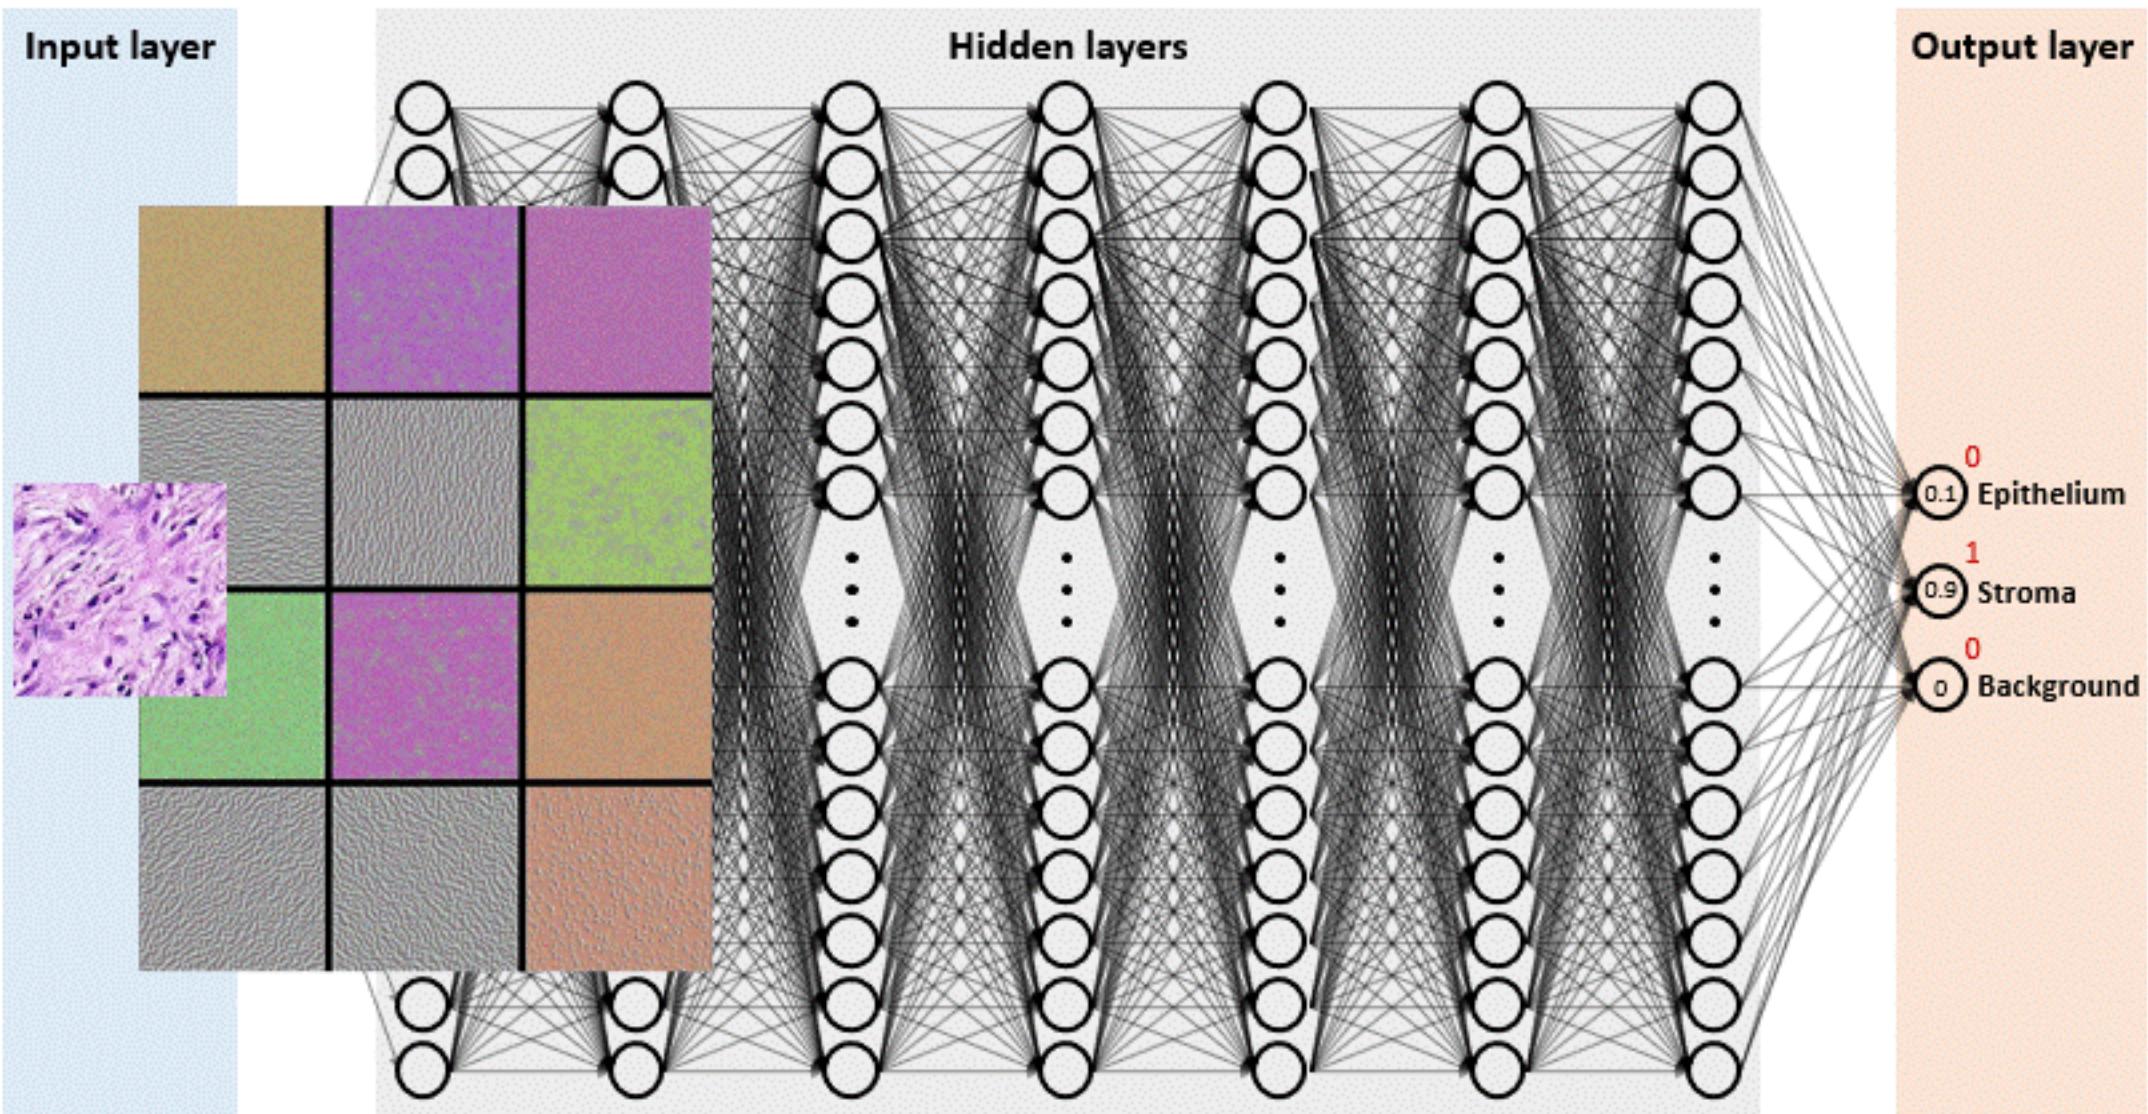
Hidden layers



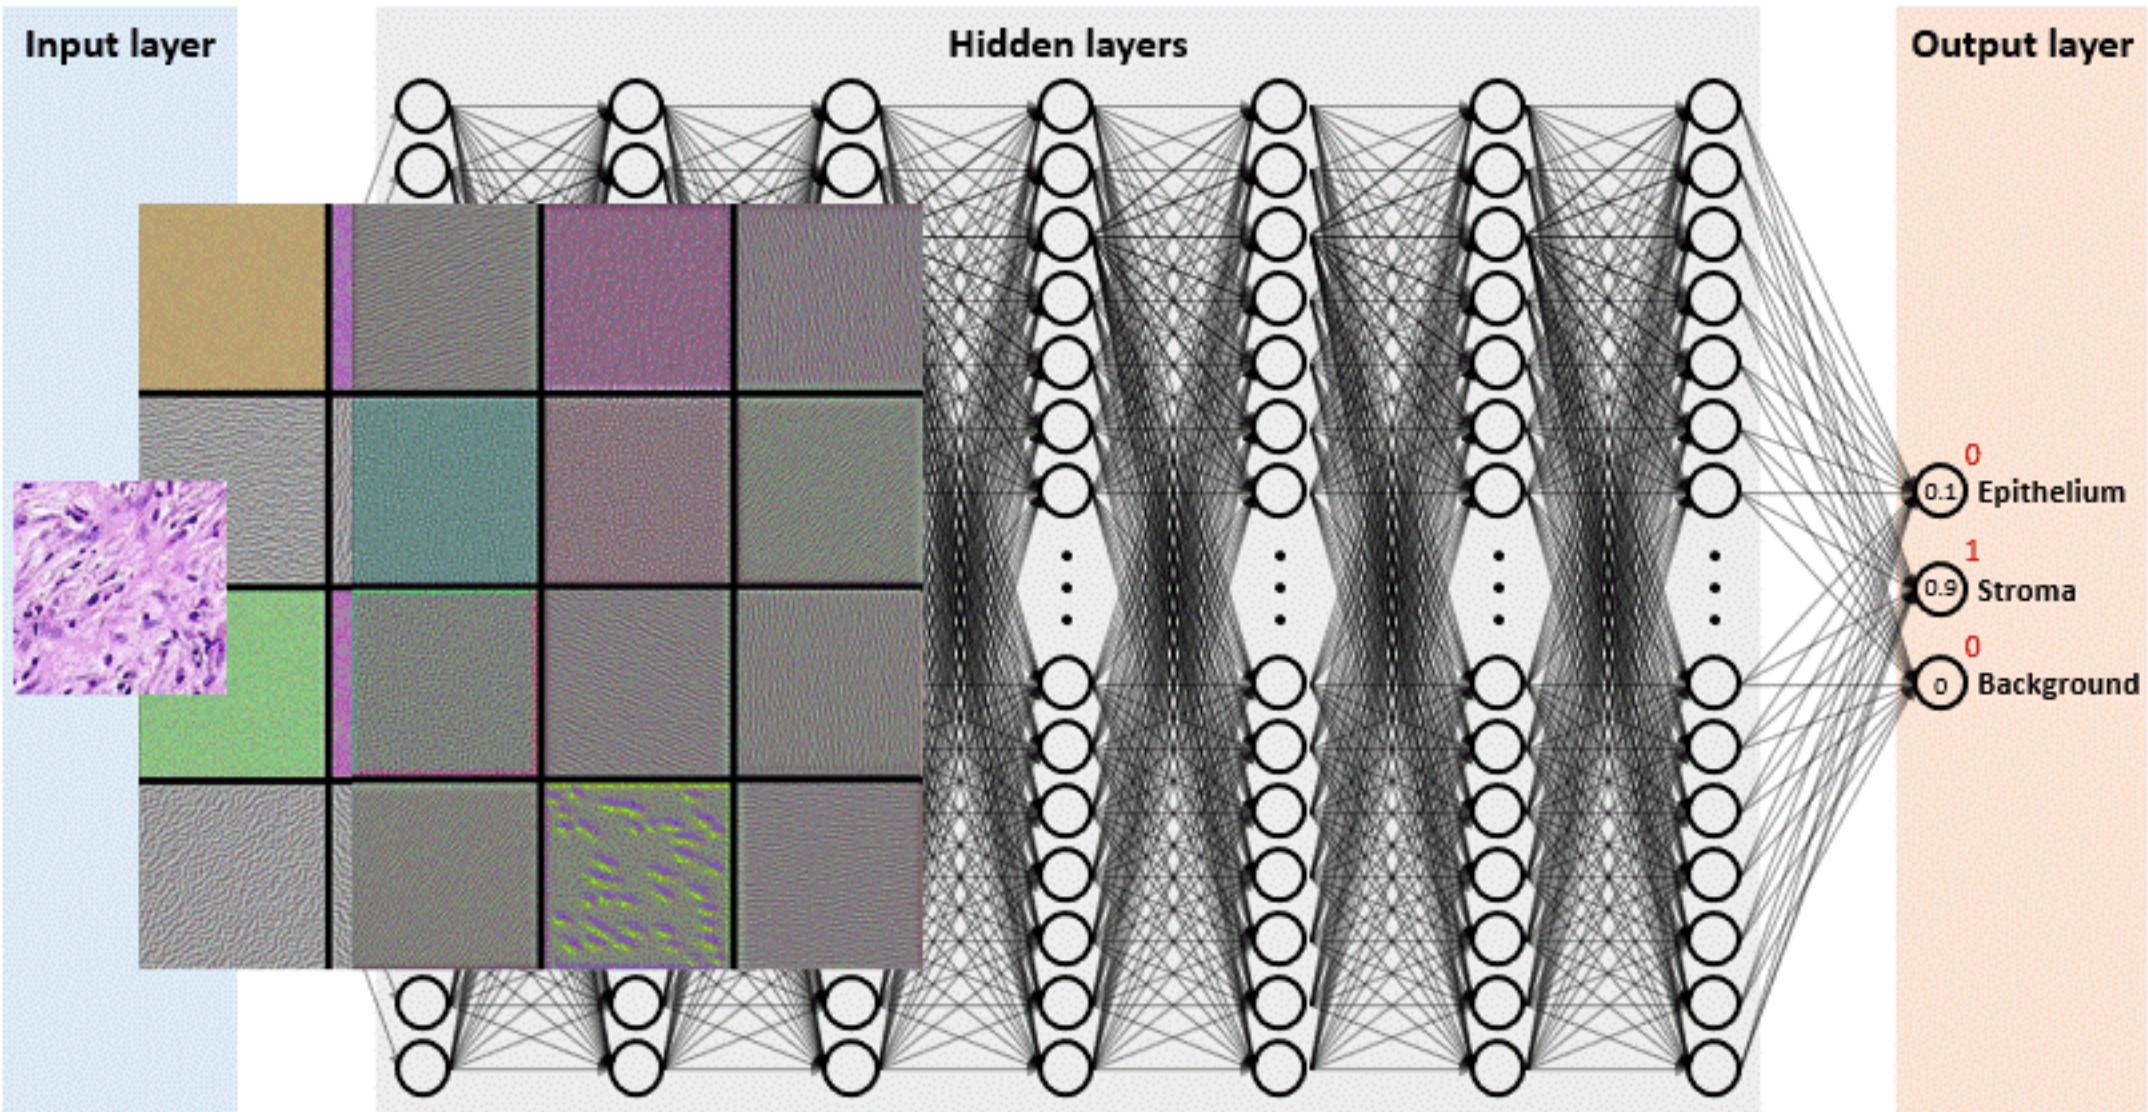
Output layer

0 Epithelium
1 Stroma
0 Background

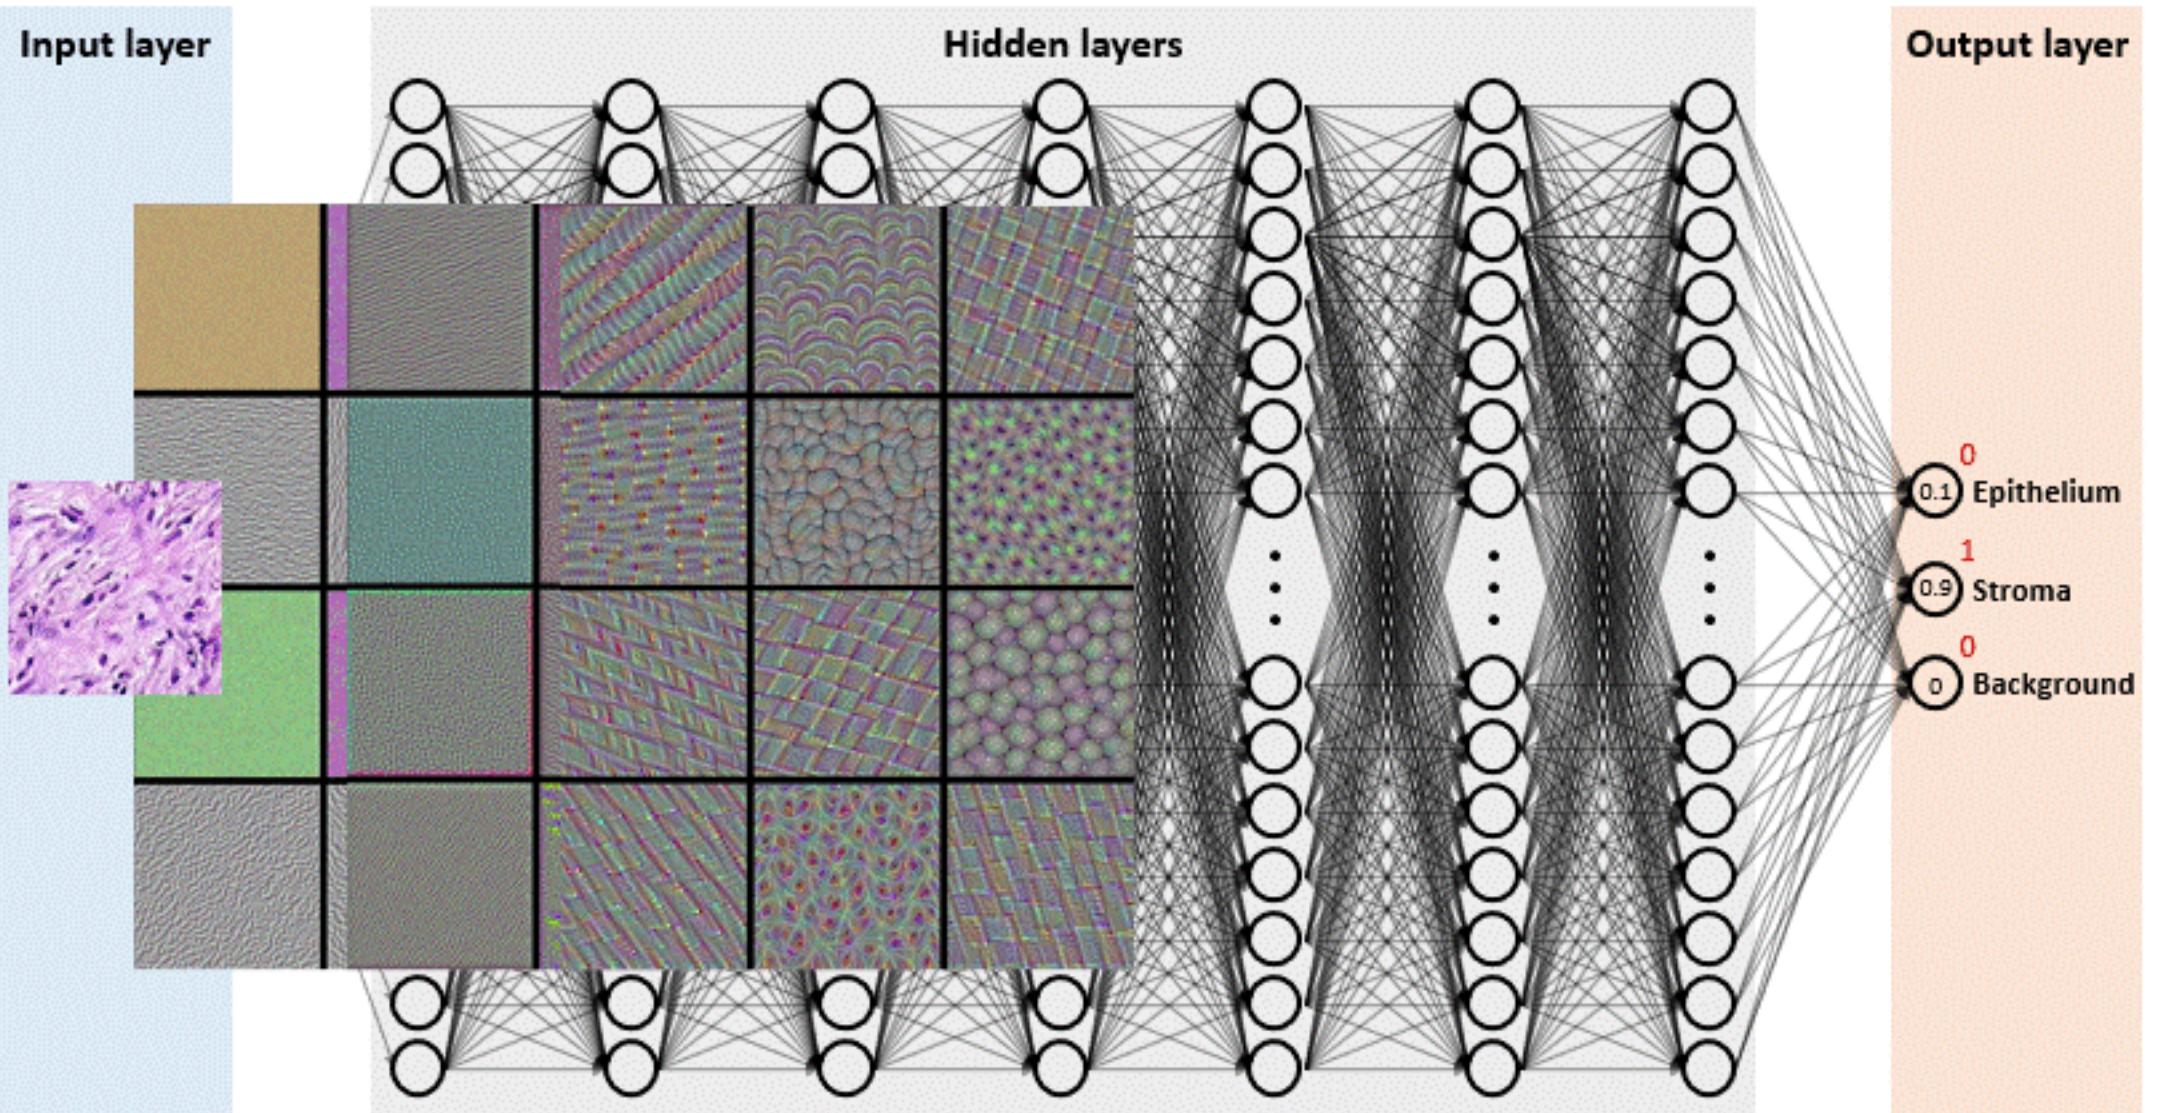
DEEP CONVOLUTIONAL NEURAL NETWORKS



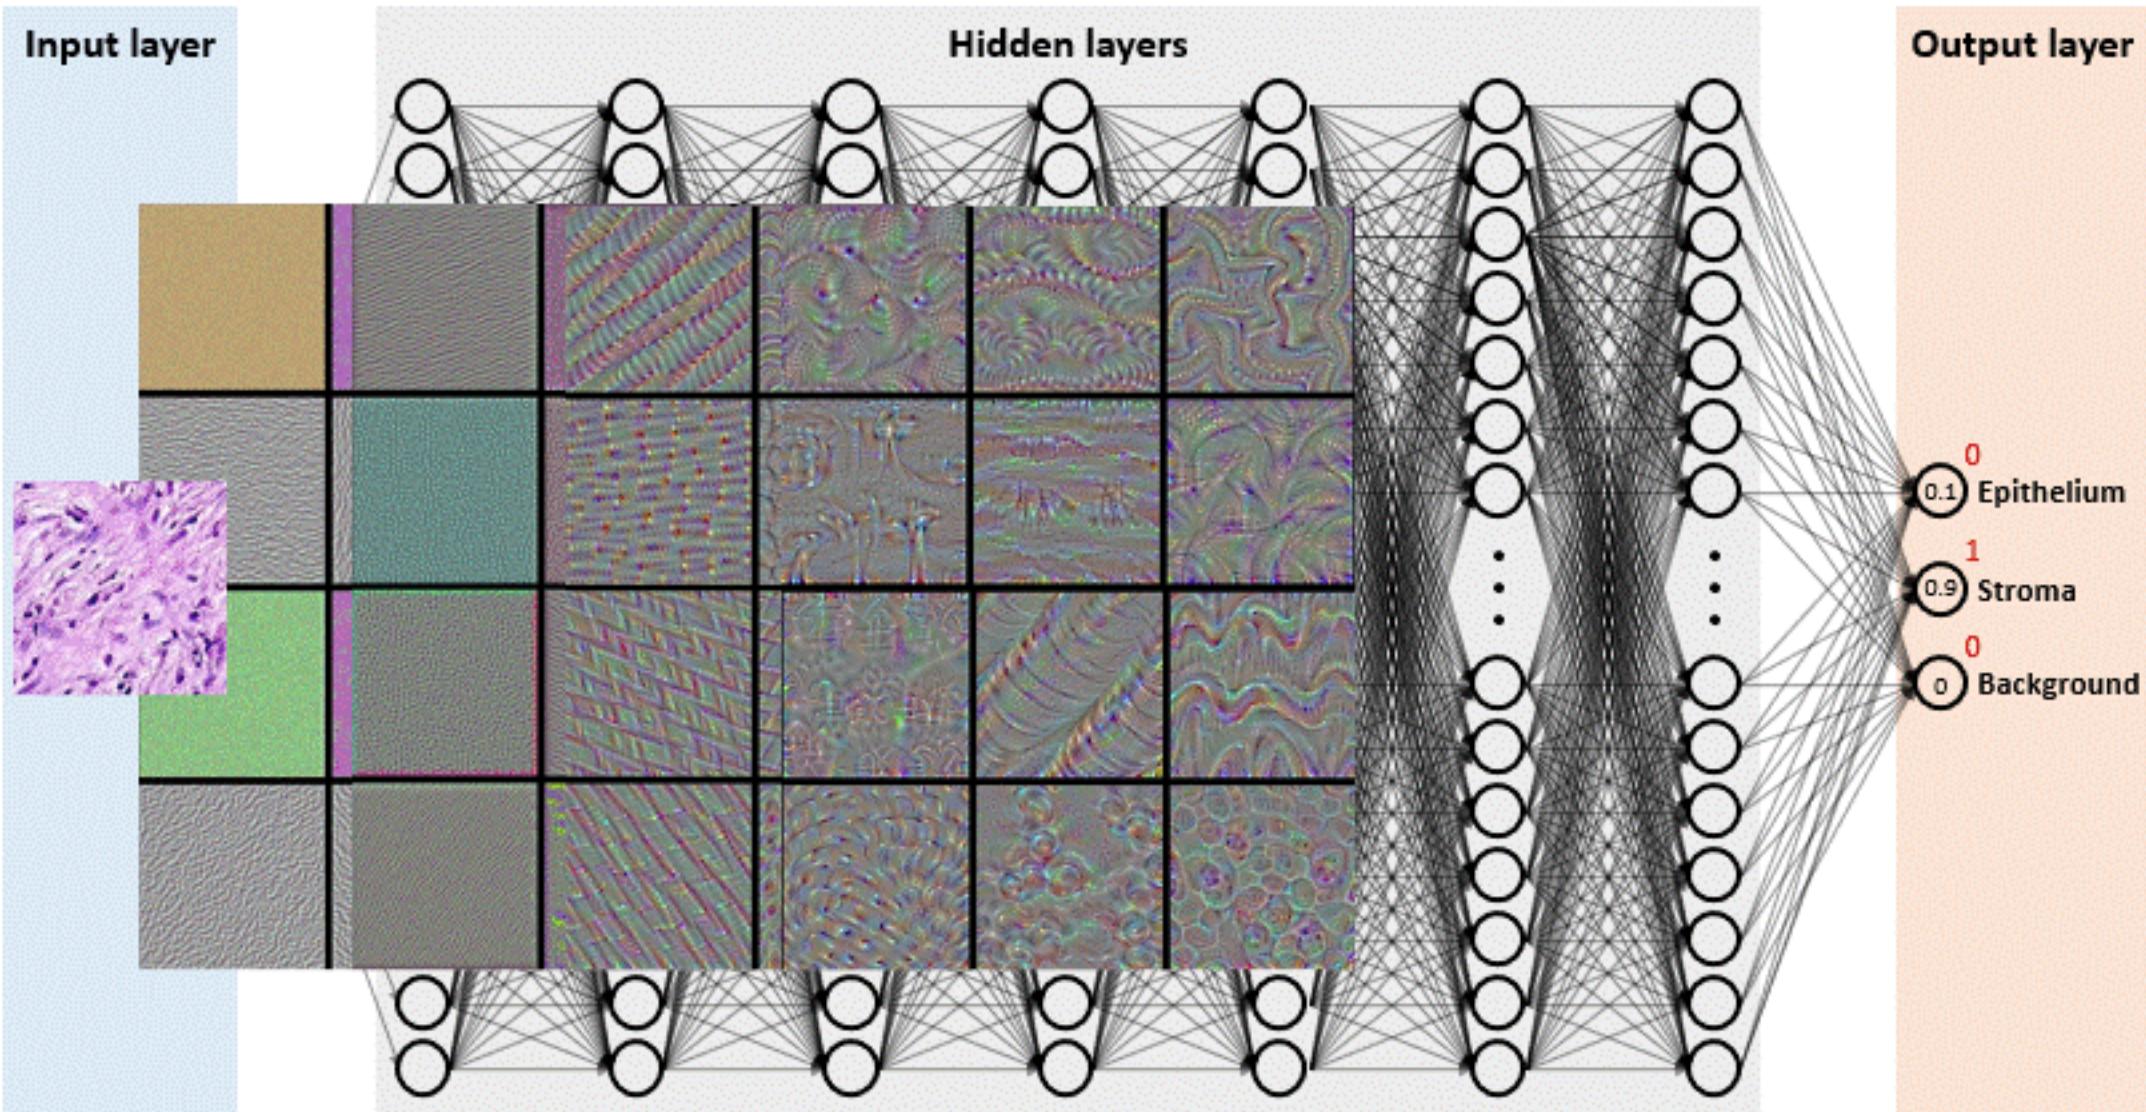
DEEP CONVOLUTIONAL NEURAL NETWORKS



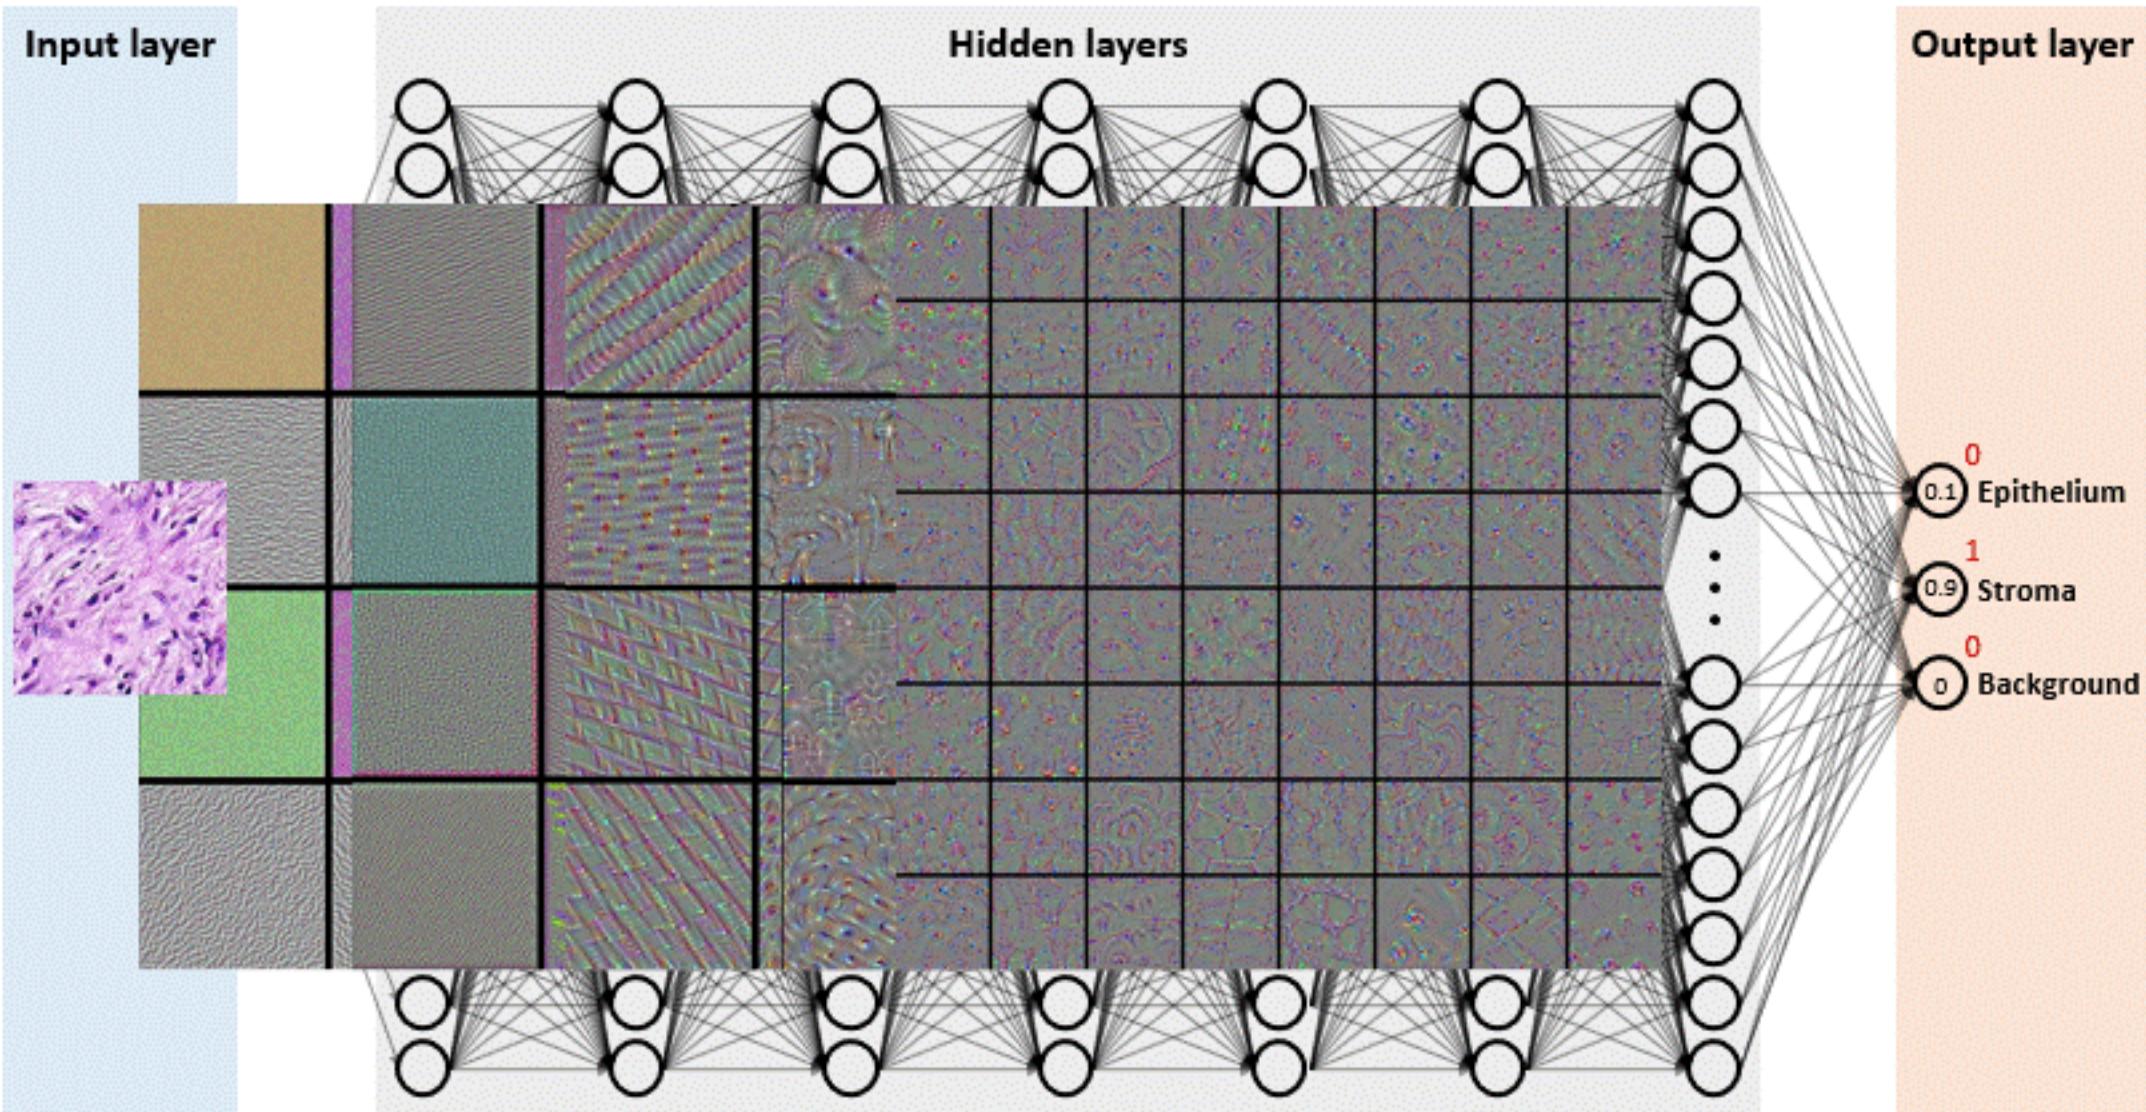
DEEP CONVOLUTIONAL NEURAL NETWORKS



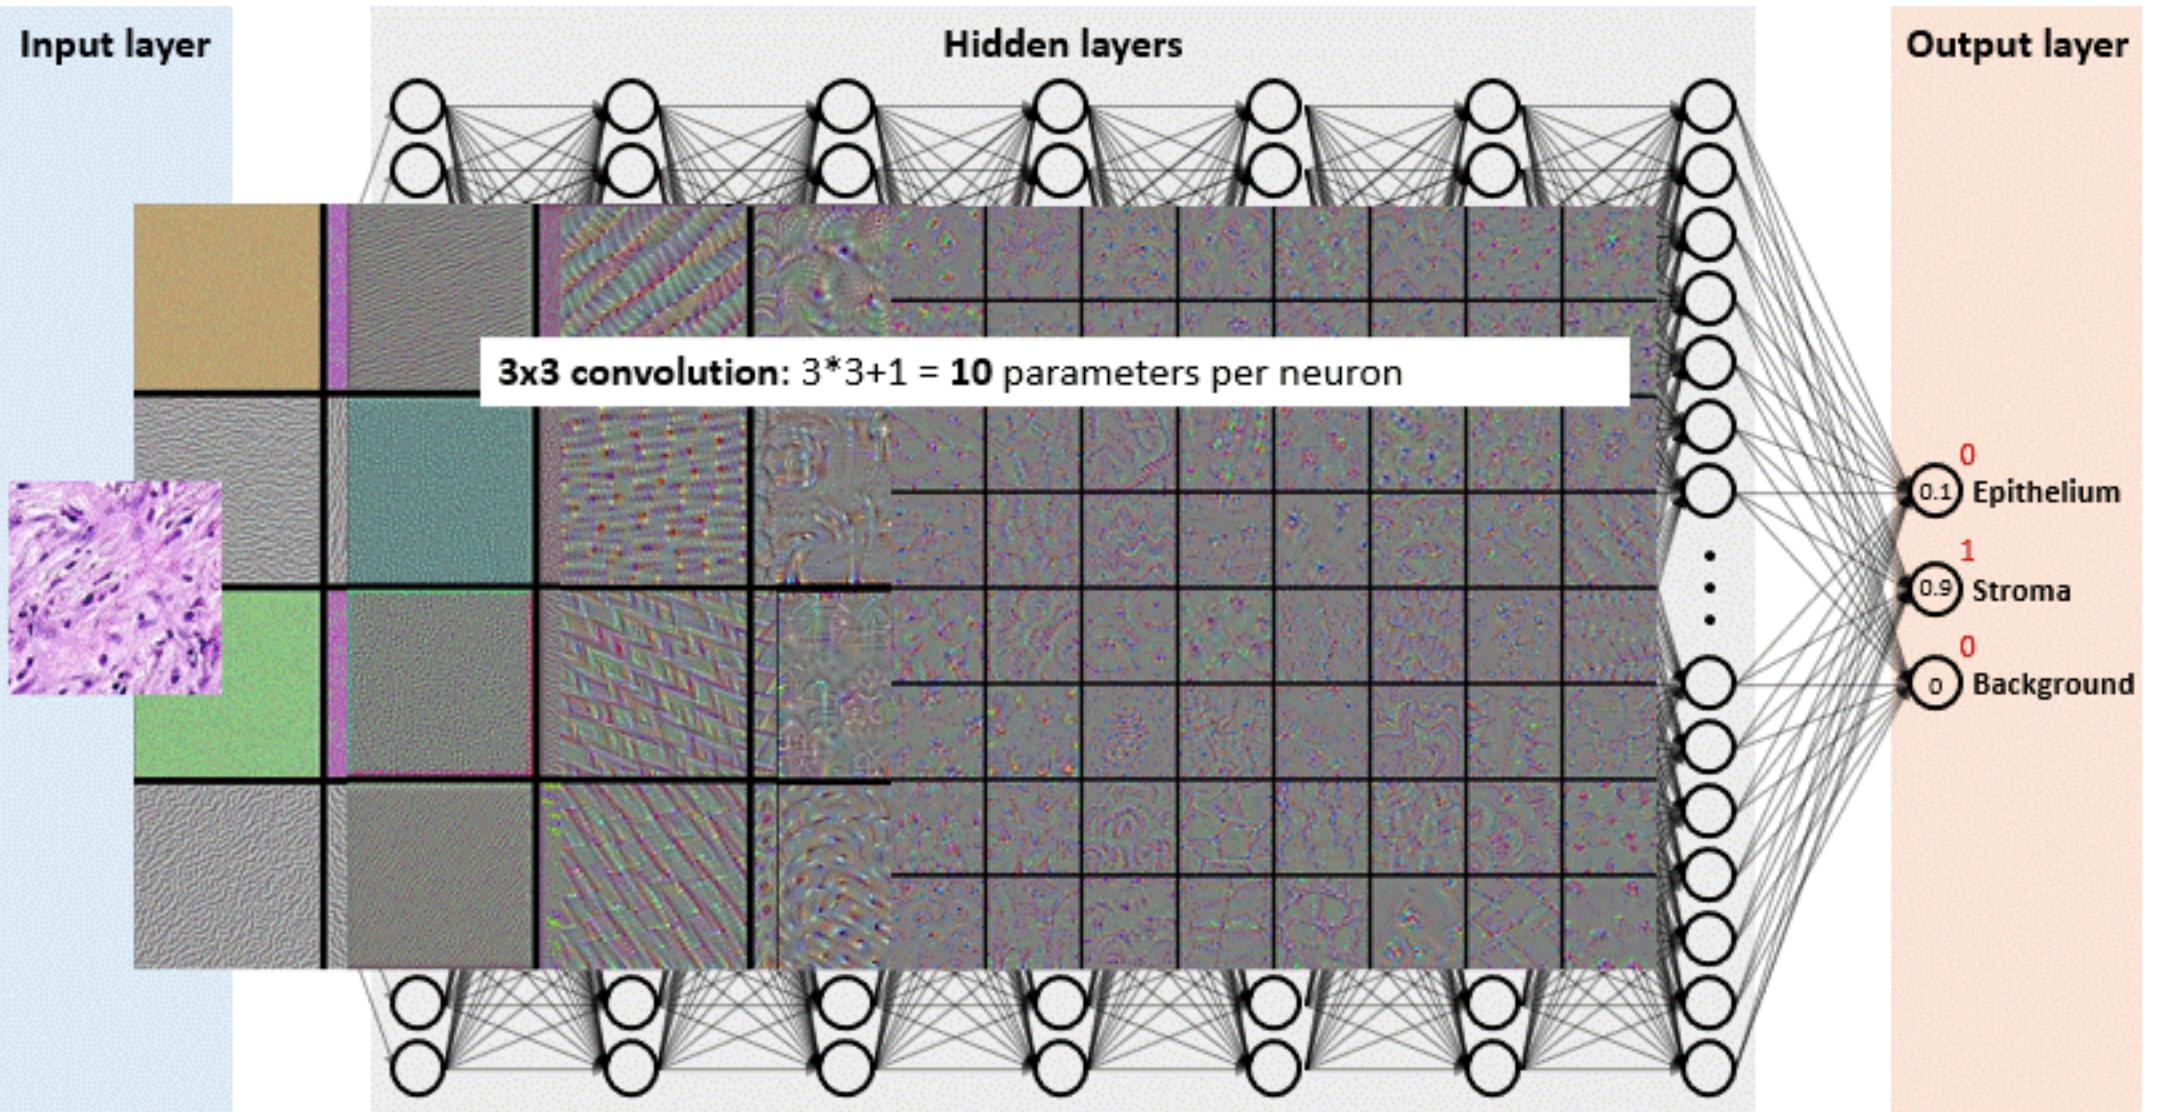
DEEP CONVOLUTIONAL NEURAL NETWORKS



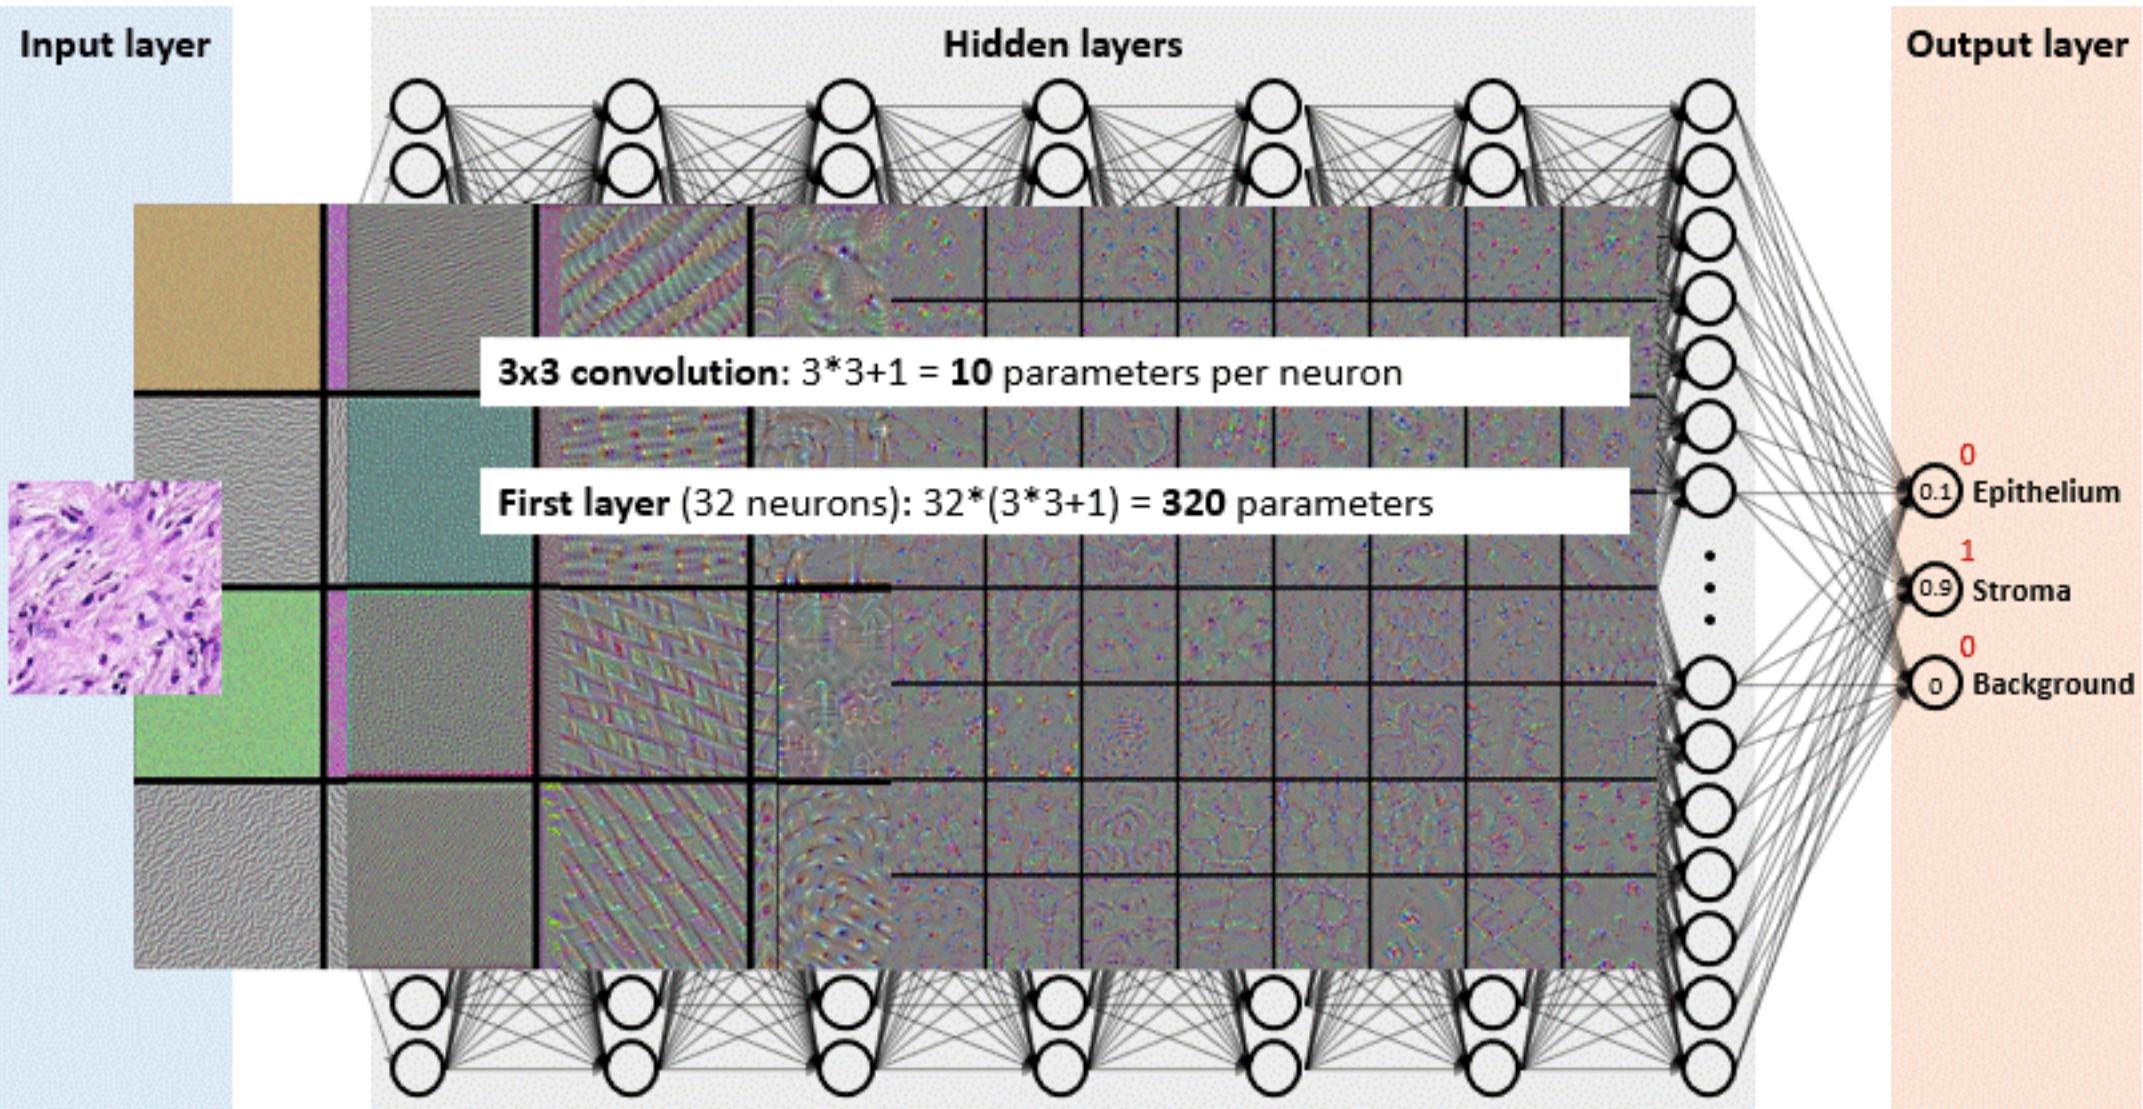
DEEP CONVOLUTIONAL NEURAL NETWORKS



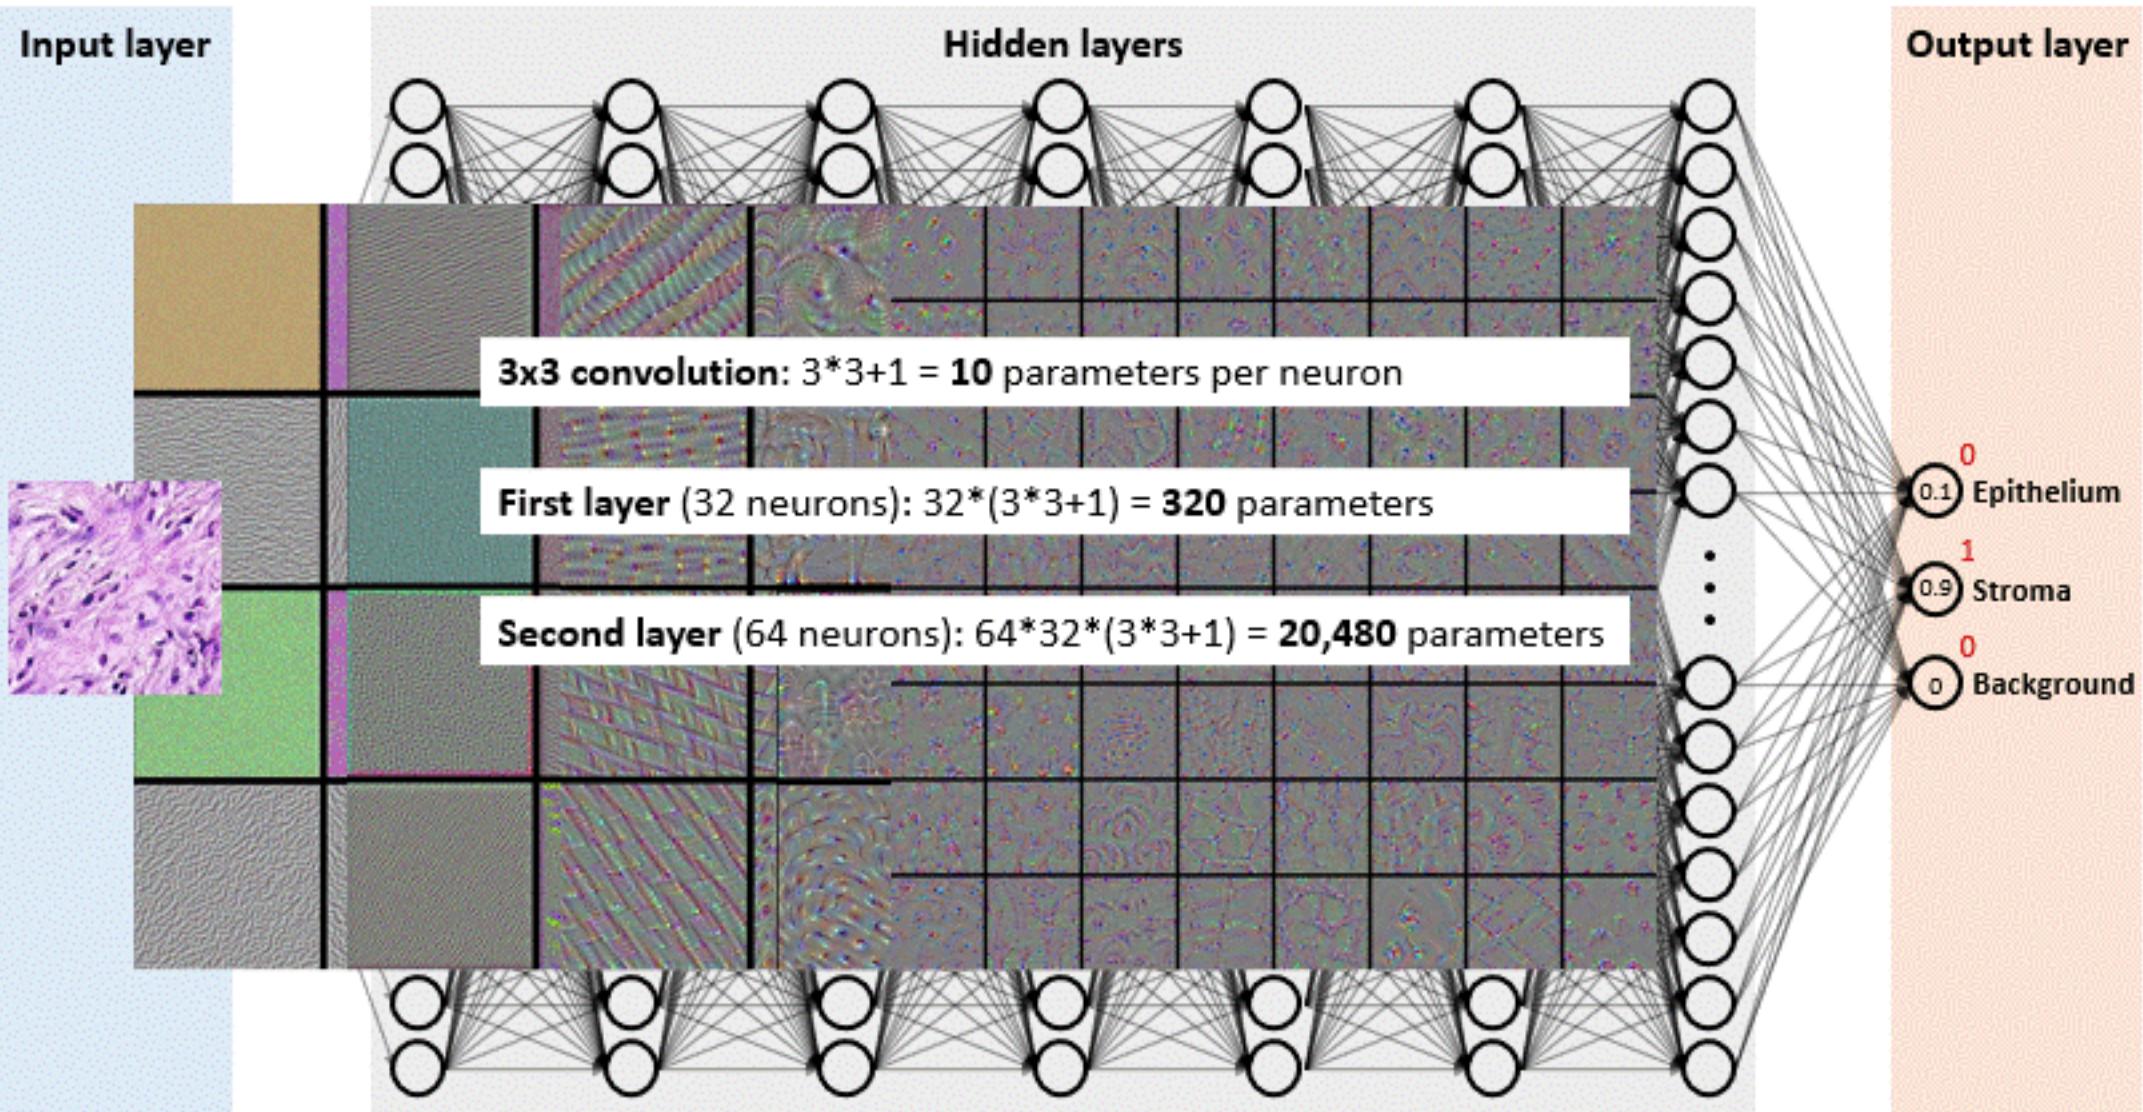
DEEP CONVOLUTIONAL NEURAL NETWORKS



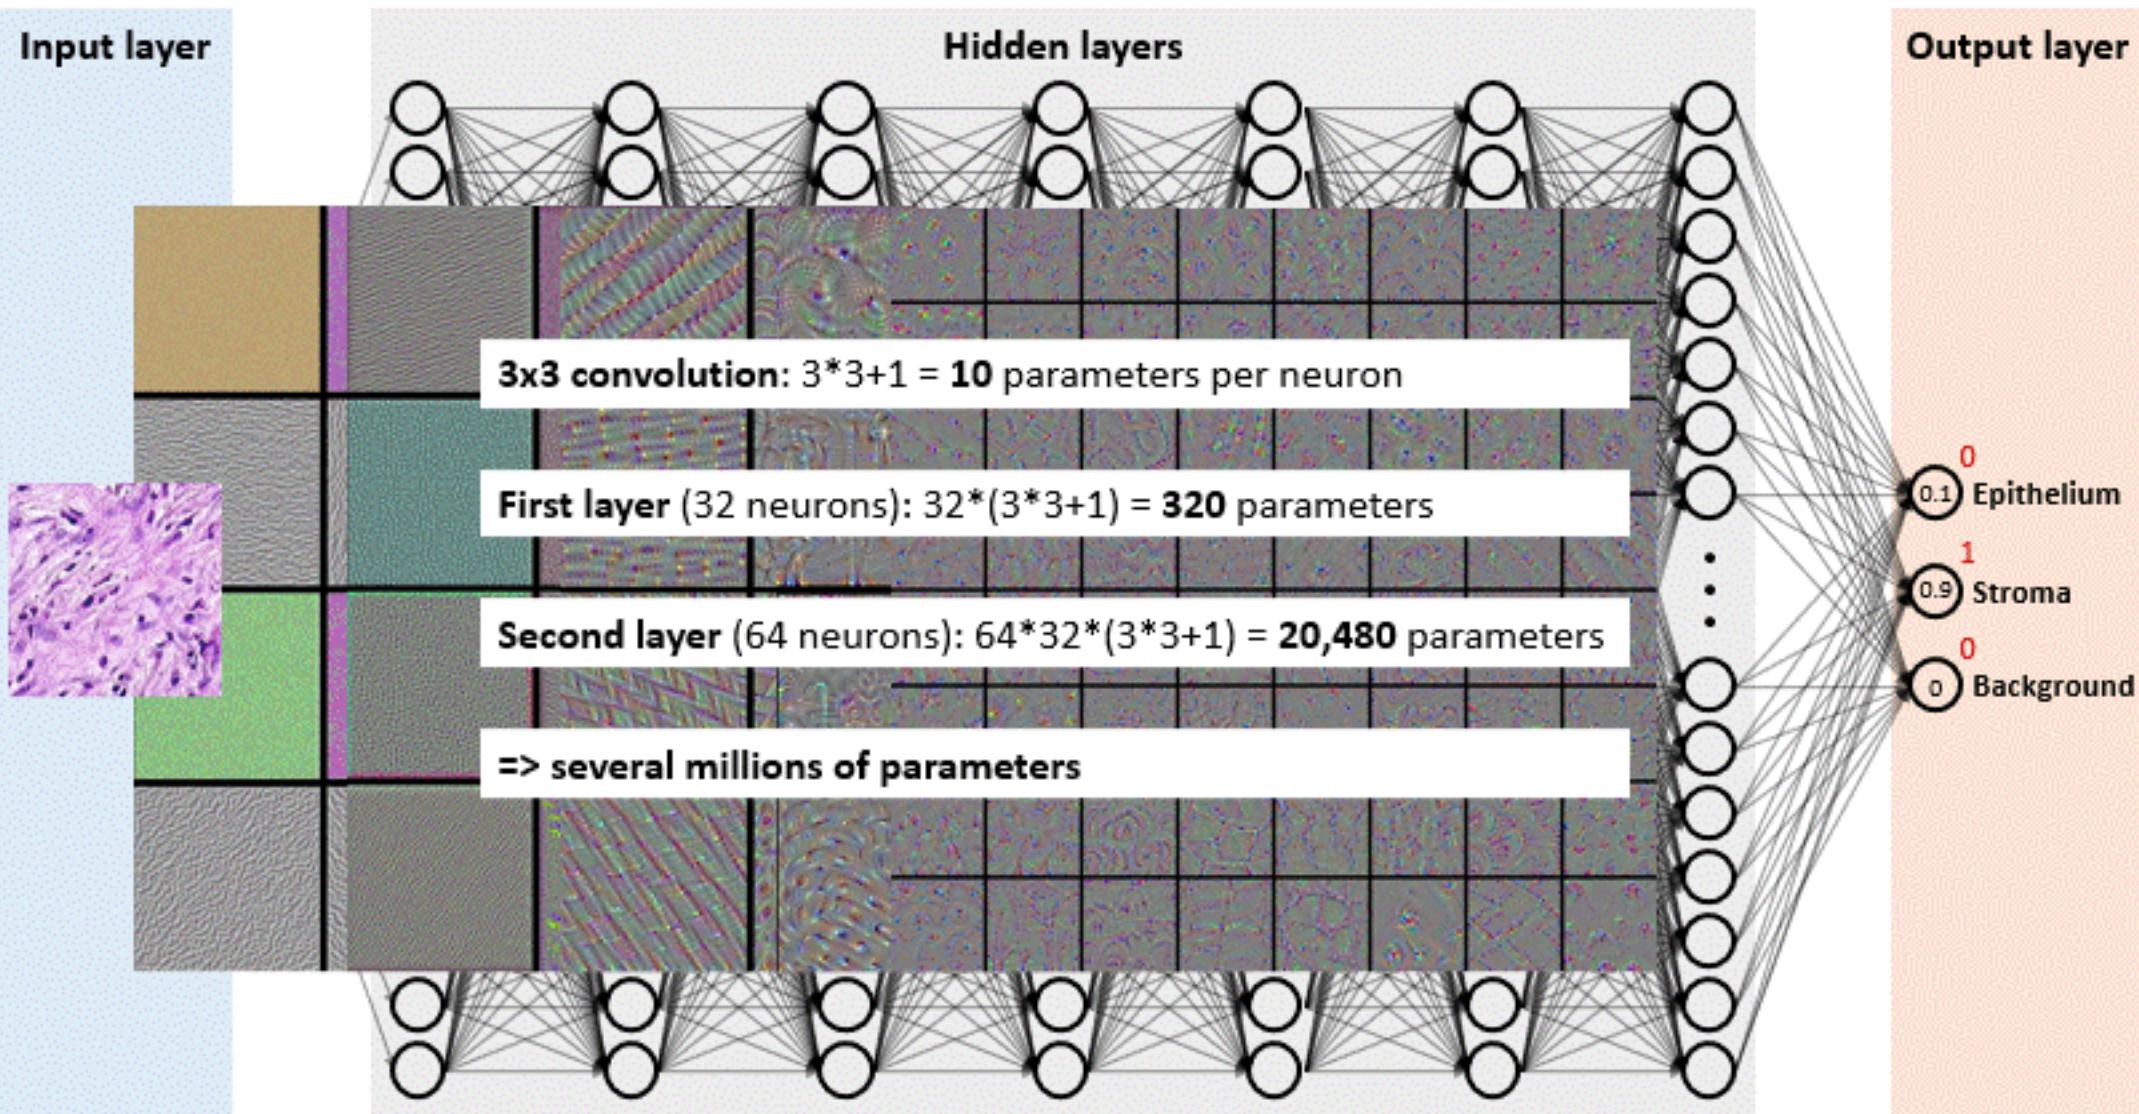
DEEP CONVOLUTIONAL NEURAL NETWORKS



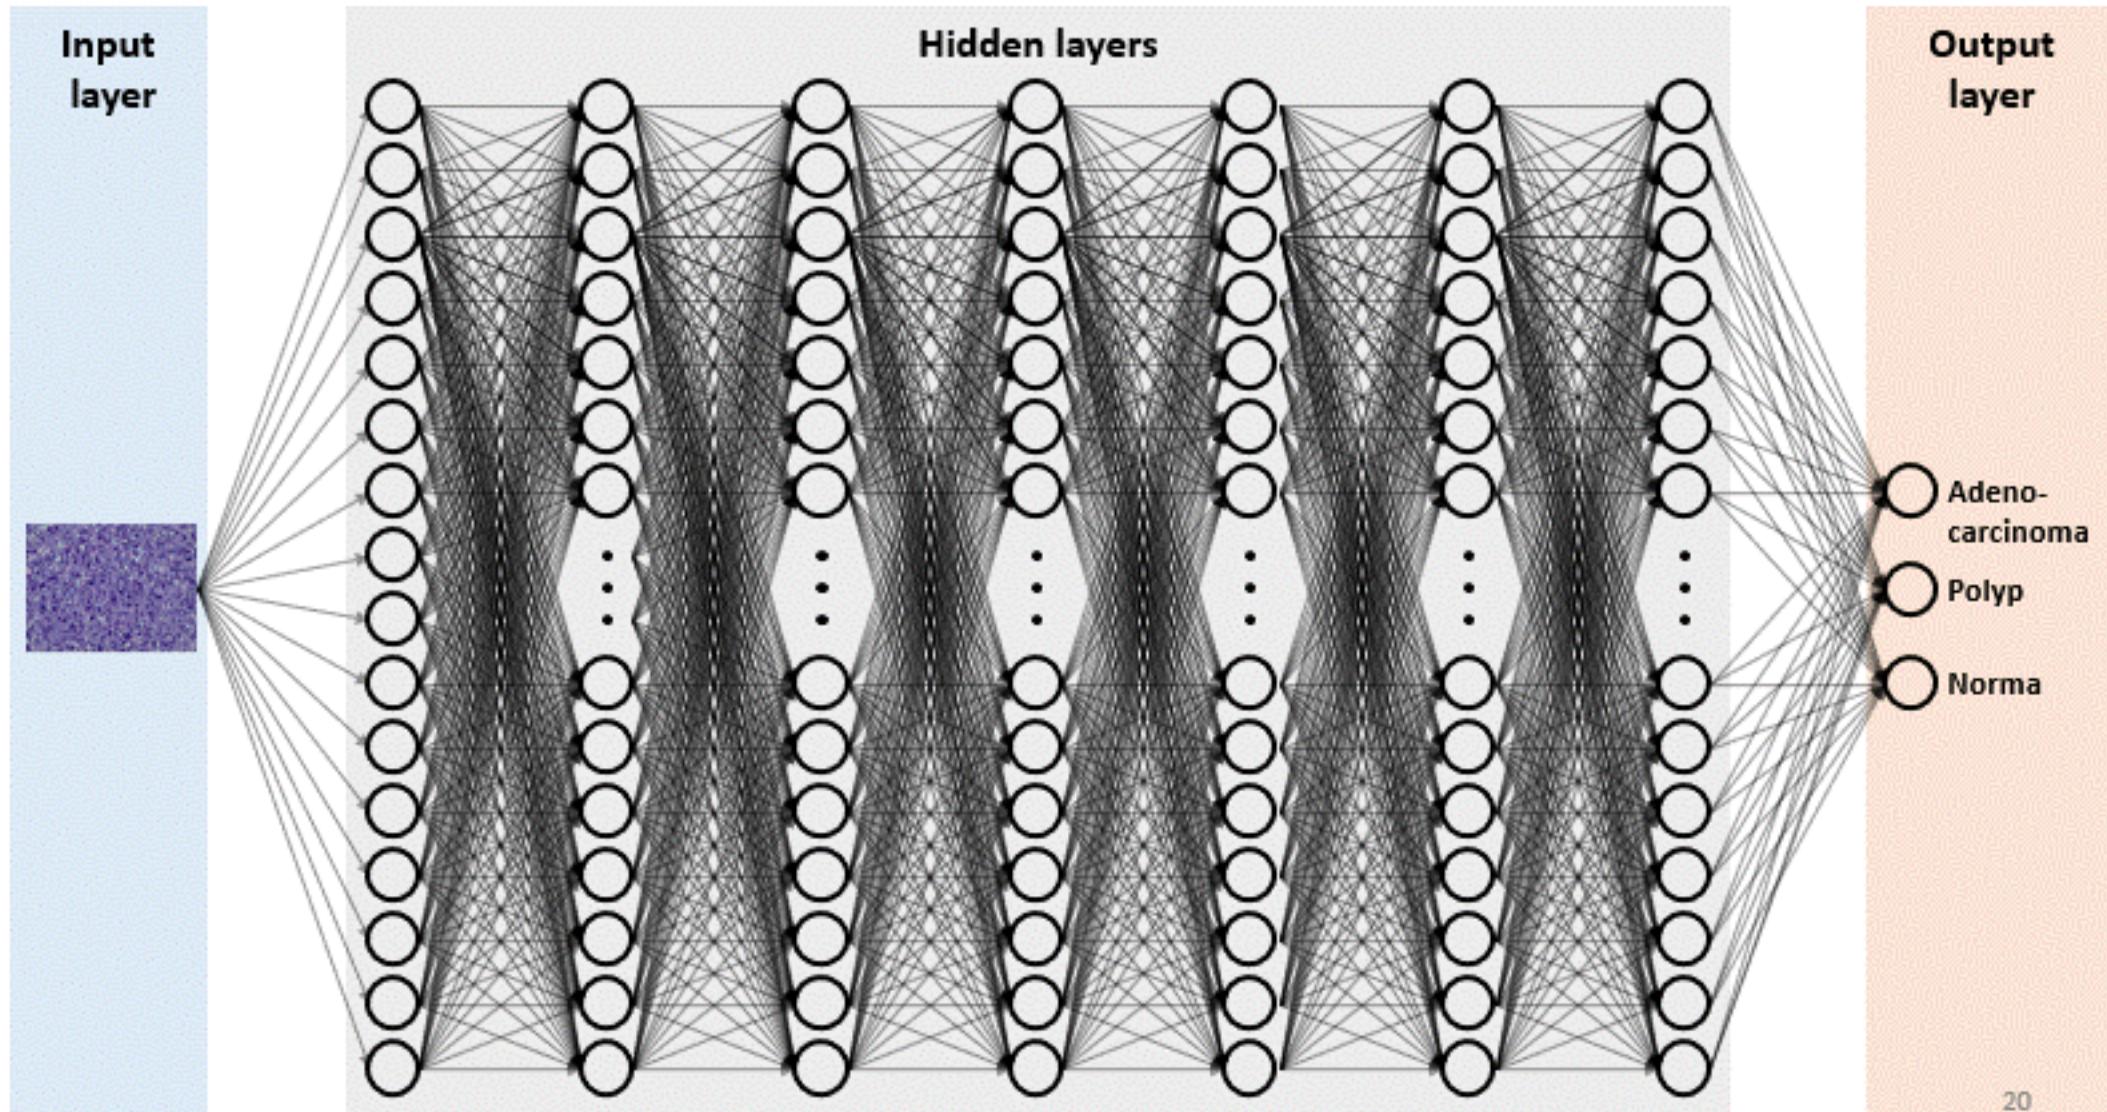
DEEP CONVOLUTIONAL NEURAL NETWORKS



DEEP CONVOLUTIONAL NEURAL NETWORKS

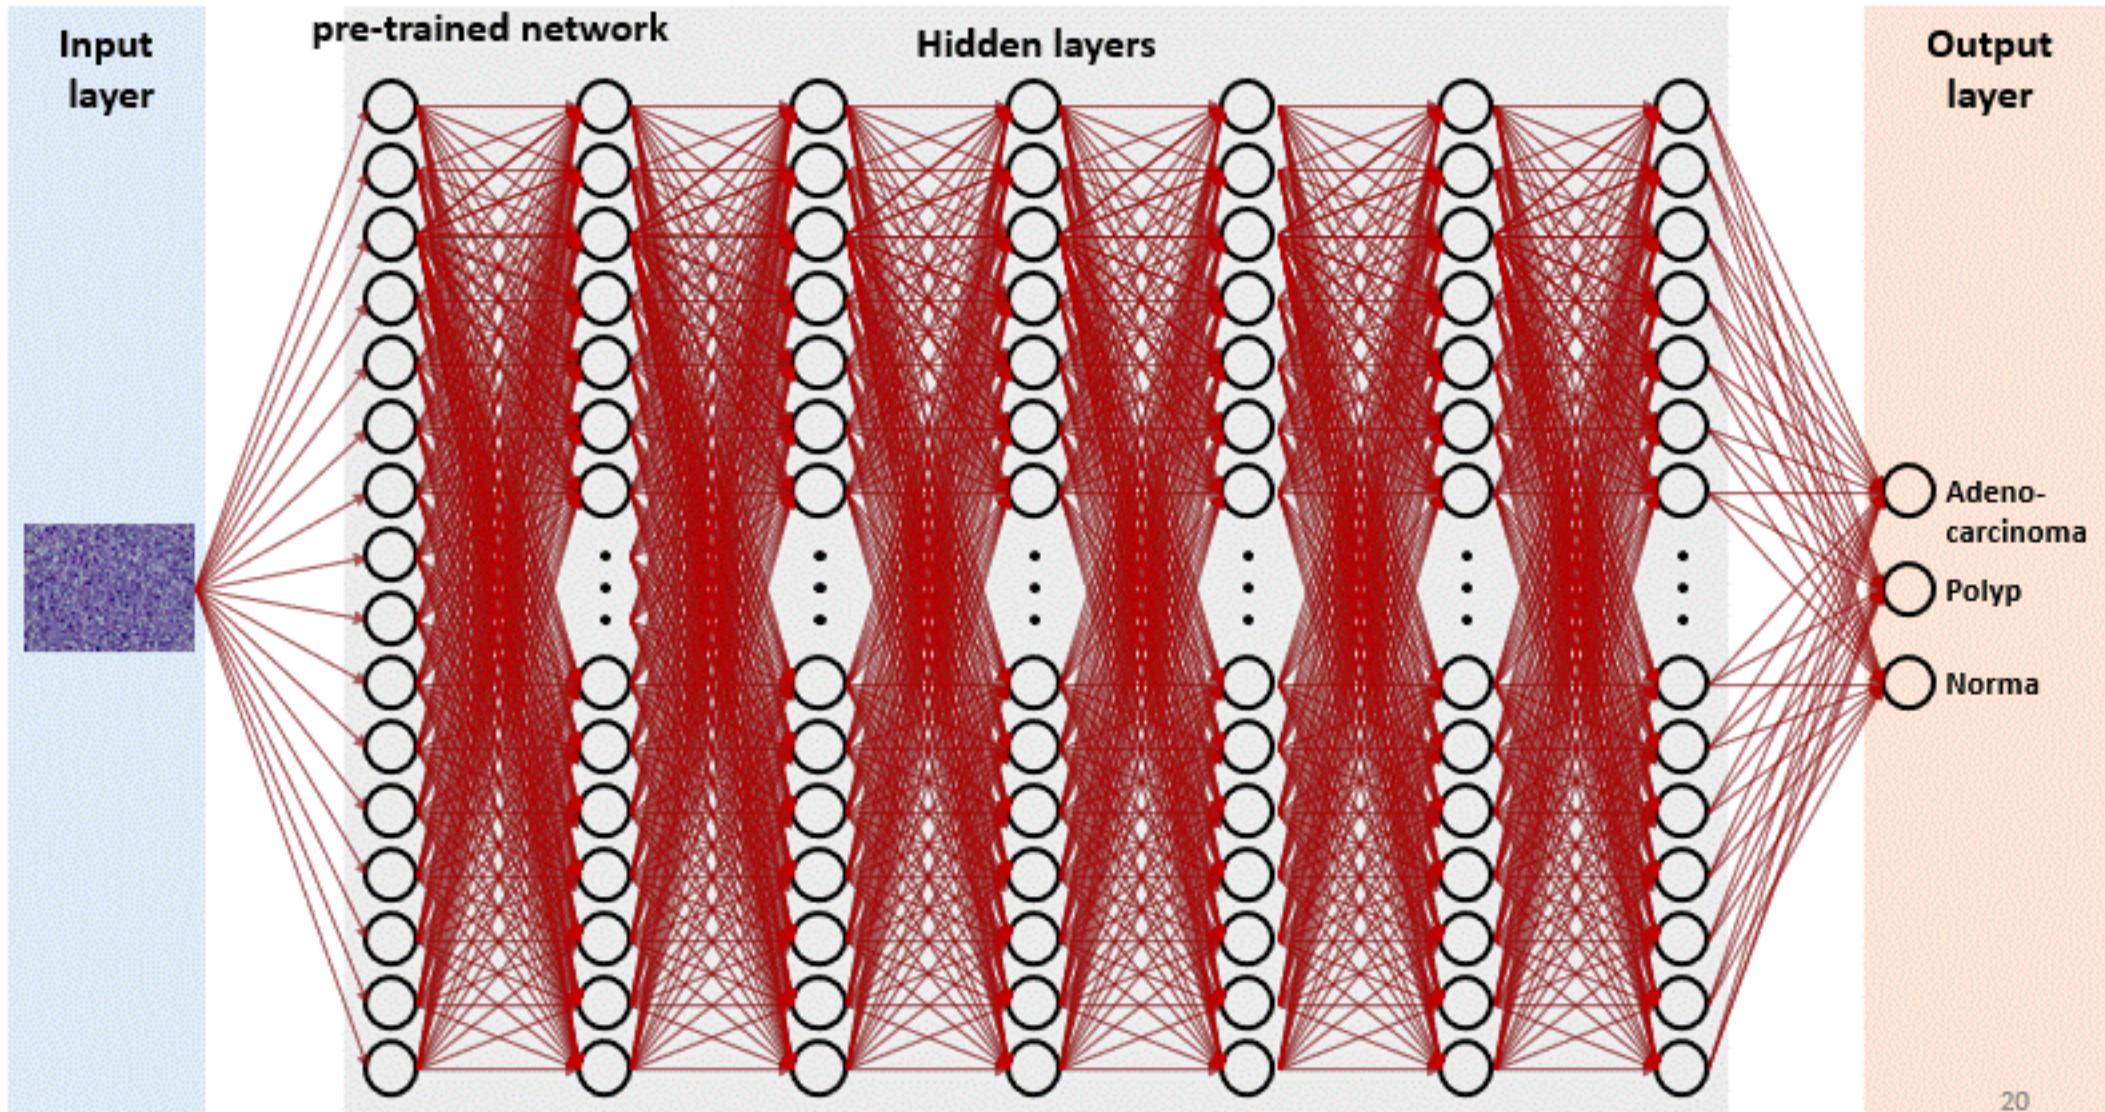


TRANSFER LEARNING



TRANSFER LEARNING

Weight initialization with a
pre-trained network



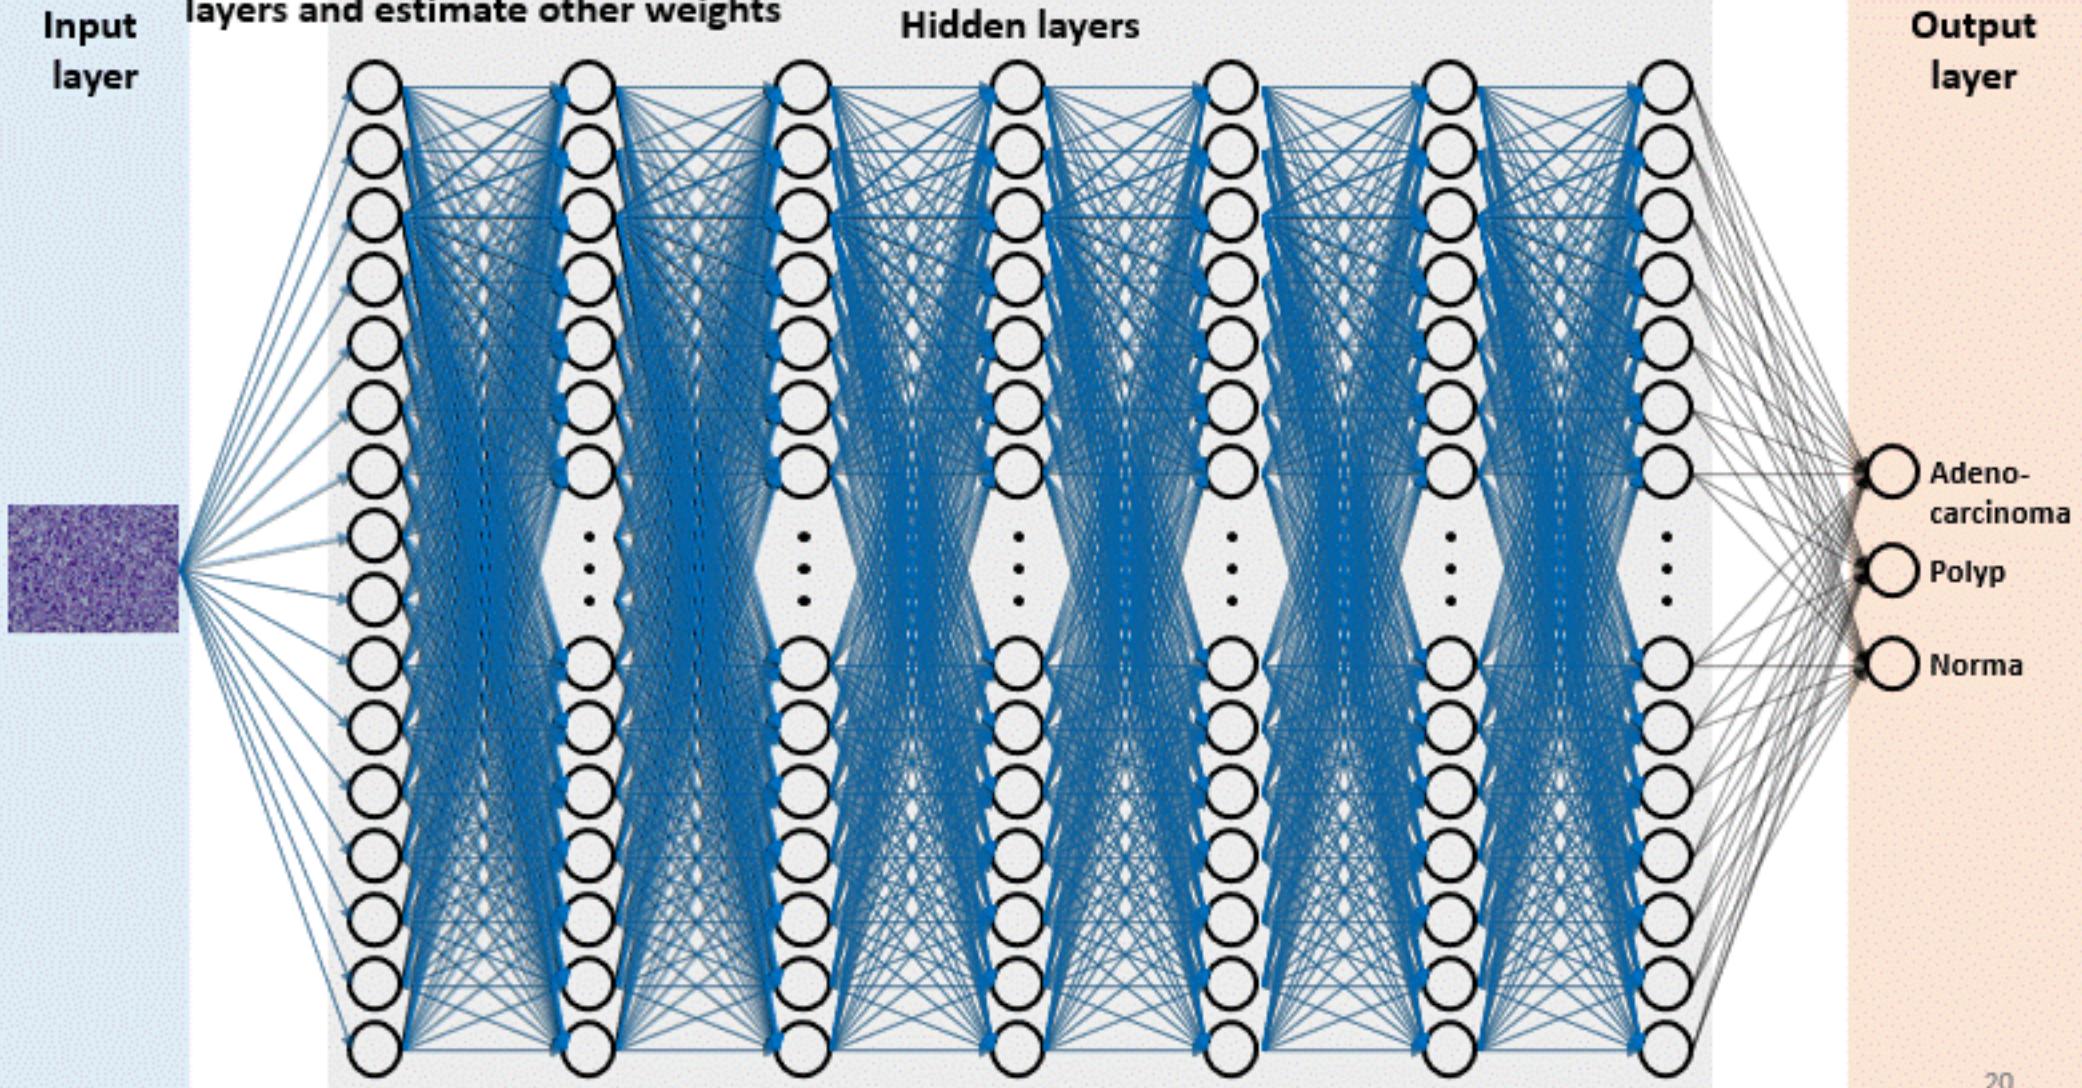
TRANSFER LEARNING

Freeze weight estimation for first layers and estimate other weights

Hidden layers

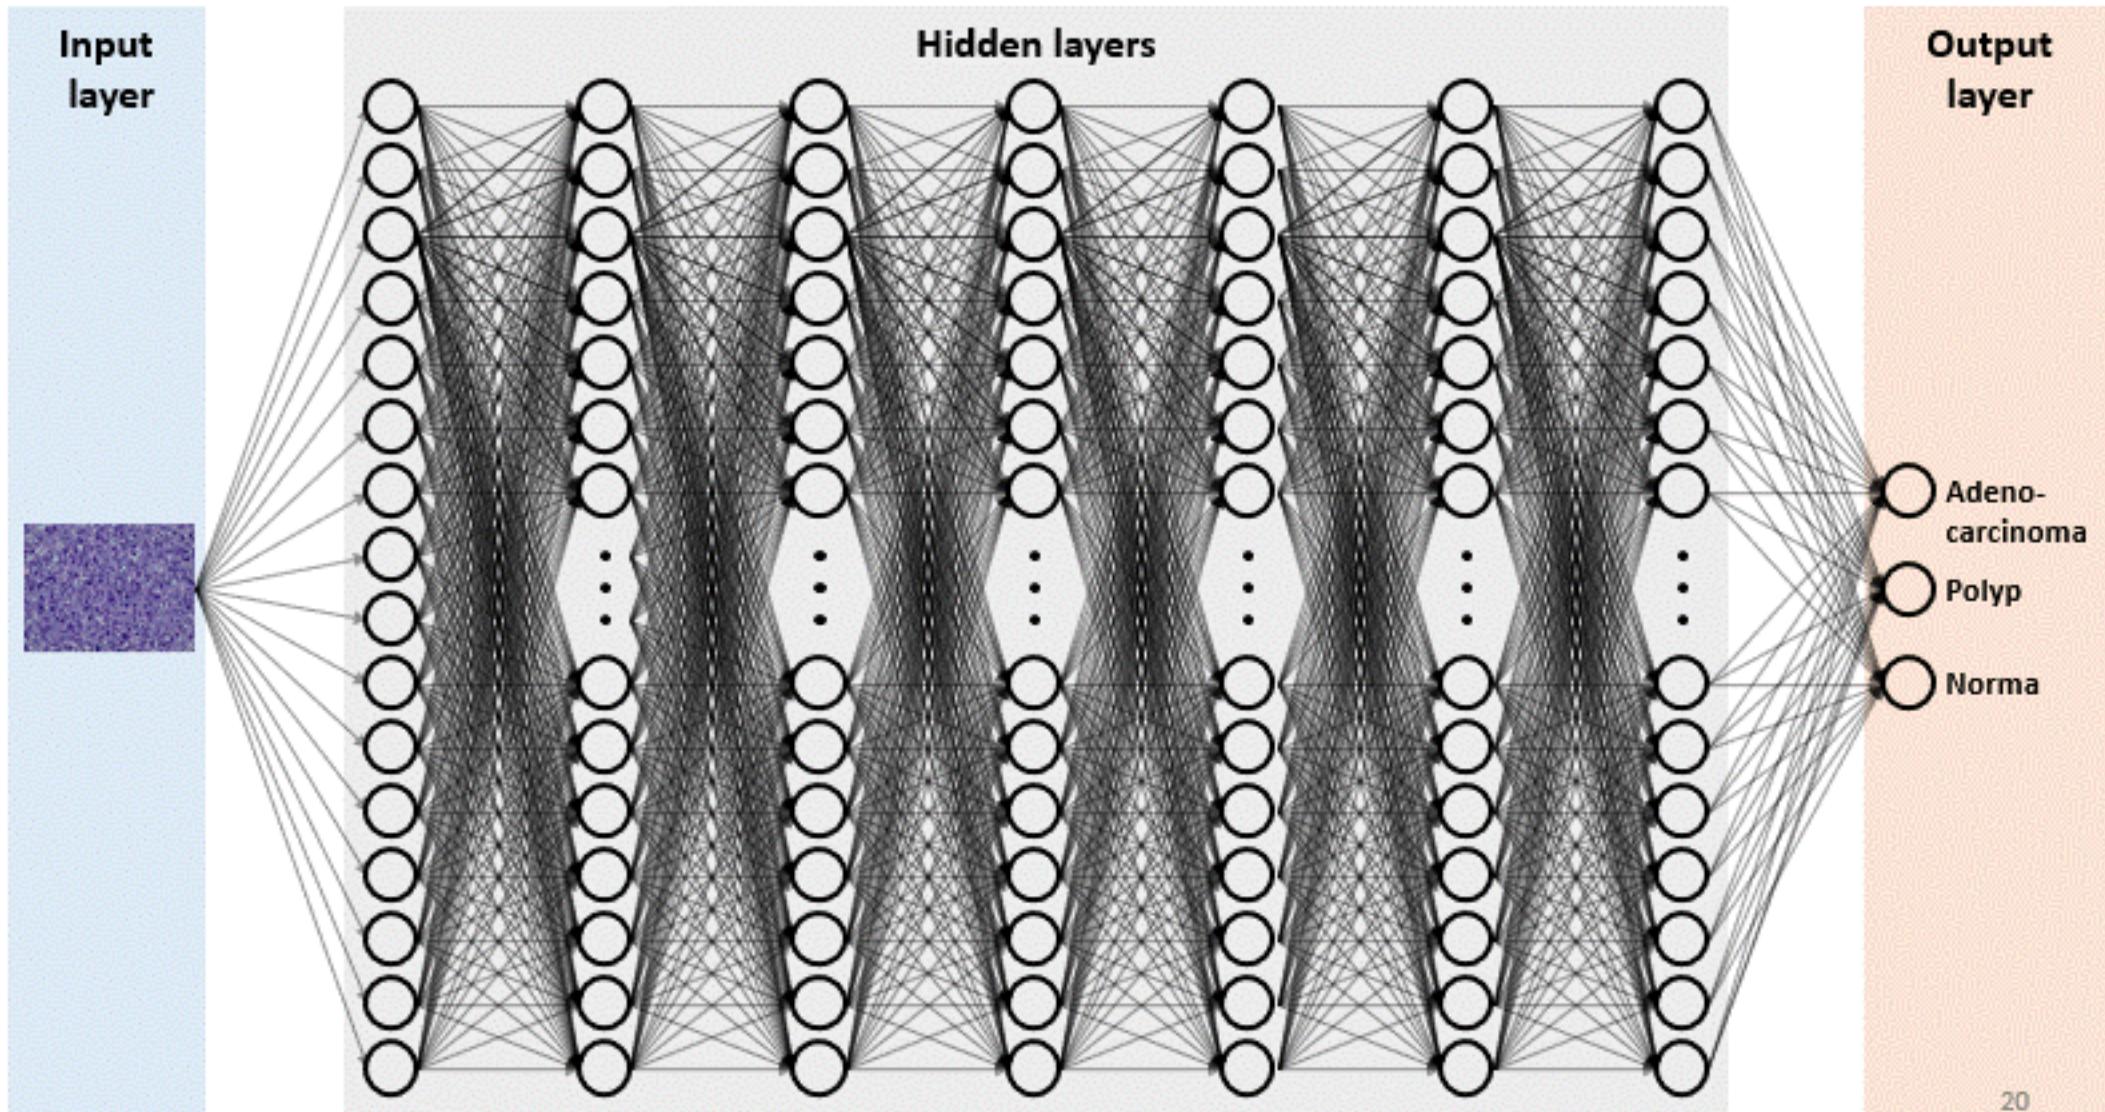
Output layer

Input layer



TRANSFER LEARNING

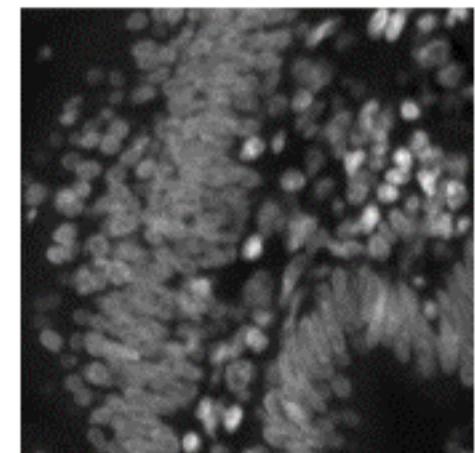
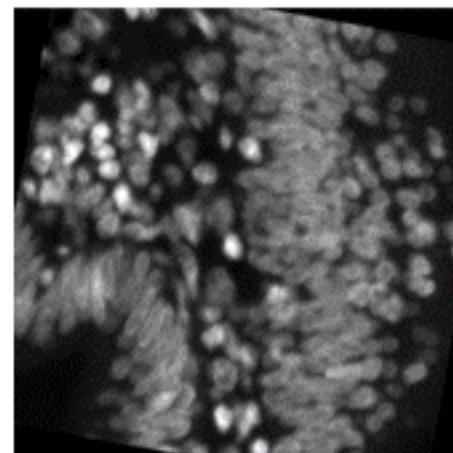
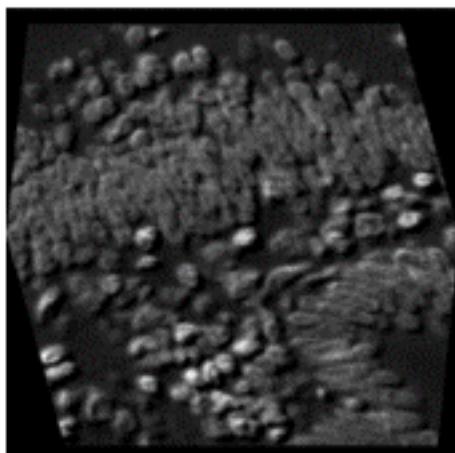
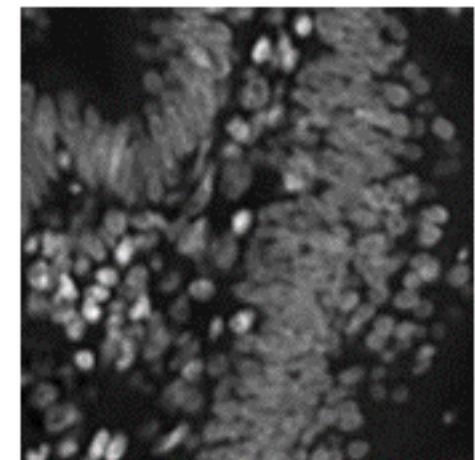
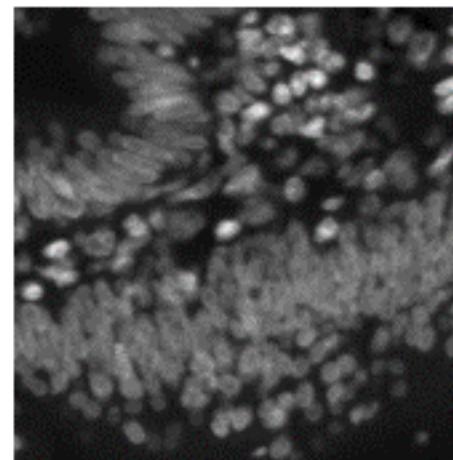
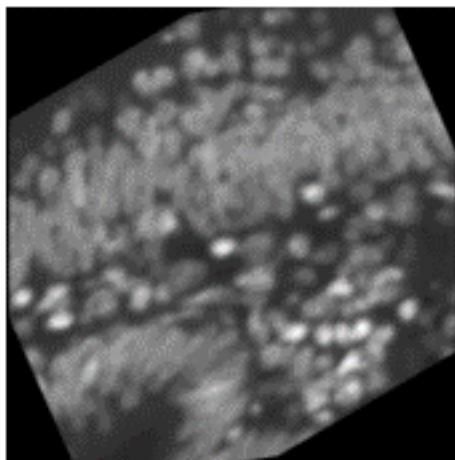
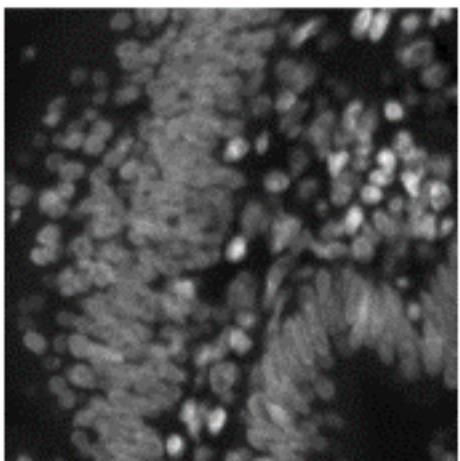
Estimate all weights



DATA AUGMENTATION

Augmented images

Original image



PRE-TRAINED MODELS



Figure 4: Example cellpose segmentations for 36 test images. The ground truth masks segmented by a Human operator are shown in yellow, and the predicted masks are shown in dotted red line. Compare to StarCell and Mask-RCNN in Figure 5d and Figure 5e.

The homepage of the Biolimage Model Zoo features a central banner with the text:

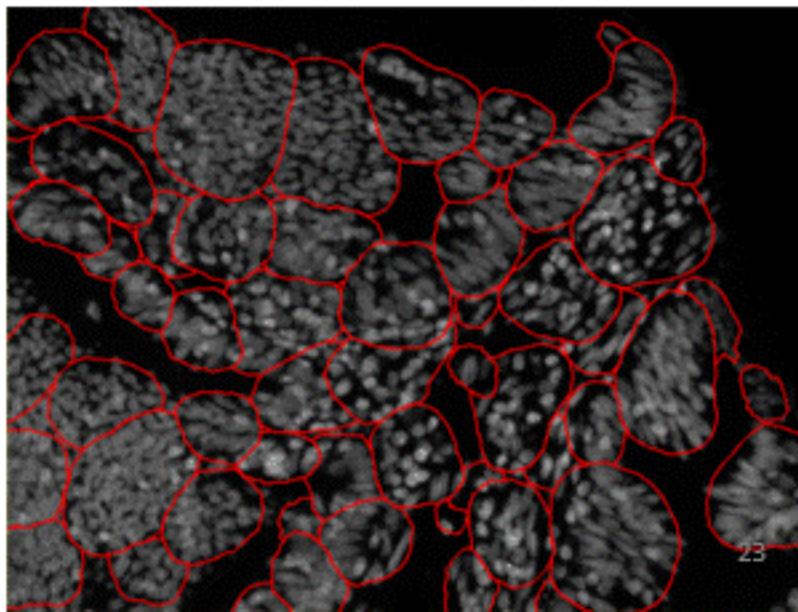
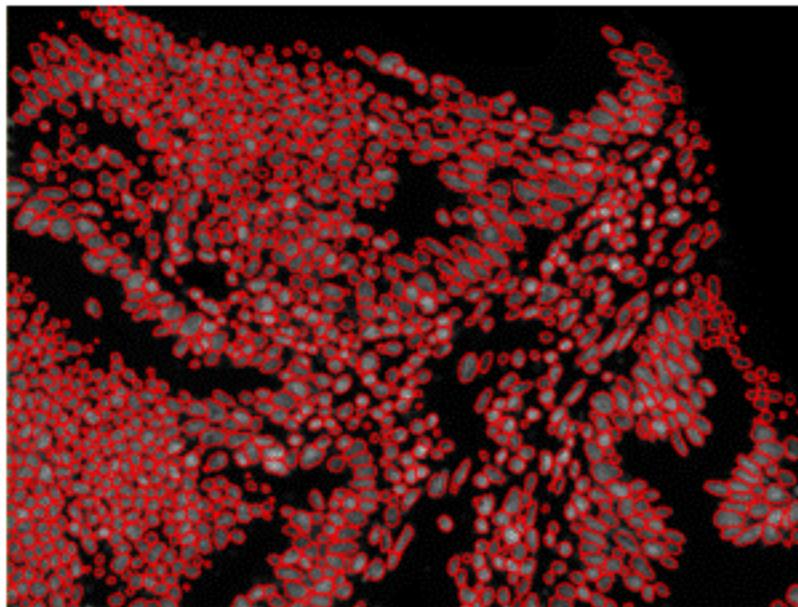
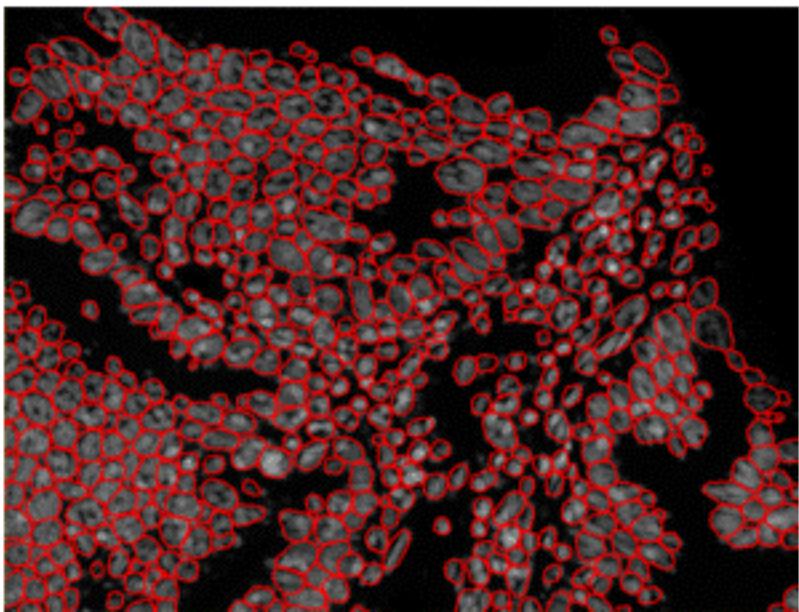
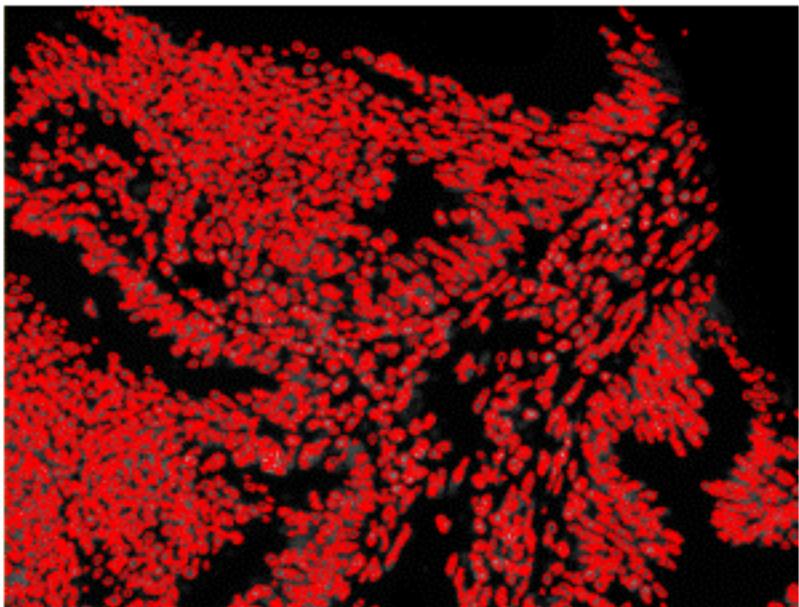
- Biolimage Model Zoo**
- Advanced AI models in one-click

Below the banner are several calls to action:

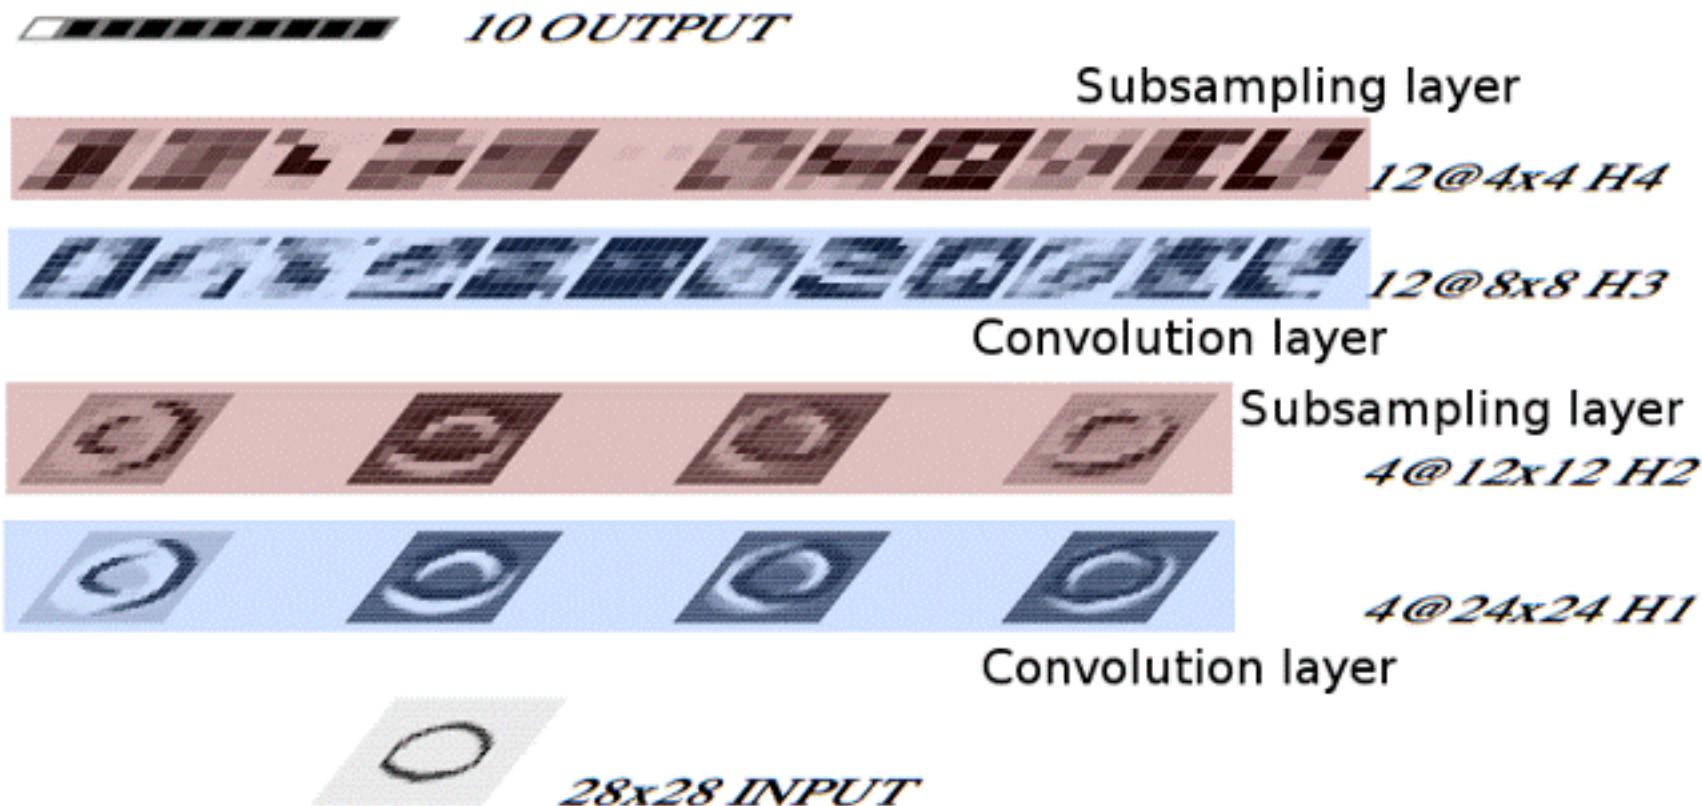
- Integrate with Fiji, ilastik, ImJoy and more
- Try model instantly with BioEngine
- Contribute your models via Github
- Link models to datasets and applications

A prominent "Explore the Zoo" button is located below the text. To the right of the text, there is a cartoon illustration of a giraffe, elephant, and tree. At the bottom, there is a "Community Partners" section featuring logos for various organizations, followed by navigation links for "All", "models", "applications", "datasets", and search/filter options.

WARNINGS



CONVOLUTIONAL NEURAL NETWORK



LeCun et al., Handwritten digit recognition with a back-propagation network, *Neural Information Processing Systems*, 1989

4,635 pixels and 2,578 parameters (much lower than fully connected network)

MOBILENETV2

Input	Operator	t	c	n	s
$224^2 \times 3$	conv2d	-	32	1	2
$112^2 \times 32$	bottleneck	1	16	1	1
$112^2 \times 16$	bottleneck	6	24	2	2
$56^2 \times 24$	bottleneck	6	32	3	2
$28^2 \times 32$	bottleneck	6	64	4	2
$14^2 \times 64$	bottleneck	6	96	3	1
$14^2 \times 96$	bottleneck	6	160	3	2
$7^2 \times 160$	bottleneck	6	320	1	1
$7^2 \times 320$	conv2d 1x1	-	1280	1	1
$7^2 \times 1280$	avgpool 7x7	-	-	1	-
$1 \times 1 \times 1280$	conv2d 1x1	-	k	-	-

Table 2: MobileNetV2 : Each line describes a sequence of 1 or more identical (modulo stride) layers, repeated n times. All layers in the same sequence have the same number c of output channels. The first layer of each sequence has a stride s and all others use stride 1. All spatial convolutions use 3×3 kernels. The expansion factor t is always applied to the input size as described in Table 1.

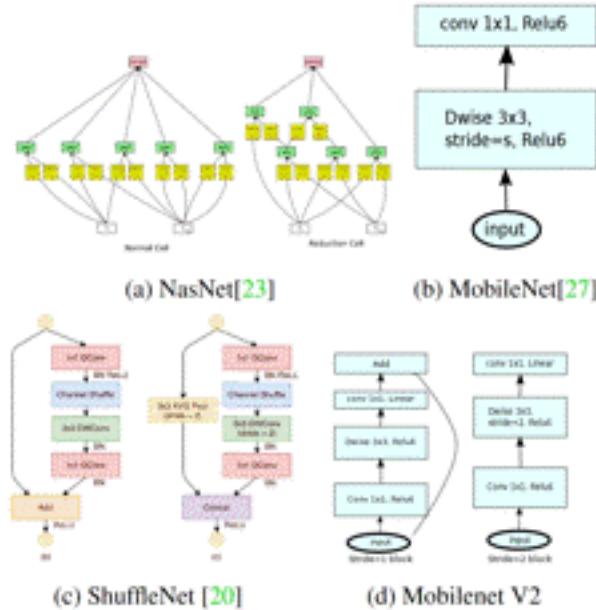


Figure 4: Comparison of convolutional blocks for different architectures. ShuffleNet uses Group Convolutions [20] and shuffling, it also uses conventional residual approach where inner blocks are narrower than output. ShuffleNet and NasNet illustrations are from respective papers.

UNET

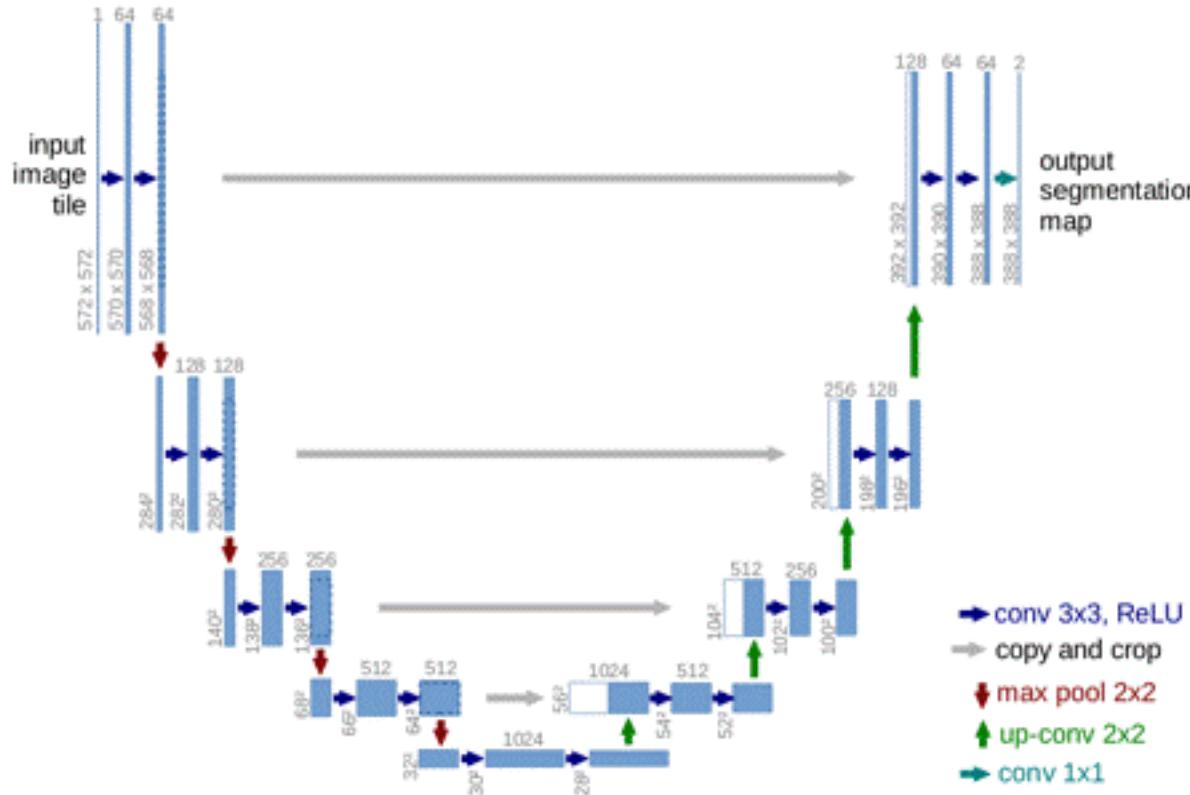


Fig. 1. U-net architecture (example for 32x32 pixels in the lowest resolution). Each blue box corresponds to a multi-channel feature map. The number of channels is denoted on top of the box. The x-y-size is provided at the lower left edge of the box. White boxes represent copied feature maps. The arrows denote the different operations.

UNET

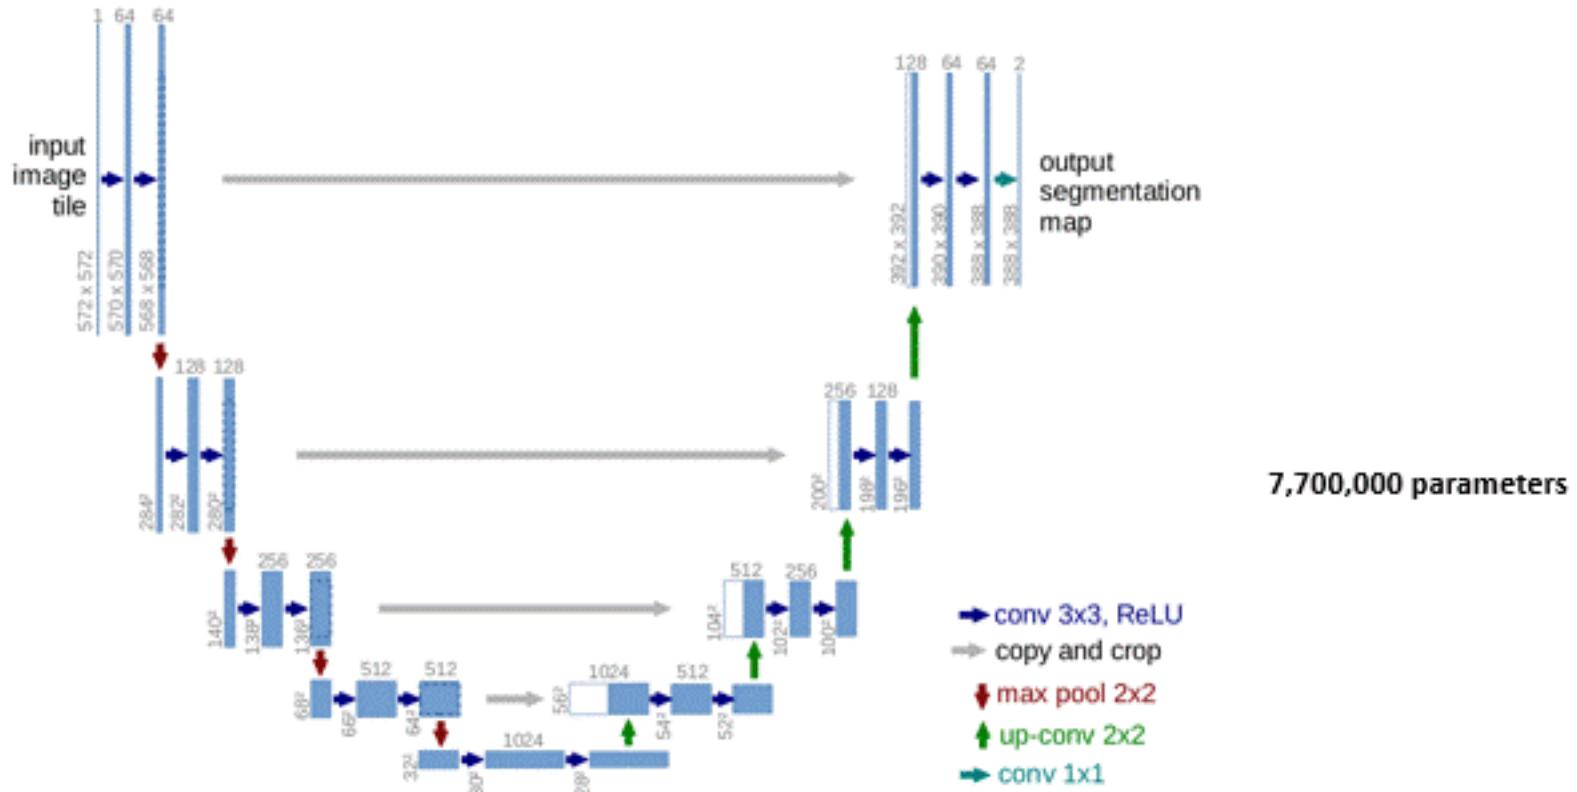


Fig. 1. U-net architecture (example for 32x32 pixels in the lowest resolution). Each blue box corresponds to a multi-channel feature map. The number of channels is denoted on top of the box. The x-y-size is provided at the lower left edge of the box. White boxes represent copied feature maps. The arrows denote the different operations.

STAR DIST – 2D

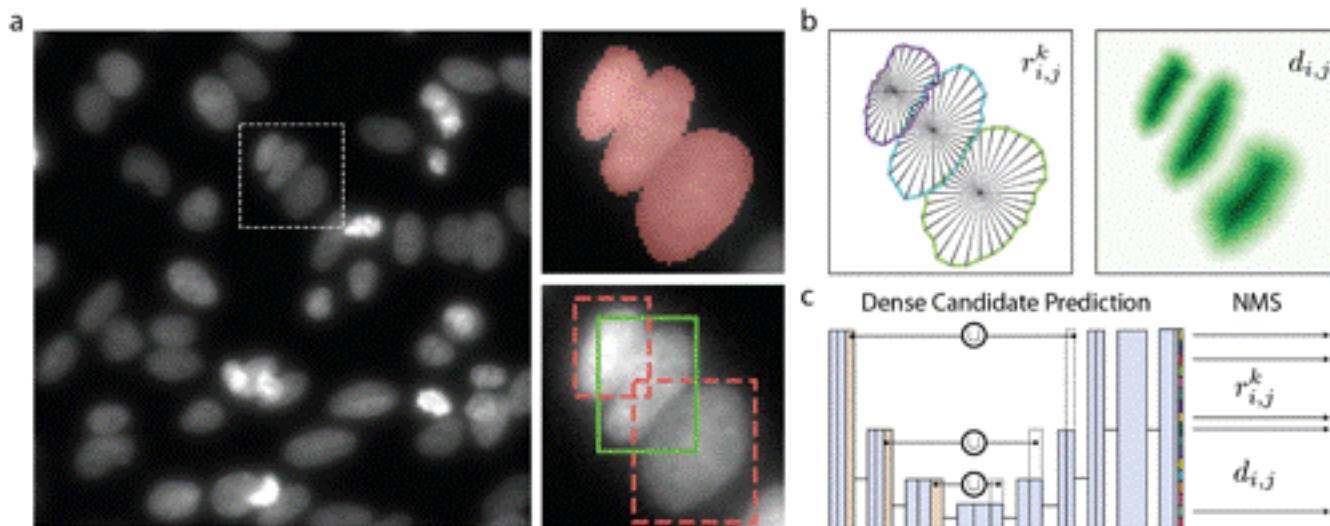
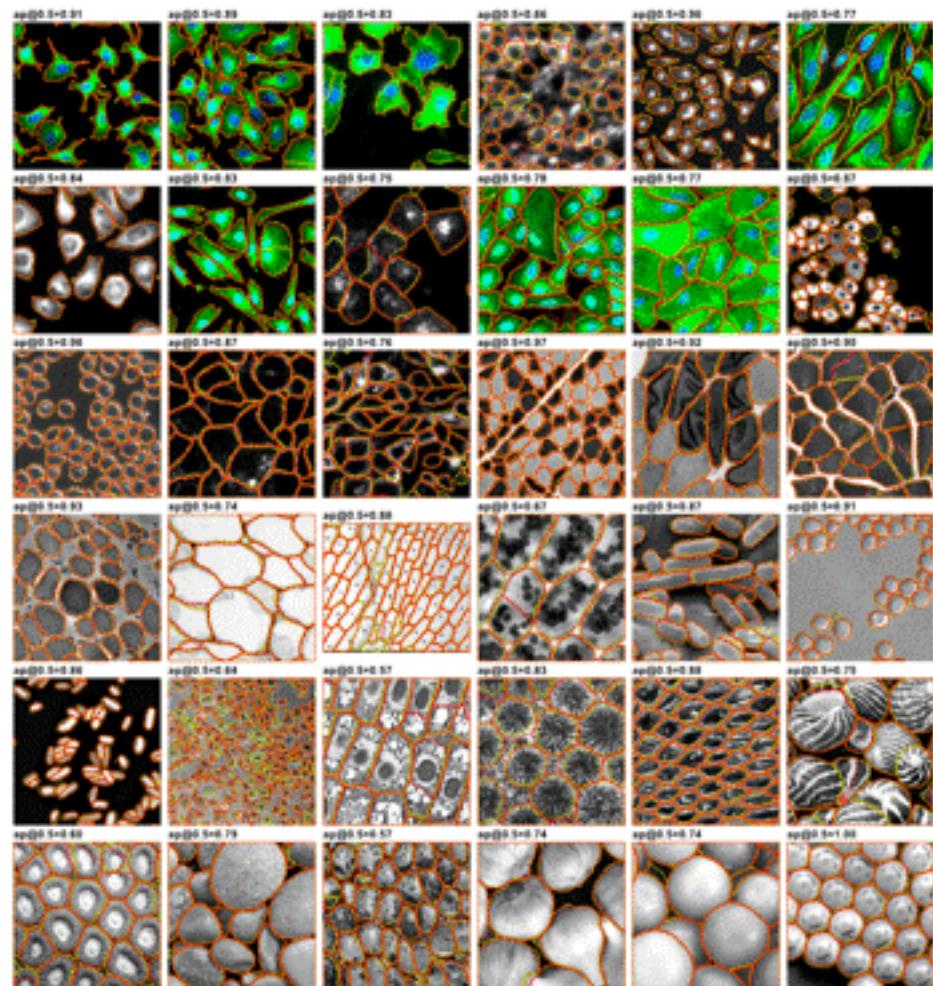
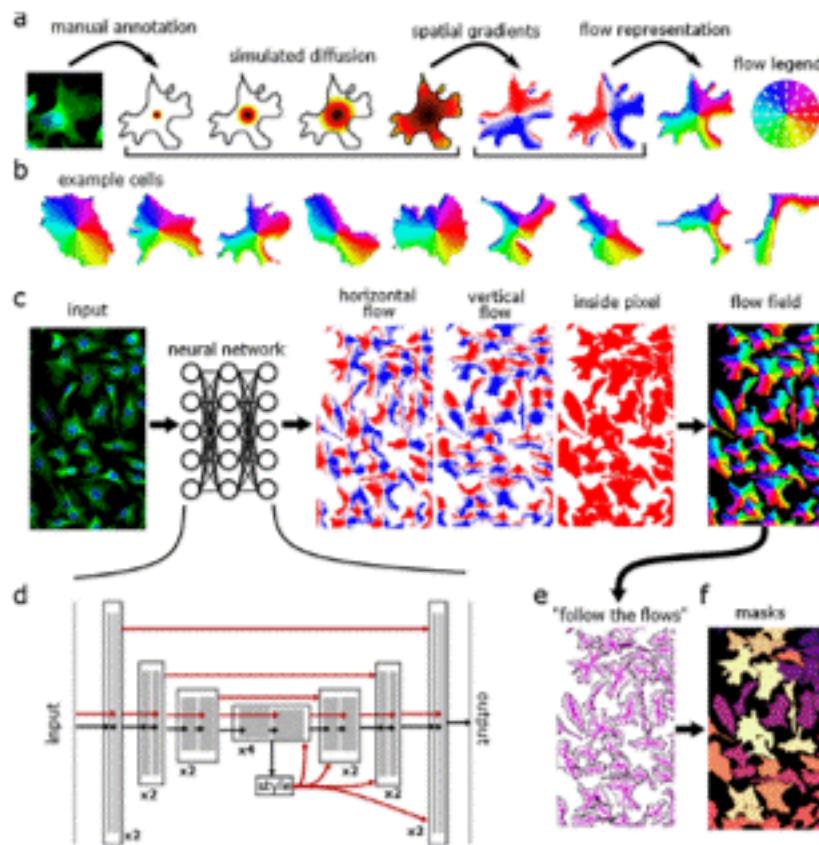


Fig. 1: (a) Potential segmentation errors for images with crowded nuclei: Merging of touching cells (upper right) or suppression of valid cell instances due to large overlap of bounding box localization (lower right). (b) The proposed STAR DIST method predicts object probabilities $d_{i,j}$ and star-convex polygons parameterized by the radial distances $r_{i,j}^k$. (c) We densely predict $r_{i,j}^k$ and $d_{i,j}$ using a simple U-Net architecture [15] and then select the final instances via non-maximum suppression (NMS).

CELLPOSE



Stringer *et al.*, Cellpose: a generalist algorithm for cellular segmentation. *Nature Methods*, 2021

OMNIPOSE

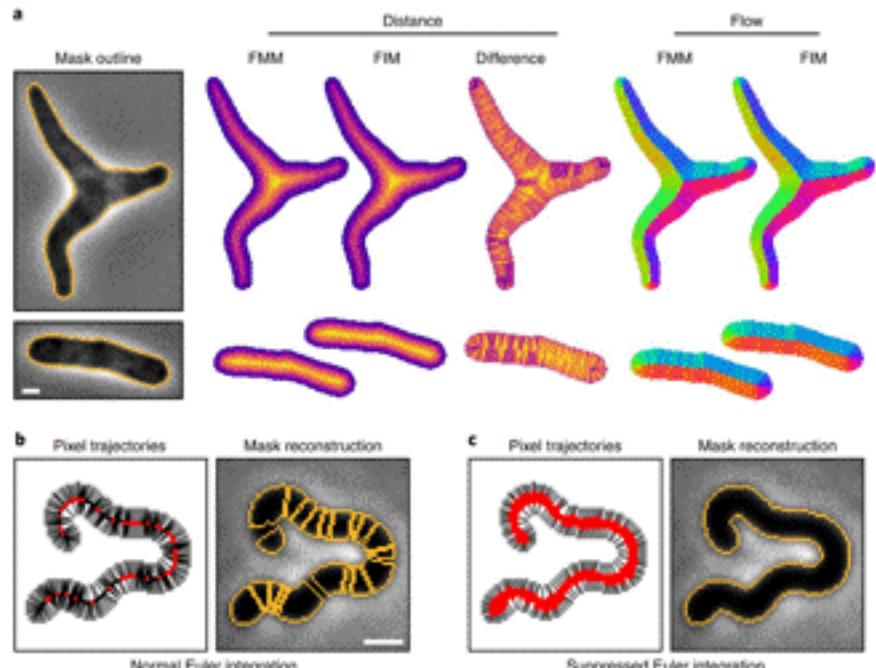


Fig. 2 | Core innovations of Omnipose. **a**, Comparison of distance field algorithms and corresponding flow fields on ground-truth masks. FMM produces ridges in the distance field resulting from pixelation on the cell mask boundary. Our smooth FIM algorithm minimizes these features. The difference image ($\text{OFIM} - \text{FMM}$) highlights artifacts in the FMM method. Flow fields are calculated as the normalized gradient of the distance field. Boundary pixelation affects the FMM flow field deep into the cell, regardless of cell size. **b,e**, Comparison of mask

reconstruction algorithms on a smooth flow field. Boundary pixel trajectories and resulting mask outlines from standard Euler integration (**b**). Trajectories and mask outlines under suppressed Euler integration (**c**). Red dots indicate the final positions of all cell pixels, not only the boundary pixels for which trajectories are displayed. Bacteria displayed are *E. coli* CS703-1 (**a**) and *H. pylori* (**b,c**) both treated with aztreonam. Scale bars, 1 μm . Images are representative of 1,299 *E. coli* and 701 *H. pylori* cells in the total ground-truth dataset, respectively.

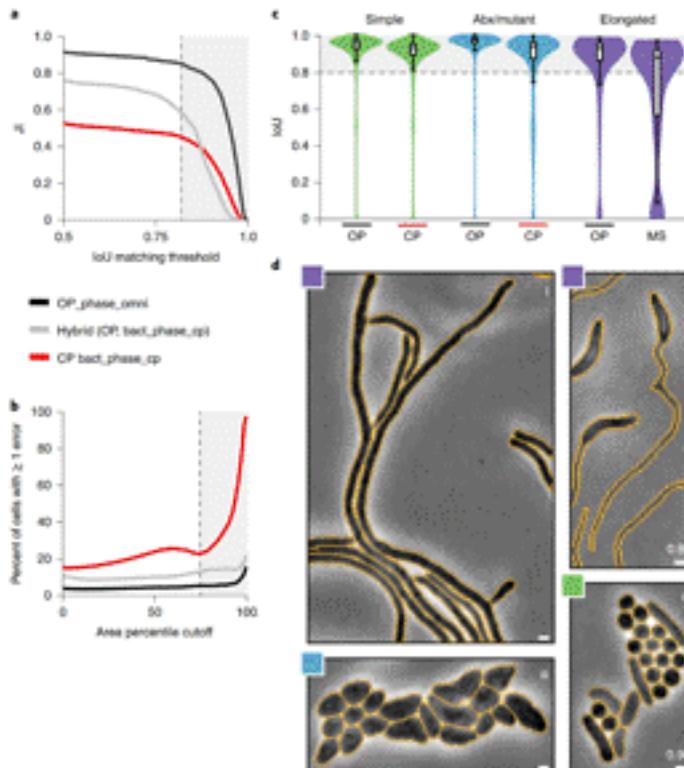


Fig. 3 | Omnipose substantially outperforms Cellpose on elongated cells. **a**, Overall performance of Omnipose (OP) (*bact_phase_omni*) and Cellpose (CP) (*bact_phase_cp*) measured by JI. The hybrid method (gray) uses the original center-seeking flow output of *bact_phase_cp* and the mask reconstruction of Omnipose. Gray box represents *bact_phase_cp* and the mask reconstruction of Omnipose. Gray box represents the top quartile. **c**, Omnipose JI distribution on the *bact_phase* dataset compared to the next highest performing

algorithm in each of three cell categories (simple, $n = 12,869$; Abn/mutant, $n = 6,138$; and elongated, $n = 530$). Boxes centered on median from Q1 to Q3, whiskers from $Q1 - 1.5 \times \text{IQR}$ to $Q3 + 1.5 \times \text{IQR}$. **d**, Example micrographs and Omnipose segmentation. Mean matched JI values shown. Bacteria displayed are *S. pristinaespiralis* (**i**), *Caulobacter crescentus* grown in HBGG medium (**j**), *S. flexneri* treated with A22 (**k**) and a mix of *Pseudomonas aeruginosa*, *Staphylococcus aureus*, *K. pneumoniae*, and *Bacillus subtilis* (**l**). HBGG, Hulme base-midazole-buffered glucose-glutamate. Scale bars, 1 μm .

FOUNDATION MODEL FOR INSTANCE SEGMENTATION

Segment Anything

Alexander Kirillov^{1,2,4} Eric Mintun² Nikhila Ravi^{1,2} Hanzi Mao² Chloe Rolland³ Laura Gustafson³
Tete Xiao³ Spencer Whitehead Alexander C. Berg Wan-Yen Lo Piotr Dollár⁴ Ross Girshick⁴
¹project lead ²joint first author ³equal contribution ⁴directional lead

Meta AI Research, FAIR

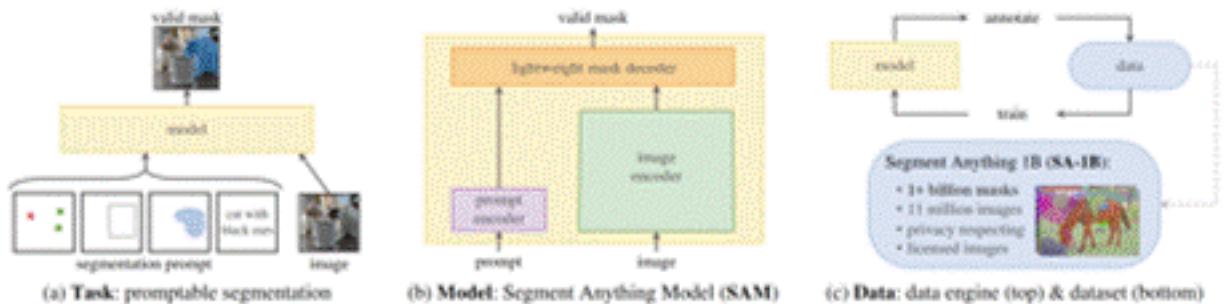


Figure 1: We aim to build a foundation model for segmentation by introducing three interconnected components: a promptable segmentation *task*, a segmentation *model* (SAM) that powers data annotation and enables zero-shot transfer to a range of tasks via prompt engineering, and a *data engine* for collecting SA-1B, our dataset of over 1 billion masks.

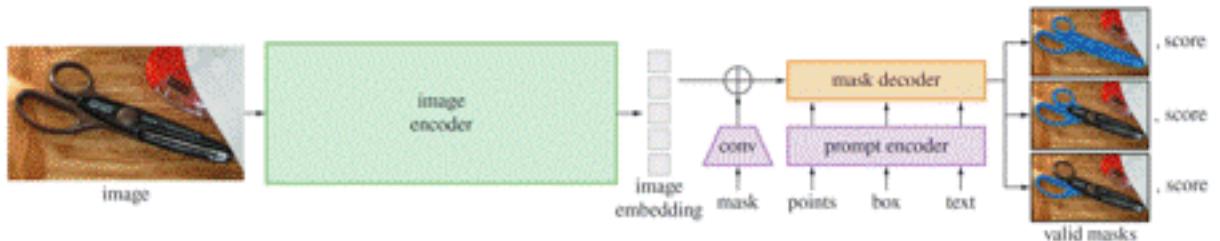


Figure 4: Segment Anything Model (SAM) overview. A heavyweight image encoder outputs an image embedding that can then be efficiently queried by a variety of input prompts to produce object masks at amortized real-time speed. For ambiguous prompts corresponding to more than one object, SAM can output multiple valid masks and associated confidence scores.

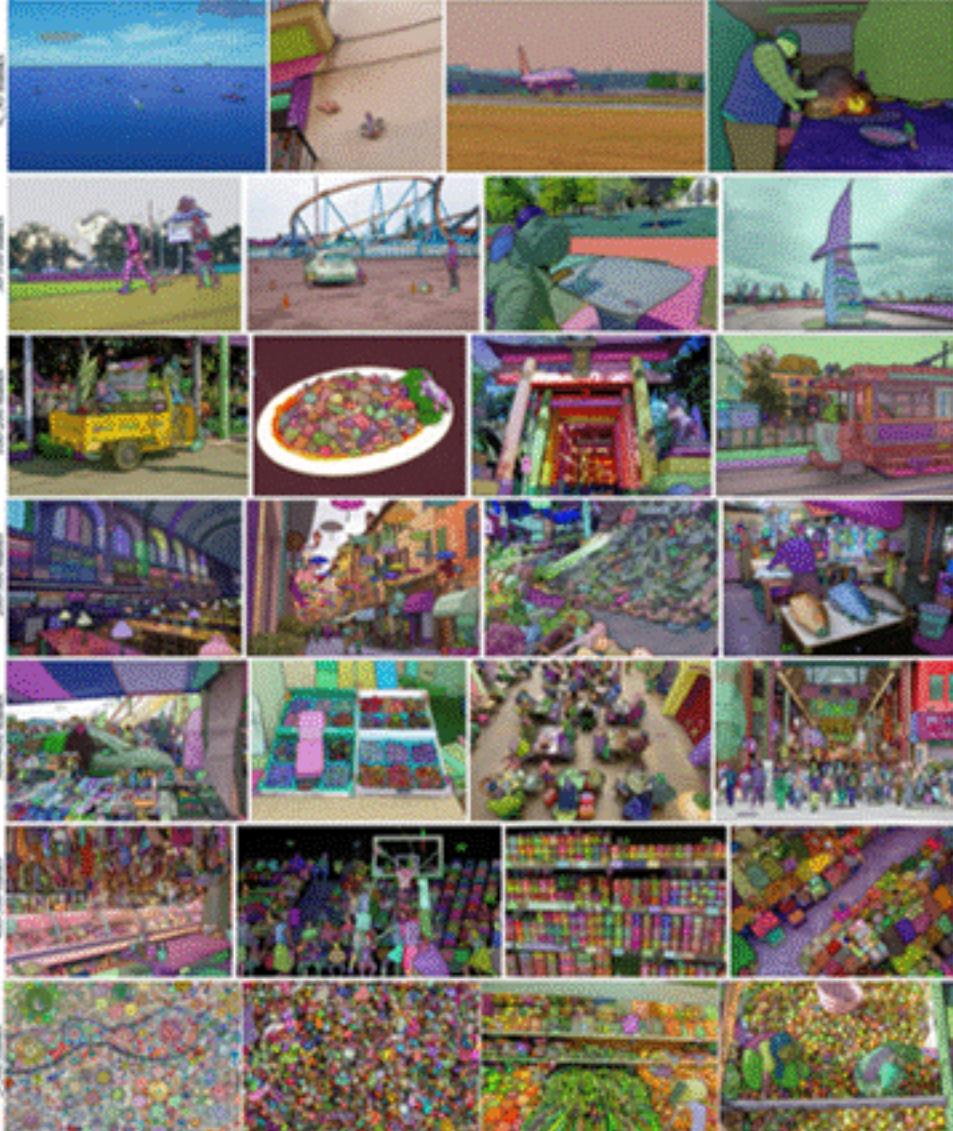


Figure 2: Example images with overlaid masks from our newly introduced dataset, SA-1B. SA-1B contains 11M diverse, high-resolution, licensed, and privacy protecting images and 1.1B high-quality segmentation masks. These masks were annotated *fully automatically* by SAM, and as we verify by human ratings and numerous experiments, are of high quality and diversity. We group images by number of masks per image for visualization (there are ~100 masks per image on average).

FOUNDATION MODEL FOR INSTANCE SEGMENTATION

Segment Anything

Alexander Kirillov^{1,2,4} Eric Mintun² Nikhila Ravi^{1,2} Hanzi Mao² Chloe Rolland³ Laura Gustafson³
Tete Xiao³ Spencer Whitehead Alexander C. Berg Wan-Yen Lo Piotr Dollár⁴ Ross Girshick⁴
¹project lead ²joint first author ³equal contribution ⁴directional lead

Meta AI Research, FAIR

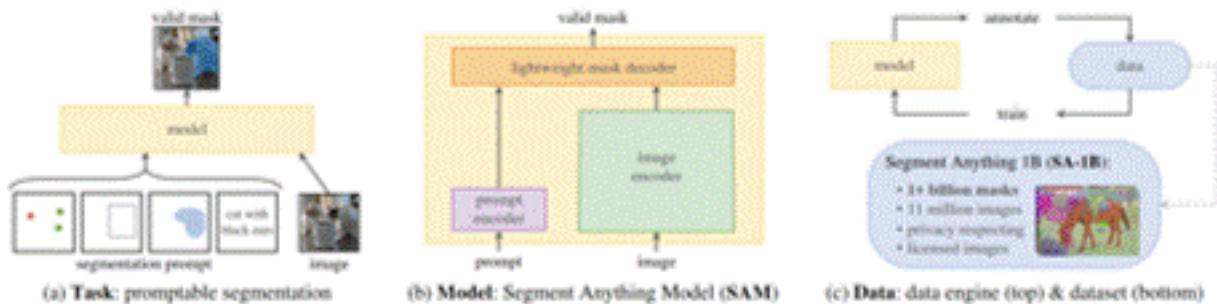


Figure 1: We aim to build a foundation model for segmentation by introducing three interconnected components: a promptable segmentation *task*, a segmentation *model* (SAM) that powers data annotation and enables zero-shot transfer to a range of tasks via prompt engineering, and a *data engine* for collecting SA-1B, our dataset of over 1 billion masks.

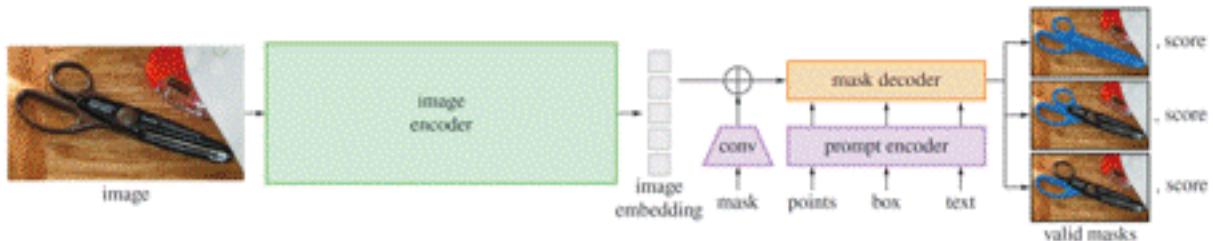


Figure 4: Segment Anything Model (SAM) overview. A heavyweight image encoder outputs an image embedding that can then be efficiently queried by a variety of input prompts to produce object masks at amortized real-time speed. For ambiguous prompts corresponding to more than one object, SAM can output multiple valid masks and associated confidence scores.



Figure 3: Each column shows 3 valid masks generated by SAM from a single ambiguous point prompt (green circle).